

Stochastic Calculus

Quasar Chunawala

July 24, 2023

Abstract

The most important results and ideas in basic mathematical finance.

1 Measure.

1.1 Null Sets.

Definition 1.1. (*Null Set*). A *null set* is a set that can be covered by a sequence of intervals of arbitrarily small total length. Given any $\epsilon > 0$, there exists a sequence of intervals $(I_n)_{n \geq 1}$ such that:

$$A \subseteq \bigcup_{n=1}^{\infty} I_n$$

and

$$\sum_{n=1}^{\infty} l(I_n) < \epsilon$$

Problem 1.1. Show that we get an equivalent notion if in the above definition we replace the word intervals by any of these: open-intervals, closed-intervals, intervals of the form $(a, b]$ or intervals of the form $[a, b)$.

Proof. Let A be a null set. Then, we can cover it by a sequence of intervals, such that total length of the cover can be made as small as we please. Mathematically,

$$\forall \epsilon > 0, \exists (I_n)_{n \geq 1}, A \subseteq \bigcup_{n=1}^{\infty} I_n, \text{ such that } \sum_{n=1}^{\infty} l(I_n) < \frac{\epsilon}{2}$$

Let $I_n = [a_n, b_n]$ and define:

$$J_n := \left(a_n - \frac{\epsilon}{2^{n+2}}, b_n + \frac{\epsilon}{2^{n+2}} \right)$$

Since, $I_n \subseteq J_n$, it follows that:

$$A \subseteq \bigcup_{n=1}^{\infty} I_n \subseteq \bigcup_{n=1}^{\infty} J_n$$

Moreover,

$$\begin{aligned} l(J_n) &= l(I_n) + \frac{\epsilon}{2^{n+1}} \\ \sum_{n=1}^{\infty} l(J_n) &= \sum_{n=1}^{\infty} l(I_n) + \frac{\epsilon}{2^2} \left(1 + \frac{1}{2} + \frac{1}{2^2} + \dots \right) \\ &< \frac{\epsilon}{2} + \frac{\epsilon}{2^2} \cdot \frac{1}{1 - 1/2} \\ &= \epsilon \end{aligned}$$

Consequently, A can be covered by a sequence of open intervals, whose total length can be made arbitrarily small. This closes the proof. \square

Theorem 1.1. *If $(N_n)_{n \geq 1}$ is a sequence of null sets, then their countable union*

$$N = \bigcup_{n=1}^{\infty} N_n$$

is also null.

Proof. Since N_1 is null, there exists a sequence of intervals (I_k^1) such that $N_1 \subseteq \bigcup_{k=1}^{\infty} I_k^1$ and $\sum_{k=1}^{\infty} l(I_k^1) < \frac{\epsilon}{2^2}$.

Since N_2 is null, there exists a sequence of intervals (I_k^2) such that $N_2 \subseteq \bigcup_{k=1}^{\infty} I_k^2$ and $\sum_{k=1}^{\infty} l(I_k^2) < \frac{\epsilon}{2^3}$.

Since N_j is null, there exists a sequence of intervals (I_k^j) such that $N_j \subseteq \bigcup_{k=1}^{\infty} I_k^j$ and $\sum_{k=1}^{\infty} l(I_k^j) < \frac{\epsilon}{2^{2+j}}$.

Clearly, we have:

$$\bigcup_{j=1}^{\infty} N_j \subseteq \bigcup_{j=1}^{\infty} \bigcup_{k=1}^{\infty} I_k^j$$

Moreover,

$$\begin{aligned} \sum_{j=1}^{\infty} \sum_{k=1}^{\infty} l(I_k^j) &< \frac{\epsilon}{2^2} + \frac{\epsilon}{2^3} + \dots \\ &= \frac{\epsilon}{2^2} \left[1 + \frac{1}{2} + \frac{1}{2^2} + \dots \right] \\ &= \frac{\epsilon}{2} < \epsilon \end{aligned}$$

Consequently, N is a null set. □

A singleton set $\{x\}$ is a null set - let $I_1 = [x - \frac{\epsilon}{4}, x + \frac{\epsilon}{4}]$, $I_n = [x, x]$ for $n \geq 2$. Thus, any countable set is a null set, and null sets appear to be closely related to countable sets - this is no surprise, as any proper interval is uncountable, so any countable subset is quite sparse when compared with an interval, hence makes no real contribution to its *length*.

However, uncountable sets can be null, provided their points are sufficiently *sparsely distributed*, as the following example due to Cantor shows:

1. Start with the interval $C_0 = [0, 1]$, remove the open middle one-third, that is the interval $(\frac{1}{3}, \frac{2}{3})$, obtaining C_1 which consists of two intervals $[0, \frac{1}{3}]$ and $[\frac{2}{3}, 1]$.
2. Next, remove the middle third of each of these two intervals leaving C_2 , consisting of four intervals $[0, \frac{1}{9}]$, $[\frac{2}{9}, \frac{3}{9}]$, $[\frac{6}{9}, \frac{7}{9}]$ and $[\frac{8}{9}, 1]$.

3. At the n th stage, we have a set C_n consisting of 2^n disjoint closed intervals, each of length $\frac{1}{3^n}$. Thus, the total length of C_n is $\left(\frac{2}{3}\right)^n$.

We call

$$C = \bigcap_{n=1}^{\infty} C_n$$

the *Cantor set*. Now, we show that C is null as promised.

Given any $\epsilon > 0$, choose n such that $\left(\frac{2}{3}\right)^n < \epsilon$. Since, $C \subseteq C_n$ and C_n is a union of disjoint intervals of total length less than ϵ , we see that C is a null set. All that remains to be checked is that C is an uncountable set.

Problem 1.2. Prove that C is uncountable.

Proof. Let $x \in C$ be an arbitrary point.

Starting with $I_0 = [0, 1]$, for all $n \in \mathbf{N}$, define the sequence of intervals (I_n) , where $I_n = [a_n, b_n]$, (L_n) and (R_n) as:

$$\begin{aligned} L_{n+1} &= \left[a_n, a_n + \frac{1}{3^{n+1}} \right] \\ R_{n+1} &= \left[b_n - \frac{1}{3^{n+1}}, b_n \right] \\ I_{n+1} &= \begin{cases} L_{n+1} & \text{if } x \in L_{n+1} \\ R_{n+1} & \text{otherwise} \end{cases} \end{aligned}$$

Clearly, the left-end point of I_{n+1} , is a_n , if $x \in L_{n+1}$, otherwise it is $b_n - \frac{1}{3^{n+1}} = a_n + \frac{1}{3^n} - \frac{1}{3^{n+1}} = a_n + \frac{2}{3^{n+1}}$. To summarize:

$$a_{n+1} = \begin{cases} a_n + \frac{0}{3^{n+1}} & \text{if } x \in L_{n+1} \\ a_n + \frac{2}{3^{n+1}} & \text{if } x \in R_{n+1} \end{cases}$$

We have that, since $C \subseteq [0, 1]$, it implies $x \in I_0$. By construction, if $x \in I_n$, then $x \in I_{n+1}$.

Hence,

$$a_{n+1} = \sum_{i=1}^{n+1} \frac{x_k}{2^k}$$

where $x_k \in \{0, 2\}$ and (by induction)

$$x \in I_n, \quad \forall n \in \mathbf{N}$$

Pick an arbitrary $\epsilon > 0$. We can choose N such that $l(I_N) = \frac{1}{3^N} < \epsilon$. Consequently, for all $n \geq N$, $|a_n - x| < l(I_n) < \epsilon$. Hence, $(a_n) \rightarrow x$.

Thus, x can be written in the ternary system as an infinite-length (non-terminating) string of 0s and 2s. That is $x = (0.x_1x_2x_3 \dots)_3$.

By Cantor's diagonal argument, the collection of all infinite-length (non-terminating) binary strings consisting of 0s and 2s is uncountable. So, C is uncountable. \square

1.2 Outer Measure.

Definition 1.2. (*Outer measure*) The *outer measure* of any set $A \subseteq \mathbf{R}$ is given by:

$$\mu^*(A) = \inf Z_A$$

where

$$Z_A = \left\{ \sum_{n=1}^{\infty} l(I_n) : I_n \text{ are intervals, } A \subseteq \bigcup_{n=1}^{\infty} I_n \right\}$$

We say that the $(I_n)_{n \geq 1}$ covers the set A . So, the outer measure is the infimum of lengths of all possible covers of A (Note again, that some of the I_n may be empty; this avoids having to worry whether the sequence (I_n) has finitely or infinitely many different members.)

Clearly, $\mu^*(A) \geq 0$ for any $A \subseteq \mathbf{R}$. For some sets A , the series $\sum_{n=1}^{\infty} l(I_n)$ may diverge for any covering of A , so $\mu^*(A)$ may be equal to ∞ . Since we wish to be able to add the outer measures of various sets we have to adopt a convention to deal with infinity. An obvious choice is $a + \infty = \infty$, $\infty + \infty = \infty$ and a less obvious but quite practical assumption is $0 \times \infty = 0$, as we have already seen.

The set Z_A is bounded from below by 0 so that the infimum always exists. If $r \in Z_A$, then $[r, +\infty] \subseteq Z_A$ (clearly we may expand the first interval of any cover to increase the total length by any number). This shows that Z_A is either $+\infty$ or the interval (x, ∞) or $[x, \infty]$ for some real number x . So, the infimum of Z_A is just x .

First, we show that the concept of a null set is consistent with that of Outer measure.

Theorem 1.2. *A set $A \subseteq \mathbf{R}$ is a null set if and only if $\mu^*(A) = 0$.*

Proof. (\Rightarrow direction).

Suppose that A is a null set. We wish to show that $\inf Z_A = 0$. To this end, our claim is, that given any $\epsilon > 0$, there exists $z \in Z_A$ such that $0 < z < \epsilon$.

By definition of a null set, we can find a sequence of intervals $(I_n)_{n \geq 1}$ covering A such that $\sum_{n=1}^{\infty} l(I_n) < \epsilon/2$ and so $\sum_{n=1}^{\infty} l(I_n)$ is an element of Z_A .

(\Leftarrow direction).

Suppose that A is a set such that $\mu^*(A) = 0$. That is, $\inf Z_A = 0$. Pick an arbitrary $\epsilon > 0$. By the definition of \inf , there exists $z \in Z_A$, such that $z < \epsilon$. But, a member of Z_A is the total length of some covering of A . That is, there exists a covering (I_n) of A , with total length smaller than ϵ . Since, $\epsilon > 0$ was arbitrary to begin with, this is true for all $\epsilon > 0$. Hence, A is a null set. \square

This combines our general *outer measure* with the special case of zero measure. Note that, $\mu^*(\emptyset) = 0$ and $\mu^*(\{x\}) = 0$ and $\mu^*(\mathbf{Q}) = 0$.

Next, we observe that μ^* is monotone: the bigger the set, the greater is its outer measure.

Proposition 1.1. *If $A \subset B$, then $\mu^*(A) \leq \mu^*(B)$.*

Proof. Let (I_n) be an arbitrary covering for B . Then, $B \subseteq \bigcup_{n=1}^{\infty} I_n$. Since $A \subset B$, it follows that $A \subset \bigcup_{n=1}^{\infty} I_n$. Thus, $(I_n)_{n \geq 1}$ covers A . So, every cover for B covers A . Consequently, $Z_B \subseteq Z_A$.

Now, $B \setminus A$ is non-empty, so let $x \in B \setminus A$.

Now, let (J_n) be a covering for A , where $J_n = (a_n, b_n)$. Define:

$$J'_n = \begin{cases} (a_n, x) \cup (x, b_n) & \text{if } x \in J_n \\ J_n & \text{otherwise} \end{cases}$$

$(J'_n)_{n \geq 1}$ covers A , but not B . Let $z = \sum_{n=1}^{\infty} l(J'_n)$. Thus, there exists $z \in Z_A$, such that $z \notin Z_B$. So, $Z_B \subset Z_A$.

By the properties of \inf , it follows that $\inf Z_A \leq \inf Z_B$. Thus, $\mu^*(A) \leq \mu^*(B)$. \square

Theorem 1.3. *The outer measure of an interval equals its length.*

If I is an interval, we have:

$$\mu^*(I) = l(I)$$

Proof. If I is unbounded then, it is clear that it cannot be covered by a system of intervals of with finite total length. This shows that $\mu^*(I) = \infty$ and so $\mu^*(I) = l(I) = \infty$.

So we restrict ourselves to bounded intervals.

Step 1. $\mu^*(I) \leq l(I)$.

Take the following sequence of intervals. $I_1 = I$, $I_n = [0, 0]$ for all $n \geq 2$. Then, $\sum_{n=1}^{\infty} l(I_n) = l(I)$. So, $l(I) \in Z_I$. But, $\mu^*(I) = \inf Z_I \leq l(I)$.

Step II. $l(I) \leq \mu^*(I)$.

(i) $I = [a, b]$. We shall show that for any $\epsilon > 0$:

$$l([a, b]) \leq \mu^*([a, b]) + \epsilon \tag{1.1}$$

Pick an arbitrary $\epsilon > 0$. By the definition of outer measure, there exists a sequence of intervals (I_n) such that :

$$\inf Z_I = \mu^*(I) \leq \sum_{n=1}^{\infty} l(I_n) < \mu^*(I) + \frac{\epsilon}{2} \tag{1.2}$$

We shall slightly increase each of the intervals to an open one. Let the endpoints of I_n be a_n, b_n and we take:

$$J_n = \left(a_n - \frac{\epsilon}{2^{n+2}}, b_n + \frac{\epsilon}{2^{n+2}} \right)$$

It is clear that

$$l(I_n) = l(J_n) - \frac{\epsilon}{2^{n+1}}$$

so that:

$$\sum_{n=1}^{\infty} l(I_n) = \sum_{n=1}^{\infty} l(J_n) - \frac{\epsilon}{2}$$

We insert this in 1.2, and we have:

$$\sum_{n=1}^{\infty} l(J_n) \leq \mu^*([a, b]) + \epsilon \quad (1.3)$$

The new sequence of intervals cover $[a, b]$, so by the Heine Borel theorem, we can choose a finite number of J_n to cover $[a, b]$ (the set $[a, b]$ is compact in \mathbf{R}). We can add some intervals to this finite family to form an initial segment of the sequence - just for the simplicity of notation. So, for some finite index m we have:

$$[a, b] \subseteq \bigcup_{n=1}^m J_n$$

Let $J_n = [c_n, d_n]$. Put $c = \min\{c_1, \dots, c_m\}$ and $d = \max\{d_1, \dots, d_m\}$. Then, the above covering means that $c < a$ and $b < d$ and hence $l([a, b]) < d - c$.

Next, the number $d - c$ is certainly smaller than the total length of $J_n, n = 1, 2, 3, \dots, m$ (some overlapping takes place) and

$$l(a, b) < d - c < \sum_{j=1}^m l(J_n) \quad (1.4)$$

Now, it is sufficient to put (1.3) and (1.4) together to deduce (1.1). (The finite sum is less than equal to the sum of the series, since all terms are non-negative) Letting $\epsilon \rightarrow 0$, we have the desired result. $l([a, b]) \leq \mu([a, b])$.

(ii) What if $I = (a, b)$?

Fix an arbitrary $\epsilon > 0$ as before. As before it is sufficient to show (1.1). We have:

$$\begin{aligned} l((a, b)) &= l\left(\left[a + \frac{\epsilon}{2}, b - \frac{\epsilon}{2}\right]\right) + \epsilon \\ &= \mu^*\left(\left[a + \frac{\epsilon}{2}, b - \frac{\epsilon}{2}\right]\right) + \epsilon \\ &\quad \{ \text{From part I} \} \\ &\leq \mu^*((a, b)) + \epsilon \\ &\quad \{ \text{By monotonicity of outer measure (1.1)} \} \end{aligned}$$

(iii) $I = (a, b]$ or $I = [a, b)$.

$$\begin{aligned} l(I) &= l((a, b)) \leq \mu^*((a, b)) && \{ \text{From part II} \} \\ &\leq \mu^*(I) && \{ \text{Monotonicity of Lebesgue Measure} \} \end{aligned}$$

This closes the proof. □

Theorem 1.4. (Countable Subadditivity). *The outer measure is countably subadditive.*

For all sequences of sets (E_n) , we have:

$$\mu^*\left(\bigcup_{n=1}^{\infty} E_n\right) \leq \sum_{n=1}^{\infty} \mu^*(E_n)$$

(Note that both sides might be infinite here.)

Proof. (A warm-up)

Let's first prove a simpler statement:

$$\mu^*(E_1 \cup E_2) \leq \mu^*(E_1) + \mu^*(E_2)$$

Take an $\epsilon > 0$ and we show an even easier inequality:

$$\mu^*(E_1 \cup E_2) \leq \mu^*(E_1) + \mu^*(E_2) + \epsilon$$

By the definition of outer measure,

There exists a sequence of intervals (I_n^1) covering E_1 such that:

$$\mu^*(E_1) < \sum_{n=1}^{\infty} l(I_n^1) < \mu^*(E_1) + \frac{\epsilon}{2}$$

There exists a sequence of intervals (I_n^2) covering E_2 such that:

$$\mu^*(E_2) < \sum_{n=1}^{\infty} l(I_n^2) < \mu^*(E_2) + \frac{\epsilon}{2}$$

Now, the sequence of intervals $I_1^1, I_1^2, I_2^1, I_2^2, \dots$ covers $E_1 \cup E_2$. Hence,

$$\begin{aligned} \mu^*(E_1 \cup E_2) &\leq \sum_{n=1}^{\infty} (l(I_n^1) + l(I_n^2)) \\ &\leq \mu^*(E_1) + \frac{\epsilon}{2} + \mu^*(E_2) + \frac{\epsilon}{2} \\ &= \mu^*(E_1) + \mu^*(E_2) + \epsilon \end{aligned}$$

Since ϵ was arbitrary, this is true for all $\epsilon > 0$.

Choosing $\epsilon = \frac{1}{n}$, passing to the limit as $n \rightarrow \infty$, we have:

$$\mu^*(E_1 \cup E_2) \leq \mu^*(E_1) + \mu^*(E_2)$$

(*Proof of the theorem.*)

If the right-hand side is infinite, then inequality is of course true. So, suppose that $\sum_{k=1}^{\infty} \mu^*(E_k) < \infty$. For each given $\epsilon > 0$ and $k \geq 1$, find a covering sequence (I_n^k) of E_k with :

$$\sum_{n=1}^{\infty} l(I_n^k) < \mu^*(E_k) + \frac{\epsilon}{2^k}$$

The iterated series

$$\sum_{k=1}^{\infty} \left(\sum_{n=1}^{\infty} l(I_n^k) \right) < \sum_{k=1}^{\infty} \mu^*(E_k) + \epsilon < \infty$$

Now, $I_1^1, I_2^1, I_1^2, I_3^2, I_2^3, I_1^3, \dots$ is a countable sequence (since $\mathbf{N} \times \mathbf{N}$ is countable) that covers $\bigcup_{k=1}^{\infty} E_k$. So,

$$\mu^* \left(\bigcup_{k=1}^{\infty} E_k \right) \leq \sum_{k=1}^{\infty} \left(\sum_{n=1}^{\infty} l(I_n^k) \right) < \sum_{k=1}^{\infty} \mu^*(E_k) + \epsilon$$

To complete the proof, we simply let $\epsilon \rightarrow 0$. □

Problem 1.3. Prove that if $\mu^*(A) = 0$ then for each B , $\mu^*(A \cup B) = \mu^*(B)$.

Proof. Let B be an arbitrary set. By countable additivity of outer-measure, we have:

$$\begin{aligned} \mu^*(A \cup B) &\leq \mu^*(A) + \mu^*(B) \\ &= \mu^*(B) \end{aligned}$$

Since $B \subseteq A \cup B$, by the monotonicity of outer-measure,

$$\mu^*(B) \leq \mu^*(A \cup B)$$

From the above discussion, it follows that, $\mu^*(A \cup B) = \mu^*(B)$. Since, B was arbitrary, this must be true for all sets B . This closes the proof. \square

Problem 1.4. Prove that if $\mu^*(A \triangle B) = 0$, then $\mu^*(A) = \mu^*(B)$.

Proof. We know that, $A \subseteq B \cup (A \triangle B)$. Hence:

$$\begin{aligned} \mu^*(A) &\leq \mu^*(B \cup (A \triangle B)) \\ &\quad \{ \text{Monotonicity of Outer Measure} \} \\ &\leq \mu^*(B) + \mu^*(A \triangle B) \\ &\quad \{ \text{Countable Subadditivity} \} \\ &= \mu^*(B) \end{aligned}$$

On the other hand, $B \subseteq A \cup (A \triangle B)$. Hence, $\mu^*(B) \leq \mu^*(A)$. Consequently, it follows that

$$\mu^*(A) = \mu^*(B)$$

\square

Proposition 1.2. *The outer measure is translation invariant.*

$$\mu^*(A) = \mu^*(A + t)$$

for each A and t .

Proof. Let $A \subset \mathbf{R}$ and t be a fixed real.

Let (I_n) be any sequence of intervals covering A . Then, $I'_n = [a_n + t, b_n + t]$ is a covering for $A + t$. Now, $l(I'_n) = l(I_n)$ for all $n \in \mathbf{N}$. So, $\sum_{n=1}^{\infty} l(I'_n) = \sum_{n=1}^{\infty} l(I_n)$. Hence, if $z \in Z_A$, it follows that $z \in Z_{A+t}$ and vice-versa. Consequently, $Z_A = Z_{A+t}$. So, $\inf Z_A = \inf Z_{A+t}$. Therefore, $\mu^*(A) = \mu^*(A + t)$. \square

1.3 Lebesgue measurable sets and Lebesgue measure.

With the outer measure, subadditivity as in Theorem (1.4) is as far as we can get. We wish to however, ensure, that, if the sets (E_n) are pairwise disjoint (that is $E_i \cap E_j = \emptyset$, $i \neq j$) then the inequality in Theorem (1.4) becomes an equality. It turns out that this will not in general be true for the outer-measure. But our wish is entirely a reasonable one: any length function should atleast be finitely additive, since decomposing a set into finitely many disjoint pieces, should not alter it's length. Moreover, since we constructed our length function via the approximation of complicated sets by simpler sets (that is intervals), it seems fair to demand a *continuity property* : if pairwise disjoint E_n have union E , then the lengths of sets $B_n = E \setminus \bigcup_{k=1}^n E_k$ may be expected to decrease to 0 as $n \rightarrow \infty$. Combining this with finite additivity leads quite naturally to demand that length be countably additive, that is:

$$\mu^* \left(\bigcup_{n=1}^{\infty} E_n \right) = \sum_{n=1}^{\infty} \mu^*(E_n) \quad \text{when } E_i \cap E_j = \emptyset \text{ for } i \neq j \quad (1.5)$$

We therefore turn to the task of finding the class of sets in \mathbf{R} which have this property.

Definition 1.3. A set E is Lebesgue *measurable* if for every set $A \subseteq \mathbf{R}$ we have:

$$\mu^*(A) = \mu^*(A \cap E) + \mu^*(A \cap E^C) \quad (1.6)$$

We write $E \in \mathcal{F}$.

We obviously have $A = (A \cap E) \cup (A \cap E^C)$, hence by countable subadditivity (1.4), we have:

$$\mu^*(A) \leq \mu^*(A \cap E) + \mu^*(A \cap E^C)$$

for any A and E . So, our future task of verifying countable additivity property (1.5) has simplified: $E \in \mathcal{F}$ if and only the following inequality holds:

$$\mu^*(A) \geq \mu^*(A \cap E) + \mu^*(A \cap E^C)$$

for all $A \subseteq \mathbf{R}$.

Now, we give examples of measurable sets.

Theorem 1.5. (i) Any null set is measurable.

(ii) Any interval is measurable.

Proof. (i) If N is a null set, then by Theorem (1.2), the null set has outer measure zero, so $\mu^*(N) = 0$. For all $A \subseteq \mathbf{R}$, since $A \cap N \subseteq N$ and $A \cap N^C \subseteq A$. Thus,

$$\begin{aligned}\mu^*(A \cap N) + \mu^*(A \cap N^C) &\leq \mu^*(N) + \mu^*(A) \\ \mu^*(A \cap N) + \mu^*(A \cap N^C) &\leq \mu^*(A)\end{aligned}$$

(ii) Let $E = I$ be an interval. Suppose, for example, $I = [a, b]$. Take any $A \subseteq \mathbf{R}$ and $\epsilon > 0$. Find a covering of A with:

$$\mu^*(A) \leq \sum_{n=1}^{\infty} l(I_n) \leq \mu^*(A) + \epsilon$$

Clearly, the intervals $I'_n = I_n \cap [a, b]$ cover $A \cap [a, b]$ and hence $\sum l(I'_n) \in Z_{A \cap [a, b]}$, that is,

$$\mu^*(A \cap [a, b]) \leq \sum_{n=1}^{\infty} l(I'_n)$$

The intervals $I''_n = I_n \cap (-\infty, a)$ and $I'''_n = I_n \cap (b, +\infty)$ cover $A \cap [a, b]^c$, so:

$$\mu^*(A \cap [a, b]^c) \leq \sum_{n=1}^{\infty} l(I''_n) + l(I'''_n)$$

Since, the intervals $I'_n \cup I''_n \cup I'''_n$ cover A , it follows that:

$$\begin{aligned}\mu^*(A \cap [a, b]) + \mu^*(A \cap [a, b]^c) &\leq \sum_{n=1}^{\infty} l(I'_n) + l(I''_n) + l(I'''_n) \\ &= \sum_{n=1}^{\infty} l(I_n) \\ &\leq \mu^*(A) + \epsilon\end{aligned}$$

Letting $\epsilon \rightarrow 0$, we have the desired result. \square

The fundamental properties of the class \mathcal{F} of all Lebesgue measurable subsets of \mathbf{R} can now be proved. They fall into two categories: first we show that certain set operations on \mathcal{F} produce sets in \mathcal{F} (these are what we call closure properties) and second we prove that for sets in \mathcal{F} the outer measure μ^* has the property of countable additivity announced above.

Theorem 1.6. (*Closure properties of \mathcal{F}*)

(i) $\mathbf{R} \in \mathcal{F}$.

(ii) If $E \in \mathcal{F}$, then $E^C \in \mathcal{F}$.

(iii) If $E_n \in \mathcal{F}$, for all $n = 1, 2, 3, \dots$ then $\bigcup_{n=1}^{\infty} E_n \in \mathcal{F}$.

Moreover, if $E_n \in \mathcal{F}$, for all $n = 1, 2, 3, \dots$ and $E_i \cap E_j = \emptyset$ for $i \neq j$, then:

$$\mu^*\left(\bigcup_{n=1}^{\infty} E_n\right) = \sum_{n=1}^{\infty} \mu^*(E_n) \quad (1.7)$$

Remark. This result is the most important theorem in this chapter and provides the basis for all that follows. It also allows us to give names to the quantities under discussion.

Conditions (i)-(iii) mean that \mathcal{F} is a sigma-algebra. In other words, we say that a family of sets is a sigma-algebra, if it contains the base set and is closed under countable unions, and complements. A $[0, \infty)$ -valued function defined on a sigma-algebra is called a measure if it satisfies countable additivity (1.7) for pairwise disjoint sets.

An alternative, rather more abstract and general approach to measure theory is to begin with the above properties as axioms, i.e. to call the triple $(\Omega, \mathcal{F}, \mu)$ a measure space, if Ω is an abstractly given set, \mathcal{F} is a sigma-algebra of the subsets of Ω and $\mu : \mathcal{F} \rightarrow [0, \infty]$ is a function satisfying countable additivity. The task of defining the Lebesgue measure on \mathbf{R} then becomes that of verifying, with \mathcal{F} and $\mu = \mu^*$ on \mathcal{F} defined above, that the triple $(\Omega, \mathcal{F}, \mu)$ satisfies these axioms.

Although the requirements of probability theory will mean that we have to consider such general measure spaces in due course, we have chosen our more concrete approach to the fundamental example of Lebesgue measure in order to demonstrate how this important measure space arises quite naturally from the considerations of the *lengths* of sets in \mathbf{R} and leads to a theory of integration which greatly extends that of

Riemann. It is also sufficient to allow us to develop most of the important examples of probability distributions.

Proof. (1) Let $A \subseteq \mathbf{R}$. Note that $A \cap \mathbf{R} = A$, $\mathbf{R}^C = \emptyset$, so that $A \cap \mathbf{R}^C = \emptyset$. Thus, the equation (1.6) now reads, $\mu^*(A) = \mu^*(A) + \mu^*(\emptyset)$ which is obviously true, since \emptyset is a null set and $\mu^*(\emptyset) = 0$.

(2) Suppose $E \in \mathcal{F}$ and take any arbitrary $A \subseteq \mathbf{R}$. We have to show (1.6) for E^C , that is:

$$\mu^*(A) = \mu^*(A \cap E^C) + \mu^*(A \cap (E^C)^C) \quad (1.8)$$

but since $(E^C)^C = E$, this reduces to the condition for E which holds by hypothesis.

(3) We split the proof (iii) into several steps. But first:

A warm up. Suppose that $E_1 \cap E_2 = \emptyset$, $E_1, E_2 \in \mathcal{F}$. We shall show that $E_1 \cup E_2 \in \mathcal{F}$ and $\mu^*(E_1 \cup E_2) = \mu^*(E_1) + \mu^*(E_2)$.

Let $A \subseteq \mathbf{R}$. We have the condition for E_1 :

$$\mu^*(A) = \mu^*(A \cap E_1) + \mu^*(A \cap E_1^C) \quad (1.9)$$

Now, we apply (1.6) for E_2 with $A \cap E_1^C$ in place of A .

$$\begin{aligned} \mu^*(A \cap E_1^C) &= \mu^*(A \cap E_1^C \cap E_2) + \mu^*(A \cap E_1^C \cap E_2^C) \\ &= \mu^*(A \cap (E_1^C \cap E_2)) + \mu^*(A \cap (E_1^C \cap E_2^C)) \end{aligned}$$

The situation is depicted in the figure below.




Figure. The sets A , E_1 and E_2 .

Since E_1 and E_2 are disjoint, $E_1^C \cap E_2 = E_2$. By De-Morgan's laws, $E_1^C \cap E_2^C = (E_1 \cup E_2)^C$. We substitute and we have:

$$\mu^*(A \cap E_1^C) = \mu^*(A \cap E_2) + \mu^*(A \cap (E_1 \cup E_2)^C)$$

Substituting this into (1.9), we get:

$$\mu^*(A) = \mu^*(A \cap E_1) + \mu^*(A \cap E_2) + \mu^*(A \cap (E_1 \cup E_2)^C) \quad (1.10)$$

Now, by the subadditivity property of μ^* , we have:

$$\begin{aligned} \mu^*(A \cap E_1) + \mu^*(A \cap E_2) &\geq \mu^*((A \cap E_1) \cup (A \cap E_2)) \\ &= \mu^*(A \cap (E_1 \cup E_2)) \end{aligned}$$

So, (1.10) gives:

$$\mu^*(A) \geq \mu^*(A \cap (E_1 \cup E_2)) + \mu^*(A \cap (E_1 \cup E_2)^C)$$

which is sufficient for $E_1 \cup E_2$ to belong to \mathcal{F} .

Finally, let $A = E_1 \cup E_2$. Then, the equation (1.10) yields:

$$\mu^*(E_1 \cup E_2) = \mu^*(E_1) + \mu^*(E_2)$$

We return to the main proof of (iii).

Step 1. Our claim is: if pairwise disjoint $E_k, k = 1, 2, \dots$ are in \mathcal{F} then their countable union is in \mathcal{F} and countable additivity (1.5) holds.

We begin as in the proof of the **Warm Up** and we have:

$$\begin{aligned} \mu^*(A) &= \mu^*(A \cap E_1) + \mu^*(A \cap E_1^C) \\ \mu^*(A) &= \mu^*(A \cap E_1) + \mu^*(A \cap E_2) + \mu^*(A \cap (E_1 \cup E_2)^C) \end{aligned} \quad (1.11)$$

(See (1.10)).

E_3 is also measurable. Let $A = A \cap E_1^C \cap E_2^C$. Then:

$$\mu^*(A \cap E_1^C \cap E_2^C) = \mu^*(A \cap E_1^C \cap E_2^C \cap E_3) + \mu^*(A \cap E_1^C \cap E_2^C \cap E_3^C)$$

But, $E_1^C \cap E_2^C \cap E_3 = E_3$ since they are pairwise disjoint. So,

$$\mu^*(A \cap (E_1 \cup E_2)^C) = \mu^*(A \cap E_3) + \mu^*(A \cap (E_1 \cup E_2 \cup E_3)^C) \quad (1.12)$$

Substituting the value of (1.12) in equation (1.11), we get after $n = 3$ steps:

$$\mu^*(A) = \sum_{k=1}^3 \mu^*(A \cap E_k) + \mu^* \left(A \cap \left(\bigcup_{k=1}^3 E_k \right)^C \right) \quad (1.13)$$

We proceed by mathematical induction. We induct on k . Our hypothesis is, that after n steps, we expect:

$$\mu^*(A) = \sum_{k=1}^n \mu^*(A \cap E_k) + \mu^* \left(A \cap \left(\bigcup_{k=1}^n E_k \right)^C \right) \quad (1.14)$$

Let's assume that

$$\mu^*(A) = \sum_{k=1}^{n-1} \mu^*(A \cap E_k) + \mu^* \left(A \cap \left(\bigcup_{k=1}^{n-1} E_k \right)^C \right) \quad (1.15)$$

is true.

Since $E_n \in \mathcal{F}$, we may apply the definition (1.6) with $A = A \cap \left(\bigcup_{k=1}^{n-1} E_k \right)^C$:

$$\mu^* \left(A \cap \left(\bigcup_{k=1}^{n-1} E_k \right)^C \right) = \mu^* \left(A \cap \left(\bigcup_{k=1}^{n-1} E_k \right)^C \cap E_n \right) + \mu^* \left(A \cap \left(\bigcup_{k=1}^{n-1} E_k \right)^C \cap E_n^C \right) \quad (1.16)$$

Now we make the same observations as in the **Warm Up**:

$$\begin{aligned} \left(\bigcup_{k=1}^{n-1} E_k \right)^C \cap E_n &= E_n \quad \{E_i \text{ are pairwise disjoint}\} \\ \left(\bigcup_{k=1}^{n-1} E_k \right)^C \cap E_n^C &= \left(\bigcup_{k=1}^n E_k \right)^C \quad \{\text{De-Morgan's laws}\} \end{aligned}$$

Inserting these into equation (1.16), we get:

$$\mu^* \left(A \cap \left(\bigcup_{k=1}^{n-1} E_k \right)^C \right) = \mu^*(A \cap E_n) + \mu^* \left(A \cap \left(\bigcup_{k=1}^n E_k \right)^C \right)$$

and inserting this into (1.15), we get :

$$\mu^*(A) = \sum_{k=1}^{n-1} \mu^*(A \cap E_k) + \mu^*(A \cap E_n) + \mu^* \left(A \cap \left(\bigcup_{k=1}^n E_k \right)^C \right)$$

This proves the induction hypothesis.

As will be seen at the next step, the fact that E_k are pairwise disjoint is not necessary in order to ensure that their union belongs to \mathcal{F} . However, with this assumption we have equality in (1.14) which does not hold otherwise. This equality will allow us to prove countable additivity (1.7).

Since:

$$\left(\bigcup_{k=1}^n E_k \right)^C \supseteq \left(\bigcup_{k=1}^{\infty} E_k \right)^C$$

from (1.14) by monotonicity of measure, we get:

$$\begin{aligned}
\mu^*(A) &= \sum_{k=1}^n \mu^*(A \cap E_k) + \mu^* \left(A \cap \left(\bigcup_{k=1}^n E_k \right)^C \right) \\
&\geq \sum_{k=1}^n \mu^*(A \cap E_k) + \mu^* \left(A \cap \left(\bigcup_{k=1}^{\infty} E_k \right)^C \right)
\end{aligned}$$

By the Order Limit Theorem, the inequality remains true, if we pass to the limit, as $n \rightarrow \infty$:

$$\mu^*(A) \geq \sum_{k=1}^{\infty} \mu^*(A \cap E_k) + \mu^* \left(A \cap \left(\bigcup_{k=1}^{\infty} E_k \right)^C \right) \quad (1.17)$$

By countable sub-additivity of μ^* (Theorem (1.4)) :

$$\sum_{k=1}^{\infty} \mu^*(A \cap E_k) \geq \mu^* \left(A \cap \left(\bigcup_{k=1}^{\infty} E_k \right) \right)$$

and so:

$$\mu^*(A) \geq \mu^* \left(A \cap \left(\bigcup_{k=1}^{\infty} E_k \right) \right) + \mu^* \left(A \cap \left(\bigcup_{k=1}^{\infty} E_k \right)^C \right) \quad (1.18)$$

So, we have shown that $\bigcup_{k=1}^{\infty} E_k \in \mathcal{F}$ and hence the two sides of (1.18) are equal.

The right hand side of (1.17) is squeezed between the left and right of (1.18). That is:

$$\begin{aligned}
\mu^*(A) &= \mu^* \left(A \cap \left(\bigcup_{k=1}^{\infty} E_k \right) \right) + \mu^* \left(A \cap \left(\bigcup_{k=1}^{\infty} E_k \right)^C \right) \\
&\leq \sum_{k=1}^{\infty} \mu^*(A \cap E_k) + \mu^* \left(A \cap \left(\bigcup_{k=1}^{\infty} E_k \right)^C \right) \\
&\leq \mu^*(A)
\end{aligned}$$

Consequently,

$$\mu^*(A) = \sum_{k=1}^{\infty} \mu^*(A \cap E_k) + \mu^* \left(A \cap \left(\bigcup_{k=1}^{\infty} E_k \right)^C \right) \quad (1.19)$$

The equality here is a consequence of the assumption that E_k are pairwise disjoint. It holds for any set A so we may insert $A = \bigcup_{j=1}^{\infty} E_j$. The last term on the right is zero, because the length of the empty set, $\mu^*(\emptyset) = 0$. And, since the E_i are disjoint, $\left(\bigcup_{j=1}^{\infty} E_j \right) \cap E_k = E_k$. As a result, we have:

$$\mu^* \left(\bigcup_{j=1}^{\infty} E_j \right) = \sum_{j=1}^{\infty} \mu^*(E_j)$$

Step 2. Our claim is, if $E_1, E_2 \in \mathcal{F}$, then $E_1 \cup E_2 \in \mathcal{F}$ (not necessarily disjoint).

Again we begin as in the **Warm Up**:

$$\mu^*(A) = \mu^*(A \cap E_1) + \mu^*(A \cap E_1^C) \quad (1.20)$$

Next applying the definition (1.6) to E_2 and with $A \cap E_1^C$ in place of A we get:

$$\mu^*(A \cap E_1^C) = \mu^*(A \cap E_1^C \cap E_2) + \mu^*(A \cap E_1^C \cap E_2^C)$$

We insert this into (1.20) to get:

$$\mu^*(A) = \mu^*(A \cap E_1) + \mu^*(A \cap E_1^C \cap E_2) + \mu^*(A \cap E_1^C \cap E_2^C) \quad (1.21)$$

By DeMorgan's law, $E_1^C \cap E_2^C = (E_1 \cup E_2)^C$ so as before:

$$\mu^*(A \cap E_1^C \cap E_2^C) = \mu^*(A \cap (E_1 \cup E_2)^C)$$

Now, consider the set $(A \cap E_1) \cup (A \cap E_1^C \cap E_2)$. This, can be written as: $A \cap (E_1 \cup (E_1^C \cap E_2)) = A \cap (E_1 \cup E_1^C) \cap (E_1 \cup E_2) = A \cap (E_1 \cup E_2)$.

By Countable Subadditivity of μ^* , we have:

$$\mu^*(A \cap E_1) + \mu^*(A \cap E_1^C \cap E_2) \geq \mu^*(A \cap (E_1 \cup E_2))$$

Inserting these facts into (1.21), we get:

$$\mu^*(A) \geq \mu^*(A \cap (E_1 \cup E_2)) + \mu^*(A \cap (E_1 \cup E_2)^C)$$

as required.

Step 3. Our claim is, if $E_k \in \mathcal{F}$, $k = 1, 2, \dots, n$, then the finite union $E_1 \cup E_2 \cup \dots \cup E_n \in \mathcal{F}$. (not necessarily disjoint)

We argue by induction. Suppose that the claim is true for $n - 1$. Then,

$$E_1 \cup E_2 \cup \dots \cup E_n = (E_1 \cup \dots \cup E_{n-1}) \cup E_n$$

so that the result follows from **Step 2**.

Step 4. If $E_1, E_2 \in \mathcal{F}$, then $E_1 \cap E_2 \in \mathcal{F}$.

We have $E_1^C, E_2^C \in \mathcal{F}$ by (ii), $E_1^C \cup E_2^C \in \mathcal{F}$ by step 2, and $(E_1^C \cup E_2^C)^C \in \mathcal{F}$ by (ii) again. But, by De-Morgan's laws, this is $(E_1^C \cup E_2^C)^C = E_1 \cap E_2$.

Step 5. The general case: if E_1, E_2, \dots are in \mathcal{F} , then so is the countably infinite union $\bigcup_{k=1}^{\infty} E_k$.

Let $E_k \in \mathcal{F}$, $k = 1, 2, \dots$. We define the auxiliary sequence of pairwise disjoint sets F_k with the same union as E_k :

$$\begin{aligned} F_1 &= E_1 \\ F_2 &= E_2 \setminus E_1 = E_2 \cap E_1^C \\ F_3 &= E_3 \setminus (E_1 \cup E_2) = E_3 \cap (E_1 \cup E_2)^C \\ &\vdots \\ F_k &= E_k \setminus (E_1 \cup E_2 \cup \dots \cup E_{k-1}) = E_k \cap (E_1 \cup \dots \cup E_{k-1})^C \end{aligned}$$

By steps 3 and 4, we know that all F_k are in \mathcal{F} . By the very construction, they are pairwise disjoint, so by step 1, their union is in \mathcal{F} . We shall show that:

$$\bigcup_{k=1}^{\infty} F_k = \bigcup_{k=1}^{\infty} E_k$$

The inclusion:

$$\bigcup_{k=1}^{\infty} F_k \subseteq \bigcup_{k=1}^{\infty} E_k$$

is obvious since for each k , $F_k \subseteq E_k$ by definition. For the inverse, let $a \in \bigcup_{k=1}^{\infty} E_k$. Put $S = \{n \in \mathbf{N} : a \in E_n\}$ which is non-empty since a belongs to the union. Let $n_0 = \min S \in S$. If $n_0 = 1$, then $a \in E_1 = F_1$. Suppose $n_0 > 1$. So, $a \in E_{n_0}$ and by definition of n_0 , $a \notin E_1, \dots, a \notin E_{n_0-1}$. By the definition of F_{n_0} , this means that $a \in F_{n_0}$ so a is in $\bigcup_{k=1}^{\infty} F_k$. This closes the proof. \square

Using De-Morgan's laws, we can easily verify an additional property of \mathcal{F} .

Proposition 1.3. *If $E_k \in \mathcal{F}$, $k = 1, 2, \dots$, then*

$$E = \bigcap_{k=1}^{\infty} E_k \in \mathcal{F}$$

Proof. \mathcal{F} is closed under complementation. Thus, $E_k \in \mathcal{F} \implies E_k^C \in \mathcal{F}$. Since, \mathcal{F} is closed under countable unions, $\bigcup_{k=1}^{\infty} E_k^C \in \mathcal{F}$. And it follows that, $(\bigcup_{k=1}^{\infty} E_k^C)^C \in \mathcal{F}$. By De-Morgan's laws, $(\bigcup_{k=1}^{\infty} E_k^C)^C = \bigcap_{k=1}^{\infty} E_k$. This closes the proof. \square

We can therefore summarize the properties of the family \mathcal{F} of Lebesgue measurable sets as follows:

\mathcal{F} is closed under countable unions, countable intersections and complements. It contains intervals and null sets.

Definition 1.4. (*Lebesgue Measure*). We shall write $\mu(E)$ instead of $\mu^*(E)$ for any E in \mathcal{F} and call $\mu(E)$ the Lebesgue measure of the set E .

The Lebesgue measure $\mu : \mathcal{F} \rightarrow [0, \infty]$ is a countably additive set function defined on the sigma-algebra \mathcal{F} of measurable sets. The Lebesgue measure of an interval is equal to its length. The Lebesgue measure of a null-set is zero.

1.4 Basic Properties of Lebesgue Measure.

Since Lebesgue measure is nothing else than the outer measure restricted to a special class of sets \mathcal{F} , some properties of the outer measure are automatically inherited by the Lebesgue measure.

Proposition 1.4. *Suppose that $A, B \in \mathcal{F}$.*

- (1) *If $A \subset B$, then $\mu(A) \leq \mu(B)$.*
- (2) *If $A \subset B$ and $\mu(A)$ is finite, then $\mu(B \setminus A) = \mu(B) - \mu(A)$.*
- (3) *μ is translation invariant.*

Since the empty set $\emptyset \in \mathcal{F}$, we can take $E_i = \emptyset$ for all $i > n$ in (1.7) to conclude that Lebesgue measure is finitely additive: if $E_i \in \mathcal{F}$ are pairwise disjoint, then:

$$\mu \left(\bigcup_{i=1}^n E_i \right) = \sum_{i=1}^n \mu(E_i)$$

Remark. Property (2) is derived as follows. Since $B = (B \setminus A) \cup A$ and $B \setminus A$ and A are disjoint, $\mu(B) = \mu(B \setminus A) + \mu(A)$. Consequently, $\mu(B \setminus A) = \mu(B) - \mu(A)$.

Problem 1.5. Find a formula describing $\mu(A \cup B)$ and $\mu(A \cup B \cup C)$ in terms of measures of the individual sets and their intersections (we do not assume that the sets are pairwise disjoint).

Proof. We have:

$$A \cup B = (A \cap (A \cap B)^C) \cup (B \cap (A \cap B)^C) \cup (A \cap B)$$

The three sets $A \setminus (A \cap B)$, $B \setminus (A \cap B)$ and $A \cap B$ are pairwise disjoint. Consequently, by finite additivity of the Lebesgue measure:

$$\begin{aligned}\mu(A \cup B) &= \mu(A \setminus (A \cap B)) + \mu(B \setminus (A \cap B)) + \mu(A \cap B) \\ &= \mu(A) - \mu(A \cap B) + \mu(B) - \mu(A \cap B) + \mu(A \cap B) \\ &\quad \{\because (A \cap B) \subseteq A, \mu(A \setminus (A \cap B)) = \mu(A) - \mu(A \cap B)\} \\ &= \mu(A) + \mu(B) - \mu(A \cap B)\end{aligned}$$

Let $B = B \cup C$

$$\begin{aligned}\mu(A \cup (B \cup C)) &= \mu(A) + \mu(B \cup C) - \mu(A \cap (B \cup C)) \\ &= \mu(A) + \mu(B) + \mu(C) - \mu(B \cap C) - \mu((A \cap B) \cup (A \cap C)) \\ &= \mu(A) + \mu(B) + \mu(C) - \mu(B \cap C) - \mu(A \cap B) - \mu(A \cap C) \\ &\quad + \mu(A \cap B \cap C)\end{aligned}$$

□

Recalling that the *symmetric difference* $A \Delta B$ of two sets is defined by $A \Delta B = (A \setminus B) \cup (B \setminus A)$ the following result is also easy to check:

Proposition 1.5. *If $A \in \mathcal{F}$, and $\mu(A \Delta B) = 0$, then $B \in \mathcal{F}$ and $\mu(A) = \mu(B)$.*

Proof. Null sets belong to \mathcal{F} . Since $A \Delta B$ is a null set, it belongs to \mathcal{F} . Now, $A \cap B^C$ and $A^C \cap B$ are subsets of $A \Delta B$, they are also null sets and belong to \mathcal{F} . We have:

$$\begin{aligned}\mu(A) &= \mu(A \cap (B \cup B^C)) \\ &= \mu(A \cap B) + \mu(A \cap B^C) \\ &= \mu(A \cap B)\end{aligned}$$

And likewise, $\mu(B) = \mu(A \cap B)$. Hence, $\mu(A) = \mu(B)$. □

Lemma 1.1. *Let $(A_n)_{n=1}^\infty, (B_n)_{n=1}^\infty$ be a sequence of sets. The difference of the union of sets is contained in the union of the difference of sets. We have:*

$$\left(\bigcup_{n \geq 1} A_n \right) - \left(\bigcup_{n \geq 1} B_n \right) \subset \bigcup_{n \geq 1} (A_n - B_n)$$

Proof. We have:

$$\begin{aligned}
A - \left(\bigcup_{n \geq 1} B_n \right) &= A \cap \left(\bigcup_{n \geq 1} B_n \right)^C \\
&= A \cap \left(\bigcap_{n \geq 1} B_n^C \right) \\
&= \bigcap_{n \geq 1} (A \cap B_n^C) \\
&= \bigcap_{n \geq 1} (A - B_n)
\end{aligned}$$

Thus,

$$\begin{aligned}
\left(\bigcup_{n \geq 1} A_n \right) - \left(\bigcup_{n \geq 1} B_n \right) &= \left(\bigcup_{n \geq 1} A_n \right) \cap \left(\bigcup_{n \geq 1} B_n \right)^C \\
&= \left(\bigcup_{m \geq 1} A_m \right) \cap \left(\bigcap_{n \geq 1} B_n^C \right) \\
&= \bigcup_{m \geq 1} \left(A_m \cap \left(\bigcap_{n \geq 1} B_n^C \right) \right) \\
&= \bigcup_{m \geq 1} \bigcap_{n \geq 1} (A_m \cap B_n^C) \\
&= \bigcup_{m \geq 1} \left\{ \bigcap_{n \geq 1} (A_m - B_n) \right\} \\
&\subset \bigcup_{m \geq 1} \{A_m - B_m\}
\end{aligned}$$

□

Every open set in \mathbf{R} can be expressed as the union of a countable number of open intervals. This ensures that open sets in \mathbf{R} are Lebesgue measurable, since \mathcal{F} contains

intervals and is closed under countable unions. We can approximate the measure of any $A \in \mathcal{F}$ from the above by the measures of a sequence of open sets containing A . This is clear from the below result:

Theorem 1.7. (i) For any $\epsilon > 0$, $A \in \mathbf{R}$, we can find an open set O such that :

$$A \subset O, \quad \mu(O) \leq \mu^*(A) + \epsilon$$

Consequently, for any $E \in \mathcal{F}$ we can find an open set O containing E such that $\mu(O \setminus E) < \epsilon$.

(ii) For any $A \subset \mathbf{R}$, we can find a sequence of open sets O_n , such that:

$$A \subset \bigcap_{n=1}^{\infty} O_n, \quad \mu \left(\bigcap_{n=1}^{\infty} O_n \right) = \mu^*(A)$$

Proof. (i) By definition of $\mu^*(A)$ we can find a sequence (I_n) of intervals with $A \subset \bigcup_{n=1}^{\infty} I_n$ and $\sum_{n=1}^{\infty} l(I_n) \leq \mu^*(A) + \epsilon/2$. That is,

$$\exists (I_n)_{n=1}^{\infty}, \quad A \subset \bigcup_n I_n, \quad \sum_{n=1}^{\infty} l(I_n) - \frac{\epsilon}{2} \leq \mu^*(A)$$

Each I_n is contained in an open interval whose length is very close to that of I_n ; if the left and right end-points of I_n are a_n and b_n respectively, let $J_n = (a_n - \frac{\epsilon}{2^{n+2}}, b_n + \frac{\epsilon}{2^{n+2}})$. Set $O = \bigcup_{n=1}^{\infty} J_n$, which is open. Remember, that J_n 's are overlapping. Then, $A \subset O$ and

$$\mu(O) \leq \sum_{n=1}^{\infty} l(J_n) = \sum_{n=1}^{\infty} l(I_n) + \frac{\epsilon}{2} \leq \mu^*(A) + \epsilon$$

When $\mu(E) < \infty$ the final statement follows at once from (ii) in proposition (1.4), since $\mu(O \setminus E) = \mu(O) - \mu(E) \leq \epsilon$.

When $\mu(E) = \infty$ we first write \mathbf{R} as the countable union of the finite intervals: $\mathbf{R} = \bigcup_n (-n, n)$. Now, $E_n = E \cap (-n, n)$ has finite measure, so we can find an

open set $O_n \supset E_n$ with $\mu(O_n \setminus E_n) \leq \frac{\epsilon}{2^n}$. The set $O = \bigcup_n O_n$ is open and contains E . Now,

$$\begin{aligned} O \setminus E &= \left(\bigcup_n O_n \right) \setminus \left(\bigcup_n E_n \right) \\ &\subset \bigcup_n (O_n \setminus E_n) \end{aligned}$$

so that $\mu(O \setminus E) \leq \sum_n \mu(O_n \setminus E_n) \leq \epsilon$.

(ii) In (i) use $\epsilon = \frac{1}{n}$ and let O_n be the open set so obtained. With $E = \bigcap_n O_n$ we obtain a measurable set containing A such that $\mu(E) < \mu(O_n) \leq \mu^*(A) + \frac{1}{n}$ for each n , hence the result follows. \square

Remark. Theorem (1.7) shows how the freedom of movement allowed by the closure properties of the sigma-field \mathcal{F} can be exploited by producing, for any set $A \subset \mathbf{R}$, a measurable set $O \supset A$ which is obtained from open intervals using two operations and whose measure(length) equals the outer measure of A .

Theorem 1.8. (Continuity Property of the Lebesgue measure) *The Lebesgue measure μ preserves limits.*

(1) Suppose that $(A_n)_{n=1}^\infty$ is a sequence of measurable sets in \mathcal{F} . Then, we have:

$$\lim_{m \rightarrow \infty} \mu \left(\bigcup_{i=1}^m A_i \right) = \mu \left(\lim_{m \rightarrow \infty} \bigcup_{i=1}^m A_i \right) = \mu \left(\bigcup_{i=1}^\infty A_i \right) \quad (1.22)$$

(2) If $A_n \subset A_{n+1}$ is a monotonically increasing sequence of sets in \mathcal{F} , then we have:

$$\lim_{m \rightarrow \infty} \mu(A_m) = \mu \left(\bigcup_{m=1}^\infty A_m \right) \quad (1.23)$$

(3) If $A_n \supset A_{n+1}$ is a monotonically decreasing sequence of sets in \mathcal{F} , then we have:

$$\lim_{m \rightarrow \infty} \mu(A_m) = \mu \left(\bigcap_{m=1}^\infty A_m \right) \quad (1.24)$$

Proof. (1) Define a new family of sets $B_1 = A_1, B_2 = A_2 \setminus A_1, \dots, B_n = A_n \setminus \bigcup_{i=1}^{n-1} A_i$ and so forth. Then, we make the following claims:

Claim I. $B_i \cap B_j = \emptyset$, for all $i \neq j$.

We proceed by contradiction. Let $m < n$. Assume that there exists an element $x \in B_m \cap B_n$. It follows that:

$$\begin{aligned} x \in (B_m \cap B_n) &\iff (x \in B_m) \wedge (x \in B_n) \\ &\iff \left(x \in \left(A_m \setminus \bigcup_{i=1}^{m-1} A_i \right) \right) \wedge \left(x \in \left(A_n \setminus \bigcup_{j=1}^{n-1} A_j \right) \right) \end{aligned}$$

In words, x belongs to both A_m and the set $\left(\bigcup_{j=1}^{n-1} A_j \right)^C$. Since, $m, n \in \mathbf{Z}_+$, and $m < n$, we must have $m \leq n-1$. If $x \in A_m$, then it must belong to $\bigcup_{j=1}^{n-1} A_j$. This is a contradiction. Hence, our initial assumption is false. $B_m \cap B_n$ is disjoint.

Claim II. $\bigcup_{i=1}^{\infty} A_i = \bigcup_{i=1}^{\infty} B_i$.

We proceed by mathematical induction. The claim is vacuously true for $n = 1$, since $B_1 = A_1$ by construction. For $n = 2$, we have:

$$\begin{aligned} A_1 \cup A_2 &= (A_2 \cup A_1) \cap (A_1 \cup A_1^C) \\ &= ((A_2 \cup A_1) \cap A_1) \cup ((A_2 \cup A_1) \cap A_1^C) \\ &= A_1 \cup ((A_2 \cap A_1^C) \cup \emptyset) \\ &= A_1 \cup (A_2 \setminus A_1) \\ &= B_1 \cup B_2 \end{aligned}$$

Assume that the claim is true for $n-1$. Define $S = \left(\bigcup_{i=1}^{n-1} A_i \right)$ We have:

$$\begin{aligned}
\bigcup_{i=1}^n A_i &= (A_n \cup S) \cap (S \cup S^C) \\
&= S \cup (A_n \setminus S) \\
&= \left(\bigcup_{i=1}^{n-1} A_i \right) \cup \left(A_n \setminus \left(\bigcup_{i=1}^{n-1} A_i \right) \right) \\
&= \left(\bigcup_{i=1}^{n-1} B_i \right) \cup B_n \\
&\quad \{\text{since the claim holds for } n-1\} \\
&= \bigcup_{i=1}^n B_i
\end{aligned}$$

Hence, the proposition holds true for all n . Passing to the limit as $n \rightarrow \infty$, we have the desired result. This closes the proof.

Since $\{B_i, i \geq 1\}$ is a disjoint sequence of events, and using the above claims, we get:

$$\begin{aligned}
\mu \left(\bigcup_{i=1}^{\infty} A_i \right) &= \mu \left(\bigcup_{i=1}^{\infty} B_i \right) \\
&= \sum_{i=1}^{\infty} \mu(B_i) \\
&\quad \{\text{Countable additivity}\}
\end{aligned}$$

Therefore:

$$\begin{aligned}
\mu\left(\bigcup_{i=1}^{\infty} A_i\right) &= \sum_{i=1}^{\infty} \mu(B_i) \\
&= \lim_{m \rightarrow \infty} \sum_{i=1}^m \mu(B_i) \\
&\quad \{\text{An infinite series converges to the limit of the sequence of partial sums.}\} \\
&= \lim_{m \rightarrow \infty} \mu\left(\bigcup_{i=1}^m B_i\right) \\
&\quad \{\text{Finite additivity}\} \\
&= \lim_{m \rightarrow \infty} \mu\left(\bigcup_{i=1}^m A_i\right) \\
&\quad \{\text{By construction}\}
\end{aligned}$$

(2) If $A_n \subset A_{n+1}$, then $\bigcup_{i=1}^m A_i = A_m$. Consequently,

$$\mu\left(\bigcup_{i=1}^{\infty} A_i\right) = \lim_{m \rightarrow \infty} \mu\left(\bigcup_{i=1}^m A_i\right) = \lim_{m \rightarrow \infty} \mu(A_m)$$

(3) If $A_n \supset A_{n+1}$, then $A_1 \setminus A_n \subset A_1 \setminus A_{n+1}$. Thus, $\{A_1 \setminus A_n, n \geq 1\}$ is an increasing sequence of sets. From (2), it follows that:

$$\begin{aligned}
\lim_{m \rightarrow \infty} \mu(A_1 \setminus A_m) &= \mu \left(\lim_{m \rightarrow \infty} \bigcup_{i=1}^m A_1 \setminus A_i \right) \\
&= \mu \left(\lim_{m \rightarrow \infty} \bigcup_{i=1}^m (A_1 \cap A_i^C) \right) \\
&= \mu \left(\lim_{m \rightarrow \infty} A_1 \cap \left(\bigcup_{i=1}^m A_i^C \right) \right) \\
&= \mu \left(\lim_{m \rightarrow \infty} A_1 \cap \left(\bigcap_{i=1}^m A_i \right)^C \right) \\
\lim_{m \rightarrow \infty} \mu(A_1) - \lim_{m \rightarrow \infty} \mu(A_m) &= \mu(A_1) - \mu \left(\lim_{m \rightarrow \infty} \bigcap_{i=1}^m A_i \right) \\
\Rightarrow \lim_{m \rightarrow \infty} \mu(A_m) &= \mu \left(\lim_{m \rightarrow \infty} \bigcap_{i=1}^m A_i \right)
\end{aligned}$$

□

Remark. The proof of theorem (1.8) simply relies on countable additivity of μ and on the definition of the sum of an infinite series in $[0, \infty]$, i.e. that:

$$\sum_{i=1}^{\infty} \mu(A_i) = \lim_{n \rightarrow \infty} \sum_{i=1}^n \mu(A_i)$$

Consequently, this result is true not only for the set function μ , but any countably additive set function defined on a sigma-field. It also leads us to the following claim, which, though, we consider it here only for μ , actually characterizes countably additive set functions.




Figure. The sets B_n and A_n (light-gray).

Theorem 1.9. *The set function μ satisfies:*

(1) μ is finitely additive, that is, for pairwise disjoint sets (A_i) we have:

$$\mu \left(\bigcup_{i=1}^n A_i \right) = \sum_{i=1}^n \mu(A_i)$$

for each n ;

(ii) μ is continuous at \emptyset , that is, if (B_n) decrease to \emptyset , $\mu(B_n)$ decreases to 0.

Proof. To prove this claim, recall that $\mu : \mathcal{F} \rightarrow [0, \infty]$ is countably additive. This implies (i), as we have already seen. To prove (ii), consider a sequence (B_n) in \mathcal{F} which decreases to \emptyset . Then, $A_n = B_n \setminus B_{n+1}$ defines a disjoint sequence in \mathcal{F} and $\bigcup_{n=1}^{\infty} A_n = B_1$. We may assume that B_1 is bounded, so that $\mu(B_n)$ is finite for all n , so that, $\mu(A_n) = \mu(B_n \setminus B_{n+1}) = \mu(B_n) - \mu(B_{n+1}) \geq 0$ and hence we have:

$$\begin{aligned}
\mu(B_1) &= \sum_{n=1}^{\infty} \mu(A_n) \\
&= \sum_{n=1}^{\infty} \mu(B_n) - \mu(B_{n+1}) \\
&= \lim_{n \rightarrow \infty} (\mu(B_1) - \mu(B_n))
\end{aligned}$$

which shows that $\lim_{n \rightarrow \infty} \mu(B_n) \rightarrow 0$. □

1.5 Borel Sets.

The definition of \mathcal{F} does not lend itself easily to the verification that a particular set belongs to \mathcal{F} ; in our proofs we have had to work quite hard to show that \mathcal{F} is closed under various operations. It is therefore useful to add another construction to our armoury; one which shows more directly, how open sets (and indeed open intervals) and the structure of sigma-fields lie at the heart of many of the concepts we have developed. We begin with an auxiliary construction enabling us to produce new sigma-fields.

Theorem 1.10. *The intersection of a family of σ -fields is a σ -field.*

Proof. Let \mathcal{F}_α be σ -fields for $\alpha \in \Lambda$ (the index set Λ can be arbitrary). Put

$$\mathcal{F} = \bigcap_{\alpha \in \Lambda} \mathcal{F}_\alpha$$

We verify the conditions of the definition.

1. $\mathbf{R} \in \mathcal{F}_\alpha$ for all $\alpha \in \Lambda$ so $\mathbf{R} \in \mathcal{F}$.
2. If $E \in \mathcal{F}$, then $E \in \mathcal{F}_\alpha$ for all $\alpha \in \Lambda$. Since \mathcal{F}_α is a σ -field, it is closed under complementation, so E^C belongs to \mathcal{F}_α for all $\alpha \in \Lambda$. Hence, $E^C \in \mathcal{F}$.
3. If E_k belongs to \mathcal{F} for $k = 1, 2, 3, \dots$, then $E_k \in \mathcal{F}_\alpha$ for all α, k hence, $\bigcup_{k=1}^{\infty} E_k \in \mathcal{F}_\alpha$ for all α and so $\bigcup_{k=1}^{\infty} E_k \in \mathcal{F}$. □

Definition 1.5. Put

$$\mathcal{B} = \bigcap \{ \mathcal{F} : \mathcal{F} \text{ is a sigma-field containing all intervals} \} \quad (1.25)$$

We say that \mathcal{B} is the σ -field generated by all the intervals and we call the elements of \mathcal{B} - *Borel sets* (after Emile Borel 1871-1956). It is obviously the smallest σ -field containing all the intervals. In general, we say that \mathcal{G} is the σ -field generated by a family of sets \mathcal{A} if $\mathcal{G} = \bigcap \{ \mathcal{F} : \mathcal{F} \text{ is a sigma-field such that } \mathcal{F} \supset \mathcal{A} \}$.

Example 1.1. (Borel Sets) The following examples illustrate how the closure properties of the σ -field \mathcal{B} may be used to verify that most familiar sets in \mathbf{R} belong to \mathcal{B} .

- (1) By construction, all intervals belong to \mathcal{B} and since \mathcal{B} is a σ -field, all open sets must belong to \mathcal{B} , as any open set is the countable union of open intervals.
- (2) Countable sets are Borel sets, since each set is a countable union of closed intervals of the form $[a, a]$; in particular \mathbf{N} and \mathbf{Q} are Borel sets. Since, \mathcal{B} is a σ -field, it is closed under complementation. So, $\mathbf{R} \setminus \mathbf{Q}$ - the set of irrational numbers belongs to \mathcal{B} and it is a Borel set. Similarly, finite sets are also Borel sets.

The definition of \mathcal{B} is also very flexible - as long as we start with all intervals of a particular type, these collections generate the same Borel σ -field:

Theorem 1.11. *If instead of all intervals, we take all open intervals, all closed intervals, all intervals of the form (a, ∞) (or of the form $[a, \infty)$, $(-\infty, b)$ or $(-\infty, b]$), all open sets, or all closed sets, then the σ -field generated by them is the same as \mathcal{B} .*

Proof. Let I be the set of all intervals and O be the set of all open intervals. Consider for example the σ -field generated by the family of open intervals O and denote it by \mathcal{C} :

$$\mathcal{C} = \bigcap \{ \mathcal{F} \supset O, \mathcal{F} \text{ is a sigma-field} \}$$

We have to show that $\mathcal{B} = \mathcal{C}$. Since open intervals are intervals, $O \subset I$ (the family of all intervals), then :

$$\{ \mathcal{F} \supset I \} \subset \{ \mathcal{F} \supset O \}$$

that is the collection of all σ -fields \mathcal{F} which contain I is smaller than the collection of all σ -fields which contain the smaller family O , since it is a more demanding requirement to contain a bigger family, so there are fewer such objects. The inclusion is reversed after we take the intersection on both sides, thus $\mathcal{C} \subset \mathcal{B}$ (the intersection of a smaller family is bigger, as the requirement of belong to each of its members is a less stringent one).

We shall show that \mathcal{C} contains all the intervals. This will be sufficient, since \mathcal{B} is the intersection of such σ -fields, so it is contained in each, and therefore $\mathcal{B} \subset \mathcal{C}$.

To this end, consider the intervals $[a, b)$, $[a, b]$, (a, b) (the intervals of the form (a, b) are in \mathcal{C} by definition):

$$[a, b) = \bigcap_{n=1}^{\infty} \left(a - \frac{1}{n}, b \right)$$

$$[a, b] = \bigcap_{n=1}^{\infty} \left(a - \frac{1}{n}, b + \frac{1}{n} \right)$$

$$(a, b] = \bigcap_{n=1}^{\infty} \left(a, b + \frac{1}{n} \right)$$

\mathcal{C} as a σ -field is closed with respect to countable intersection, so it contains the sets on the right. The argument for unbounded intervals is similar:

$$(a, \infty) = \bigcup_{n=1}^{\infty} (a, n)$$

and

$$(-\infty, b) = \bigcup_{n=1}^{\infty} (-n, b)$$

The proof is complete. □

Remark. Since \mathcal{F} is a σ -field containing all the intervals, and \mathcal{B} is the smallest such σ -field, we have the inclusion $\mathcal{B} \subset \mathcal{F}$, that is every Borel set in \mathbf{R} is Lebesgue measurable. The question therefore arises whether these σ -fields might be the same. In fact, the inclusion is proper. It is not altogether straightforward to construct a set in $\mathcal{F} \setminus \mathcal{B}$. However, by theorem 1.7 (ii), given any $E \in \mathcal{F}$, we can find a Borel set $B \supset E$ of the form $B = \bigcap_n O_n$, where the O_n are open sets, such that $\mu(E) = \mu(B)$. In particular,

$$\mu(B \Delta E) = \mu(B \setminus E) = 0$$

Hence, μ cannot distinguish between the measurable set E and the Borel set B we have constructed.

Thus, given a Lebesgue measurable set E we can find a Borel set B such that their symmetric difference $E \Delta B$ is a null set. Now, we know that $E \Delta B \in \mathcal{F}$, and it is obvious that subsets of null sets are also null, and hence in \mathcal{F} . However, we cannot conclude that every null set will be a Borel set (if \mathcal{B} did contain all the null sets then by theorem 1.7 (ii), we should have

2 Expectation.

The goal of this section is to define the expectation of random variables and establish it's basic properties.

2.1 Lebesgue-measurable functions.

Integration is concerned with the process of approximation. In the Riemann integral, we split the interval $I = [a, b]$ over which we integrate into a partition $\{x_0 = a < x_1 < x_2 < \dots < x_n = b\}$. Define $I_n := [x_{n-1}, x_n]$. Then, we construct approximating sums by multiplying the lengths of small subintervals by certain numbers a_n (related to the values of the function in question; for example $a_n = \sup_{I_n} f(x)$, $a_n = \inf_{I_n} f(x)$):

$$\sum_{n=1}^{\infty} a_n l(I_n) \tag{2.1}$$

For large n , this sum is close to the Riemann integral $\int_a^b f(x)dx$.




Figure. Riemann sums.

The approach to the Lebesgue integral is similar but there is a crucial difference. Instead of splitting the integration domain into various parts, we decompose the range of the function. Again, a simple way is to introduce short intervals J_n of equal length.




Figure. Lebesgue sums.

To build the approximating sums, we first take the inverse images of J_n by f , that is by $f^{-1}(J_n)$. These may be complicated sets, not necessarily intervals. Here the theory of measure developed previously comes into its own. We are able to measure sets provided they are measurable i.e. they are in \mathcal{F} . Given that, we compute:

$$\sum_{n=1}^N a_n \mu(f^{-1}(J_n)) \quad (2.2)$$

where $a_n \in J_n$ or $a_n = \inf J_n = y_{n-1}$ for example. The following definition guarantees that the above procedure makes sense.

Definition 2.1. Suppose that E is a measurable set. We say that a function $f : E \rightarrow \mathbf{R}$ is (*Lebesgue*)-measurable if for any interval $I \subseteq \mathbf{R}$

$$f^{-1}(I) = \{x \in \mathbf{R} : f(x) \in I\} \in \mathcal{F}$$

In what follows, the term *measurable* (without qualification) will refer to Lebesgue-measurable functions.

If all the sets $f^{-1}(I) \in \mathcal{B}$, that is, if they are Borel sets, we call f *Borel-measurable*, or simply a Borel function.

The underlying philosophy is one which is common for various mathematical notions : the inverse image of a *nice set* is *nice*. Remember continuous functions, for example, where the inverse image of an open set is open. The actual meaning of the word nice depends on the particular branch of mathematics.

Remark. The terminology is unfortunate. *Measurable objects* should be measured (as with measurable sets). However, measurable functions will be integrated. The confusion here stems from the fact that the word *integrable* which would probably best fit here, carries a more restricted meaning as we shall see later.

2.2 Simple Random Variables.

In the special case of probability spaces we use the phrase *random variable* to mean a measurable function. That is, if $(\Omega, \mathcal{F}, \mathbb{P})$ is a probability space, then $X : \Omega \rightarrow \mathbf{R}$ is a random variable if for all $x \in \mathbf{R}$, the set $X^{-1}((-\infty, x])$ is in \mathcal{F} :

$$\{\omega \in \Omega : X(\omega) \leq x\} \in \mathcal{F}$$

A function $f : \Omega \rightarrow \mathbf{R}$ is called *simple* if its image $f(\Omega)$ is a finite-set. That is, f can be written as a finite linear-combination of indicator functions. We can write:

$$f(\omega) = \sum_{i=1}^n a_i I_{A_i}(\omega)$$

for all $\omega \in \Omega$, for some distinct $a_1, \dots, a_n \geq 0$ (values) and sets A_1, \dots, A_n which form a partition of Ω .

A random variable $X : \Omega \rightarrow \mathbf{R}$ is called simple, if its image $X(\Omega)$ takes a finite set of values. That is, X can be written as a finite linear-combination of indicator random variables. We can write:

$$X(\omega) = \sum_{i=1}^n a_i I_{A_i}(\omega)$$

for all $\omega \in \Omega$, for some distinct $a_1, \dots, a_n \geq 0$ and events A_1, \dots, A_n which form a partition of Ω . Note that: $X \geq 0$.

The abstract(Lebesgue) integral of a simple function f (with respect to the measure μ), denoted $\int f d\mu$ is defined as:

$$\int f d\mu = \sum_{k=1}^n a_k \mu(A_k)$$

The **expectation** of the simple random variable X , denoted by $\mathbb{E}X$ is defined as :

$$\int X d\mathbb{P} = \mathbb{E}X = \sum_{k=1}^n x_k \mathbb{P}(A_k)$$

This equates to discretising the y -axis.

The expectation of a non-negative random variable X is defined as :

$$\mathbb{E}X = \sup\{\mathbb{E}Z : Z \text{ is simple and } Z \leq X\}$$

Note that, we can always take $Z = 0$, so that, $\mathbb{E}Z = 0$ and therefore $\mathbb{E}X$ is bounded below by 0. That is, $\mathbb{E}X \geq 0$.

The abstract(Lebesgue) integral of a non-negative function f is defined as:

$$\int f d\mu = \sup\left\{\int q d\mu : q \text{ is simple and } q \leq f\right\}$$

Again, we can always take $q = 0$, so that $\int q d\mu = 0 \cdot I_\Omega = 0$. Therefore, $\int f d\mu$ is bounded below by zero and $\int f d\mu \geq 0$.

For an arbitrary random variable X , we can always write :

$$X = X^+ - X^-$$

where $X^+ = \max\{X, 0\} = X \cdot I_{\{X \geq 0\}}$ and $X^- = \max\{-X, 0\} = -X \cdot I_{\{X \leq 0\}}$.

These are non-negative random variables and the expectation of X is defined as:

$$\mathbb{E}X = \mathbb{E}X^+ - \mathbb{E}X^-$$

Theorem 2.1. Let X and Y be simple random variables. Then, $\mathbb{E}(X + Y) = \mathbb{E}X + \mathbb{E}Y$.

Proof. Let $X = \sum_{k=1}^m x_k I_{A_k}$ and $Y = \sum_{l=1}^n y_l I_{B_l}$ for some non-negative numbers x_k, y_l and events A_k and B_l are such that the A_k and B_l partition Ω . Then, the events $A_k \cap B_l$ partition Ω and

$$\begin{aligned}
\mathbb{E}(X + Y) &= \sum_{k \leq m, l \leq n} (x_k + y_l) \mathbb{P}(A_k \cap B_l) \\
&= \sum_{k \leq m, l \leq n} x_k \mathbb{P}(A_k \cap B_l) + \sum_{k \leq m, l \leq n} y_l \mathbb{P}(A_k \cap B_l) \\
&= \sum_{k \leq m} x_k \sum_{l \leq n} \mathbb{P}(A_k \cap B_l) + \sum_{l \leq n} y_l \sum_{k \leq m} \mathbb{P}(A_k \cap B_l) \\
&= \sum_{k \leq m} x_k (\mathbb{P}(A_k \cap B_1) + \mathbb{P}(A_k \cap B_2) + \dots + \mathbb{P}(A_k \cap B_n)) \\
&\quad + \sum_{l \leq n} y_l (\mathbb{P}(A_1 \cap B_l) + \mathbb{P}(A_2 \cap B_l) + \dots + \mathbb{P}(A_m \cap B_l)) \\
&= \sum_{k \leq m} x_k \mathbb{P}(A_k) + \sum_{l \leq n} y_l \mathbb{P}(B_l) \\
&= \mathbb{E}X + \mathbb{E}Y
\end{aligned}$$

□

2.3 Non-negative Random Variables.

Our main goal is to prove the linearity of expectation. We first establish a few basic properties of expectation for non-negative random variables.

Theorem 2.2. Let X and Y be non-negative random variables. We have:

- (a) If $A \in \mathcal{F}$, then $\mathbb{E}I_A = \int I_A \cdot d\mathbb{P} = \mathbb{P}(A)$.
- (b) (Monotonicity). If $X \leq Y$, then $\mathbb{E}X \leq \mathbb{E}Y$.
- (c) (Translation and Scaling) For $a \geq 0$, $\mathbb{E}(a + X) = a + \mathbb{E}X$ and $\mathbb{E}(aX) = a\mathbb{E}X$.
- (d) If $\mathbb{E}X = 0$, then $X = 0$ almost surely (that is $\mathbb{P}\{X = 0\} = 1$).
- (e) If A and B are events such that $A \subset B$, then $\mathbb{E}X1_A \leq \mathbb{E}X1_B$.

Proof. (a) I_A is a simple random variable. Then, by the definition of the Lebesgue integral, $\mathbb{E}I_A = \int I_A d\mathbb{P} = 1 \cdot \mathbb{P}(A)$.

(b) Let S_X, S_Y be the set of all simple random variables which are less than or equal to X, Y respectively. Since $X \leq Y$, every simple random variable which is less than or equal to X is also less than or equal to Y . But, there exists simple random variables that are less than or equal to Y but greater than X . Consequently, $S_X \subseteq S_Y$. Thus, $\{\mathbb{E}Z : Z \text{ is simple and } Z \leq X\} \subseteq \{\mathbb{E}Z : Z \text{ is simple and } Z \leq Y\}$. Therefore, it follows that $\sup\{\mathbb{E}Z : Z \text{ is simple and } Z \leq X\} \leq \sup\{\mathbb{E}Z : Z \text{ is simple and } Z \leq Y\}$. Consequently, $\mathbb{E}X \leq \mathbb{E}Y$.

(c) Let Z be an arbitrary simple random variable which is less than or equal to X . Then, $Z = \sum_{k=1}^m x_k I_{A_k}$ where $x_k \geq 0$. We have:

$$\begin{aligned}\mathbb{E}(a + Z) &= \sum_{k=1}^m (a + x_k) \mathbb{P}(A_k) \\ &= a \sum_{k=1}^m \mathbb{P}(A_k) + \sum_{k=1}^m x_k \mathbb{P}(A_k) \\ &= a + \mathbb{E}Z\end{aligned}$$

Note that, for all simple random variables $Z \leq X \iff a + Z \leq a + X$.

$$\begin{aligned}\mathbb{E}(a + X) &= \sup\{\mathbb{E}(a + Z) : a + Z \text{ is a simple random variable and } a + Z \leq a + X\} \\ &= \sup\{a + \mathbb{E}Z : Z \text{ is a simple random variable and } Z \leq X\} \\ &= a + \sup\{\mathbb{E}Z : Z \text{ is a simple random variable and } Z \leq X\} \\ &= a + \mathbb{E}X\end{aligned}$$

Also,

$$\begin{aligned}\mathbb{E}(aZ) &= \sum_{k=1}^m ax_k \mathbb{P}(A_k) \\ &= a \sum_{k=1}^m x_k \mathbb{P}(A_k) \\ &= a\mathbb{E}Z\end{aligned}$$

(\forall simple random variables Z)($Z \leq X$) $\iff aZ \leq aX$.

$$\begin{aligned}\mathbb{E}(aX) &= \sup\{\mathbb{E}aZ : aZ \text{ is a simple random variable and } aZ \leq aX\} \\ &= \sup\{a\mathbb{E}Z : Z \text{ is a simple random variable and } Z \leq X\} \\ &= a \sup\{\mathbb{E}Z : Z \text{ is a simple random variable and } Z \leq X\} \\ &= a\mathbb{E}X\end{aligned}$$

(d) For $n \geq 1$, we have $X \geq XI_{\{X \geq \frac{1}{n}\}} \geq \frac{1}{n}I_{\{X \geq \frac{1}{n}\}}$. So, by (a) and (b), we have:

$$0 = \mathbb{E}X \geq \frac{1}{n}\mathbb{E}I_{\{X \geq \frac{1}{n}\}} = \frac{1}{n}\mathbb{P}\{X \geq \frac{1}{n}\}$$

But since $\mathbb{P}\{X \geq \frac{1}{n}\} \geq 0$, we conclude that $\mathbb{P}\{X \geq \frac{1}{n}\} = 0$.

Now,

$$\mathbb{P}(X > 0) = \mathbb{P}\left(\bigcup_{n=1}^{\infty} \{X \geq \frac{1}{n}\}\right) = \mathbb{P}(\lim\{X \geq \frac{1}{n}\}) = \lim \left(\mathbb{P}\left\{X \geq \frac{1}{n}\right\} \right) = 0$$

(e) Clearly, if $A \subset B$, then $X \cdot I_A \leq X \cdot I_B$. Thus, by the monotonicity property, $\mathbb{E}X I_A \leq \mathbb{E}X I_B$. \square

The following lemma gives a way to approximate non-negative random variables with monotone sequences of simple ones.

Lemma 2.1. *If X is a random variable, then there is a sequence (Z_n) of non-negative simple random variables such that for every $\omega \in \Omega$, $Z_n(\omega) \leq Z_{n+1}(\omega)$ and $Z_n(\omega) \rightarrow X(\omega)$ pointwise.*

Proof. For each positive integer n , define

$$Z_n = \sum_{k=1}^{n \cdot 2^n} \frac{k-1}{2^n} \mathbf{1}_{\{\frac{k-1}{2^n} < X < \frac{k}{2^n}\}} + n \cdot \mathbf{1}_{\{X \geq n\}}$$

Essentially, we are dividing the interval $(0, n)$ on the y -axis into $n \cdot 2^n$ strips, each of size $1/2^n$. Beyond the point $X \geq n$, Z_n takes the constant value n .

If $n = 2$, this is what $Z_2(\omega)$ looks like. It chops the interval $[0, 2]$ on the Y -axis into 8 sub-intervals.




Figure. The step function $Z_2(\omega)$.

As n increases, $Z_n(\omega)$ better approximates $X(\omega)$.

Pick any arbitrary $\omega \in \Omega$. Let $\epsilon > 0$.

By the Archimedean property, there exists a natural number $N_1 \in \mathbf{N}$, such that $N_1 > X(\omega)$.

We have that $X(\omega)$ lies in an interval I_n , that is $\frac{k-1}{2^n} < X(\omega) < \frac{k}{2^n}$ for some $1 \leq k \leq n \cdot 2^n$, $k \in \mathbf{Z}^+$, for all $n \geq N_1$.

There exists $N_2 \in \mathbf{N}$, such that $l(I_n) = \frac{1}{2^n} < \epsilon$ for all $n \geq N_2$.

Pick $N = \max\{N_1, N_2\}$. Then, for all $n \geq N$, $|Z_n(\omega) - X(\omega)| < \epsilon$.

Thus, $(Z_n(\omega))$ converges pointwise to $X(\omega)$ for all $\omega \in \Omega$.

Note that, the partition points at stage $(n + 1)$ include the partition points at stage n and new partition points at the mid-points of the old ones. Because of this, (Z_n) is a monotonically increasing sequence. \square

Lemma 2.2. *If X is a positive random variable, and if $(X_n)_{n=1}^\infty$ is any sequence of positive simple random variables increasing to X , then $\mathbf{E}[X_n]$ increases to $\mathbf{E}[X]$.*

Proof. Recall the definition of $\mathbf{E}[X]$. We have:

$$\mathbf{E}[X] = \sup \{ \mathbf{E}Z : Z \text{ is a simple random variable and } Z \leq X \}$$

Therefore, observe that if the sequence $\mathbf{E}X_n$ increases to some limit a , then $X_n \leq X$ implies $\mathbf{E}X_n \leq \mathbf{E}X$, so $\lim \mathbf{E}X_n \leq \mathbf{E}X$, and so, $a \leq \mathbf{E}[X]$.

To obtain, that indeed $a = \mathbf{E}[X]$, it should be enough to prove that if Y is a simple random variable such that $0 \leq Y \leq X$, then $\mathbf{E}[Y] \leq a$.

Suppose the random variable Y takes on m different values, say, a_1, \dots, a_m and let the set $A_k = \{\omega : Y(\omega) = a_k\}$. Choose $\epsilon \in (0, 1]$. The random variable $Y_{n,\epsilon} = (1 - \epsilon)Y\mathbf{1}_{\{(1-\epsilon)Y \leq X_n\}}$ takes the value $(1 - \epsilon)a_k$ on the set $A_{k,n,\epsilon} = A_k \cap \{(1 - \epsilon)Y \leq X_n\}$ and 0 on the set $\{(1 - \epsilon)Y \geq X_n\}$.

At this point, I think a sketch is pretty helpful. Essentially, $Y_{n,\epsilon}$ is a shifted step function.




Figure. The shifted step function $Y_{n,\epsilon} = (1 - \epsilon)Y\mathbf{1}_{(1-\epsilon)Y \leq X_n}$.

Furthermore, it is obvious that $Y_{n,\epsilon} \leq X_n$.

So, the expectation of this random variable is:

$$\mathbf{E}Y_{n,\epsilon} = (1 - \epsilon) \sum_{k=1}^m a_k \mathbb{P}(A_{n,k,\epsilon}) \quad (2.3)$$

and by monotonicity, $\mathbf{E}Y_{n,\epsilon} \leq \mathbf{E}X_n = a$.

Now, recall that $Y \leq \lim_n X_n$, and hence $(1 - \epsilon)Y < \lim_n X_n$. It should be clear from the diagram, intuitively, that, as $n \rightarrow \infty$, the set $A_{k,n,\epsilon} = A_k \cap \{(1-\epsilon)Y \leq X_n\}$ approaches A_k . Consequently, by the continuity of probability measure, $\mathbb{P}(A_{n,k,\epsilon}) \rightarrow \mathbb{P}(A_k)$. Hence, taking the limit in (2.3) gives:

$$(1 - \epsilon) \sum_{k=1}^m a_k \mathbb{P}(A_k) = (1 - \epsilon) \mathbf{E}Y \leq a$$

Letting $\epsilon \rightarrow 0$ in the above, we deduce that $\mathbf{E}[Y] \leq a$, as required. \square

Theorem 2.3. (*Monotone Convergence Theorem*). *Let X_1, X_2, \dots, X_n be a sequence of random variables converging almost surely to another random variable X . That is,*

$$0 \leq X_1 \leq X_2 \leq X_3 \leq \dots \leq$$

almost surely, then

$$\lim_{n \rightarrow \infty} \mathbf{E}X_n = \mathbf{E}(\lim_{n \rightarrow \infty} X_n) = \mathbf{E}X$$

That is, expectation preserves limits.

Proof. Using lemma (2.1), for each n , we can choose an increasing sequence $Y_{n,k}$, $k = 1, 2, 3, \dots$ of positive simple random variables increasing to X_n and set:

$$Z_k = \max_{n \leq k} Y_{n,k}$$

Essentially, we have these sequences of positive increasing random variables $(Y_{1,k}) \rightarrow X_1$, $(Y_{2,k}) \rightarrow X_2$, \dots , $(Y_{n,k}) \rightarrow X_n$. And now, we construct a sequence (Z_k) by taking Z_k as the maximum of the r.v.'s $\{Y_{k,1}, Y_{k,2}, \dots, Y_{k,k}, \dots\}$. Thus, $(Z_k : k \geq 1)$ is a non-decreasing sequence of positive simple random variables and thus it has a limit $Z = \lim_{k \rightarrow \infty} Z_k$. Also,

$$Y_{n,k} \leq Z_k \leq X_k \leq X \quad \text{almost surely } \forall n \leq k$$

Hence,

$$\lim_{k \rightarrow \infty} Y_{n,k} \leq \lim_{k \rightarrow \infty} Z_k \leq X \quad \text{almost surely}$$

In other words,

$$X_n \leq Z \leq X \quad \text{almost surely}$$

Next, if we let $n \rightarrow \infty$, we have:

$$X = Z \quad \text{almost surely}$$

Since the expectation is a positive operator, we have:

$$\mathbf{E}[Y_{n,k}] \leq \mathbf{E}[Z_k] \leq \mathbf{E}[X_k] \quad \text{for } n \leq k$$

Fix n and let $k \rightarrow \infty$. Using lemma (2.2), we obtain:

$$\mathbf{E}[X_n] \leq \mathbf{E}[Z] \leq \lim_{k \rightarrow \infty} \mathbf{E}[X_k]$$

Now, let $n \rightarrow \infty$ to obtain:

$$\lim_{n \rightarrow \infty} \mathbf{E}[X_n] \leq \mathbf{E}[Z] \leq \lim_{k \rightarrow \infty} \mathbf{E}[X_k]$$

By the squeeze theorem,

$$\lim_{n \rightarrow \infty} \mathbf{E}[X_n] = \mathbf{E}[Z] = \mathbf{E}[\lim_{n \rightarrow \infty} X_n]$$

□

Theorem 2.4. (*Linearity of Expectations*) Let X and Y be non-negative random variables. Then,

$$\mathbb{E}(X + Y) = \mathbb{E}X + \mathbb{E}Y$$

Proof. By lemma 2.1, there exists monotonic sequences of non-negative random variables $(X_n)_{n=1}^\infty$ and $(Y_n)_{n=1}^\infty$ such that $(X_n) \rightarrow X$ and $(Y_n) \rightarrow Y$. Then, the sequence $X_n + Y_n$ is also monotone, and by the Algebraic limit theorem for sequences, $X_n + Y_n \rightarrow X + Y$. By theorem 2.1,

$$\mathbb{E}(X_n + Y_n) = \mathbb{E}X_n + \mathbb{E}Y_n$$

Passing to the limits, we get:

$$\lim \mathbb{E}(X_n + Y_n) = \lim \mathbb{E}X_n + \lim \mathbb{E}Y_n$$

By the Monotone convergence theorem, \mathbb{E} preserves limits, so,

$$\mathbb{E}(X + Y) = \mathbb{E}X + \mathbb{E}Y$$

□

2.4 Fatou's Lemma.

Theorem 2.5. (Fatou's Lemma) *Let Y be a random variable that satisfies $\mathbb{E}[|Y|] < \infty$. Let $(X_n)_{n=1}^\infty$ be a sequence of random variables. Then the following holds:*

- If $Y \leq X_n$, for all n , then $\mathbb{E}[\liminf_{n \rightarrow \infty} X_n] \leq \liminf_{n \rightarrow \infty} \mathbb{E}[X_n]$.
- If $Y \geq X_n$, for all n , then $\mathbb{E}[\limsup_{n \rightarrow \infty} X_n] \geq \limsup_{n \rightarrow \infty} \mathbb{E}[X_n]$.

Proof. Firstly, if $X_n \geq Y$, that is, (X_1, X_2, X_3, \dots) is any sequence of random variables bounded below, analogous to a sequence of real numbers, the point-wise limit, $\liminf_{n \rightarrow \infty} X_n$ always exists and therefore \liminf random variable is defined. Similarly, if $X_n \leq Y$, that is, (X_1, X_2, X_3, \dots) is any sequence of random variables bounded above, then $\limsup_{n \rightarrow \infty} X_n$ always exists and therefore \limsup random variable is defined.

Fix some $n \in \mathbf{N}$. From the definition of infimum, we have:

$$\inf_{k \geq n} X_k - Y \leq X_m - Y, \quad \forall m \geq n$$

By the monotonicity property, it follows that:

$$\mathbb{E} \left[\inf_{k \geq n} X_k - Y \right] \leq \mathbb{E} [X_m - Y] \quad \forall m \geq n$$

The left-hand side is a constant real number. The right-hand side is indexed by m . So, this inequality holds for a sequence of real numbers (a_m) , $m \geq n$, where $a_m = X_m(\omega) - Y(\omega)$.

Consider the set:

$$\{a_m, a_{m+1}, a_{m+2}, \dots\}$$

This set is bounded below for all $m \geq n$. Hence, its infimum exists. I can take infimum with respect to m , on both sides. By the order limit theorem, we have:

$$\inf_{m \geq n} \mathbb{E} \left[\inf_{k \geq n} X_k - Y \right] \leq \inf_{m \geq n} \mathbb{E} [X_m - Y] \quad \forall m \geq n$$

Thus,

$$\mathbb{E} \left[\inf_{k \geq n} X_k - Y \right] \leq \inf_{m \geq n} \mathbb{E} [X_m - Y] \quad \forall m \geq n$$

Define $Z_n = \inf_{k \geq n} X_k - Y$ and $S_n = \inf_{m \geq n} \mathbb{E} [X_m - Y]$. So, we can write:

$$\mathbb{E} Z_n \leq S_n$$

Passing to the limit as $n \rightarrow \infty$, by the Order limit theorem, we have:

$$\lim_{n \rightarrow \infty} \mathbb{E} Z_n \leq \lim_{n \rightarrow \infty} S_n$$

Note that, $Z_n \geq 0$, since $X_k \geq Y$. And Z_n is a sequence of monotonically increasing random variables. Thus, $\lim Z_n$ exists. By the Monotone Convergence theorem,

$$\lim_{n \rightarrow \infty} \mathbb{E} Z_n = \mathbb{E} \left[\lim_{n \rightarrow \infty} Z_n \right] = \mathbb{E} \left[\liminf_{n \rightarrow \infty} X_n - Y \right] \leq \lim_{n \rightarrow \infty} S_n = \liminf_{n \rightarrow \infty} \mathbb{E} [X_n - Y]$$

and so, it follows that:

$$\mathbb{E} \left[\liminf_{n \rightarrow \infty} X_n \right] \leq \liminf_{n \rightarrow \infty} \mathbb{E} [X_n]$$

□

The DCT is an important result which asserts a sufficient condition under which we can interchange a limit and expectation.

Theorem 2.6. (*Dominated Convergence Theorem*). *Consider a sequence of random variables that converges almost surely to X . Suppose that there exists a random variable Y , such that $|X_n| \leq Y$ almost surely for all n and $\mathbb{E}[Y] < \infty$. Then, we have:*

$$\lim_{n \rightarrow \infty} \mathbb{E}[X_n] = \mathbb{E}[X]$$

Proof. Since $-Y \leq X_n \leq Y$ for all $n \in \mathbf{N}$, we can invoke both sides of Fatou's Lemma:

$$\mathbb{E} \left[\liminf_{n \rightarrow \infty} X_n \right] \leq \liminf_{n \rightarrow \infty} \mathbb{E} X_n$$

and

$$\mathbb{E} \left[\limsup_{n \rightarrow \infty} X_n \right] \geq \limsup_{n \rightarrow \infty} \mathbb{E} X_n$$

Thus,

$$\mathbb{E} X = \mathbb{E} \left[\lim_{n \rightarrow \infty} X_n \right] = \mathbb{E} \left[\liminf_{n \rightarrow \infty} X_n \right] \leq \liminf_{n \rightarrow \infty} \mathbb{E} X_n \leq \mathbb{E} X_n \leq \limsup_{n \rightarrow \infty} \mathbb{E} X_n \leq \mathbb{E} \left[\limsup_{n \rightarrow \infty} X_n \right] = \mathbb{E} X$$

This implies that:

$$\liminf_{n \rightarrow \infty} \mathbb{E} X_n = \limsup_{n \rightarrow \infty} \mathbb{E} X_n$$

so

$$\lim_{n \rightarrow \infty} \mathbb{E}X_n$$

exists and further

$$\lim_{n \rightarrow \infty} \mathbb{E}X_n = \mathbb{E}X$$

□

3 Gaussian Processes.

3.1 Random Vectors.

Consider a probability space $(\Omega, \mathcal{F}, \mathbb{P})$. We can define several random variables on Ω . A n -tuple of random variables on this space is called a random vector. For example, if X_1, X_2, \dots, X_n are random variables on $(\Omega, \mathcal{F}, \mathbb{P})$, then the n -tuple (X_1, X_2, \dots, X_n) is a random vector on $(\Omega, \mathcal{F}, \mathbb{P})$. The vector is said to be n -dimensional because it contains n -variables. We will sometimes denote a random vector by X .

A good point of view is to think of a random vector $X = (X_1, \dots, X_n)$ as a random variable (point) in \mathbf{R}^n . In other words, for an outcome $\omega \in \Omega$, $X(\omega)$ is a point sampled in \mathbf{R}^n , where $X_j(\omega)$ represents the j -th coordinate of the point. The distribution of X , denoted μ_X is the probability on \mathbf{R}^n defined by the events related to the values of X :

$$\mathbb{P}\{X \in A\} = \mu_X(A) \quad \text{for a subset } A \text{ in } \mathbf{R}^n$$

In other words, $\mathbb{P}(X \in A) = \mu_X(A)$ is the probability that the random point X falls in A . The distribution of the vector X is also called the joint distribution of (X_1, \dots, X_n) .

Definition 3.1. The **joint distribution function** of $\mathbf{X} = (X, Y)$ is the function $F : \mathbf{R}^2 \rightarrow [0, 1]$ given by:

$$F_{\mathbf{X}}(x, y) = \mathbb{P}(X \leq x, Y \leq y) \tag{3.1}$$

Definition 3.2. The joint **PDF** $f_{\mathbf{X}}(x_1, \dots, x_n)$ of a random vector \mathbf{X} is a function $f_{\mathbf{X}} : \mathbf{R}^n \rightarrow \mathbf{R}$ such that the probability that X falls in a subset A of \mathbf{R}^n and is expressed as the multiple integral of $f(x_1, x_2, \dots, x_n)$ over A :

$$\mathbb{P}(X \in A) = \int_A f(x_1, x_2, \dots, x_n) dx_1 dx_2 \dots dx_n$$

Note that: we must have that the integral of f over the whole of \mathbf{R}^n is 1.

If F is differentiable at the point (x, y) , then we usually specify:

$$f(x, y) = \frac{\partial^2}{\partial x \partial y} F(x, y) \quad (3.2)$$

Theorem 3.1. Let (X, Y) be the random variables with joint density function $f_{X,Y}(x, y)$. The marginal density function $f_X(x)$ and $f_Y(y)$ of the random variables X and Y respectively is given by:

$$\begin{aligned} f_X(x) &= \int_{-\infty}^{+\infty} f_{(X,Y)}(x, y) dy \\ f_Y(y) &= \int_{-\infty}^{+\infty} f_{(X,Y)}(x, y) dx \end{aligned}$$

Proof. We have:

$$\begin{aligned} F_X(x) &= P(X \leq x) \\ &= \int_{-\infty}^x \int_{y=-\infty}^{y=+\infty} f(x, y) dy dx \end{aligned}$$

Differentiating both sides with respect to x ,

$$f_X(x) = \int_{y=-\infty}^{y=+\infty} f(x, y) dy$$

□

Definition 3.3. For continuous random variables X and Y with the joint density function $f_{(X,Y)}$, the conditional density of Y given $X = x$ is:

$$f_{Y|X}(y|x) = \frac{f_{(X,Y)}(x, y)}{f_X(x)}$$

for all x with $f_X(x) > 0$. This is considered as a function of y for a fixed x . As a convention, in order to make $f_{Y|X}(y|x)$ well-defined for all real x , let $f_{Y|X}(y|x) = 0$ for all x with $f_X(x) = 0$.

We are essentially slicing the the joint density function of $f_{(X,Y)}(x, y)$ by a thin plane $X = x$. How can we speak of conditioning on $X = x$ for X being a continuous random variable, considering that this event has probability zero. Rigorously speaking, we are actually conditioning on the event that X falls within a small interval containing x , say $X \in (x - \epsilon, x + \epsilon)$ and then taking the limit as ϵ approaches zero from the right.

We can recover the joint PDF $f_{(X,Y)}$ if we have the conditional PDF $f_{Y|X}$ and the corresponding marginal f_X :

$$f_{(X,Y)}(x, y) = f_{Y|X}(y|x) \cdot f_X(x)$$

Theorem 3.2. (*Bayes rule and LOTP*) Let (X, Y) be continuous random variables. We have the following continuous form of the Bayes rule:

$$f_{Y|X}(y|x) = \frac{f_{X|Y}(x|y) \cdot f_Y(y)}{f_X(x)} \quad (3.3)$$

And we have the following continuous form of the law of total probability:

$$f_X(x) = \int_{y=-\infty}^{y=+\infty} f_{X|Y}(x|y) \cdot f_Y(y) dy$$

Proof. By the definition of conditional PDFs, we have:

$$f_{X|Y}(x|y) \cdot f_Y(y) = f_{(X,Y)}(x, y) = f_{Y|X}(y|x) \cdot f_X(x)$$

Dividing throughout by $f_X(x)$, we have:

$$f_{Y|X}(x) = \frac{f_{X|Y}(x|y) \cdot f_Y(y)}{f_X(x)} = \frac{f_{(X,Y)}(x, y)}{f_X(x)}$$

□

Example 3.1. (Sampling uniformly in the unit disc). Consider the random vector $\mathbf{X} = (X, Y)$ corresponding to a random point chosen uniformly in the unit disc $\{(x, y) : x^2 + y^2 \leq 1\}$. \mathbf{X} is said to have uniform on the unit circle distribution. In this case the PDF is 0 outside the disc and $\frac{1}{\pi}$ inside the disc:

$$f(x, y) = \frac{1}{\pi} \quad \text{if } x^2 + y^2 \leq 1$$

The random point (X, Y) has x -coordinate X and Y coordinate Y . Each of these are random variables and their PDFs and CDFs can be computed. This is a valid PDF, because:

$$\begin{aligned} \int \int_D f(x, y) dy dx &= \int_{-1}^1 \int_{-\sqrt{1-x^2}}^{\sqrt{1-x^2}} \frac{1}{\pi} dy dx \\ &= \frac{1}{\pi} \int_{-1}^1 [y]_{-\sqrt{1-x^2}}^{+\sqrt{1-x^2}} dx \\ &= \frac{2}{\pi} \int_{-1}^1 \sqrt{1-x^2} dx \end{aligned}$$

Substituting $x = \sin \theta$, we have: $dx = \cos \theta d\theta$ and $\sqrt{1-x^2} = \cos \theta$. The limits of integration are $\theta = -\pi/2$ to $\theta = \pi/2$. Thus,

$$\begin{aligned}
\int \int_D f(x, y) dy dx &= \frac{2}{\pi} \int_{-\pi/2}^{\pi/2} \cos^2 \theta d\theta \\
&= \frac{1}{\pi} \int_{-\pi/2}^{\pi/2} (1 + \cos 2\theta) d\theta \\
&= \frac{1}{\pi} \left[\theta + \frac{1}{2} \sin 2\theta \right]_{-\pi/2}^{\pi/2} \\
&= \frac{1}{\pi} \cdot \pi \\
&= 1
\end{aligned}$$

The CDF of X is given by:

$$\begin{aligned}
F_X(a) &= \int_{-1}^a \int_{-\sqrt{1-x^2}}^{\sqrt{1-x^2}} \frac{1}{\pi} dy dx \\
&= \frac{2}{\pi} \int_{-1}^a \sqrt{1-x^2} dx
\end{aligned}$$

I leave it in this integral form. The PDF of X is obtained by differentiating the CDF, so it is:

$$f_X(x) = \frac{2}{\pi} \sqrt{1-x^2} \quad (3.4)$$

Let's quickly plot the density of X over the domain of the definition $-1 \leq x \leq 1$.




Figure. The PDF of the random variable X .

Not suprisingly the distribution of the x -coordinate is no longer uniform!

If (X_1, X_2, \dots, X_n) is a random vector, the distribution of a single coordinate, say X_1 is called the *marginal distribution*. In the example 3.1, the marginal distribution of X is determined by the PDF 3.4.

Random variables X_1, X_2, \dots, X_n defined on the same probability space are said to be independent if for any intervals A_1, A_2, \dots, A_n in \mathbf{R} , the probability factors:

$$\mathbb{P}(X_1 \in A_1, X_2 \in A_2, \dots, X_n \in A_n) = \mathbb{P}(X_1 \in A_1) \times \mathbb{P}(X_2 \in A_2) \times \dots \times \mathbb{P}(X_n \in A_n)$$

We say that the random variables are independent and identically distributed (IID) if they are independent and their marginal distributions are the same.

When the random vector (X_1, X_2, \dots, X_n) has a joint PDF $f(x_1, x_2, \dots, x_n)$, the independence of random variables is equivalent to saying that the joint PDF is given by the product of the marginal PDFs:

$$f(x_1, x_2, \dots, x_n) = f_1(x_1) \times f_2(x_2) \times \dots \times f_n(x_n) \quad (3.5)$$

3.2 Basic Probabilistic Inequalities.

Inequalities are extremely useful tools in the theoretical development of probability theory.

3.2.1 Jensen's inequality.

Theorem 3.3. *If g is a convex function, and $a > 0$, $b > 0$, with $p \in [0, 1]$, it follows that:*

$$g(pa + (1 - p)b) \leq pg(a) + (1 - p)g(b) \quad (3.6)$$

Proof. This directly follows from the definition of convex functions. \square

3.2.2 Jensen's inequality for Random variables.

Theorem 3.4. *If g is a convex function, then it follows that:*

$$\mathbb{E}(g(X)) \geq g(\mathbb{E}X) \quad (3.7)$$

Proof. Another way to express the idea, that a function is convex is to observe that the tangent line at an arbitrary point $(t, g(t))$ always lies below the curve. Let $y = a + bx$ be the tangent to g at the point t . Then, it follows that:

$$\begin{aligned} a + bt &= g(t) \\ a + bx &\leq g(x) \end{aligned}$$

for all x .

Thus, it follows that, for any point t , there exists b such that:

$$g(x) - g(t) \geq b(x - t)$$

for all x . Set $t = \mathbb{E}X$ and $x = X$. Then,

$$g(X) - g(\mathbb{E}X) \geq b(X - \mathbb{E}X)$$

Taking expectations on both sides and simplifying:

$$\begin{aligned} \mathbb{E}(g(X)) - g(\mathbb{E}X) &\geq b(\mathbb{E}X - \mathbb{E}X) = 0 \\ \mathbb{E}g(X) &\geq g(\mathbb{E}X) \end{aligned}$$

\square

3.2.3 Young's Inequality.

Theorem 3.5. *If $a \geq 0$ and $b \geq 0$ are non-negative real numbers and if $p > 1$ and $q > 1$ are real numbers such that $\frac{1}{p} + \frac{1}{q} = 1$, then:*

$$ab \leq \frac{a^p}{p} + \frac{b^q}{q} \quad (3.8)$$

Proof. Consider $g(x) = \log x$. Being a concave function, Jensen's inequality can be reversed. We have:

$$\begin{aligned} g\left(\frac{1}{p}a^p + \frac{1}{q}b^q\right) &\geq \frac{1}{p}g(a^p) + \frac{1}{q}g(b^q) \\ \log\left(\frac{1}{p}a^p + \frac{1}{q}b^q\right) &\geq \frac{1}{p}\log(a^p) + \frac{1}{q}\log(b^q) \\ \log\left(\frac{1}{p}a^p + \frac{1}{q}b^q\right) &\geq \frac{1}{p} \cdot p\log(a) + \frac{1}{q} \cdot q\log(b) \\ \log\left(\frac{1}{p}a^p + \frac{1}{q}b^q\right) &\geq \log ab \end{aligned}$$

By the Monotonicity of the $\log x$ function, it follows that :

$$ab \leq \frac{a^p}{p} + \frac{b^q}{q}$$

□

3.2.4 Chebyshev's inequality.

One of the simplest and very useful probabilistic inequalities is a tail bound by expectation: the so called Chebyshev's inequality.

Theorem 3.6. *(Chebyshev's inequality) If X is a non-negative random variable, then for every $t \geq 0$:*

$$\mathbb{P}(X \geq t) \leq \frac{1}{t} \mathbb{E}X \quad (3.9)$$

Proof. We have:

$$t \cdot \mathbf{1}_{\{X \geq t\}} \leq \mathbf{1}_{\{X \geq t\}} \cdot X$$

By the monotonicity of expectations, we have:

$$\begin{aligned} \mathbb{E} \mathbf{1}_{\{X \geq t\}} &\leq \frac{1}{t} \mathbb{E} X \\ \Rightarrow \mathbb{P}\{X \geq t\} &\leq \frac{1}{t} \mathbb{E} X \end{aligned}$$

This closes the proof. □

There are several variants, easily deduced from Chebyshev's inequality monotonicity of several functions. For a non-negative random variable X and $t > 0$, using the power function x^p , $p > 0$, we get:

$$\mathbb{P}(X \geq t) = \mathbb{P}(X^p \geq t^p) \leq \frac{1}{t^p} \mathbb{E} X^p \quad (3.10)$$

For a real valued random variable X , every $t \in \mathbf{R}$, using the square function x^2 and variance, we have:

$$\mathbb{P}(|X - \mathbb{E}X| \geq t) \leq \frac{1}{t^2} \mathbb{E}|X - \mathbb{E}X|^2 = \frac{1}{t^2} \text{Var}(X) \quad (3.11)$$

For a real-valued random variable X , every $t \in \mathbf{R}$ and $\lambda > 0$, using the exponential function $e^{\lambda x}$ (which is monotonic), we have:

$$\mathbb{P}(X \geq t) = \mathbb{P}(\lambda X \geq \lambda t) = \mathbb{P}(e^{\lambda X} \geq e^{\lambda t}) \leq \frac{1}{e^{\lambda t}} \mathbb{E} e^{\lambda X} \quad (3.12)$$

Our next inequality, the so-called Holder's inequality is a very effective inequality to factor out the expectation of a product.

3.2.5 Holder's inequality.

Theorem 3.7. *Let $p, q \geq 1$ be such that $\frac{1}{p} + \frac{1}{q} = 1$, For random variables X and Y , we have:*

$$\mathbb{E}|XY| \leq (\mathbb{E}|X^p|)^{1/p} (\mathbb{E}|Y^q|)^{1/q}$$

Proof. From the Young's inequality, for any $a, b \in \mathbf{R}$, $p, q \geq 1$, we have:

$$ab \leq \frac{a^p}{p} + \frac{b^q}{q}$$

Setting $a = \frac{|X|}{(\mathbb{E}|X^p|)^{1/p}}$ and $b = \frac{|Y|}{(\mathbb{E}|Y^q|)^{1/q}}$, we get:

$$\frac{|XY|}{(\mathbb{E}|X^p|)^{1/p} (\mathbb{E}|Y^q|)^{1/q}} \leq \frac{1}{p} \cdot \frac{|X|^p}{\mathbb{E}|X^p|} + \frac{1}{q} \cdot \frac{|Y|^q}{\mathbb{E}|Y^q|}$$

Taking expectations on both sides, and using the monotonicity of expectation property, we get:

$$\frac{\mathbb{E}|XY|}{(\mathbb{E}|X^p|)^{1/p} (\mathbb{E}|Y^q|)^{1/q}} \leq \frac{1}{p} \cdot \frac{\mathbb{E}|X|^p}{\mathbb{E}|X^p|} + \frac{1}{q} \cdot \frac{\mathbb{E}|Y|^q}{\mathbb{E}|Y^q|} = \frac{1}{p} + \frac{1}{q} = 1$$

Consequently,

$$\mathbb{E}|XY| \leq (\mathbb{E}|X^p|)^{1/p} (\mathbb{E}|Y^q|)^{1/q}$$

Let $p = 2$ and $q = 2$. Then, we get the Cauchy-Schwarz inequality:

$$\mathbb{E}|XY| \leq [\mathbb{E}(X^2)]^{1/2} [\mathbb{E}(Y^2)]^{1/2}$$

In some ways, the p -th moment of a random variable can be thought of as its length or p -norm.

Define:

$$\|X\|_p = (\mathbb{E}|X|^p)^{1/p}$$

□

3.2.6 Minkowski's Inequality.

Theorem 3.8. *For random variables X and Y , and for all $p \geq 1$ we have:*

$$\|X + Y\|_p \leq \|X\|_p + \|Y\|_p \quad (3.13)$$

Proof. The basic idea of the proof is to use Holder's inequality. Let $\frac{1}{q} = 1 - \frac{1}{p}$ or in other words, $q = \frac{p}{p-1}$. We have:

$$\mathbb{E}|X||X + Y|^{p-1} \leq (\mathbb{E}|X|^p)^{1/p} (\mathbb{E}|X + Y|^{(p-1)q})^{1/q} \quad (a)$$

$$\mathbb{E}|Y||X + Y|^{p-1} \leq (\mathbb{E}|Y|^p)^{1/p} (\mathbb{E}|X + Y|^{(p-1)q})^{1/q} \quad (b)$$

Adding the above two equations, we get:

$$\mathbb{E}(|X + Y||X + Y|^{p-1}) \leq \mathbb{E}(|X| + |Y|)(|X + Y|^{p-1}) \leq \left\{ (\mathbb{E}|X|^p)^{1/p} + (\mathbb{E}|Y|^p)^{1/p} \right\} (\mathbb{E}|X + Y|^{(p-1)q})^{1/q}$$

$$\mathbb{E}|X + Y|^p \leq \left\{ \|X\|_p + \|Y\|_p \right\} (\mathbb{E}|X + Y|^p)^{1/q}$$

$$(\mathbb{E}|X + Y|^p)^{1/p} \leq \|X\|_p + \|Y\|_p$$

$$\|X + Y\|_p \leq \|X\|_p + \|Y\|_p$$

□

3.3 A quick refresher of linear algebra.

Many of the concepts in this chapter have very elegant interpretations, if we think of real-valued random variables on a probability space as vectors in a vector space. In particular, variance is related to the concept of norm and distance, while covariance is related to inner-products. These concepts can help unify some of the ideas in this chapter from a geometric point of view. Of course, real-valued random variables are simply measurable, real-valued functions on the abstract space Ω .

Definition 3.4. (Vector Space).

By a vector space, we mean a non-empty set V with two operations:

- Vector addition: $+: (\mathbf{x}, \mathbf{y}) \rightarrow \mathbf{x} + \mathbf{y}$
- Scalar multiplication: $\cdot : (\alpha, \mathbf{x}) \rightarrow \alpha \mathbf{x}$

such that the following conditions are satisfied:

(A1) Commutativity. $\mathbf{x} + \mathbf{y} = \mathbf{y} + \mathbf{x}$ for all $\mathbf{x}, \mathbf{y} \in V$

(A2) Associativity: $(\mathbf{x} + \mathbf{y}) + \mathbf{z} = \mathbf{x} + (\mathbf{y} + \mathbf{z})$ for all $\mathbf{x}, \mathbf{y}, \mathbf{z} \in V$

(A3) Zero Element: There exists a zero element, denoted $\mathbf{0}$ in V , for all $\mathbf{x} \in V$, such that $\mathbf{x} + \mathbf{0} = \mathbf{x}$.

(A4) Additive Inverse: For all $\mathbf{x} \in V$, there exists an additive inverse (negative element) denoted $-\mathbf{x}$ in V , such that $\mathbf{x} + (-\mathbf{x}) = \mathbf{0}$.

(M1) Scalar multiplication by identity element in F : For all $\mathbf{x} \in V$, $1 \cdot \mathbf{x} = \mathbf{x}$, where 1 denotes the multiplicative identity in F .

(M2) Scalar multiplication and field multiplication mix well: For all $\alpha, \beta \in F$ and $\mathbf{v} \in V$, $(\alpha\beta)\mathbf{v} = \alpha(\beta\mathbf{v})$.

(D1) Distribution of scalar multiplication over vector addition: For all $\alpha \in F$, and $\mathbf{u}, \mathbf{v} \in V$, $\alpha(\mathbf{u} + \mathbf{v}) = \alpha\mathbf{u} + \alpha\mathbf{v}$.

(D2) Distribution of field addition over scalar multiplication: For all $\alpha, \beta \in F$, and $\mathbf{v} \in V$, $(\alpha + \beta)\mathbf{v} = \alpha\mathbf{v} + \beta\mathbf{v}$.

As usual, our starting point is a random experiment modeled by a probability space $(\Omega, \mathcal{F}, \mathbb{P})$, so that Ω is the set of outcomes, \mathcal{F} is the σ -algebra of events and \mathbb{P} is the probability measure on the measurable space (Ω, \mathcal{F}) . Our basic vector space V consists of all real-valued random variables defined on $(\Omega, \mathcal{F}, \mathbb{P})$. We define vector addition and scalar multiplication in the usual way point-wise.

- Vector addition: $(X + Y)(\omega) = X(\omega) + Y(\omega)$.
- Scalar multiplication: $(\alpha X)(\omega) = \alpha X(\omega)$

Clearly, any function g of a random variable $X(\omega)$ is also a random variable on the same probability space and any linear combination of random variables on $(\Omega, \mathcal{F}, \mathbb{P})$ also define a new random variable on the same probability space. Thus, V is closed under vector addition and scalar-multiplication. Since vector-addition and scalar multiplication is defined point-wise, it is easy to see that - all the axioms of a vector space (A1)-(A4), (M1-M2), (D1), (D2) are satisfied. The constantly zero random variable $0(\omega) = 0$ and the indicator random variable $I_\Omega(\omega)$ can be thought of as the zero and identity vectors in this vector space.

3.3.1 Inner Products.

In Euclidean geometry, the angle between two vectors is specified by their dot product, which is itself formalized by the abstract concept of inner products.

Definition 3.5. (Inner Product). An inner product on the real vector space V is a pairing that takes two vector $\mathbf{v}, \mathbf{w} \in V$ and produces a real number $\langle \mathbf{v}, \mathbf{w} \rangle \in \mathbf{R}$. The inner product is required to satisfy the following three axioms for all $\mathbf{u}, \mathbf{v}, \mathbf{w} \in V$ and scalars $c, d \in \mathbf{R}$.

(i) Bilinearity:

$$\langle c\mathbf{u} + d\mathbf{v}, \mathbf{w} \rangle = c \langle \mathbf{u}, \mathbf{w} \rangle + d \langle \mathbf{v}, \mathbf{w} \rangle \quad (3.14)$$

$$\langle \mathbf{u}, c\mathbf{v} + d\mathbf{w} \rangle = c \langle \mathbf{u}, \mathbf{v} \rangle + d \langle \mathbf{u}, \mathbf{w} \rangle \quad (3.15)$$

(ii) Symmetry:

$$\langle \mathbf{v}, \mathbf{w} \rangle = \langle \mathbf{w}, \mathbf{v} \rangle \quad (3.16)$$

(iii) Positive Definiteness:

$$\langle \mathbf{v}, \mathbf{v} \rangle > 0 \quad \text{whenever } \mathbf{v} \neq \mathbf{0} \quad (3.17)$$

$$\langle \mathbf{v}, \mathbf{v} \rangle = 0 \quad \text{whenever } \mathbf{v} = \mathbf{0} \quad (3.18)$$

Definition 3.6. (Norm). A norm on a real vector space V is a function $\|\cdot\| : V \rightarrow \mathbf{R}$ satisfying :

(i) Positive Definiteness.

$$\|\mathbf{v}\| \geq 0 \quad (3.19)$$

and

$$\|\mathbf{v}\| = 0 \quad \text{if and only if } \mathbf{v} = \mathbf{0} \quad (3.20)$$

(ii) Scalar multiplication.

$$\|\alpha \mathbf{v}\| = |\alpha| \|\mathbf{v}\| \quad (3.21)$$

(iii) Triangle Inequality.

$$\|\mathbf{x} + \mathbf{y}\| \leq \|\mathbf{x}\| + \|\mathbf{y}\| \quad (3.22)$$

As mentioned earlier, we can define the p -norm of a random variable as :

$$\|X\|_p = (\mathbb{E}|X|^p)^{1/p}$$

(i) Positive semi-definiteness: Since $|X|$ is a non-negative random variable, $|X|^p \geq 0$ and the expectation of a non-negative random variable is also non-negative. Hence, $(\mathbb{E}|X|^p)^{1/p} \geq 0$. Moreover, $\|X\|_p = 0$ implies that $\mathbb{E}|X|^p = 0$. From property (iv) of expectations, $X = 0$.

(ii) Scalar-multiplication: We have:

$$\begin{aligned} \|cX\|_p &= (\mathbb{E}|cX|^p)^{1/p} \\ &= (|c|^p)^{1/p} (\mathbb{E}|X|^p)^{1/p} \\ &= |c| \cdot \|X\|_p \end{aligned}$$

(iii) Triangle Inequality. This followed from the Minkowski's inequality.

The space of all random variables defined on $(\Omega, \mathcal{F}, \mathbb{P})$ such that $\|X\|_p < \infty$ is finite is called the L^p space.

3.3.2 Orthogonal Matrices.

Definition 3.7. (Orthogonal Matrix). Let A be an $n \times n$ square matrix. We say that the matrix A is orthogonal, if its transpose is equal to its inverse.

$$A' = A^{-1}$$

This may seem like an odd property to study, but the following theorem explains why it is so useful. Essentially, an orthogonal matrix rotates (or reflects) vectors without distorting angles or distances.

Proposition 3.1. For an $n \times n$ square matrix A , the following are equivalent:

- (1) A is orthogonal. That is, $A'A = I$.
- (2) A preserves norms. That is, for all \mathbf{x} ,

$$\|A\mathbf{x}\| = \|\mathbf{x}\|$$

- (3) A preserves inner products, that is, for every $\mathbf{x}, \mathbf{y} \in \mathbf{R}^n$:

$$(A\mathbf{x}) \cdot (A\mathbf{y}) = \mathbf{x} \cdot \mathbf{y}$$

Proof. We have:

$$\begin{aligned}\|A\mathbf{x}\|^2 &= (A\mathbf{x})' (A\mathbf{x}) \\ &= \mathbf{x}' (A'A) \mathbf{x} \\ &= \mathbf{x}' I \mathbf{x} \\ &= \mathbf{x}' \mathbf{x} \\ &= \|\mathbf{x}\|^2\end{aligned}$$

Consequently, $\|A\mathbf{x}\| = \|\mathbf{x}\|$. The matrix A preserves norms. Thus, (1) implies (2). Moreover, consider

$$\begin{aligned}
\|A(\mathbf{x} + \mathbf{y})\|^2 &= (A\mathbf{x} + A\mathbf{y}) \cdot (A\mathbf{x} + A\mathbf{y}) \\
&= (A\mathbf{x}) \cdot (A\mathbf{x}) + (A\mathbf{x}) \cdot (A\mathbf{y}) + (A\mathbf{y}) \cdot (A\mathbf{x}) + (A\mathbf{y}) \cdot (A\mathbf{y}) \\
&= \|A\mathbf{x}\|^2 + 2(A\mathbf{x}) \cdot (A\mathbf{y}) + \|A\mathbf{y}\|^2 && \{\mathbf{x} \cdot \mathbf{y} = \mathbf{y} \cdot \mathbf{x}\} \\
&= \|\mathbf{x}\|^2 + 2(A\mathbf{x}) \cdot (A\mathbf{y}) + \|\mathbf{y}\|^2 && \{A \text{ preserves norms}\}
\end{aligned}$$

But, $\|A(\mathbf{x} + \mathbf{y})\|^2 = \|\mathbf{x} + \mathbf{y}\|^2 = \|\mathbf{x}\|^2 + 2\mathbf{x} \cdot \mathbf{y} + \|\mathbf{y}\|^2$. Equating the two expressions, we have the desired result. Hence, (2) implies (3).

Lastly, if A preserves inner products, we may write:

$$\begin{aligned}
\langle A\mathbf{x}, A\mathbf{x} \rangle &= \langle \mathbf{x}, \mathbf{x} \rangle \\
(A\mathbf{x})' (A\mathbf{x}) &= \mathbf{x}' \mathbf{x} \\
\mathbf{x}' A' A \mathbf{x} &= 0
\end{aligned}$$

Since $\mathbf{x} \neq \mathbf{0}$, it must be true that $\mathbf{x}' A' A - \mathbf{x}' = 0$. Again, since $\mathbf{x}' \neq \mathbf{0}$, it follows that $A' A - I = 0$. \square

Theorem 3.9. *If $\mathbf{q}_1, \mathbf{q}_2, \dots, \mathbf{q}_k \in V$ be mutually orthogonal elements, such that $\mathbf{q}_i \neq \mathbf{0}$ for all i , then $\mathbf{q}_1, \mathbf{q}_2, \dots, \mathbf{q}_k$ are linearly independent.*

Proof. Let

$$c_1 \mathbf{q}_1 + c_2 \mathbf{q}_2 + \dots + c_k \mathbf{q}_k = \mathbf{0}$$

Since $\langle \mathbf{q}_i, \mathbf{q}_i \rangle = 1$ and $\langle \mathbf{q}_i, \mathbf{q}_j \rangle = 0$ where $i \neq j$, we can take the inner product of the vector $(c_1 \mathbf{q}_1 + c_2 \mathbf{q}_2 + \dots + c_i \mathbf{q}_i + \dots + c_k \mathbf{q}_k)$ with \mathbf{q}_i for each $i = 1, 2, 3, \dots, k$. It results in $c_i \|\mathbf{q}_i\|^2 = 0$. Since $\mathbf{q}_i \neq \mathbf{0}$, $\|\mathbf{q}_i\|^2 > 0$. So, $c_i = 0$. We conclude that $c_1 = c_2 = \dots = c_k = 0$. Consequently, $\mathbf{q}_1, \mathbf{q}_2, \dots, \mathbf{q}_k$ are linearly independent. \square

Theorem 3.10. *Let $Q = \begin{bmatrix} \mathbf{q}_1 & \mathbf{q}_2 & \dots & \mathbf{q}_n \end{bmatrix}$ be an $n \times n$ orthogonal matrix. Then, $\{\mathbf{q}_1, \dots, \mathbf{q}_n\}$ form an orthonormal basis for \mathbf{R}^n .*

Proof. We have $Q\mathbf{e}_i = \mathbf{q}_i$. Consequently,

$$\begin{aligned}
 \langle \mathbf{q}_i, \mathbf{q}_i \rangle &= \mathbf{q}_i' \mathbf{q}_i \\
 &= (Q\mathbf{e}_i)'(Q\mathbf{e}_i) \\
 &= \mathbf{e}_i' Q' Q \mathbf{e}_i \\
 &= \mathbf{e}_i' I \mathbf{e}_i \\
 &= \mathbf{e}_i' \mathbf{e}_i \\
 &= 1
 \end{aligned}$$

Assume that $i \neq j$. We have:

$$\begin{aligned}
 \langle \mathbf{q}_i, \mathbf{q}_j \rangle &= \mathbf{q}_i' \mathbf{q}_j \\
 &= \mathbf{e}_i' Q' Q \mathbf{e}_j \\
 &= \mathbf{e}_i' \mathbf{e}_j \\
 &= 0
 \end{aligned}$$

From theorem (3.9), $\{\mathbf{q}_1, \dots, \mathbf{q}_n\}$ are linearly independent and hence form an orthonormal basis for \mathbf{R}^n . \square

3.3.3 Quadratic Forms.

An expression of the form:

$$\mathbf{x}' A \mathbf{x}$$

where \mathbf{x} is a $n \times 1$ column vector and A is an $n \times n$ matrix is called a quadratic form in \mathbf{x} and

$$\mathbf{x}' A \mathbf{x} = \sum_{i=1}^n \sum_{j=1}^n a_{ij} x_i x_j$$

If A and B are $n \times n$ and \mathbf{x}, \mathbf{y} are n -vectors, then

$$\mathbf{x}'(A + B)\mathbf{y} = \mathbf{x}'A\mathbf{y} + \mathbf{x}'B\mathbf{y}$$

The quadratic form of the matrix A is called positive definite if:

$$\mathbf{x}'A\mathbf{x} > 0 \quad \text{whenever } \mathbf{x} \neq \mathbf{0}$$

and positive semidefinite if:

$$\mathbf{x}'A\mathbf{x} \geq 0 \quad \text{whenever } \mathbf{x} \neq \mathbf{0}$$

Letting \mathbf{e}_i be the unit vector with its i th coordinate vector 1, we have:

$$\mathbf{e}_i' A \mathbf{e}_i = [a_{i1} a_{i2} \dots a_{ii} \dots a_{ni}] \begin{bmatrix} 0 \\ 0 \\ \vdots \\ 1 \\ \vdots \\ 0 \end{bmatrix} = a_{ii}$$

3.3.4 Eigenthings and diagonalizability.

Let V and W be finite dimensional vector spaces with $\dim(V) = n$ and $\dim(W) = m$. A linear transformation $T : V \rightarrow W$, is defined by its action on the basis vectors. Suppose:

$$T(\mathbf{v}_i) = \sum_{j=1}^n a_{ij} \mathbf{w}_j$$

for all $1 \leq i \leq m$.

Then, the matrix $A = [T]_{\mathcal{B}_V}^{\mathcal{B}_W}$ of the linear transformation is defined as:

$$A = \begin{bmatrix} a_{11} & a_{12} & \dots & a_{1n} \\ a_{21} & a_{22} & \dots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \dots & a_{mn} \end{bmatrix}$$

Definition 3.8. A linear transformation $T : V \rightarrow V$ is **diagonalizable** if there exists an ordered basis $\mathcal{B} = \{\mathbf{v}_1, \dots, \mathbf{v}_n\}$ for V so that the matrix for T with respect to \mathcal{B} is diagonal. This means precisely that, for some scalars $\lambda_1, \lambda_2, \dots, \lambda_n$, we have:

$$\begin{aligned} T(\mathbf{v}_1) &= \lambda_1 \mathbf{v}_1 \\ T(\mathbf{v}_2) &= \lambda_2 \mathbf{v}_2 \\ &\vdots \\ T(\mathbf{v}_n) &= \lambda_n \mathbf{v}_n \end{aligned}$$

In other words, if $A = [T]_{\mathcal{B}}$, then we have:

$$A\mathbf{v}_i = \lambda_i \mathbf{v}_i$$

Thus, if we let P be the $n \times n$ matrix whose columns are the vectors $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_n$ and Λ be the $n \times n$ diagonal matrix with diagonal entries $\lambda_1, \lambda_2, \dots, \lambda_n$, then we have:

$$A \begin{bmatrix} \mathbf{v}_1 & \mathbf{v}_2 & \dots & \mathbf{v}_n \end{bmatrix} = \begin{bmatrix} \mathbf{v}_1 & \mathbf{v}_2 & \dots & \mathbf{v}_n \end{bmatrix} \begin{bmatrix} \lambda_1 & & & \\ & \lambda_2 & & \\ & & \ddots & \\ & & & \lambda_n \end{bmatrix}$$

$$AP = P\Lambda$$

$$A = P\Lambda P^{-1}$$

There exists a large class of diagonalizable matrices - the symmetric matrices. A square matrix A is symmetric, if $A = A'$.

Definition 3.9. Let $T : V \rightarrow V$ be a linear transformation. A **non-zero** vector $\mathbf{v} \in V$ is called the eigenvector of T , if there is a scalar λ so that $T(\mathbf{v}) = \lambda\mathbf{v}$. The scalar λ is called the eigenvalue of T .

Lemma 3.1. Let A be an $n \times n$ matrix, and let λ be any scalar. Then,

$$E(\lambda) = \{\mathbf{x} \in \mathbf{R}^n : A\mathbf{x} = \lambda\mathbf{x}\} = \ker(A - \lambda I)$$

is a subspace of \mathbf{R}^n . Moreover, if $E(\lambda) \neq \{\mathbf{0}\}$ if and only if λ is an eigenvalue, in which case we call $E(\lambda)$ the λ -eigenspace of the matrix A .

Proof. We know that, $E(\lambda)$ is a subset of \mathbf{R}^n . Moreover, if $\mathbf{u}, \mathbf{v} \in E(\lambda)$, then $A(c_1\mathbf{u} + c_2\mathbf{v}) = c_1A\mathbf{u} + c_2A\mathbf{v} = \lambda(c_1\mathbf{u} + c_2\mathbf{v})$. Consequently, $c_1\mathbf{u} + c_2\mathbf{v} \in E(\lambda)$. Thus, $E(\lambda)$ is a subspace of \mathbf{R}^n .

Moreover, by definition, λ is an eigenvalue of A precisely when $\mathbf{x} \neq \mathbf{0}$ vector in $E(\lambda)$. This closes the proof. \square

Theorem 3.11. Let A be a $n \times n$ square matrix. If A is a singular matrix, then $\det A = 0$.

Proof. By definition, a square matrix is said to be non-singular, if it can be reduced to an upper triangular form with all non-zero elements on the diagonal - the pivots, by elementary row operations. A singular matrix is such that its echelon form has a row of zeroes, and its row vectors are linearly dependent and $\det A = 0$. \square

Theorem 3.12. Let A be a $n \times n$ square matrix. Then, λ is an eigenvalue of A if and only if $\det(A - \lambda I) = 0$.

Proof. λ is an eigenvalue of A , if and only, the homogenous system of linear equations $(A - \lambda I)\mathbf{x} = \mathbf{0}$ has non-trivial solutions. Consequently, the only possibility is that there are one more free variables (more variables than the number of equations). In other words, $(A - \lambda I)$ must be a singular matrix and $\det(A - \lambda I) = 0$. \square

Example 3.2. Let's find the eigenvalues and eigenvectors of the matrix

$$A = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 1 & 0 \\ 1 & 3 & 1 \end{bmatrix}$$

We begin by computing

$$\begin{aligned}
 \det(A - \lambda I) &= \begin{vmatrix} 1 - \lambda & 2 & 1 \\ 0 & 1 - \lambda & 0 \\ 1 & 2 & 1 - \lambda \end{vmatrix} \\
 &= (1 - \lambda)(1 - \lambda)^2 - (1 - \lambda) \\
 &= (1 - \lambda)[(1 - \lambda)^2 - 1] \\
 &= -\lambda(1 - \lambda)(2 - \lambda)
 \end{aligned}$$

Thus, the eigenvalues of A are $\lambda = 0$, $\lambda = 1$ and $\lambda = 2$.

We find the respective eigenspaces:

1) Fix $\lambda = 0$. We see that:

$$\begin{bmatrix} 1 & 2 & 1 \\ 0 & 1 & 0 \\ 1 & 3 & 1 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}$$

The augmented matrix $[A|b]$ is :

$$\left[\begin{array}{ccc|c} 1 & 2 & 1 & 0 \\ 0 & 1 & 0 & 0 \\ 1 & 3 & 1 & 0 \end{array} \right]$$

$R_3 - R_1$, $R_3 - R_2$ and $R_1 - 2R_2$ leaves us with:

$$\left[\begin{array}{ccc|c} 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 \end{array} \right]$$

So, $x_1 + x_3 = 0$ and $x_2 = 0$. Here, x_3 is a free variable. Thus,

$$E(0) = \{\alpha(1, 0, -1) | \alpha \in \mathbf{R}\}$$

2) Fix $\lambda = 1$. We see that:

$$\begin{bmatrix} 0 & 2 & 1 \\ 0 & 0 & 0 \\ 1 & 3 & 0 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}$$

Thus, $2x_2 + x_3 = 0$ and $x_1 + 3x_2 = 0$. Here x_3 is a free variable. Let $x_3 = -2\alpha$. Then, $x_2 = \alpha$ and $x_1 = -3\alpha$. Consequently,

$$E(1) = \{\alpha(-3, 1, -2) | \alpha \in \mathbf{R}\}$$

3) Fix $\lambda = 3$. We see that:

$$\begin{bmatrix} -1 & 2 & 1 \\ 0 & -1 & 0 \\ 1 & 3 & -1 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}$$

The augmented matrix $[A|b]$ is :

$$\left[\begin{array}{ccc|c} -1 & 2 & 1 & 0 \\ 0 & -1 & 0 & 0 \\ 1 & 3 & -1 & 0 \end{array} \right]$$

$R_3 + R_1$, $R_3 + 5R_2$ followed by $R_1 + 2R_2$ gives:

$$\left[\begin{array}{ccc|c} -1 & 0 & 1 & 0 \\ 0 & -1 & 0 & 0 \\ 0 & 0 & 0 & 0 \end{array} \right]$$

Thus, $x_2 = 0$ and $x_1 - x_3 = 0$. Here x_3 is the free variable. Hence,

$$E(2) = \{\alpha(1, 0, 1) : \alpha \in \mathbf{R}\}$$

Clearly, there exists a basis $\mathcal{B} = \{(1, 0, -1), (-3, 1, -2), (1, 0, 1)\}$ with respect to which the matrix of T is diagonal. Hence, A is diagonalizable.

Judging from the previous example, it appears that when an $n \times n$ square matrix has n distinct eigen values, the corresponding eigenvectors form a linearly independent set and will therefore give a *diagonalizing basis*. Let's begin with a slightly stronger statement.

Theorem 3.13. *Let $T : V \rightarrow V$ be a linear transformation. Suppose $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_k$ are eigenvectors of T corresponding to the distinct eigenvalues $\lambda_1, \lambda_2, \dots, \lambda_k$. Then, $\{\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_k\}$ is a linearly independent set of vectors.*

Proof. Let m be the largest number between 1 and k (inclusive) so that $\{\mathbf{v}_1, \dots, \mathbf{v}_m\}$ is linearly independent. We proceed by contradiction. We want to see $m = k$. Assume that $m < k$. Then, we know that $\{\mathbf{v}_1, \dots, \mathbf{v}_m\}$ is linearly independent and $\{\mathbf{v}_1, \dots, \mathbf{v}_m, \mathbf{v}_{m+1}\}$ is linearly dependent. Thus, $\mathbf{v}_{m+1} = c_1\mathbf{v}_1 + c_2\mathbf{v}_2 + \dots + c_m\mathbf{v}_m$ such that atleast one of c_1, c_2, \dots, c_m are non-zero. Then, using repeatedly the fact that $T(\mathbf{v}_i) = \lambda_i\mathbf{v}_i$:

$$\begin{aligned} \mathbf{0} &= (T - \lambda_{m+1}I)\mathbf{v}_{m+1} = (T - \lambda_{m+1}I)(c_1\mathbf{v}_1 + \dots + c_m\mathbf{v}_m) \\ &= c_1(T\mathbf{v}_1 - \lambda_{m+1}I\mathbf{v}_1) + c_2(T\mathbf{v}_2 - \lambda_{m+1}I\mathbf{v}_2) + \dots + c_m(T\mathbf{v}_m - \lambda_{m+1}I\mathbf{v}_m) \\ &= c_1(\lambda_1 - \lambda_{m+1})\mathbf{v}_1 + c_2(\lambda_2 - \lambda_{m+1})\mathbf{v}_2 + \dots + c_m(\lambda_m - \lambda_{m+1})\mathbf{v}_m \end{aligned}$$

Since $\lambda_i \neq \lambda_{m+1}$ for $i = 1, 2, 3, \dots, m$ and since $\{\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_m\}$ is linearly independent, the only other possibility is $c_1 = c_2 = \dots = c_m = 0$. But, this contradicts the fact that \mathbf{v}_{m+1} is an eigenvector and $\mathbf{v}_{m+1} \neq \mathbf{0}$. Thus, it cannot happen that $m < k$. Consequently, $m = k$. \square

What is underlying this formal argument is the observation that: if $\mathbf{v} \in E(\lambda) \cap E(\mu)$, then $T\mathbf{v} = \lambda\mathbf{v}$ and $T\mathbf{v} = \mu\mathbf{v}$. Hence, if $\lambda \neq \mu$, then $\mathbf{v} = \mathbf{0}$. That is, if $\lambda \neq \mu$, we have $E(\lambda) \cap E(\mu) = \{\mathbf{0}\}$.

Corollary 3.1. *Suppose V is an n -dimensional vector space and $T : V \rightarrow V$ has n distinct eigenvalues. Then T is diagonalizable.*

Proof. The set of n corresponding eigenvectors must be linearly independent and hence form a basis for V . The matrix of T with respect to the eigenbasis is always diagonal. \square

The converse of this statement is not true. There are many diagonalizable matrices with repeated eigen-values.

Definition 3.10. Let λ be an eigenvalue of a linear transformation. The algebraic multiplicity of λ is its multiplicity as a root of the characteristic polynomial $p(t)$ that is, the highest power of $t - \lambda$ dividing $p(t)$. The geometric multiplicity of λ is the dimension of the eigenspace $E(\lambda)$.

Proposition 3.2. Let λ be an eigenvalue of algebraic multiplicity m and geometric multiplicity d . Then, the geometric multiplicity is always bounded by the algebraic multiplicity, and $1 \leq d \leq m$.

Proof. Suppose λ is the eigenvalue of the linear transformation T . Then, $d = \dim E(\lambda) \geq 1$ by definition. Now, choose a basis $\{\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_d\}$ for $E(\lambda)$ and extend it to a basis $\mathcal{B} = \{\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_n\}$ for V . Then, the matrix of T with respect to \mathcal{B} is of the form

$$A = \begin{bmatrix} \lambda I_d & B \\ 0_{(n-d) \times d} & C \end{bmatrix}$$

The characteristic polynomial $p(t)$ of the matrix A is given by:

$$\begin{aligned} p(t) &= \det(A - tI) \\ &= \det((\lambda - t)I_d) \cdot \det(C - tI) \\ &= (\lambda - t)^d \cdot \det(C - tI) \end{aligned}$$

Since the characteristic polynomial does not depend on the choice of basis, the algebraic multiplicity of λ is at least d . \square

Lemma 3.2. (*Lagrange Multipliers*) Suppose $f, g : \mathbf{R}^n \rightarrow \mathbf{R}$ are scalar-valued C^1 functions - that is partial derivatives ∂_{x_i} in all variables are continuous. Let $S = \{\mathbf{x} \in \mathbf{R}^n | g(\mathbf{x}) = c\}$ denote the level set of g at height c . Then if $f|_S$ (the restriction of f to S) has an extremum point \mathbf{x}_0 in S such that $\nabla g(\mathbf{x}_0) \neq \mathbf{0}$, there exists a scalar λ such that

$$\nabla f(\mathbf{x}_0) = \lambda \nabla g(\mathbf{x}_0) \tag{3.23}$$

Proof. Let's visualize the situation for the case $n = 3$, where the constraint equation $g(x, y, z) = c$ defines a surface S in \mathbf{R}^3 .

Thus, suppose that \mathbf{x}_0 is an extremum of f restricted to S . We consider a further restriction of f - to a curve lying in S and passing through \mathbf{x}_0 . Let $\mathbf{x}(t) = (x(t), y(t), z(t))$ be the parametric equation of one such arbitrary path $\mathbf{x} : I \subseteq \mathbf{R} \rightarrow \mathbf{R}^3$ lying in S with $\mathbf{x}(t_0) = \mathbf{x}_0$ for some $t_0 \in I$. Then, the restriction of f to \mathbf{x} can be written as a function of a single variable t . That is:

$$F(t) := f(\mathbf{x}(t))$$

Because \mathbf{x}_0 is an extremum of f on the whole of S , it is also an extremum on the path \mathbf{x} . Since F is a differentiable function of t , by the interior-extremum theorem, it follows that $F'(t_0) = 0$. The chain rule implies that:

$$F'(t) = \nabla f(\mathbf{x}) \cdot \mathbf{x}'(t)$$

Evaluating at $t = t_0$, we have:

$$F'(t_0) = 0 = \nabla f(\mathbf{x}(t_0)) \cdot \mathbf{x}'(t_0)$$

Thus, $\nabla f(\mathbf{x}(t_0))$ is perpendicular to any curve in S passing through \mathbf{x}_0 ; that is $\nabla f(\mathbf{x}_0)$ is normal to S at \mathbf{x}_0 . We've already seen previously that the gradient vector $\nabla g(\mathbf{x}_0)$ is also normal to S at \mathbf{x}_0 . Since the normal direction to the level S is uniquely determined, we must conclude that $\nabla f(\mathbf{x}_0)$ and $\nabla g(\mathbf{x}_0)$ are parallel vectors. Therefore, there exists a scalar λ such that:

$$\nabla f(\mathbf{x}_0) = \lambda \nabla g(\mathbf{x}_0)$$

□

3.3.5 The Gram-Schmidt Process.

The advantage of using an orthonormal basis is, that the coordinates of any vector are explicitly given as inner products. Let $\{\mathbf{u}_1, \mathbf{u}_2, \dots, \mathbf{u}_n\}$ be an orthonormal basis of \mathbf{R}^n . And let $\mathbf{v} = c_1\mathbf{u}_1 + \dots + c_n\mathbf{u}_n$ be an arbitrary vector. Then we have:

$$c_i = \mathbf{v} \cdot \mathbf{u}_i$$

Moreover, the magnitude (norm) of the vector is given by the Pythagorean formula:

$$\begin{aligned}\|\mathbf{v}\|_2^2 &= \langle \mathbf{v}, \mathbf{v} \rangle \\ &= c_1^2 + c_2^2 + \dots + c_n^2\end{aligned}$$

Once we are convinced of the utility of orthogonal and orthonormal bases, a natural question arises: how can we construct them? A practical algorithm was discovered Pierre-Simon Laplace in the eighteenth century. Today, the algorithm is known as the *Gram-Schmidt process*, after its rediscovery by Gram and twentieth century mathematician Schmidt.

Let W be a finite dimensional vector space, such that $\dim W = n$. We assume that, we already know some basis $\{\mathbf{w}_1, \dots, \mathbf{w}_n\}$ of W , where $n = \dim W$. Our goal is to use this information to construct an orthogonal basis $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_n$.

We will construct the orthogonal basis one-by-one. Since initially, we are not worrying about normality, there are no conditions on the first orthogonal basis element \mathbf{v}_1 , so there is no harm in choosing :

$$\mathbf{v}_1 = \mathbf{w}_1$$

Note that, $\mathbf{v}_1 \neq \mathbf{0}$, since \mathbf{w}_1 appears in the original basis. Starting with \mathbf{w}_2 , the second basis vector \mathbf{v}_2 must be orthogonal to the first: $\langle \mathbf{v}_2, \mathbf{v}_1 \rangle = 0$.




Figure. Resolving the vector \mathbf{w}_2 into two components (1) along \mathbf{u}_1 and (2) perpendicular to \mathbf{u}_1 .

Let us try to arrange this, by subtracting a suitable multiple of \mathbf{v}_1 , and set:

$$\mathbf{v}_2 = \mathbf{w}_2 - c\mathbf{v}_1$$

The orthogonality condition

$$\begin{aligned} 0 &= \langle \mathbf{v}_2, \mathbf{v}_1 \rangle \\ &= (\mathbf{w}_2 - c\mathbf{v}_1) \cdot \mathbf{v}_1 \\ &= \mathbf{w}_2 \cdot \mathbf{v}_1 - c \|\mathbf{v}_1\|^2 \\ c &= \frac{\mathbf{w}_2 \cdot \mathbf{v}_1}{\|\mathbf{v}_1\|^2} \end{aligned}$$

and therefore

$$\mathbf{v}_2 = \mathbf{w}_2 - \left(\frac{\mathbf{w}_2 \cdot \mathbf{v}_1}{\|\mathbf{v}_1\|^2} \right) \mathbf{v}_1$$

The linear independence of $\mathbf{v}_1 = \mathbf{w}_1$ and \mathbf{w}_2 ensures that $\mathbf{v}_2 \neq \mathbf{0}$.

Next, we construct:

$$\mathbf{v}_3 = \mathbf{w}_3 - c_1\mathbf{v}_1 - c_2\mathbf{v}_2$$

by subtracting suitable multiples of the first two orthogonal basis elements from \mathbf{w}_3 . We want \mathbf{v}_3 to be orthogonal to both \mathbf{v}_1 and \mathbf{v}_2 . Since we already arranged that $\mathbf{v}_1 \cdot \mathbf{v}_2 = 0$, this requires:

$$\begin{aligned} 0 &= \mathbf{v}_3 \cdot \mathbf{v}_1 = (\mathbf{w}_3 \cdot \mathbf{v}_1) - c_1 \|\mathbf{v}_1\|^2 \\ 0 &= \mathbf{v}_3 \cdot \mathbf{v}_2 = (\mathbf{w}_3 \cdot \mathbf{v}_2) - c_2 \|\mathbf{v}_2\|^2 \end{aligned}$$

And hence:

$$c_1 = \frac{\mathbf{w}_3 \cdot \mathbf{v}_1}{\|\mathbf{v}_1\|^2}$$

$$c_2 = \frac{\mathbf{w}_3 \cdot \mathbf{v}_2}{\|\mathbf{v}_2\|^2}$$

Therefore the next orthogonal basis vector isw given by the formula:

$$\mathbf{v}_3 = \mathbf{w}_3 - \frac{\mathbf{w}_3 \cdot \mathbf{v}_1}{\|\mathbf{v}_1\|^2} \mathbf{v}_1 - \frac{\mathbf{w}_3 \cdot \mathbf{v}_2}{\|\mathbf{v}_2\|^2} \mathbf{v}_2$$

Since \mathbf{v}_1 and \mathbf{v}_2 are linear combinations of \mathbf{w}_1 and \mathbf{w}_2 , we must have that $\mathbf{v}_3 \neq \mathbf{0}$, since otherwise this would imply that \mathbf{w}_3 can be written as a linear combination of \mathbf{w}_1 and \mathbf{w}_2 making them linearly dependent.

Continuing in the same manner, suppose we have already constructed the mutually orthogonal vectors $\mathbf{v}_1, \dots, \mathbf{v}_{k-1}$ as linear combinations of $\mathbf{w}_1, \dots, \mathbf{w}_{k-1}$. The next orthogonal basis element \mathbf{v}_k will be obtained from \mathbf{w}_k by subtracting a suitable linear combination of the previous orthogonal basis elements. In this fashion we establish the general *Gram-Schmidt* formula -

$$\mathbf{v}_k = \mathbf{w}_k - \sum_{j=1}^{k-1} \frac{\mathbf{w}_k \cdot \mathbf{v}_j}{\|\mathbf{v}_j\|^2} \mathbf{v}_j \quad (3.24)$$

If we are after an orthonormal basis $\mathbf{u}_1, \dots, \mathbf{u}_n$ we merely normalize the resulting orthogonal basis vectors, setting $\mathbf{u}_k = \frac{\mathbf{v}_k}{\|\mathbf{v}_k\|}$.

3.3.6 Modifications of the Gram-Schmidt process.

With the basic Gram-Schmidt algorithm now in hand, it is worth looking at a couple of reformulations that have both practical and theoretical advantages. The first can be used to construct orthonormal basis vectors $\mathbf{u}_1, \mathbf{u}_2, \dots, \mathbf{u}_n$ directly from the basis $\mathbf{w}_1, \dots, \mathbf{w}_n$.

We begin by replacing each orthogonal basis vector in the basic Gram-Schmidt formula (3.24) by its normalized version $\mathbf{u}_j = \mathbf{v}_j / \|\mathbf{v}_j\|$. The original basis vectors can be expressed in terms of the orthonormal basis via a triangular system.

$$\begin{aligned}
\mathbf{w}_1 &= r_{11}\mathbf{u}_1 \\
\mathbf{w}_2 &= r_{12}\mathbf{u}_1 + r_{22}\mathbf{u}_2 \\
\mathbf{w}_3 &= r_{13}\mathbf{u}_1 + r_{23}\mathbf{u}_2 + r_{33}\mathbf{u}_3 \\
&\vdots \\
\mathbf{w}_n &= r_{1n}\mathbf{u}_1 + r_{2n}\mathbf{u}_2 + r_{3n}\mathbf{u}_3 + \dots + r_{nn}\mathbf{u}_n
\end{aligned} \tag{3.25}$$

The coefficients r_{ij} can, in fact, be computed directly from these formulas. Indeed taking, the inner product of the equation for \mathbf{w}_j with the orthonormal basis vector \mathbf{u}_i for $i \leq j$, we obtain in view of the orthonormality constraints:

$$\begin{aligned}
\mathbf{w}_j \cdot \mathbf{u}_i &= r_{1j}\mathbf{u}_1 \cdot \mathbf{u}_i + \dots + r_{ij}\mathbf{u}_i \cdot \mathbf{u}_i + \dots + r_{jj}\mathbf{u}_j \cdot \mathbf{u}_i \\
&= r_{ij}
\end{aligned}$$

and hence:

$$r_{ij} = \langle \mathbf{w}_j, \mathbf{u}_i \rangle \tag{3.26}$$

On the other hand, we have:

$$\begin{aligned}
\|\mathbf{w}_j\|^2 &= \|r_{1j}\mathbf{u}_1 + r_{2j}\mathbf{u}_2 + \dots + r_{jj}\mathbf{u}_j\|^2 \\
&= r_{1j}^2 + r_{2j}^2 + \dots + r_{jj}^2
\end{aligned} \tag{3.27}$$

The pair of equations (3.26) and (3.27) can be rearranged to devise a recursive procedure to compute the orthonormal basis. We begin by setting $r_{11} = \|\mathbf{w}_1\|$ and so $\mathbf{u}_1 = \mathbf{w}_1/r_{11}$. At each subsequent stage, $j \geq 2$, we assume that we have already constructed $\mathbf{u}_1, \dots, \mathbf{u}_{j-1}$. We then compute

$$r_{ij} = \langle \mathbf{w}_j, \mathbf{u}_i \rangle \quad \text{for each } i = 1, 2, \dots, j-1 \tag{3.28}$$

We obtain next the orthonormal basis vector \mathbf{u}_j by computing

$$\begin{aligned}
r_{jj} &= \sqrt{\|\mathbf{w}_j\|^2 - r_{1j}^2 - r_{2j}^2 - \dots - r_{j-1,j}^2} \\
\mathbf{u}_j &= \frac{\mathbf{w}_j - r_{1j}\mathbf{u}_1 - r_{2j}\mathbf{u}_2 - \dots - r_{j-1,j}\mathbf{u}_{j-1}}{r_{jj}}
\end{aligned} \tag{3.29}$$

3.3.7 The QR Factorization.

The Gram-Schmidt procedure for orthonormalizing bases of \mathbf{R}^n can be reinterpreted as a matrix factorization.

Let $\mathbf{w}_1, \dots, \mathbf{w}_n$ be a basis of \mathbf{R}^n , and let $\mathbf{u}_1, \dots, \mathbf{u}_n$ be the corresponding orthonormal basis that results from any one of the implementations of the Gram-Schmidt process. We assemble both sets of column vectors to form non-singular $n \times n$ matrices:

$$A = [\mathbf{w}_1 \ \mathbf{w}_2 \ \dots \ \mathbf{w}_n], \quad Q = [\mathbf{u}_1 \ \mathbf{u}_2 \ \dots \ \mathbf{u}_n]$$

Since the \mathbf{u}_i form an orthonormal basis, Q is an orthogonal matrix. In view of the matrix multiplication formula, the Gram-Schmidt equations (3.25) can be recast into an equivalent matrix form:

$$A = [\mathbf{u}_1 \ \mathbf{u}_2 \ \dots \ \mathbf{u}_n] \begin{bmatrix} r_{11} & r_{12} & r_{13} & \dots & r_{1n} \\ & r_{22} & r_{23} & \dots & r_{2n} \\ & & r_{33} & \dots & r_{3n} \\ & & & \ddots & \\ & & & & r_{nn} \end{bmatrix}$$

$$= QR$$

Since the Gram-Schmidt algorithm works on any basis, the only requirement on the matrix A is that its columns are linearly-independent and form a basis of \mathbf{R}^n , and hence A can be any non-singular matrix. We have therefore established the celebrated QR -factorization of non-singular matrices.

Theorem 3.14. *Every non-singular matrix A can be factored, $A = QR$ into the product of an orthogonal matrix Q and an upper triangular matrix R .*

3.3.8 Numerically stable implementation of QR-Factorization.

We can treat all vectors simultaneously instead of sequentially and compute in the $j = 1$ st iteration:

$$\begin{aligned}
\mathbf{u}_1 &= \mathbf{w}_1 / r_{11} \\
\mathbf{w}_2^{(2)} &= \left(\mathbf{w}_2^{(1)} - \langle \mathbf{w}_2^{(1)}, \mathbf{u}_1 \rangle \mathbf{u}_1 \right) \\
\mathbf{w}_3^{(2)} &= \left(\mathbf{w}_3^{(1)} - \langle \mathbf{w}_3^{(1)}, \mathbf{u}_1 \rangle \mathbf{u}_1 \right) \\
&\vdots \\
\mathbf{w}_n^{(2)} &= \left(\mathbf{w}_n^{(1)} - \langle \mathbf{w}_n^{(1)}, \mathbf{u}_1 \rangle \mathbf{u}_1 \right)
\end{aligned}$$

Note that, the vectors $\mathbf{w}_2^{(2)}, \mathbf{w}_3^{(2)}, \dots, \mathbf{w}_n^{(2)}$ are orthogonal to \mathbf{u}_1 .

In the $j = 2$ nd iteration, we compute:

$$\begin{aligned}
\mathbf{u}_2 &= \mathbf{w}_2^{(2)} / r_{22} \\
\mathbf{w}_3^{(3)} &= \left(\mathbf{w}_3^{(2)} - \langle \mathbf{w}_3^{(2)}, \mathbf{u}_2 \rangle \mathbf{u}_2 \right) \\
\mathbf{w}_4^{(3)} &= \left(\mathbf{w}_4^{(2)} - \langle \mathbf{w}_4^{(2)}, \mathbf{u}_2 \rangle \mathbf{u}_2 \right) \\
&\vdots \\
\mathbf{w}_n^{(3)} &= \left(\mathbf{w}_n^{(2)} - \langle \mathbf{w}_n^{(2)}, \mathbf{u}_2 \rangle \mathbf{u}_2 \right)
\end{aligned}$$

Since $\mathbf{w}_2^{(2)}$ was orthogonal to \mathbf{u}_1 , \mathbf{u}_2 must also be orthogonal to \mathbf{u}_1 . Further, $\mathbf{w}_3^{(3)}, \dots, \mathbf{w}_n^{(3)}$ are orthogonal to both $\mathbf{u}_1, \mathbf{u}_2$.

In particular, in the j th iteration we compute:

$$\begin{aligned}
\mathbf{u}_j &= \mathbf{w}_j^{(j)} / r_{jj} \\
\mathbf{w}_{j+1}^{(j+1)} &= \left(\mathbf{w}_{j+1}^{(j)} - \langle \mathbf{w}_{j+1}^{(j)}, \mathbf{u}_j \rangle \mathbf{u}_j \right) \\
\mathbf{w}_{j+2}^{(j+1)} &= \left(\mathbf{w}_{j+2}^{(j)} - \langle \mathbf{w}_{j+2}^{(j)}, \mathbf{u}_j \rangle \mathbf{u}_j \right) \\
&\vdots \\
\mathbf{w}_n^{(j+1)} &= \left(\mathbf{w}_n^{(j)} - \langle \mathbf{w}_n^{(j)}, \mathbf{u}_j \rangle \mathbf{u}_j \right)
\end{aligned}$$

We can summarize the above steps as follows. We iterate $j = 1$ to n . For $j = 1$, we start with the initial basis $\mathbf{w}_k^{(1)} = \mathbf{w}_k$, and set $\mathbf{u}_1 = \mathbf{w}_1^{(1)} / r_{11}$.

In the j th iteration, we set $\mathbf{u}_j = \mathbf{w}_j^{(j)} / r_{jj}$ and for all $k = j + 1$ to n , we let $\mathbf{w}_k^{(j+1)} = \mathbf{w}_k^{(j)} - \langle \mathbf{w}_k^{(j)}, \mathbf{u}_j \rangle \mathbf{u}_j$. Also, we set $r_{jk} = \langle \mathbf{w}_k^{(j)}, \mathbf{u}_j \rangle$.

Listing 1: QR Factorization

```
#include <iostream>
#include <Eigen/Dense>
#include <cmath>

using Eigen::MatrixXd;

MatrixXd QRFactorization(MatrixXd& A)
{
    const int dimSize{ A.rows() };

    MatrixXd R(dimSize, dimSize);

    //We proceed column-wise and iteratively build the
    orthonormal vectors u_0, u_1, ..., u_{n-1}
    for (int j{ 0 }; j < dimSize; ++j)
    {
        // The scalar r_jj = ||w_j^(j)||
        for (int i{ 0 }; i < dimSize; ++i)
        {
            R(j, j) += A(i, j) * A(i, j);
        }
        R(j, j) = sqrt(R(j, j));

        // The vector u_j = w_j^(j)/r_jj
        for (int i{ 0 }; i < dimSize; ++i)
        {
            A(i, j) = A(i, j) / R(j, j);
        }

        // for all k=j+1 to n-1, this loop computes the vectors:
        // w_k^(j+1) = w_k^(j) - <w_k^(j),u_j>u_j
        for (int k{ j + 1 }; k < dimSize; ++k)
        {
            //this loop computes the inner product of the vector
            w_k with u_j
            double sum{ 0 };
            for (int i{ 0 }; i < dimSize; ++i)
```

```

        {
            sum += A(i, k) * A(i, j);
        }

        R(j, k) = sum;

        for (int i{ 0 }; i < dimSize; ++i)
        {
            A(i, k) = A(i, k) - sum * A(i, j);
        }
    }
}

```

3.3.9 Gram Matrices.

Symmetric matrices whose entries are given by the inner products of elements of an inner product space are called *Gram matrices*, after the Danish mathematician *Jorgen Gram*.

Definition 3.11. Let V be an inner product space, and let $\mathbf{v}_1, \dots, \mathbf{v}_n \in V$. The associated *Gram matrix*

$$K = \begin{bmatrix} \langle \mathbf{v}_1, \mathbf{v}_1 \rangle & \langle \mathbf{v}_1, \mathbf{v}_2 \rangle & \dots & \langle \mathbf{v}_1, \mathbf{v}_n \rangle \\ \langle \mathbf{v}_2, \mathbf{v}_1 \rangle & \langle \mathbf{v}_2, \mathbf{v}_2 \rangle & \dots & \langle \mathbf{v}_2, \mathbf{v}_n \rangle \\ \vdots & \vdots & \ddots & \vdots \\ \langle \mathbf{v}_n, \mathbf{v}_1 \rangle & \langle \mathbf{v}_n, \mathbf{v}_2 \rangle & \dots & \langle \mathbf{v}_n, \mathbf{v}_n \rangle \end{bmatrix}$$

is the $n \times n$ symmetric matrix whose entries are the inner-products between the selected vector space elements.

Theorem 3.15. *All Gram matrices are positive semi-definite.*

Proof. Let K be an arbitrary Gram matrix. To prove the positive semi-definiteness of K , we need to examine the associated quadratic form:

$$\begin{aligned} q(\mathbf{x}) &= \mathbf{x}' K \mathbf{x} \\ &= \sum_{i=1}^n \sum_{j=1}^n k_{ij} x_i x_j \end{aligned}$$

But, $k_{ij} = \langle \mathbf{v}_i, \mathbf{v}_j \rangle$. Substituting the values for the matrix entries, we obtain:

$$q(\mathbf{x}) = \sum_{i=1}^n \sum_{j=1}^n \langle \mathbf{v}_i, \mathbf{v}_j \rangle x_i x_j$$

For intuition, let's choose $n = 2$. The quadratic form becomes:

$$\begin{aligned} q(\mathbf{x}) &= \langle \mathbf{v}_1, \mathbf{v}_1 \rangle x_1^2 + \langle \mathbf{v}_1, \mathbf{v}_2 \rangle x_1 x_2 + \langle \mathbf{v}_2, \mathbf{v}_1 \rangle x_2 x_1 + \langle \mathbf{v}_2, \mathbf{v}_2 \rangle x_2^2 \\ &= \langle x_1 \mathbf{v}_1 + x_2 \mathbf{v}_2, x_1 \mathbf{v}_1 + x_2 \mathbf{v}_2 \rangle && \{\text{Bi-linearity of inner products}\} \\ &= \|x_1 \mathbf{v}_1 + x_2 \mathbf{v}_2\|^2 \end{aligned}$$

Therefore, we can write the original quadratic form as a single inner product:

$$\begin{aligned} q(\mathbf{x}) &= \left\langle \sum_{i=1}^n x_i \mathbf{v}_i, \sum_{j=1}^n x_j \mathbf{v}_j \right\rangle \\ &= \left\| \sum_{i=1}^n x_i \mathbf{v}_i \right\|^2 \\ &= \|\mathbf{v}\|^2 && \{\text{Norm } \|\cdot\| \text{ is positive semi-definite}\} \\ &\geq 0 \end{aligned}$$

Moreover,

□

3.3.10 Positive Definiteness.

Gram matrices furnish us with an almost inexhaustible supply of positive semi-definite matrices. However, we still do not know how to test whether a given symmetric matrix is positive definite.

From elementary school, we recall the algebraic technique known as *completing the square*, first arising in the derivation of the formula for the solution to the quadratic equation

$$q(x) = ax^2 + 2bx + c = 0 \quad (3.30)$$

The idea is to combine the first two terms in the equation (3.30) to form a perfect square and thereby rewrite the quadratic function in the form :

$$\begin{aligned} q(x) &= a \left[x^2 + 2\frac{b}{a}x + \frac{c}{a} \right] \\ &= a \left[x^2 + 2x \cdot \frac{b}{a} + \left(\frac{b}{a} \right)^2 + \frac{c}{a} - \left(\frac{b}{a} \right)^2 \right] \\ &= a \left[\left(x + \frac{b}{a} \right)^2 + \frac{ac - b^2}{a^2} \right] \end{aligned}$$

As a consequence,

$$\left(x + \frac{b}{a} \right)^2 = \frac{b^2 - ac}{a^2}$$

The familiar *quadratic formula*:

$$x = \frac{-b \pm \sqrt{b^2 - ac}}{a}$$

follows by taking the square root on both sides and then solving for x .

We can perform the same kind of manipulation on a homogenous quadratic form:

$$q(x_1, x_2) = ax_1^2 + 2bx_1x_2 + cx_2^2 \quad (3.31)$$

In this case, provided $a \neq 0$, completing the square amounts to writing:

$$\begin{aligned}
q(x_1, x_2) &= ax_1^2 + 2bx_1x_2 + cx_2^2 \\
&= a \left[x_1^2 + 2x_1 \cdot \frac{b}{a}x_2 + \left(\frac{b}{a}x_2 \right)^2 + cx_2^2 - \frac{b^2}{a^2}x_2^2 \right] \\
&= a \left[\left(x_1 + \frac{b}{a}x_2 \right)^2 + \frac{ac - b^2}{a^2}x_2^2 \right] \\
&= ay_1^2 + \frac{ac - b^2}{a}y_2^2
\end{aligned} \tag{3.32}$$

The net result is to re-express $q(x_1, x_2)$ as a simpler sum of squares of the new variables:

$$y_1 = x_1 + \frac{b}{a}x_2, \quad y_2 = x_2 \tag{3.33}$$

It is not hard to see that the final expression in (3.32) is positive definite, as a function of y_1 and y_2 if and only if both coefficients are positive:

$$a > 0, \quad \frac{ac - b^2}{a} > 0 \tag{3.34}$$

Our goal is to adapt this simple idea to analyse the positive semi-definiteness of quadratic forms depending on more than two variables. To this end, let us write the quadratic form identity in the matrix form. The original quadratic form in (3.31) can be written as:

$$\begin{aligned}
q(\mathbf{x}) &= \mathbf{x}'K\mathbf{x} \\
&= \begin{bmatrix} x_1 & x_2 \end{bmatrix} \begin{bmatrix} a & b \\ b & c \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}
\end{aligned}$$

Similarly, the right hand side of (3.32) can be written as:

$$\hat{q}(\mathbf{y}) = \mathbf{y}'D\mathbf{y}, \quad \text{where} \quad D = \begin{bmatrix} a & 0 \\ 0 & \frac{ac-b^2}{a} \end{bmatrix}, \quad \mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \end{bmatrix} \tag{3.35}$$

Anticipating the final result, the equations (3.33) connecting \mathbf{x} and \mathbf{y} can themselves be written in the matrix form as:

$$\mathbf{y} = L'\mathbf{x} \quad \text{or} \quad \begin{bmatrix} y_1 \\ y_2 \end{bmatrix} = \begin{bmatrix} x_1 + \frac{b}{a}x_2 \\ x_2 \end{bmatrix}, \quad \text{where} \quad L' = \begin{bmatrix} 1 & 0 \\ b/a & 1 \end{bmatrix} \quad (3.36)$$

Substituting \mathbf{y} into (3.35), we obtain:

$$\mathbf{y}'D\mathbf{y} = (L'\mathbf{x})'D(L'\mathbf{x}) = \mathbf{x}'LDL'\mathbf{x} = \mathbf{x}'K\mathbf{x}, \quad \text{where} \quad K = LDL' \quad (3.37)$$

We are thus led to the realization that completing the square is the same as the LDL' factorization of a symmetric matrix K .

From basic algebra, we know that, if A is a non-singular matrix, with all its pivot elements $a_{kk}^{(k)}$ non-zero in the Gaussian elimination process, then $A = LDU$ where L and U are lower and upper uni-triangular matrices and D is a diagonal matrix consisting of the pivots of A . If the matrix is symmetric, then it admits the unique factorization LDL' .

The identity (3.37) is therefore valid for all real symmetric matrices that are non-singular and can be reduced to an upper triangular matrix by performing elementary row operations (without row interchanges). It also shows how to write the associated quadratic form as a sum of squares:

$$q(\mathbf{x}) = \mathbf{x}'K\mathbf{x} = \mathbf{y}'D\mathbf{y} = d_1y_1^2 + d_2y_2^2 + \dots + d_ny_n^2 \quad \text{where} \quad \mathbf{y} = L'\mathbf{x} \quad (3.38)$$

The coefficients d_i are the diagonal entries of D , which are the pivots of K . The diagonal quadratic form is positive definite, $\mathbf{y}'D\mathbf{y} > 0$ for all $\mathbf{y} \neq \mathbf{0}$ if and only if, when performing the Gaussian elimination process, all the pivots are positive. We can now add this to our list of standard results.

Theorem 3.16. (Positive Definiteness) *Let K be a $n \times n$ real symmetric positive definite (SPD) matrix. Then the following statements are equivalent.*

- (i) K is non-singular and can be reduced to an upper triangular matrix by performing elementary row operations (without row permutations), and it has positive pivot elements when performing Gaussian elimination.
- (ii) K admits a factorization $K = LDL'$, where $D = \text{diag}(d_1, \dots, d_n)$ such that $d_i > 0$ for all $i = 1, 2, 3, \dots, n$.

3.3.11 Cholesky Factorization.

The identity (3.37) shows us how to write an arbitrary regular quadratic form $q(\mathbf{x})$ as linear combination of squares. We can push this result slightly further in the positive definite case. Since each pivot d_i is positive, we can write the quadratic form as a sum of squares:

$$\begin{aligned} d_1 y_1^2 + d_2 y_2^2 + \dots + d_n y_n^2 &= (\sqrt{d_1} y_1)^2 + (\sqrt{d_2} y_2)^2 + \dots + (\sqrt{d_n} y_n)^2 \\ &= z_1^2 + z_2^2 + \dots + z_n^2 \end{aligned}$$

where $z_i = \sqrt{d_i} y_i$. In the matrix form, we are writing:

$$\begin{aligned} \hat{q}(\mathbf{y}) &= \mathbf{y}' D \mathbf{y} \\ &= \mathbf{z}' \mathbf{z} \\ &= \|\mathbf{z}\|^2 \end{aligned}$$

where $\mathbf{z} = S\mathbf{y}$, with $S = \text{diag}(\sqrt{d_1}, \sqrt{d_2}, \dots, \sqrt{d_n})$. Since $D = S^2$, the matrix S can be thought of as a square root of the diagonal matrix D . Substituting back into the equation $K = LDL'$, we deduce the *Cholesky factorization*:

$$\begin{aligned} K &= LDL' \\ &= LSS'L' \\ &= LS(LS)' \\ &= MM' \end{aligned}$$

of a positive definite matrix, first proposed by the early twentieth-century French geographer Andrew Louis Cholesky for solving problems in geodetic surveying. Note that, M is a lower triangular matrix with all positive diagonal entries, namely the square roots of the pivots: $m_{ii} = \sqrt{d_i}$.

Example 3.3. Let the matrix $K = \begin{bmatrix} 1 & 2 & -1 \\ 2 & 6 & 0 \\ -1 & 0 & 9 \end{bmatrix}$. Let $KX = I$. We consider the augmented matrix $[K \mid I]$. Performing Gaussian elimination, we have:

$$\left[\begin{array}{ccc|ccc} 1 & 2 & -1 & 1 & 0 & 0 \\ 2 & 6 & 0 & 0 & 1 & 0 \\ -1 & 0 & 9 & 0 & 0 & 1 \end{array} \right]$$

The pivot element $a_{11}^{(1)} = 1$. Performing $R_2 = R_2 - 2R_1$ and $R_3 = R_3 + R_1$, the above system is row-equivalent to:

$$\left[\begin{array}{ccc|ccc} 1 & 2 & -1 & 1 & 0 & 0 \\ 0 & 2 & 2 & -2 & 1 & 0 \\ 0 & 2 & 8 & 1 & 0 & 1 \end{array} \right]$$

The pivot element $a_{22}^{(2)} = 2$. Performing $R_3 = R_3 - R_2$, the above system is row-equivalent to:

$$\left[\begin{array}{ccc|ccc} 1 & 2 & -1 & 1 & 0 & 0 \\ 0 & 2 & 2 & -2 & 1 & 0 \\ 0 & 0 & 6 & 3 & -1 & 1 \end{array} \right]$$

The pivot element $a_{33}^{(3)} = 6$. We have now reduced the system to the form $[DU \mid C]$, where U is an upper uni-triangular matrix. Thus, Gaussian Elimination produces the factors:

$$L = \begin{bmatrix} 1 & 0 & 0 \\ 2 & 1 & 0 \\ -1 & 1 & 1 \end{bmatrix}, \quad D = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 6 \end{bmatrix}, \quad L^T = \begin{bmatrix} 1 & 2 & -1 \\ 0 & 1 & 1 \\ 0 & 0 & 1 \end{bmatrix}$$

Thus,

$$M = LS = \begin{bmatrix} 1 & 0 & 0 \\ 2 & 1 & 0 \\ -1 & 1 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 \\ 0 & \sqrt{2} & 0 \\ 0 & 0 & \sqrt{6} \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 \\ 2 & \sqrt{2} & 0 \\ -1 & \sqrt{2} & \sqrt{6} \end{bmatrix}$$

and $K = MM^T$.

We conclude our discussion by observing the following:

Lemma 3.3. *If a square matrix K is SPD, it admits a Cholesky factorization of the form $K = MM^T$.*

Example 3.4. Prove that, if K is real SPD (symmetric positive definite matrix), then the diagonal elements of K are positive.

Proof. Since K is real SPD, K admits a factorization $K = LL^T$. Since the diagonal element (j, j) is the inner product of the j -th row of L and the j -th column of L^T , we have:

$$k_{jj} = \sum_{m=1}^n l_{jm} l'_{mj}$$

But, $l_{jm} = l'_{mj}$, since $L = (L^T)^T$. Hence, k_{jj} is a sum of squares. Further, since the diagonal elements of L , that is, all elements l_{jj} are strictly positive, the sum $k_{jj} = l_{j1}^2 + \dots + l_{jn}^2 > 0$. Consequently, the diagonal elements of K are positive. \square

3.3.12 Cholesky Factorization Algorithm.

We adopt the commonly used notation where Greek lower-case letters refer to scalars, lower-case letters refer to (column) vectors and upper case letters refer to matrices. The \star refers to a part of A that is neither stored nor updated. By substituting these partitioned matrices into $A = LL^T$ we find that:

$$\begin{bmatrix} \alpha_{11} & a_{21}^T \\ a_{21} & A_{22} \end{bmatrix} = \begin{bmatrix} l_{21} & L_{22} \end{bmatrix} \begin{bmatrix} \lambda_{11} & l_{21}^T \\ 0 & L_{22}^T \end{bmatrix} = \begin{bmatrix} \lambda_{11}^2 & \star \\ \lambda_{11} l_{21} & l_{21} l_{21}^T + L_{22} L_{22}^T \end{bmatrix}$$

so that :

$$\frac{a_{11} = \lambda_{11}^2 \quad \star}{a_{21} = \lambda_{11} l_{21} \quad A_{22} = l_{21} l_{21}^T + L_{22} L_{22}^T}$$

and hence:

$$\frac{\lambda_{11} = \sqrt{a_{11}} \quad \star}{l_{21} = a_{21} / \lambda_{11} \quad L_{22} = \text{Cholesky}(A_{22} - l_{21} l_{21}^T)}$$

The last equality is clever. Essentially, if $A_{22} = l_{21} l_{21}^T - L_{22} L_{22}^T$, we must have: $L_{22} L_{22}^T = A_{22} - l_{21} l_{21}^T$. So, to find L_{22} , we recursively perform the cholesky factorization of the matrix $A_{22} - l_{21} l_{21}^T$. These equalities motivate the following block algorithm:

1. Partition $A = \left[\begin{array}{c|c} \alpha_{11} & \star \\ \hline a_{21} & A_{22} \end{array} \right]$.
2. Overwrite $\alpha_{11} := \lambda_{11} = \sqrt{\alpha_{11}}$.
3. Overwrite $a_{21} := l_{21} = a_{21}/\lambda_{11}$.
4. Overwrite $A_{22} := A_{22} - l_{21}l_{21}^T$.
5. Continue with $A = A_{22}$.

We can also implement a serial algorithm by multiplying out the matrices:

$$\begin{aligned}
 \begin{bmatrix} a_{11} & a_{21} & a_{31} & a_{41} \\ a_{21} & a_{22} & a_{32} & a_{42} \\ a_{31} & a_{32} & a_{33} & a_{43} \\ a_{41} & a_{42} & a_{43} & a_{44} \end{bmatrix} &= \begin{bmatrix} l_{11} & 0 & 0 & 0 \\ l_{21} & l_{22} & 0 & 0 \\ l_{31} & l_{32} & l_{33} & 0 \\ l_{41} & l_{42} & l_{43} & l_{44} \end{bmatrix} \begin{bmatrix} l_{11} & l_{21} & l_{31} & l_{41} \\ 0 & l_{22} & l_{32} & l_{42} \\ 0 & 0 & l_{33} & l_{43} \\ 0 & 0 & 0 & l_{44} \end{bmatrix} \\
 &= \begin{bmatrix} l_{11}^2 & 0 & 0 & 0 \\ l_{21}l_{11} & l_{21}^2 + l_{22}^2 & 0 & 0 \\ l_{31}l_{11} & l_{31}l_{21} + l_{32}l_{22} & l_{31}^2 + l_{32}^2 + l_{33}^2 & 0 \\ l_{41}l_{11} & l_{41}l_{31} + l_{42}l_{32} & l_{41}l_{31} + l_{42}l_{32} + l_{43}l_{33} & l_{41}^2 + l_{42}^2 + l_{43}^2 + l_{44}^2 \end{bmatrix}
 \end{aligned}$$

We can thus solve for the elements of the matrix L , column-by-column. The expressions for l_{jj} and l_{ij} in general, are given by:

$$\begin{aligned}
 l_{jj} &= \sqrt{a_{jj} - \sum_{k=1}^{j-1} l_{jk}^2} \\
 l_{ij} &= \frac{1}{l_{jj}} \left(a_{ij} - \sum_{k=1}^{j-1} l_{ik} \cdot l_{jk} \right), \quad \forall i > j
 \end{aligned}$$

Listing 2: Cholesky Factorization

```

#include <iostream>
#include <Eigen/Dense>
#include <cmath>

using Eigen::MatrixXd;

```

```

// Cholesky-Crout algorithm starts from the upper-left corner of
// the matrix L and proceeds
// to calculate matrix column by column
MatrixXd choleskyDecomposition(const MatrixXd& A)
{
    MatrixXd L = MatrixXd::Zero(A.rows(), A.cols());

    for (int j{ 0 }; j < A.cols(); ++j)
    {
        double sum{ 0.0 };
        for (int k{ 0 }; k < j; ++k)
        {
            sum += L(j, k) * L(j, k);
        }
        L(j, j) = sqrt(A(j, j) - sum);

        for (int i{ j + 1 }; i < A.rows(); ++i)
        {
            double sum{ 0.0 };
            for (int k{ 0 }; k < j; ++k) {
                sum += L(i, k) * L(j, k);
            }

            L(i, j) = (A(i, j) - sum)/L(j,j);
        }
    }

    return L;
}

int main()
{
    MatrixXd K(3, 3);

    K <<      4, 12, -16,
              12, 37, -43,
              -16, -43, 98;

    MatrixXd L = choleskyDecomposition(K);

    std::cout << "The SPD(Symmetric Positive Definite) matrix K
is : " << std::endl;
    std::cout << K << std::endl;
    std::cout << "The Cholesky Decomposition of K into K=LL\'
yields L : " << std::endl;

```

```

std::cout << L << std::endl;

return 0;
}

```

3.3.13 Eigen-decomposition of real symmetric matrices.

We review couple of lemmas from basic algebra, which we shall need in the main result.

Lemma 3.4. *Every linearly independent sequence can be extended to a basis.*

Let V be a finite-dimensional vector space and let $\mathbf{l}_1, \mathbf{l}_2, \dots, \mathbf{l}_n$ be linearly independent. Then, there exists a basis of V containing $\mathbf{l}_1, \mathbf{l}_2, \dots, \mathbf{l}_n$.

Proof. Let $\mathcal{L} = \mathbf{l}_1, \mathbf{l}_2, \dots, \mathbf{l}_n$. Since V is finite-dimensional, there exist elements $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_m$ of V such that they span V .

Define a sequence of sequences of the elements of V as follows. Set $\mathcal{L}_0 = \mathcal{L}$ and for $i \geq 0$, define:

$$\mathcal{L}_{i+1} = \begin{cases} \mathcal{L}_i & \text{if } \mathbf{v}_i \in \text{span}(\mathcal{L}_i) \\ \mathcal{L}_i, \mathbf{v}_{i+1} & \text{otherwise} \end{cases}$$

Here, $\mathcal{L}_i, \mathbf{v}_{i+1}$ just means take the sequence \mathcal{L}_i and add \mathbf{v}_{i+1} on to the end.

Note that in either case, $\mathbf{v}_{i+1} \in \text{span}(\mathcal{L}_{i+1})$ and also that $\mathcal{L}_0 \subseteq \mathcal{L}_1 \subseteq \dots \subseteq \mathcal{L}_m$.

By construction, each sequence \mathcal{L}_i is linearly independent and in particular \mathcal{L}_m is linearly independent. Furthermore, $\text{span}(\mathcal{L}_m)$ contains $\{\mathbf{v}_1, \dots, \mathbf{v}_m\}$ and therefore contains $\text{span}(\mathcal{L}_m) = V$. Therefore, \mathcal{L}_m is a basis for V containing \mathcal{L} . This completes the proof. \square

Lemma 3.5. (EMHE) *Every matrix has an (atleast one) eigenvalue, and a corresponding eigenvector.*

Proof. This is just the Fundamental Theorem of Algebra(FTA), but it's still worth enumerating as a theorem.

Let $A \subseteq \mathbf{C}^{n \times n}$ and the scalar field $\mathbf{F} = \mathbf{C}$.

Let \mathbf{v} be any non-zero vector in \mathbf{C}^n . Consider the list $\mathcal{L} = \mathbf{v}, A\mathbf{v}, A^2\mathbf{v}, \dots, A^n\mathbf{v}$. There are $n + 1$ vectors in the list, so they must be linearly dependent. There exists scalars a_0, a_1, \dots, a_n from \mathbf{C} not all zero, such that:

$$a_0\mathbf{v} + a_1A\mathbf{v} + a_2A^2\mathbf{v} + \dots + a_nA^n\mathbf{v} = \mathbf{0}$$

By FTA, the polynomial equation of degree n :

$$p(x) = a_0 + a_1x + a_2x^2 + \dots + a_nx^n = 0$$

has n linear factors

$$p(x) = (x - \lambda_1)(x - \lambda_2) \cdots (x - \lambda_n) = 0$$

where $\lambda_i \in \mathbf{C}, i = 1, 2, \dots, n$.

Putting it all together,

$$\begin{aligned} p(A)\mathbf{v} &= \mathbf{0} = a_0\mathbf{v} + a_1A\mathbf{v} + a_2A^2\mathbf{v} + \dots + a_nA^n\mathbf{v} \\ &= (a_0 + a_1A + a_2A^2 + \dots + a_nA^n)\mathbf{v} \\ &= (A - \lambda_1I)(A - \lambda_2I) \cdots (A - \lambda_nI)\mathbf{v} \end{aligned}$$

This shows that the composition of the factors has a non-trivial nullspace. $\ker((A - \lambda_1I)(A - \lambda_2I) \cdots (A - \lambda_nI)) \neq \{\mathbf{0}\}$. So, atleast one of the factors must fail to be injective. There exists λ_i , such that $(A - \lambda_i)\mathbf{v} = \mathbf{0}$ such that $\mathbf{v} \neq \mathbf{0}$. Thus, A has atleast one eigenvalue and a corresponding eigenvector. \square

Theorem 3.17. (*Spectral Theorem*) Every real symmetric matrix is diagonalizable.

Let A be a symmetric $n \times n$ real matrix. Then,

- 1) The eigenvalues of A are real.
- 2) There exists an orthonormal basis $\{\mathbf{q}_1, \mathbf{q}_2, \dots, \mathbf{q}_n\}$ for \mathbf{R}^n consisting of the eigenvectors of A . That is, there is an orthogonal matrix Q so that $Q^{-1}AQ = \Lambda$ is diagonal.

Proof. (I) Before we get to the proof, note that for any square matrix A , we have:

$$\begin{aligned}\langle A\mathbf{x}, \mathbf{y} \rangle &= \mathbf{x}' A' \mathbf{y} \\ &= \langle \mathbf{x}, A' \mathbf{y} \rangle\end{aligned}$$

Since for a symmetric matrix A , we have, $A = A'$, it follows that:

$$\langle A\mathbf{x}, \mathbf{y} \rangle = \langle \mathbf{x}, A\mathbf{y} \rangle$$

Or using the dot-product notation, we could write:

$$(A\mathbf{x}) \cdot \mathbf{y} = \mathbf{x} \cdot (A\mathbf{y})$$

Suppose $\mathbf{v} \neq \mathbf{0}$ be a non-zero vector in \mathbf{R}^n such that there exists a complex scalar λ , satisfying:

$$A\mathbf{v} = \lambda\mathbf{v} \tag{3.39}$$

We can therefore write:

$$(A\mathbf{v}) \cdot \mathbf{v} = (\lambda\mathbf{v}) \cdot \mathbf{v} = \lambda(\mathbf{v} \cdot \mathbf{v}) \tag{3.40}$$

Alternatively,

$$(A\mathbf{v}) \cdot \mathbf{v} = \mathbf{v} \cdot (A\mathbf{v}) \tag{3.41}$$

We can now take the complex conjugate of the (3.39) equation. Remember that A is a real matrix so $\overline{A} = A$. Thus, we have the conjugated version of the eigen-value equation:

$$\overline{(A\mathbf{v})} = \overline{A\mathbf{v}} = A\overline{\mathbf{v}} = \overline{(\lambda\mathbf{v})} = \overline{\lambda}\overline{\mathbf{v}}$$

In equation (3.40), if we replace the second vector \mathbf{v} with its conjugate, $\bar{\mathbf{v}}$, we get:

$$(A\mathbf{v}) \cdot \bar{\mathbf{v}} = \lambda(\mathbf{v} \cdot \bar{\mathbf{v}}) \quad (3.42)$$

In equation (3.41), if we replace the second vector \mathbf{v} with its conjugate, $\bar{\mathbf{v}}$, we get:

$$(A\mathbf{v}) \cdot \bar{\mathbf{v}} = \mathbf{v} \cdot (A\bar{\mathbf{v}}) = \mathbf{v} \cdot (\bar{\lambda}\bar{\mathbf{v}}) = \bar{\lambda}(\mathbf{v} \cdot \bar{\mathbf{v}}) \quad (3.43)$$

Now, since \mathbf{v} is an eigenvector, it cannot be the zero vector.

Without loss of generality, if $\mathbf{v} = (v_1, \dots, v_n)$, then $\mathbf{v} \cdot \bar{\mathbf{v}} = |v_1|^2 + \dots + |v_n|^2 \neq 0$, so $\mathbf{v} \cdot \bar{\mathbf{v}} \neq 0$.

The two expressions for $(A\mathbf{v}) \cdot \bar{\mathbf{v}}$ are equal, so $(\lambda - \bar{\lambda})(\mathbf{v} \cdot \bar{\mathbf{v}}) = 0$. But, $(\mathbf{v} \cdot \bar{\mathbf{v}}) \neq 0$, so $\lambda = \bar{\lambda}$. Therefore, $\lambda \in \mathbf{R}$.

(II) We proceed by mathematical induction on n .

For $n = 1$, any 1×1 symmetric matrix is already diagonal. Since A and $v \in V$ are both scalars, $Av = \lambda v$ where $\lambda = A$. Thus, we can pick any non-zero scalar v to form a basis of \mathbf{R} . And we can write, $A = P^{-1}\Lambda P$, where $P = I$ and $\Lambda = A$.

Induction hypothesis: Every $k \times k$ symmetric matrix is diagonalizable for $k = 1, 2, 3, \dots, n-1$. If C is a real symmetric matrix of size $k \times k$, then there exists an orthogonal matrix R such that $R^{-1}CR$ is diagonal.

By lemma (3.5), the square matrix A has atleast one eigenvalue. Suppose λ_1 is an eigenvalue of the matrix A . By part (I), we know that $\lambda_1 \in \mathbf{R}$. Choose a unit vector \mathbf{q}_1 that is an eigenvector with eigenvalue λ_1 . (Obviously, this is no problem. We can pick an eigenvector and then make it a unit vector by dividing by it's length.)

By lemma (3.4), we can extend this to a basis $\{\mathbf{q}_1, \mathbf{w}_2, \dots, \mathbf{w}_n\}$ of V . By the Gram-Schmidt orthogonalization algorithm, given the basis $\{\mathbf{q}_1, \mathbf{w}_2, \dots, \mathbf{w}_n\}$, we can find a corresponding orthonormal basis $\{\mathbf{q}_1, \mathbf{q}_2, \dots, \mathbf{q}_n\}$ of V .

Now, we huddle these basis vectors together as column-vectors of a matrix and formulate the matrix P .

$$P = \begin{bmatrix} \mathbf{q}_1 & \mathbf{q}_2 & \dots & \mathbf{q}_n \end{bmatrix}$$

By definition, P is an orthogonal matrix.

Let

$$B = P^{-1}AP$$

We are interested to show that B is diagonal.

Step I. B is symmetric.

We have:

$$\begin{aligned} B^T &= (P^{-1}AP)^T \\ &= (P^T AP)^T && \{P^{-1} = P^T\} \\ &= P^T A^T (P^T)^T \\ &= P^T A^T P \\ &= P^T AP && \{A \text{ is symmetric}\} \\ &= B \end{aligned}$$

We are now going to try and write B in the block form to try to see the structure that this matrix must have and hope that it looks like, it is going to be diagonal.

Step II. The structure of B .

The way we do this, is to consider the matrix B post-multiplied by \mathbf{e}_1 . Consider $B\mathbf{e}_1$. This should actually give us the first column of B . Now, we also know that $B = P^T AP$. So, we could actually say, well,

$$P^T AP\mathbf{e}_1 = P^T A\mathbf{q}_1$$

Now, remember that \mathbf{q}_1 is the normalized eigenvector corresponding to the eigenvalue λ_1 . So, $A\mathbf{q}_1 = \lambda_1\mathbf{q}_1$. That means, this is equal to:

$$\begin{aligned}
P^T A \mathbf{q}_1 &= P^T \lambda_1 \mathbf{q}_1 \\
&= \lambda_1 P^T \mathbf{q}_1 \\
&= \lambda_1 \begin{bmatrix} \mathbf{q}_1^T \\ \mathbf{q}_2^T \\ \vdots \\ \mathbf{q}_n^T \end{bmatrix} \mathbf{q}_1 \\
&= \lambda_1 \begin{bmatrix} \mathbf{q}_1^T \mathbf{q}_1 \\ \mathbf{q}_2^T \mathbf{q}_1 \\ \vdots \\ \mathbf{q}_n^T \mathbf{q}_1 \end{bmatrix} \\
&= \lambda_1 \begin{bmatrix} 1 \\ 0 \\ \vdots \\ 0 \end{bmatrix} \\
&= \begin{bmatrix} \lambda_1 \\ 0 \\ 0 \\ 0 \end{bmatrix}
\end{aligned}$$

This is the first column of the matrix B . Since $B = B^T$, the first row should also be

$$[\lambda_1 \quad 0 \quad 0 \quad 0]$$

So, we can write the matrix B in the form:

$$B = \begin{bmatrix} \lambda_1 & O \\ O & C \end{bmatrix}$$

The first row and the first column are satisfying the need to be diagonal.

Step III.

We know that C is a $(n-1) \times (n-1)$ symmetric matrix. By the inductive hypothesis, C is diagonalizable and further there exists an orthogonal matrix R , such that $R^{-1}CR = D$ where D is diagonal.

Define the matrix Q as:

$$Q := P \begin{bmatrix} 1 & 0_{1 \times n-1} \\ 0_{n-1 \times 1} & R \end{bmatrix} \quad (3.44)$$

Our claim is that Q is orthogonal and $Q^{-1}AQ$ is diagonal.

(i) We have:

$$\begin{aligned} Q^{-1} &= \begin{bmatrix} 1 & 0_{1 \times n-1} \\ 0_{n-1 \times 1} & R^{-1} \end{bmatrix} P^{-1} && \{\text{Reverse order law}\} \\ &= \begin{bmatrix} 1 & 0_{1 \times n-1} \\ 0_{n-1 \times 1} & R^T \end{bmatrix} P^T && \{P \text{ and } R \text{ are orthogonal}\} \end{aligned}$$

But,

$$Q^T = \begin{bmatrix} 1 & 0_{1 \times n-1} \\ 0_{n-1 \times 1} & R^T \end{bmatrix} P^T$$

So,

$$Q^T = Q^{-1}$$

Thus, Q is orthogonal.

(ii) Well, let's compute $Q^{-1}AQ$.

$$\begin{aligned}
Q^{-1}AQ &= Q^T AQ && \{Q \text{ is orthogonal}\} \\
&= \begin{bmatrix} 1 & 0_{1 \times n-1} \\ 0_{n-1 \times 1} & R^T \end{bmatrix} P^T A P \begin{bmatrix} 1 & 0_{1 \times n-1} \\ 0_{n-1 \times 1} & R \end{bmatrix} \\
&= \begin{bmatrix} 1 & 0_{1 \times n-1} \\ 0_{n-1 \times 1} & R^T \end{bmatrix} B \begin{bmatrix} 1 & 0_{1 \times n-1} \\ 0_{n-1 \times 1} & R \end{bmatrix} \\
&= \begin{bmatrix} 1 & 0_{1 \times n-1} \\ 0_{n-1 \times 1} & R^T \end{bmatrix} \begin{bmatrix} \lambda_1 & 0_{1 \times n-1} \\ 0_{n-1 \times 1} & C \end{bmatrix} \begin{bmatrix} 1 & 0_{1 \times n-1} \\ 0_{n-1 \times 1} & R \end{bmatrix} \\
&= \begin{bmatrix} \lambda_1 & 0_{1 \times n-1} \\ 0_{n-1 \times 1} & R^T C \end{bmatrix} \begin{bmatrix} 1 & 0_{1 \times n-1} \\ 0_{n-1 \times 1} & R \end{bmatrix} \\
&= \begin{bmatrix} \lambda_1 & 0_{1 \times n-1} \\ 0_{n-1 \times 1} & R^T C R \end{bmatrix}
\end{aligned}$$

Since $R^T C R$ is diagonal, it follows that $Q^{-1}AQ$ is diagonal. This closes the proof. \square

3.4 Covariance and MGF of random variables.

Definition 3.12. If (X, Y) is a random vector, then the covariance of (X, Y) is given by:

$$\text{Cov}(X, Y) = \mathbb{E}[(X - \mathbb{E}X)(Y - \mathbb{E}Y)] = \mathbb{E}[XY] - \mathbb{E}[X] \cdot \mathbb{E}[Y] \quad (3.45)$$

3.4.1 Expected value of a random matrix.

Suppose our random experiment is modeled by the probability space $(\Omega, \mathcal{F}, \mathbb{P})$. We can define the expected value of a random matrix in a component-wise manner.

Suppose that \mathbf{X} is an $m \times n$ matrix of real-valued random variables, whose (i, j) entry is denoted by X_{ij} . Equivalently, \mathbf{X} is a random $m \times n$ matrix. The expected value $\mathbb{E}(\mathbf{X})$ is defined to be the $m \times n$ matrix whose (i, j) entry is $\mathbb{E}X_{ij}$, the expected value of X_{ij} .

Many of the basic properties of expected value of random variables have analogous results for expected values of random matrices/vectors. If \mathbf{X} and \mathbf{Y} are random $m \times n$

matrices, the linearity property holds: $\mathbb{E}(\mathbf{X} + \mathbf{Y}) = \mathbb{E}\mathbf{X} + \mathbb{E}\mathbf{Y}$. Similarly, if \mathbf{X} is a $n \times p$ random matrix and \mathbf{a} is a constant $m \times n$ matrix, the constant factor can be pulled out of the expectation. $\mathbb{E}[\mathbf{aX}] = \mathbf{a}\mathbb{E}[\mathbf{X}]$.

3.4.2 Covariance Matrices.

Definition 3.13. Suppose that \mathbf{X} is a random vector in \mathbf{R}^m and \mathbf{Y} is a random vector in \mathbf{R}^n . The covariance matrix of \mathbf{X} and \mathbf{Y} is the $m \times n$ matrix $\text{Cov}(\mathbf{X}, \mathbf{Y})$ whose (i, j) entry is $\text{Cov}(X_i, Y_j)$.

Definition 3.14. Let $\mathbf{X} = (X_1, \dots, X_n)$ be a random vector in \mathbf{R}^n . Then the covariance matrix of \mathbf{X} , denoted by Σ is the $n \times n$ matrix, whose (i, j) entry is $\text{Cov}(X_i, X_j)$.

Theorem 3.18. Let (X, Y) be random variables. $\text{Cov}(X, Y)$ has the following properties:

(i) $\text{Cov}(X, X) = \text{Var}(X)$

(ii) $\text{Cov}(X, Y) = \text{Cov}(Y, X)$

(iii) $\text{Cov}(X, c) = 0$

(iv) *Scaling property:* $\text{Cov}(aX, Y) = a\text{Cov}(X, Y)$

(v) *Bi-linearity:*

$$\text{Cov}(aX + bY, Z) = a\text{Cov}(X, Z) + b\text{Cov}(Y, Z)$$

$$\text{Cov}(X, cY + dZ) = c\text{Cov}(X, Y) + d\text{Cov}(X, Z)$$

(vi) $\text{Var}(X + Y) = \text{Var}(X) + \text{Var}(Y) + 2\text{Cov}(X, Y)$

Since $\text{Cov}(X, -Y) = -\text{Cov}(X, Y)$, it follows that $\text{Var}(X - Y) = \text{Var}(X) + \text{Var}(-Y) + 2\text{Cov}(X, -Y) = \text{Var}(X) + \text{Var}(Y) - 2\text{Cov}(X, Y)$

$$\text{Var}(X_1 + X_2 + \dots + X_n) = \sum_{i=1}^n \text{Var}(X_i) + \sum_{i=1}^n \sum_{j=1}^n \text{Cov}(X_i, X_j)$$

Theorem 3.19. Let $\mathbf{X} = (X_1, X_2, \dots, X_n)$ be a random vector with mean vector $\boldsymbol{\mu} = (\mu_1, \mu_2, \dots, \mu_n)$ and $n \times n$ covariance matrix Σ . Then, Σ is positive semi-definite.

Proof. We have:

$$\begin{aligned}
\Sigma &= \begin{bmatrix} \mathbb{E}(X_1 - \mu_1)(X_1 - \mu_1) & \mathbb{E}(X_1 - \mu_1)(X_2 - \mu_2) & \dots & \mathbb{E}(X_1 - \mu_1)(X_n - \mu_n) \\ \mathbb{E}(X_2 - \mu_2)(X_1 - \mu_1) & \mathbb{E}(X_2 - \mu_2)(X_2 - \mu_2) & \dots & \mathbb{E}(X_2 - \mu_2)(X_n - \mu_n) \\ \vdots & \vdots & \ddots & \vdots \\ \mathbb{E}(X_n - \mu_n)(X_1 - \mu_1) & \mathbb{E}(X_n - \mu_n)(X_2 - \mu_2) & \dots & \mathbb{E}(X_n - \mu_n)(X_n - \mu_n) \end{bmatrix} \\
&= \mathbb{E} \left[\begin{bmatrix} X_1 - \mu_1 \\ X_2 - \mu_2 \\ \vdots \\ X_n - \mu_n \end{bmatrix} \begin{bmatrix} X_1 - \mu_1 & X_2 - \mu_2 & \dots & X_n - \mu_n \end{bmatrix} \right] \\
&= \mathbb{E} [(\mathbf{X} - \boldsymbol{\mu})(\mathbf{X} - \boldsymbol{\mu})']
\end{aligned}$$

Let \mathbf{a} be an arbitrary(not random) vector in \mathbf{R}^n . Then,

$$\begin{aligned}
\mathbf{a}'\Sigma\mathbf{a} &= \mathbf{a}'\mathbb{E} [(\mathbf{X} - \boldsymbol{\mu})(\mathbf{X} - \boldsymbol{\mu})'] \mathbf{a} \\
&= \mathbb{E} [\mathbf{a}'(\mathbf{X} - \boldsymbol{\mu})(\mathbf{X} - \boldsymbol{\mu})'\mathbf{a}] \\
&= \mathbb{E} [((\mathbf{X} - \boldsymbol{\mu})'\mathbf{a})'((\mathbf{X} - \boldsymbol{\mu})'\mathbf{a})] \\
&= \mathbb{E} [(\mathbf{X} - \boldsymbol{\mu})'\mathbf{a}]^2 \\
&\geq 0
\end{aligned}$$

Consequently, Σ is a positive semi-definite matrix. \square

Definition 3.15. The MGF of a random variable X on $(\Omega, \mathcal{F}, \mathbb{P})$ is the function on \mathbf{R} defined by:

$$M_X(t) = \mathbb{E} [e^{tX}]$$

Example 3.5. The MGF of a standard Gaussian random variable given by:

$$\begin{aligned}
M_Z(t) &= \mathbb{E} [e^{tZ}] \\
&= \int_{-\infty}^{\infty} e^{tz} \phi(z) dz \\
&= \frac{1}{\sqrt{2\pi}} \int_{-\infty}^{\infty} e^{tz} e^{-z^2/2} dz
\end{aligned}$$

We can complete the square in the exponent as follows:

$$\begin{aligned}\exp\left(tz - \frac{z^2}{2}\right) &= \exp\left[-\frac{1}{2}(z^2 - 2tz + t^2 - t^2)\right] \\ &= \exp\left[-\frac{1}{2}(z - t)^2 + \frac{t^2}{2}\right]\end{aligned}$$

So,

$$\begin{aligned}M_Z(t) &= \frac{e^{t^2/2}}{\sqrt{2\pi}} \int_{-\infty}^{\infty} e^{-(z-t)^2/2} dz \\ &= \frac{e^{t^2/2}}{\sqrt{2\pi}} \sqrt{2\pi} \\ &= e^{t^2/2}\end{aligned}$$

Differentiating with respect to t , we have:

$$\begin{aligned}M'_Z(t) &= te^{t^2/2} \\ M''_Z(t) &= e^{t^2/2} + t^2e^{t^2/2} \\ M^{(3)}_Z(t) &= 3te^{t^2/2} + t^3e^{t^2/2} \\ M^{(4)}_Z(t) &= 3e^{t^2/2} + 6t^2e^{t^2/2} + t^4e^{t^2/2}\end{aligned}$$

So, the mean of the standard gaussian random variable is $M'_Z(0) = 0$, the second moment and variance of a standard gaussian random variable is $M''_Z(0) = 1$. The skewness of the standard gaussian random variable is $M^{(3)}_Z(0) = 0$, while the kurtosis of a standard gaussian random variable is $M^{(4)}_Z(0) = 3$.

Definition 3.16. (*Joint Moment Generating Function (MGF)*). The joint MGF of a random vector $\mathbf{X} = (X_1, X_2, \dots, X_n)$ on $(\Omega, \mathcal{F}, \mathbb{P})$ is the function defined on \mathbf{R}^n by:

$$M_{\mathbf{X}}(\mathbf{t}) = \mathbb{E} [\exp(\mathbf{t}^T \mathbf{X})] = \mathbb{E} [\exp(t_1 X_1 + t_2 X_2 + \dots + t_n X_n)] \quad (3.46)$$

The following result will be stated without proof. It will be useful when studying Gaussian vectors.

Proposition 3.3. *Let $(\Omega, \mathcal{F}, \mathbb{P})$ be a probability space. Two random vectors X and Y that have the same moment generating function have the same distribution.*

Example 3.6. Consider (X, Y) a random vector with value in \mathbf{R}^2 such that X and Y are IID with standard Gaussian distribution. Then, the joint PDF is:

$$f(x, y) = \frac{1}{\sqrt{2\pi}} e^{-x^2/2} \times \frac{1}{\sqrt{2\pi}} e^{-y^2/2} = \frac{1}{\sqrt{2\pi}} e^{-(x^2+y^2)/2}$$

The moment generating function is obtained by independence:

$$\begin{aligned} M_{(X,Y)}(t_1, t_2) &= \mathbb{E}[e^{t_1 X + t_2 Y}] \\ &= \mathbb{E}[e^{t_1 X} \cdot e^{t_2 Y}] \\ &= \mathbb{E}[e^{t_1 X}] \cdot \mathbb{E}[e^{t_2 Y}] \\ &= e^{t_1^2/2} \cdot e^{t_2^2/2} \\ &= e^{(t_1^2 + t_2^2)/2} \end{aligned}$$

More generally, we can consider n IID random variables with standard Gaussian distribution. We then have the joint PDF:

$$f(x_1, x_2, \dots, x_n) = \frac{e^{-(x_1^2 + x_2^2 + \dots + x_n^2)/2}}{(2\pi)^{n/2}}$$

In order to work with random vectors, we frequently use the change-of-variables theorem from vector calculus.

Theorem 3.20. *If $f : \mathbf{R}^2 \rightarrow \mathbf{R}$ and $\mathbf{T}(x_1, x_2) = (x_1(y_1, y_2), x_2(y_1, y_2))$ is a linear transformation that maps the domain D^* to D , then we have:*

$$\int \int_D f(x_1, x_2) dx_1 dx_2 = \int \int_{D^*} f(x_1(y_1, y_2), x_2(y_1, y_2)) \left| \frac{\partial(x_1, x_2)}{\partial(y_1, y_2)} \right| dy_1 dy_2$$

Corollary 3.2. *If X_1, X_2 have the joint density function f , and T is any linear transformation, then the pair $(Y_1, Y_2) = T(X_1, X_2)$ has the density function:*

$$f_{(Y_1, Y_2)}(y_1, y_2) = f(x_1(y_1, y_2), x_2(y_1, y_2)) \left| \frac{\partial(x_1, x_2)}{\partial(y_1, y_2)} \right|$$

Example 3.7. (*Computations with random vectors*). Let (X, Y) be two IID standard Gaussian random variables. We can think of (X, Y) as the random point in \mathbf{R}^2 with x -coordinate X and y -coordinate Y .

First off, let's compute the probability that the point (X, Y) is in the unit disc $D = \{(x, y) | x^2 + y^2 = 1\}$. The probability is given by the double integral:

$$\begin{aligned} P((X, Y) \in D) &= \int \int_D \frac{1}{2\pi} e^{-(x^2+y^2)/2} dx dy \\ &= \int_{-1}^{+1} \int_{-\sqrt{1-x^2}}^{+\sqrt{1-x^2}} \frac{1}{2\pi} e^{-(x^2+y^2)/2} dx dy \end{aligned}$$

We apply the linear transformation $\mathbf{T} : \mathbf{R}^2 \rightarrow \mathbf{R}^2$,

$$\mathbf{T}(r, \theta) = (x, y) = (r \cos \theta, r \sin \theta)$$

The Jacobian $\frac{\partial(x, y)}{\partial(r, \theta)}$ is given by:

$$\begin{aligned} \frac{\partial(x, y)}{\partial(r, \theta)} &= \det \begin{bmatrix} \frac{\partial x}{\partial r} & \frac{\partial x}{\partial \theta} \\ \frac{\partial y}{\partial r} & \frac{\partial y}{\partial \theta} \end{bmatrix} \\ &= \begin{vmatrix} \cos \theta & -r \sin \theta \\ \sin \theta & r \cos \theta \end{vmatrix} \\ &= r(\cos^2 \theta + \sin^2 \theta) \\ &= r \end{aligned}$$

We need to identify the region D^* that T maps in a one-to-one fashion to D . We have:

$$D^* = \{(\theta, r) | 0 \leq \theta \leq 2\pi, 0 \leq r \leq 1\}$$

Thus, D^* is a rectangular region. We can write our double integral as:

$$\begin{aligned} P((X, Y) \in D) &= \frac{1}{2\pi} \int_0^{2\pi} \int_0^1 e^{-r^2/2} r dr d\theta \\ &= \frac{1}{2\pi} \int_0^{2\pi} \int_0^{1/2} e^{-u} du d\theta \\ &= \frac{1}{2\pi} \int_0^{2\pi} -[e^{-u}]_0^{1/2} d\theta \\ &= \frac{1}{2\pi} \int_0^{2\pi} (1 - e^{-1/2}) d\theta \\ &= (1 - e^{-1/2}) \frac{1}{2\pi} \int_0^{2\pi} d\theta \\ &= (1 - e^{-1/2}) \end{aligned}$$

Consider now the random variable $R = (X^2 + Y^2)^{1/2}$ giving the random distance of the point to the origin. Let's compute $\mathbb{E}[R]$. Now, R is a function of the random variables (X, Y) . Hence, by LOTUS, we must have:

$$\begin{aligned} \mathbb{E}[R] &= \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} (x^2 + y^2)^{1/2} f_{(X,Y)}(x, y) dx dy \\ &= \frac{1}{2\pi} \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} (x^2 + y^2)^{1/2} e^{-(x^2+y^2)/2} dx dy \end{aligned}$$

Again by transforming to the polar coordinates, we have:

$$\begin{aligned} \mathbb{E}[R] &= \frac{1}{2\pi} \int_0^{2\pi} \int_0^{\infty} r e^{-r^2/2} r dr d\theta \\ &= \frac{1}{2\pi} \int_0^{2\pi} \int_0^{\infty} r^2 e^{-r^2/2} dr d\theta \end{aligned}$$

By the product rule, the inner integral can be simplified as follows:

$$\begin{array}{c|c} u & dv \\ \hline r & re^{-r^2/2}dr \\ 1 & -e^{-r^2/2} \end{array}$$

We have:

$$\begin{aligned} \int_0^\infty u dv &= uv|_0^\infty - \int_0^\infty v du \\ &= -re^{-r^2/2} \Big|_0^\infty + \int_0^\infty e^{-r^2/2} dr \\ &= 0 + \frac{\sqrt{2\pi}}{2} \end{aligned}$$

So, the desired expectation is:

$$\mathbb{E}[R] = \frac{\sqrt{2\pi}}{2} \cdot \frac{1}{2\pi} \int_0^{2\pi} d\theta = \sqrt{\frac{\pi}{2}}$$

More generally, the CDF of R is given by:

$$\mathbb{P}(R \leq r) = \mathbb{P}((X, Y) \in D)$$

where $D = \{(x, y) | x^2 + y^2 \leq r^2\}$. We know that, the probability of a random vector lying in a domain D is given by:

$$\begin{aligned} \mathbb{P}((X, Y) \in D) &= \int \int_D f_{X,Y}(x, y) dx dy \\ &= \int \int_D \frac{1}{2\pi} e^{-(x^2+y^2)/2} dx dy \end{aligned}$$

Again transforming to the Polar coordinates, we have:

$$\begin{aligned}
\mathbb{P}((X, Y) \in D) &= \frac{1}{2\pi} \int_0^{2\pi} \int_0^r r e^{-r^2/2} dr d\theta \\
&= \frac{1}{2\pi} \int_0^{2\pi} - \left[e^{-r^2/2} \right]_0^r d\theta \\
&= (1 - e^{-r^2/2})
\end{aligned}$$

Taking the derivative with respect to r , we get that the PDF of the R is:

$$f_R(r) = r e^{-r^2/2}$$

Consider now the random angle that the point (X, Y) makes with the x -axis. That is, the random variable $\Theta = \arctan \frac{Y}{X}$. It is not hard to compute the joint PDF of (R, Θ) . Define $D = \{(x, y) | x^2 + y^2 \leq r^2, \frac{y}{x} \leq \tan \theta\}$.

We have:

$$\begin{aligned}
\mathbb{P}(R \leq r, \Theta \leq \theta) &= \int \int_D \frac{1}{2\pi} e^{-(x^2+y^2)/2} dx dy \\
&= \int_0^\theta \int_0^r \frac{1}{2\pi} r e^{-r^2/2} dr d\theta \\
&= \frac{(1 - e^{-r^2/2})\theta}{2\pi}
\end{aligned}$$

And the joint PDF is just $f_{(R, \Theta)}(r, \theta) = \frac{1}{2\pi} r e^{-r^2/2}$. In particular, the variables (R, Θ) are independent since the joint the PDF is the product of the marginals. Θ is uniformly distributed on $[0, 2\pi]$ and has PDF $f_\Theta(\theta) = \frac{1}{2\pi}$.

3.4.3 The Box-Mueller Method.

The above example gives an interesting method to generate a pair of IID standard Gaussian random variables. This is called the Box-Mueller method. Let U_1 and U_2 be two independent uniform random variables on $[0, 1]$. Define the random variables (Z_1, Z_2) as follows:

$$\begin{aligned} Z_1 &= \sqrt{-2 \log U_1} \cos(2\pi U_2) \\ Z_2 &= \sqrt{-2 \log U_1} \sin(2\pi U_2) \end{aligned}$$

The CDF of the random variable R defined above is:

$$u = 1 - e^{-r^2/2}$$

Expressing r in terms of u , we have:

$$\begin{aligned} e^{-r^2/2} &= 1 - u \\ \frac{-r^2}{2} &= \log(1 - u) \\ r^2 &= -2 \log(1 - u) \\ r &= \sqrt{-2 \log(1 - u)} \end{aligned}$$

By probability integral transform, we know that if U'_1 is a Uniform $[0, 1]$ random variable, then the random variable $F_X^{-1}(U'_1)$ has the CDF F_X . By symmetry, $U_1 := 1 - U'_1$ is also uniformly distributed on $[0, 1]$. Thus, the random variable $\sqrt{-2 \log U_1}$ has the same distribution as R .

The CDF of the random variable Θ defined above is:

$$F_\Theta(\theta) = \frac{\theta}{2\pi}$$

So, if U_2 is a uniform random variable, then the random variable $2\pi U_2$ has the same distribution as Θ in the discussion above.

As seen in the example above, if R and Θ are independent and their marginal CDFs are $F_R(r) = 1 - e^{-r^2/2}$ and $F_\Theta(\theta) = \frac{\theta}{2\pi}$, we know that the random variables defined by $X = R \cos \Theta$ and $Y = R \sin \Theta$ are IID standard normal random variables.

More formally, we are making the transformation:

$$T(X, Y) = (R(X, Y), \Theta(X, Y)) = \left(\sqrt{X^2 + Y^2}, \arctan\left(\frac{Y}{X}\right) \right)$$

So, the density function of the pair (X, Y) is given :

$$\begin{aligned} f_{(X,Y)}(x, y) &= f_{(R,\Theta)}(r, \theta) \cdot \left| \begin{array}{cc} \frac{\partial r}{\partial x} & \frac{\partial r}{\partial y} \\ \frac{\partial \theta}{\partial x} & \frac{\partial \theta}{\partial y} \end{array} \right| \\ &= \frac{1}{2\pi} r e^{-r^2/2} \left| \begin{array}{cc} \frac{x}{\sqrt{x^2+y^2}} & \frac{y}{\sqrt{x^2+y^2}} \\ \frac{-y}{x^2+y^2} & \frac{x}{x^2+y^2} \end{array} \right| \\ &= \frac{1}{2\pi} \sqrt{x^2 + y^2} \cdot e^{-(x^2+y^2)/2} \cdot \frac{1}{\sqrt{x^2 + y^2}} \\ &= \frac{1}{2\pi} e^{-(x^2+y^2)/2} \end{aligned}$$

Hence, X and Y are IID standard Gaussian random variables.

3.5 Gaussian Vectors.

Definition 3.17. A n -dimensional random vector $\mathbf{X} = (X_1, X_2, \dots, X_n)$ is said to be jointly Gaussian if and only if for all real vectors $\mathbf{t} = (t_1, \dots, t_n)$, the linear combination $\mathbf{t}^T \mathbf{X} = t_1 X_1 + t_2 X_2 + \dots + t_n X_n$ of (X_1, X_2, \dots, X_n) is a Gaussian random variable.

As a simple consequence of the above definition, if (X_1, \dots, X_n) is Gaussian, then setting $t_i = 1$ and $t_j = 0$ for all $i \neq j$, we have that each X_i is also Gaussian.

An equivalent definition can also be stated in terms of the joint MGF since an MGF uniquely characterizes the distribution of a random variable. Before introducing the second definition, we first make two important observations about the mean and variance of a linear combination of random variables.

First, the mean of a linear combination of random variables is:

$$\mathbb{E}[a_1 X_1 + a_2 X_2 + \dots + a_n X_n] = a_1 \mathbb{E}X_1 + \dots + a_n \mathbb{E}X_n = \mathbf{a}^T \mathbb{E}\mathbf{X}$$

where $\mathbb{E}\mathbf{X}$ is the mean vector. The variance is obtained with a short calculation using the linearity of expectations:

$$\begin{aligned} \text{Var}(a_1X_1 + \dots + a_nX_n) &= \sum_{i=1}^n \sum_{j=1}^n a_i a_j \text{Cov}(X_i, X_j) \\ &= \mathbf{a}^T \Sigma \mathbf{a} \end{aligned}$$

where Σ is the covariance matrix of \mathbf{X} .

Proposition 3.4. *A random vector $\mathbf{X} = (X_1, X_2, \dots, X_n)$ is Gaussian if and only if the moment generating function of \mathbf{X} is:*

$$\mathbb{E} [\exp \{ \mathbf{t}^T \mathbf{X} \}] = \exp \left[\mathbf{t}^T \boldsymbol{\mu} + \frac{1}{2} \mathbf{t}^T \Sigma \mathbf{t} \right] \quad (3.47)$$

where $\boldsymbol{\mu}$ is the mean vector and Σ is the covariance matrix of \mathbf{X} .

Proof. By the definition of joint MGF:

$$M_{\mathbf{X}}(\mathbf{t}) = \mathbb{E} [\exp \{ \mathbf{t}^T \mathbf{X} \}] = \mathbb{E} [\exp \{ t_1 X_1 + \dots + t_n X_n \}] \quad (3.48)$$

But, we know that $t_1 X_1 + \dots + t_n X_n$ is a Gaussian random variable with mean $\mu = \mathbf{t}^T \boldsymbol{\mu}$ and variance $\sigma^2 = \mathbf{t}^T \Sigma \mathbf{t}$.

The MGF of a univariate Gaussian random variable is :

$$M_X(s) = \mathbb{E}[\exp(sX)] = \exp\left(\mu s + \frac{\sigma^2 s^2}{2}\right)$$

At $s = 1$, we have:

$$M_X(1) = \mathbb{E}[\exp(X)] = \exp\left(\mu + \frac{\sigma^2}{2}\right) \quad (3.49)$$

Thus, if $X = t_1 X_1 + \dots + t_n X_n$ then it follows that:

$$\mathbb{E}[\exp(t_1 X_1 + \dots + t_n X_n)] = \exp\left[\mathbf{t}^T \boldsymbol{\mu} + \frac{1}{2} \mathbf{t}^T \Sigma \mathbf{t}\right]$$

But from (3.48), this is the joint MGF of \mathbf{X} . This closes the proof. \square

Proposition 3.5. *Let $\mathbf{X} = (X_1, \dots, X_n)$ be a Gaussian vector. Then, the covariance matrix is diagonal, if and only if the random variables are independent.*

Proof. (\Rightarrow) direction.

We are given that the covariance matrix is diagonal. Our proposition is that the random variables are independent.

Remember, that if X_1 and X_2 are independent random variables, $Cov(X_1, X_2) = 0$. But, the converse is not true. We use the MGF of the random vector \mathbf{X} , to prove this claim.

We have:

$$M_{\mathbf{X}}(\mathbf{t}) = \exp\left[\mathbf{t}^T \boldsymbol{\mu} + \frac{1}{2} \mathbf{t}^T \Sigma \mathbf{t}\right]$$

Since $\Sigma = \text{Diag}(\sigma_1^2, \dots, \sigma_n^2)$, we can express:

$$\mathbf{t}^T \Sigma \mathbf{t} = t_1^2 \sigma_1^2 + t_2^2 \sigma_2^2 + \dots + t_n^2 \sigma_n^2$$

So:

$$\begin{aligned} M_{\mathbf{X}}(\mathbf{t}) &= \exp\left[t_1 \mu_1 + \frac{\sigma_1^2 t_1^2}{2}\right] \cdots \exp\left[t_n \mu_n + \frac{\sigma_n^2 t_n^2}{2}\right] \\ &= M_{X_1}(t_1) \cdots M_{X_n}(t_n) \end{aligned}$$

Consequently, the MGF can be factored into a product of the MGFs of X_1, \dots, X_n . Thus, X_1, X_2, \dots, X_n are independent random variables.

(\Leftarrow) direction.

This direction is trivial. We are given that the random variables are independent. Then, $Cov(X_i, X_j) = 0$ for all $i \neq j$. So, the covariance matrix is diagonal. \square

Before writing the joint PDF of a Gaussian vector in terms of the mean vector and the covariance matrix, we need to introduce the important notion of degenerate vector. We say a Gaussian vector is *degenerate* if its covariance matrix Σ is singular, $\det \Sigma = 0$.

Example 3.8. Consider (Z_1, Z_2, Z_3) IID standard Gaussian random variables. We define $X = Z_1 + Z_2 + Z_3$, $Y = Z_1 + Z_2$ and $W = Z_3$. Clearly, (X, Y, W) is a Gaussian vector. It has 0 mean and covariance:

$$\begin{bmatrix} 3 & 2 & 1 \\ 2 & 2 & 0 \\ 1 & 0 & 1 \end{bmatrix}$$

It is easy to check that $\det \Sigma = 3 \cdot 2 - 2 \cdot 2 + 1 \cdot (-2) = 0$. Thus, (X, Y, W) is a degenerate Gaussian vector.

The above example is helpful to illustrate the notion. Note that we have the linear relation $X - Y - W = 0$ between the random variables. Therefore, the random variables are linearly dependent. In other words, one vector is redundant, say X , in the sense that its value can be recovered from others for any outcome. The relation between degeneracy and linear dependence is general.

Lemma 3.6. *Let $\mathbf{X} = (X_1, X_2, \dots, X_n)$ be a Gaussian vector. Then, X is degenerate if and only if the coordinates are linearly dependent. That is, there exists c_1, c_2, \dots, c_n , not all zero, such that $c_1 X_1 + c_2 X_2 + \dots + c_n X_n = 0$ with probability one.*

Proof. (\Rightarrow) direction.

We are given that the vector X is degenerate. This implies that $\det \Sigma = 0$ and the columns of Σ are linearly dependent. Σ is non-singular.

TODO. □

We are now ready to state the form of the PDF of Gaussian vectors.

Definition 3.18. (Joint PDF of Gaussian vectors). Let $\mathbf{X} = (X_1, X_2, \dots, X_n)$ be a non-degenerate Gaussian vector with mean vector $\boldsymbol{\mu}$ and covariance matrix Σ , written $N(\boldsymbol{\mu}, \Sigma)$. Then the joint density of X is given by the PDF:

$$f(x_1, \dots, x_n) = \frac{1}{\sqrt{(2\pi)^n |\det \Sigma|}} \exp \left(-\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu})^T \Sigma^{-1} (\mathbf{x} - \boldsymbol{\mu}) \right) \quad (3.50)$$

where $\mathbf{x} \in \mathbf{R}^n$ and Σ is PSD (Positive symmetric definite).

Example 3.9. Consider a Gaussian vector (X_1, X_2) of mean 0 and covariance matrix $\Sigma = \begin{bmatrix} 2 & 1 \\ 1 & 2 \end{bmatrix}$. The inverse of Σ can be found out as follows.

We consider the augmented matrix $[\Sigma|I]$.

$$\left[\begin{array}{cc|cc} 2 & 1 & 1 & 0 \\ 1 & 2 & 0 & 1 \end{array} \right]$$

Performing $R_1 = 1/2R_1$, the above system is row equivalent to:

$$\left[\begin{array}{cc|cc} 1 & \frac{1}{2} & \frac{1}{2} & 0 \\ 1 & 2 & 0 & 1 \end{array} \right]$$

Performing $R_2 = R_2 - R_1$, the above system is row equivalent to:

$$\left[\begin{array}{cc|cc} 1 & \frac{1}{2} & \frac{1}{2} & 0 \\ 0 & \frac{3}{2} & -\frac{1}{2} & 1 \end{array} \right]$$

Performing $R_2 = \frac{2}{3}R_2$, the above system is row equivalent to:

$$\left[\begin{array}{cc|cc} 1 & \frac{1}{2} & \frac{1}{2} & 0 \\ 0 & 1 & -\frac{1}{3} & \frac{2}{3} \end{array} \right]$$

Performing $R_1 = R_1 - \frac{1}{2}R_2$, the above system is row equivalent to

$$\left[\begin{array}{cc|cc} 1 & 0 & \frac{2}{3} & -\frac{1}{3} \\ 0 & 1 & -\frac{1}{3} & \frac{2}{3} \end{array} \right]$$

So, $\Sigma^{-1} = \begin{bmatrix} 2/3 & -1/3 \\ -1/3 & 2/3 \end{bmatrix}$ and $\det C = 3$. By doing matrix operations, the joint PDF of (X_1, X_2) is:

$$\begin{aligned} f_{(X_1, X_2)}(x_1, x_2) &= \frac{1}{\sqrt{(2\pi)^2 \cdot 3}} \exp \left(-\frac{1}{2} \begin{bmatrix} x_1 & x_2 \end{bmatrix} \begin{bmatrix} 2/3 & -1/3 \\ -1/3 & 2/3 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \right) \\ &= \frac{1}{2\pi\sqrt{3}} \exp \left(-\frac{1}{2} \begin{bmatrix} 2/3x_1 - 1/3x_2 & -1/3x_1 + 2/3x_2 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \right) \\ &= \frac{1}{2\pi\sqrt{3}} \exp \left(-\frac{1}{3}x_1^2 + \frac{1}{3}x_1x_2 - \frac{1}{3}x_2^2 \right) \end{aligned}$$

We will not prove proposition (3.18) yet. Instead, we will take a short detour and derive it from a powerful decomposition of Gaussian vectors as a linear combination of IID Gaussians. The decomposition is the generalization of making a random variable *standard*. Suppose X is Gaussian with mean 0 and variance σ^2 . Then, we can write it as $X = \sigma Z$, where Z is a standard normal random variable. (This makes sense even when X is degenerate that is $\sigma^2 = 0$). If $\sigma^2 \neq 0$, then we can reverse the relation to get:

$$Z = \frac{X}{\sigma}$$

We generalize this procedure to Gaussian vectors.

Proposition 3.6. (*Decomposition into IID*). Let $\mathbf{X} = (X_1, X_2, \dots, X_n)$ be a Gaussian vector of mean 0 and $n \times n$ covariance matrix C . If \mathbf{X} is non-degenerate, there exists n IID gaussian random variables Z_1, Z_2, \dots, Z_n and an invertible $n \times n$ matrix A such that:

$$\mathbf{X} = AZ, \quad Z = A^{-1}\mathbf{X} \quad (3.51)$$

The choice of Z s and thus the matrix A is generally not unique as the following simple example shows:

Example 3.10. Consider the Gaussian vector (X_1, X_2) given by:

$$\begin{aligned} X_1 &= Z_1 + Z_2 \\ X_2 &= Z_1 - Z_2 \end{aligned}$$

where Z_1, Z_2 are IID standard gaussians.

The matrix $A = \begin{bmatrix} 1 & 1 \\ 1 & -1 \end{bmatrix}$. The covariance matrix of \mathbf{X} is $\begin{bmatrix} 2 & 0 \\ 0 & 2 \end{bmatrix}$. Since, the covariance matrix is diagonal, by proposition (3.5), the random variables X_1 and X_2 are independent.

Another choice of decomposition is simply $W_1 = X_1/\sqrt{2}$ and $W_2 = X_2/\sqrt{2}$.

Proof of proposition 3.6.

Proof. This is done using the same Gram-Schmidt procedure as for \mathbf{R}^n . The idea is to take the variables one-by-one and subtract the components in the directions of the previous ones using covariance. The lemma (3.6) ensures that no random variables are linear combinations of the others.

To start, we take $Z_1 = \frac{X_1}{\sqrt{C_{11}}}$. Clearly, Z_1 is a standard normal random variable.

Then, we define Z'_2 as:

$$Z'_2 = X_2 - \mathbb{E}[X_2 Z_1] Z_1$$

And let

$$Z_2 = \frac{Z'_2}{\sqrt{\text{Var}(Z'_2)}}$$

Firstly, since X_2 and Z_1 are Gaussian random variables, it follows that Z_2 is also a Gaussian random variable. Moreover, $\mathbb{E}[Z_2] = \frac{1}{\sqrt{\text{Var}(Z'_2)}} \cdot \mathbb{E}[X_2] = 0$ and $\text{Var}(Z_2) = 1$. Further:

$$\begin{aligned} \text{Cov}(Z_1, Z'_2) &= \text{Cov}(Z_1, X_2 - \mathbb{E}[X_2 Z_1] Z_1) \\ &= \text{Cov}(Z_1, X_2) - \mathbb{E}[X_2 Z_1] \text{Var}(Z_1) \\ &= \mathbb{E}[X_2 Z_1] - \mathbb{E}[X_2 Z_1] (1) \\ &= 0 \end{aligned}$$

Thus, Z_2 is independent Gaussian with mean 0 and variance 1.

In the same way, we take Z_3 to be:

$$Z'_3 = X_3 - \mathbb{E}(X_3, Z_2) Z_2 - \mathbb{E}(X_3, Z_1) Z_1$$

and

$$Z_3 = \frac{Z'_3}{\sqrt{\text{Var}(Z'_3)}}$$

Again, its easy to check that Z'_3 is independent of Z_2 and Z_1 . As above, we define Z_3 to be Z'_3 divided by the square root of variance. This procedure is carried on until we run out variables. Note that since C is non-degenerate, none of the variances of the Z'_i will be zero, and therefore they can be standardized. \square

The covariance matrix C of the Gaussian vector \mathbf{X} with mean vector $\boldsymbol{\mu} = \mathbf{0}$ can be written in terms of A . Write $A = (a_{ij})$ for the (i, j) th entry of the matrix A . By the relation $X = AZ$, we have:

$$\begin{aligned} \text{Cov}(X_i, X_j) &= \mathbb{E}(X_i X_j) \\ &= \mathbb{E} \left[\left(\sum_{k=1}^n a_{ik} Z_k \right) \left(\sum_{l=1}^n a_{jl} Z_l \right) \right] \\ &= \mathbb{E} \left[\sum_{k=1}^n a_{ik} a_{jk} Z_k^2 + \sum_{k \neq l} a_{ik} a_{jl} Z_k Z_l \right] \end{aligned}$$

Now, we know that:

$$\mathbb{E}(Z_k \cdot Z_l) = \begin{cases} 1 & \text{if } k = l \\ 0 & \text{otherwise} \end{cases}$$

So, the expectation simplifies to:

$$\begin{aligned} \text{Cov}(X_i, X_j) &= \mathbb{E} \left[\sum_{k=1}^n a_{ik} a_{jk} Z_k^2 \right] \\ &= \sum_{k=1}^n a_{ik} a_{jk} \\ &= (AA^T)_{ij} \end{aligned}$$

Thus, we have:

$$C = AA^T$$

Thus, the covariance matrix C of a Gaussian vector \mathbf{X} admits a Cholesky Factorization of the form, $C = AA^T$ and therefore, C is SPD (symmetric positive definite). For applications and numerical simulations, it is important to get the matrix A from the covariance matrix C . This decomposition is an exact analogue of the decomposition of a vector in \mathbf{R}^3 written as a sum of orthonormal basis vectors. In particular, the condition of being non-degenerate is equivalent to linear independence.

Proof of proposition 3.18

Proof. Without the loss of generality assume that $\boldsymbol{\mu} = (0, \dots, 0)$. Otherwise, we just need to subtract it from \mathbf{X} . We use the decomposition in proposition (3.6). First note that, since $C = AA^T$, the determinant of C is:

$$C = AA^T$$

so the determinant of C of C is:

$$\begin{aligned} \det C &= \det A \cdot \det A^T \\ &= \det A \cdot \det A \\ &= (\det A)^2 \end{aligned}$$

In particular, since \mathbf{X} is non-degenerate, we have that $\det C \neq 0$, so $\det A \neq 0$. Thus, A is invertible. We also have by the decomposition that there exist IID Gaussian random variables Z such that $\mathbf{X} = AZ$. Now, the event $\{\mathbf{X} \in B\} = \{AZ \in B\} = \{Z \in A^{-1}B\}$. So,

$$P(X \in B) = P(Z \in A^{-1}B)$$

But we know the joint density of n IID standard normal random variables Z_1, Z_2, \dots, Z_n . Consequently, we have:

$$P(X \in B) = \int \cdots \int_{A^{-1}B} \frac{1}{(2\pi)^{n/2}} \exp \left[-\frac{1}{2} \mathbf{z}^T \mathbf{z} \right] dz_1 \cdots dz_n$$

because $\mathbf{z} = (z_1, \dots, z_n)$ and $\mathbf{z}^T \mathbf{z} = z_1^2 + \dots + z_n^2$. It remains to do the change of variable $\mathbf{x} = A\mathbf{z}$.

Let us define the map T as:

$$\mathbf{x} = A\mathbf{z}$$

Then, the inverse map T^{-1} is:

$$\mathbf{z} = A^{-1}\mathbf{x}$$

Since \mathbf{X} is non-degenerate, A^{-1} exists and the right-hand side vector is well-defined. The Jacobian $\frac{\partial(z_1, \dots, z_n)}{\partial(x_1, \dots, x_n)}$ is:

$$\frac{\partial(z_1, \dots, z_n)}{\partial(x_1, \dots, x_n)} = |\det(A^{-1})| = \frac{1}{|\det A|} = \frac{1}{\sqrt{|\det C|}}$$

Moreover, $\mathbf{z} \in A^{-1}B$ is equivalent to $\mathbf{x} \in B$. Further, $\mathbf{z}^T \mathbf{z} = (A^{-1}\mathbf{x})^T (A^{-1}\mathbf{x}) = \mathbf{x}^T (A^{-1})^T (A^{-1}) \mathbf{x}$. Now, note that if $C = AA^T$, by the reverse order law, $C^{-1} = (A^T)^{-1} (A^{-1}) = (A^{-1})^T (A^{-1})$. Consequently, $\mathbf{z}^T \mathbf{z} = \mathbf{x}^T C^{-1} \mathbf{x}$.

$$P(X \in B) = \int \cdots \int_B \frac{1}{\sqrt{(2\pi)^n |\det C|}} \exp \left[-\frac{1}{2} \mathbf{x}^T C^{-1} \mathbf{x} \right] dx_1 \cdots dx_n$$

Consequently, the joint density function of \mathbf{X} is

$$f_{\mathbf{X}}(\mathbf{x}) = \frac{1}{\sqrt{(2\pi)^n |\det C|}} \exp \left[-\frac{1}{2} \mathbf{x}^T C^{-1} \mathbf{x} \right]$$

If \mathbf{X} has a non-zero mean vector $\boldsymbol{\mu}$, then $\mathbf{X}' = \mathbf{X} - \boldsymbol{\mu}$ has a mean vector zero. Thus, the joint density function becomes:

$$f_{\mathbf{X}}(\mathbf{x}) = \frac{1}{\sqrt{(2\pi)^n |\det C|}} \exp \left[-\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu})^T C^{-1} (\mathbf{x} - \boldsymbol{\mu}) \right]$$

□

We now explore three ways to find the matrix A in the decomposition of Gaussian vectors of proposition (3.6). We proceed by example:

Example 3.11. (Cholesky by Gram-Schmidt). This is the method suggested by the proof of proposition 3.6. It suffices to successively go through the X 's by subtracting the projection of a given X_i onto the previous random variables. Consider the random vector $\mathbf{X} = (X_1, X_2)$ with mean 0 and covariance matrix :

$$C = \begin{bmatrix} 2 & 1 \\ 1 & 2 \end{bmatrix}$$

It is easy to check that X is non-degenerate, $\det C = 3$. Take:

$$Z_1 = \frac{X_1}{\sqrt{2}}$$

This is obviously a standard Gaussian random variable. For Z_2 , first consider:

$$Z'_2 = X_2 - \mathbb{E}(X_2 Z_1) Z_1$$

It is straightforward to check that Z_1 and Z'_2 are jointly Gaussian. Z'_2 is a linear combination of Z_1 and X_2 , so Z'_2 is Gaussian. Since all linear combinations of Z'_2 and Z_1 are Gaussian, by definition, (Z_1, Z'_2) is jointly Gaussian. They are also independent, because:

$$\begin{aligned}
\mathbb{E}(Z_1 Z'_2) &= \mathbb{E}[Z_1(X_2 - \mathbb{E}(X_2 Z_1)Z_1)] \\
&= \mathbb{E}[Z_1 X_2] - \mathbb{E}(X_2 Z_1)\mathbb{E}(Z_1^2) \\
&= \mathbb{E}[Z_1 X_2] - \mathbb{E}(X_2 Z_1) \cdot 1 \\
&= 0
\end{aligned}$$

Note that:

$$\begin{aligned}
Z'_2 &= X_2 - \mathbb{E}[X_2 Z_1]Z_1 \\
&= X_2 - \mathbb{E}\left[X_2 \frac{X_1}{\sqrt{2}}\right] \frac{X_1}{\sqrt{2}} \\
&= X_2 - \frac{1}{2}\mathbb{E}[X_1 X_2]X_1 \\
&= X_2 - \frac{1}{2}X_1
\end{aligned}$$

In particular, we have by linearity of expectations:

$$\begin{aligned}
\mathbb{E}[(Z'_2)^2] &= \mathbb{E}[X_2^2 - X_1 X_2 + \frac{1}{4}X_1^2] \\
&= \mathbb{E}(X_2^2) - \mathbb{E}(X_1 X_2) + \frac{1}{4}\mathbb{E}(X_1^2) \\
&= 2 - 1 + \frac{1}{4} \cdot 2 \\
&= \frac{3}{2}
\end{aligned}$$

To get a random variable of variance 1, that is a multiple of Z'_2 , we take $Z_2 = \frac{Z'_2}{\sqrt{3/2}} = \sqrt{\frac{2}{3}}Z'_2 = \sqrt{\frac{2}{3}}X_2 - \frac{1}{\sqrt{6}}X_1$. Altogether, we get :

$$\begin{aligned}
Z_1 &= \frac{X_1}{\sqrt{2}} \\
Z_2 &= -\frac{1}{\sqrt{6}}X_1 + \sqrt{\frac{2}{3}}X_2
\end{aligned}$$

We thus constructed two standard IID Gaussians from (X_1, X_2) . In particular we have:

$$A^{-1} = \begin{bmatrix} \frac{1}{\sqrt{2}} & 0 \\ -\frac{1}{\sqrt{6}} & \sqrt{\frac{2}{3}} \end{bmatrix}, \quad A = \begin{bmatrix} \sqrt{2} & 0 \\ \frac{1}{\sqrt{2}} & \sqrt{\frac{3}{2}} \end{bmatrix}$$

We can check that :

$$\begin{aligned} AA^T &= \begin{bmatrix} \sqrt{2} & 0 \\ \frac{1}{\sqrt{2}} & \sqrt{\frac{3}{2}} \end{bmatrix} \begin{bmatrix} \sqrt{2} & \frac{1}{\sqrt{2}} \\ 0 & \sqrt{\frac{3}{2}} \end{bmatrix} \\ &= \begin{bmatrix} 2 & 1 \\ 1 & 2 \end{bmatrix} \\ &= C \end{aligned}$$

The probability of the event $P(X_1 > 2, X_2 < 3)$ can be computed as follows:

$$P(X_1 > 2, X_2 < 3) = P\left(\sqrt{2}Z_1 > 2, \frac{1}{\sqrt{2}}Z_1 + \sqrt{\frac{3}{2}}Z_2 < 3\right)$$

Example 3.12. (Cholesky by solving a system of equations). Consider the same example as above. Write $A = \begin{bmatrix} a & b \\ c & d \end{bmatrix}$. Then the relation $C = AA^T$ yields:

$$\begin{bmatrix} a & b \\ c & d \end{bmatrix} \begin{bmatrix} a & c \\ b & d \end{bmatrix} = \begin{bmatrix} a^2 + b^2 & ac + bd \\ ac + bd & c^2 + d^2 \end{bmatrix} = \begin{bmatrix} 2 & 1 \\ 1 & 2 \end{bmatrix}$$

and so we have the three equations:

$$\begin{aligned} a^2 + b^2 &= 2 \\ ac + bd &= 1 \\ c^2 + d^2 &= 2 \end{aligned}$$

There are several solutions. One of them is $a = \sqrt{2}, b = 0, c = \frac{1}{\sqrt{2}}$ and $d = \sqrt{\frac{3}{2}}$.

Example 3.13. (*Cholesky by diagonalization*) This method takes advantage of the symmetry of the covariance matrix. From the spectral theorem, we know that, if C is a symmetric matrix, it is diagonalizable, it admits a factorization of the form $Q\Lambda Q^{-1}$ where Q is an orthogonal matrix. The entries of Λ are the eigenvalues of C . Furthermore, the eigenvectors are orthogonal to each other.

Since $C = AA^T$, we get:

$$\begin{aligned} C &= Q\Lambda Q^T \\ \iff AA^T &= Q\Lambda Q^T \end{aligned}$$

It suffices to take:

$$A = Q\Lambda^{1/2}$$

where Q is the matrix with the columns given by the eigenvectors of C and $\Lambda^{1/2}$ is the diagonal matrix with the square root of the eigenvalues on the diagonal.

Example 3.14. (IID Decomposition). Let $X = (X_1, X_2, X_3)$ be a Gaussian vector with mean 0 and covariance matrix:

$$C = \begin{bmatrix} 1 & 1 & 1 \\ 1 & 2 & 2 \\ 1 & 2 & 3 \end{bmatrix}$$

Let's find a matrix A such that $X = AZ$ for $Z = (Z_1, Z_2, Z_3)$ IID standard Gaussians. The vector is not degenerate since $\det C = 1 \cdot (2 - 1) = 1$. If we do a Gram-Schmidt procedure, we get:

$$\begin{aligned} Z_1 &= X_1 \\ Z_2 &= (X_2 - X_1) \\ Z_3 &= X_3 - (X_2 - X_1) - X_1 \\ &= X_3 - X_2 \end{aligned}$$

Consequently, $X_1 = Z_1$, $X_2 = Z_1 + Z_2$ and $X_3 = Z_1 + Z_2 + Z_3$. So, the matrix A is:

$$A = \begin{bmatrix} 1 & 0 & 0 \\ 1 & 1 & 0 \\ 1 & 1 & 1 \end{bmatrix}$$

As we will see in the next section, this random vector corresponds to the position of the Brownian motion at time 1, 2 and 3.

3.6 Gaussian Processes.

In general, a *stochastic process* is an infinite collection of random variables on a probability space $(\Omega, \mathcal{F}, \mathbb{P})$. The collection can be countable or uncountable. We are mostly interested in the case, where the variables are indexed by time; for example

$$X = (X_t, t \in \mathcal{J})$$

where \mathcal{J} can be a finite set, \mathbf{N} or some uncountable set such as the closed interval $[0, T]$ or $[0, \infty)$.

In the case where $\mathcal{J} = [0, \infty)$ or $[0, T]$, the realization of the process $X(\omega)$ can be thought of as a function of time for each outcome ω . This function $t \mapsto X_t(\omega)$ is sometimes called a *path* or the *trajectory* of the process. With this in mind, we can think of the process $(X)_{t \geq 0}$ as a function-valued random variable as each outcome ω produces a function.

(a) For each t , $X(t, \cdot)$ is a random variable.

(b) For each ω , $X(\cdot, \omega)$ is a *function* (called a *sample path*)

For convenience, the random variable $X(t, \cdot)$ will be written as $X(t)$ or X_t . Thus a stochastic process $X(t, \omega)$ can also be expressed as $(X(t))_{t \geq 0}$ or simply $X(t)$.

How can we compute the probabilities for a stochastic process? In other words, what object captures its distribution? The most common way (there are others) is to use finite dimensional distributions. The idea here is to describe the probabilities related

to any finite set of time. More precisely, the finite-dimensional distributions are given by:

$$\mathbb{P}(X_{t_1} \in B_1, X_{t_2} \in B_2, \dots, X_{t_n} \in B_n)$$

for any $n \in \mathbf{N}$, any choice of $t_1, \dots, t_n \in \mathcal{J}$, and any events B_1, \dots, B_n in \mathbf{R} . Of course, for any fixed choice of t 's $(X_{t_1}, \dots, X_{t_n})$ is a random vector as seen in the previous section. The fact that we can control the probabilities for the whole random function comes from the fact that we have the distributions of these vectors for any n and any choice of t 's.

Some important types of stochastic processes include Markov processes, martingales and Gaussian processes. We will encounter them along the way. Let's start with Gaussian processes.

Definition 3.19. A *Gaussian process* $(X_t)_{t \geq 0}$ is a stochastic process whose finite dimensional distributions are jointly Gaussian. In other words, for any $n \in \mathbf{N}$ and any choice of $t_1 < \dots < t_n$ we have that $(X_{t_1}, X_{t_2}, \dots, X_{t_n})$ is a Gaussian vector. In particular, its distribution is defined by the mean function $m(t) = \mathbb{E}(X_t)$ and the covariance function $C(s, t) = \text{Cov}(X_t, X_s)$.

As before, linear combinations of Gaussian processes remain Gaussian.

Lemma 3.7. Let $X^{(1)}, X^{(2)}, \dots, X^{(m)}$ be m Gaussian processes on $[0, \infty)$ defined on the same probability space. Then, any process constructed by taking linear combinations is also a Gaussian process:

$$a_1 X^{(1)} + \dots + a_m X^{(m)} = \left(a_1 X_t^{(1)} + \dots + a_m X_t^{(m)}, t \geq 0 \right)$$

Proof. It suffices to take a finite set of times $t_1 < t_2 < \dots < t_n$. Let $\mathcal{J} = \{t_1, \dots, t_n\}$

By definition, the random vector $\mathbf{X}_{\mathcal{J}}^{(i)}$ is Gaussian.

Any linear combination of Gaussian vectors is a Gaussian vector, so

$$\begin{aligned}
\mathbf{a}^T \begin{bmatrix} \mathbf{X}_{\mathcal{J}}^{(1)} \\ \mathbf{X}_{\mathcal{J}}^{(2)} \\ \vdots \\ \mathbf{X}_{\mathcal{J}}^{(m)} \end{bmatrix} &= a_1 \mathbf{X}_{\mathcal{J}}^{(1)} + \dots + a_m \mathbf{X}_{\mathcal{J}}^{(m)} \\
&= a_1 \begin{bmatrix} X_{t_1}^{(1)} \\ X_{t_2}^{(1)} \\ \vdots \\ X_{t_n}^{(1)} \end{bmatrix} + \dots + a_m \begin{bmatrix} X_{t_1}^{(m)} \\ X_{t_2}^{(m)} \\ \vdots \\ X_{t_n}^{(m)} \end{bmatrix} \\
&= \begin{bmatrix} \sum_{k=1}^m a_k X_{t_1}^{(k)} \\ \sum_{k=1}^m a_k X_{t_2}^{(k)} \\ \vdots \\ \sum_{k=1}^m a_k X_{t_n}^{(k)} \end{bmatrix}
\end{aligned}$$

is also a Gaussian vector for any t_1, \dots, t_n . Hence, any linear combination of Gaussian processes is a Gaussian process. \square

The most important example of a Gaussian process is Brownian motion.

Definition 3.20. (*Standard Brownian motion or Wiener process*). A stochastic process $B(t, \omega)$ is called a Brownian motion if it satisfies the following conditions:

- $\mathbb{P}\{\omega : B(0, \omega) = 0\} = 1$.
- For any $0 \leq s < t$, the random variable $B(t) - B(s)$ is normally distributed with mean 0 and variance $t - s$. That is for any $a < b$:

$$\mathbb{P}\{a \leq B(t) - B(s) \leq b\} = \frac{1}{\sqrt{2\pi(t-s)}} \int_a^b e^{-\frac{x^2}{2(t-s)}} dx$$

- $B(t, \omega)$ has independent increments, i.e. for any $0 \leq t_1 < t_2 < \dots < t_n$, the random variables :

$$B(t_1), B(t_2) - B(t_1), B(t_3) - B(t_2), \dots, B(t_n) - B(t_{n-1})$$

are independent.

- Almost all sample paths of $B(t, \omega)$ are continuous functions, that is,

$$P(\omega \mid B(\cdot, \omega) \text{ is continuous}) = 1$$

Example 3.15. (Sampling a Gaussian process using Cholesky decomposition). The IID decomposition of proposition 3.6 is useful for generating a sample of the Gaussian process. $(X_t)_{t \in [0, T]}$. First, we need to fix the discretization or step-size. Take for example, a step size of 0.01, meaning we approximate the process by evaluating the position at every 0.01. This is given by the Gaussian vector:

$$(X_{\frac{j}{100}}, j = 1, 2, 3, \dots, 100T)$$

This Gaussian vector has covariance matrix C and a matrix A from the IID decomposition. Note that, we start with the vector at 0.01 and not 0. This is because in some cases (like the standard Brownian motion) the value at time 0 is 0. Including it in the covariance matrix would result in a degenerate covariance matrix. You can always add position 0 at time 0 after performing the cholesky decomposition. It then suffices to sample $100T$ IID standard Gaussian random variable $Z = (Z_1, Z_2, \dots, Z_{100T})$ and to apply the deterministic matrix A to the sample vector to get :

$$(X_{\frac{j}{100}}, j = 1, 2, \dots, 100T) = AZ$$

Example 3.16. Simulating Brownian Motion. The goal of this project is to simulate 100 paths of Brownian motion on $[0, 1]$ using a step-size of 0.01 using the Cholesky decomposition.

- (a) Construct the covariance matrix of $(B_{j/100})_{1 \leq j \leq 100}$ using a for-loop. Recall that for a Brownian motion $C(s, t) = s \wedge t$ with mean 0.
- (b) The command `numpy.linalg.cholesky` in Python gives the Cholesky decomposition of the covariance matrix C . Use this to find the matrix A .
- (c) Define a function whose output is a sample of N standard Gaussian random variables and whose input is N .
- (d) Use the above to plot $n = 100$ paths of the Brownian motion on $[0, 1]$ with a step size of 0.01. Do not forget B_0 !

Solution.

Listing 3: Generating 100 paths of a standard brownian motion

```
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

sns.set_style("whitegrid")

# A generator for N standard gaussian random variables
def standardNormalGenerator(N):
    return np.random.standard_normal(N)

# Produces 1 sample (path) of a gaussian process
# N : Number of time-steps
# A : The transformation that maps IID gaussians (Z_1,Z_2,...,
      Z_N) to a gaussian vector (X_1,X_2,...,X_N)
# with covariance matrix C = AA'
def sampleGaussianProcess(A,N):

    Z = standardNormalGenerator(N)
    X = np.matmul(A,Z)
    return X

# Produces `numOfPaths` paths of a standard brownian motion
# N : Number of time-steps, 1/N : step-size
def standardBrownianMotion(numOfPaths,N):
    C = np.zeros((N,N))

    for i in range(N):
        for j in range(N):
            s = (i+1)/N
            t = (j+1)/N

            C[i][j] = np.min([s,t])

    A = np.linalg.cholesky(C)

    B = []
    for i in range(numOfPaths):
        X = sampleGaussianProcess(A,N)
        X = np.concatenate([0], X), axis=0)
        B.append(X)

    return B

if __name__ == "__main__":
```

```

T = 1.0
N = 100

C = covarMatrix(N)
B = standardBrownianMotion(numOfPaths=100, covarianceMatrix=C
                             ,N=100)


plt.xlabel(r'$t$')
plt.ylabel(r'$B(t,\omega)$')
plt.grid(True)
plt.title(r'$100$ sample paths of a standard brownian motion
')

t = np.linspace(start=0, stop=1.0, num=101)
for n in range(100):
    plt.plot(t, B[n])

plt.show()

```

100 sample paths of a standard brownian motion



Example 3.17. (*Brownian motion with a drift.*) For $\sigma > 0$ (called the volatility or diffusion coefficient) and $\mu \in \mathbf{R}$ (called the drift), we define the process:

$$X_t = \sigma B_t + \mu t$$

where $(B_t)_{t \geq 0}$ is a standard brownian motion. This is a Gaussian process, because it is a linear transformation of a brownian motion, which is itself a Gaussian process by lemma (3.7).

A straightfoward computation shows that:

$$\begin{aligned}\mathbb{E}[X_t] &= \sigma \mathbb{E}[B_t] + \mathbb{E}[\mu t] \\ &= \mu t\end{aligned}$$

and if $0 \leq s \leq t$,

$$\begin{aligned}\mathbb{E}[X_s X_t] &= \mathbb{E}[(\sigma B_s + \mu s)(\sigma B_t + \mu t)] \\ &= \mathbb{E}[\sigma^2 B_s B_t + \mu t B_s + \mu s \sigma B_t + \mu^2 st] \\ &= \sigma^2 s + \mu^2 st\end{aligned}$$

so,


$$Cov(X_s, X_t) = \sigma^2 s$$

Listing 4: Brownian motion with a drift

```
# Given a standard brownian motion, this function produces a
# brownian motion with drift = mu and diffusion coefficient=
# sigma
def brownianMotionWithDrift(mu, sigma, B_t):
    numOfPaths = len(B_t)
    N = len(B_t[0])
    t = np.linspace(start=0, stop=1.0, num=N)

    Y = []
    for omega_i in range(numOfPaths):
        X_t = sigma * B_t[omega_i] + mu * t
        Y.append(X_t)
```

Brownian motion with drift $\mu = 1.0$, diffusion coeff $\sigma = 1.0$



Example 3.18. (*Brownian Bridge*). The Brownian bridge is a Gaussian process $(Z_t)_{t \in [0,1]}$ defined by the mean $\mathbb{E}[Z_t] = 0$ and covariance $Cov(Z_t, Z_s) = s(1-t)$ if $0 \leq s \leq t$. Note that by construction, $Z_0 = Z_1 = 0$. It turns out that if $(B_t)_{t \in [0,1]}$ is a standard brownian motion on $[0, 1]$, then the process

$$Z_t = B_t - tB_1, \quad t \in [0, 1]$$

has the distribution of a Brownian bridge.

Listing 5: Brownian bridge


```
def brownianBridge(B_t):
    numOfPaths = len(B_t)
    N = len(B_t[0])

    t = np.linspace(start=0, stop=1.0, num=N)

    Z = []
    for omega_i in range(numOfPaths):
        X = B_t[omega_i] - B_t[omega_i][N-1]* t
        Z.append(X)

    return Z
```

10 sample paths of Brownian Bridge on $[0, 1]$



Example 3.19. (*Fractional Brownian Motion*). The fractional Brownian motion $(B_t^{(H)})_{t \geq 0}$ with index $0 < H < 1$ (called the Hurst Index), is the Gaussian process with mean 0 and covariance

$$\text{Cov}(Y_s, Y_t) = \mathbb{E}[B_t^{(H)}, B_s^{(H)}] = \frac{1}{2}(t^{2H} + s^{2H} - |t - s|^{2H})$$

The case of $H = 1/2$ corresponds to the Brownian motion.

Listing 6: Fractional Brownian Motion

```
def fBM(H, numOfPaths, N):
    # Initialize the covariance matrix
    C = np.zeros((N, N))

    for i in range(N):
        for j in range(N):
            s = (i+1)/N
            t = (j+1)/N

            C[i][j] = 0.50 * (s**(2*H) + t**(2*H) - (np.abs(t -
                s))**(2*H))

    A = np.linalg.cholesky(C)

    Y = []
```


```

for i in range(numOfPaths):
    X = sampleGaussianProcess(A, N)
    X = np.concatenate(([0], X), axis=0)
    Y.append(X)


return Y

```


3 paths of fractional brownian motion with $H = 0.1$ on $[0, 1]$



3 paths of fractional brownian motion with $H = 0.5$ on $[0, 1]$



3 paths of fractional brownian motion with $H = 0.9$ on $[0, 1]$



Example 3.20. (Ornstein-Uhlenbeck process). The Ornstein-Uhlenbeck process $(Y_t)_{t \geq 0}$ starting at $Y_0 = 0$ is the Gaussian process with mean $\mathbb{E}[Y_t] = 0$ and covariance:

$$\text{Cov}(Y_s, Y_t) = \frac{e^{-2(t-s)}}{2}(1 - e^{-2s}), \quad \text{for } s \leq t$$

If the starting point Y_0 is random, specifically Gaussian with mean 0 and variance $1/2$, then we have: $\mathbb{E}Y_t = 0$ and

$$\text{Cov}(Y_s, Y_t) = \frac{e^{-2(t-s)}}{2}, \quad \text{for } s \leq t$$

The covariance only depends on the difference of the time! This means that the process $(Y_t)_{t \geq 0}$ has the same distribution if we shift time by an amount a for any $a \geq 0$: $(Y_{t+a})_{t \geq 0}$. Processes with this property are called *stationary*. As can be observed from the figure below, the statistics of stationary processes do not change over time.

3.7 A Geometric Point of View.

Before turning to Gaussian processes in more detail, it is worthwhile to spend some time to further explore the analogy between random variables in $L^2(\Omega)$ and vectors

in \mathbf{R}^n . We've already seen earlier, how the space of all random variables form a vector space. We shall now observe that L^2 is a subspace of this vector space.

Definition 3.21. For a given probability space $(\Omega, \mathcal{F}, \mathbb{P})$, the space $L^2(\Omega, \mathcal{F}, \mathbb{P})$ consists of all random variables defined on $(\Omega, \mathcal{F}, \mathbb{P})$ such that:

$$[\mathbb{E}(X^2)] < \infty$$

Such random variables are called square integrable.

In the same spirit, the space of integrable random variables is denoted by $L^1(\Omega, \mathcal{F}, \mathbb{P})$. We will see that any square-integrable random variable must be integrable. In other words, $L^2(\Omega, \mathcal{F}, \mathbb{P})$ is a subset of $L^1(\Omega, \mathcal{F}, \mathbb{P})$. In particular, square integrable random variables have a well-defined expectation. This means that, we can think of L^2 as the set of all random variables on a given probability space with finite variance. Clearly, random variables on $(\Omega, \mathcal{F}, \mathbb{P})$ with the Gaussian distribution are in L^2 .

The space L^2 is a vector space.

1) If X and Y are two random variables in L^2 , then the linear combination $aX + bY$ is also a random variable in L^2 . If $u, v \in \mathbf{R}$, we know that :

$$\begin{aligned}(u - v)^2 &\geq 0 \\ u^2 - 2uv + v^2 &\geq 0 \\ 2uv &\leq u^2 + v^2\end{aligned}$$

Setting $u = aX$ and $v = bY$, we get:

$$\begin{aligned}2abXY &\leq a^2X^2 + b^2Y^2 \\ 2ab\mathbb{E}(XY) &\leq a^2\mathbb{E}X^2 + b^2\mathbb{E}Y^2\end{aligned}$$

Having established this upper bound for $2ab\mathbb{E}(XY)$, we now proceed to show that $aX + bY$ belongs to L^2 . We have:

$$\begin{aligned}\mathbb{E}[(aX + bY)^2] &= \mathbb{E}(a^2X^2 + b^2Y^2 + 2abXY) \\ &= a^2\mathbb{E}X^2 + b^2\mathbb{E}Y^2 + 2ab\mathbb{E}(XY) \\ &\leq a^2\mathbb{E}X^2 + b^2\mathbb{E}Y^2 + a^2\mathbb{E}X^2 + b^2\mathbb{E}Y^2 \\ &= 2a^2\mathbb{E}X^2 + 2b^2\mathbb{E}Y^2\end{aligned}$$

Since $X, Y \in L^2$, $\mathbb{E}X^2 < \infty$ and $\mathbb{E}Y^2 < \infty$. Hence, $\mathbb{E}[(aX + bY)^2]$ is bounded.

2) The zero element of the linear space L^2 is the constantly zero random variable $X = 0$ (with probability one).

Example 3.21. Consider $(\Omega, \mathcal{P}(\Omega), \mathbb{P})$ where $\Omega = \{0, 1\} \times \{0, 1\}$, \mathbb{P} is the equiprobability and $\mathcal{P}(\Omega)$ is the power set of Ω i.e. all the subsets of Ω . An example of a random variable is $X = 2\mathbf{1}_{\{(0,0)\}}$ where $\mathbf{1}_{\{(0,0)\}}$ is the indicator function of the event $\{(0, 0)\}$. In other words, X takes the value 2 on the outcome $(0, 0)$ and 0 for the other outcomes. Of course, we can generalize this construction by taking a linear combination of multiples of indicator random functions. Namely, consider the random variable:

$$X = a\mathbf{1}_{\{(0,0)\}} + b\mathbf{1}_{\{(1,0)\}} + c\mathbf{1}_{\{(0,1)\}} + d\mathbf{1}_{\{(1,1)\}}$$

for some fixed $a, b, c, d \in \mathbf{R}$. Clearly, any random variable on this probability space can be written in this form. Moreover, any random variable of this form will have a finite variance. Therefore, the space L^2 in this example consists of random variables of the above form. This linear space has dimension 4, since we can write any random variables as the finite linear combination of the four indicator functions. In general, if Ω is finite, the space L^2 is finite dimensional as a linear space, if Ω is infinite, the space $L^2(\Omega, \mathcal{F}, \mathbb{P})$ might be infinite dimensional.

3.7.1 Norm in $L^2(\Omega, \mathcal{F}, \mathbb{P})$.

Similar to \mathbf{R}^n , the space $L^2(\Omega, \mathcal{F}, \mathbb{P})$ has a norm or a length: for a random variable X in L^2 , its norm $\|X\|_2$ is given by:

$$\|X\|_2 = [\mathbb{E}X^2]^{1/2} \quad (3.52)$$

Note that this is very close in spirit to the length for the vector $\|\mathbf{x}\|_2 = \sqrt{x_1^2 + \dots + x_n^2}$ in \mathbf{R}^n , since the expectation is heuristically a sum over the outcomes. We have already seen that this definition satisfies the properties of a norm.

1) *Positive Semi-Definite:*

If X is a random variable, then $X^2 \geq 0$. By monotonicity of expectations, $\mathbb{E}X^2 \geq 0$. Moreover, if $\mathbb{E}X^2 = 0$, then since X^2 is a non-negative random variable, $X^2 = 0$ almost surely. It implies that $X = 0$ a.s.

2) *Scalar multiplication.*

If X is a random variable in L^2 , we have:

$$\begin{aligned}\|aX\|_2 &= (\mathbb{E}[(aX)^2])^{1/2} \\ &= (\mathbb{E}a^2 X^2)^{1/2} \\ &= |a| (\mathbb{E}X^2)^{1/2} \\ &= |a| \|X\|_2\end{aligned}$$

3) *Triangle Inequality.*

We have:

$$\begin{aligned}\|X + Y\|_2^2 &= \mathbb{E}[(X + Y)^2] \\ &= \mathbb{E}X^2 + \mathbb{E}Y^2 + 2\mathbb{E}XY \\ &= \|X\|_2^2 + \|Y\|_2^2 + 2\mathbb{E}XY \\ &\leq \|X\|_2^2 + \|Y\|_2^2 + 2\mathbb{E}|XY| \quad \{\because XY \leq |XY|\} \\ &\leq \|X\|_2^2 + \|Y\|_2^2 + 2(\mathbb{E}X^2)^{1/2}(\mathbb{E}Y^2)^{1/2} \quad \{\text{Cauchy-Schwarz inequality}\} \\ &= \|X\|_2^2 + \|Y\|_2^2 + 2\|X\|_2\|Y\|_2 \\ &= (\|X\|_2 + \|Y\|_2)^2\end{aligned}$$

Thus, $\|X + Y\|_2 \leq \|X\|_2 + \|Y\|_2$.

3.7.2 Inner-product in $L^2(\Omega, \mathcal{F}, \mathbb{P})$.

Like \mathbf{R}^n , the space L^2 has a dot-product or scalar product between two elements X, Y of the space. It is given by:

$$\langle X, Y \rangle = \mathbb{E}(XY)$$

More generally, this operation is called the inner-product. It has the same properties as the dot product in \mathbf{R}^n .

1) *Symmetric*:

We have:

$$\mathbb{E}(XY) = \mathbb{E}(YX)$$

2) *Linearity*:

$$\mathbb{E}((aX + bY)Z) = a\mathbb{E}(XZ) + b\mathbb{E}(YZ)$$

3) *Positive semi-definite*:

$$\langle X, X \rangle = \mathbb{E}X^2$$

Since X^2 is a non-negative random variable, $X^2 \geq 0$ and by the monotonicity of expectations $\mathbb{E}X^2 \geq 0$.

Let $X, Y \in L^2(\Omega, \mathcal{F}, \mathbb{P})$ and define $\hat{X} = X - \mathbb{E}X$, $\hat{Y} = Y - \mathbb{E}Y$

$$\begin{aligned} |\mathbb{E}(\hat{X}\hat{Y})| &\leq \mathbb{E}|\hat{X}\hat{Y}| \leq \left(\mathbb{E}\hat{X}^2\right)^{1/2} \left(\mathbb{E}\hat{Y}^2\right)^{1/2} \\ |\mathbb{E}(X - \mathbb{E}X)(Y - \mathbb{E}Y)| &\leq [\mathbb{E}(X - \mathbb{E}X)^2]^{1/2} [\mathbb{E}(Y - \mathbb{E}Y)^2]^{1/2} \\ |Cov(X, Y)| &\leq \sqrt{Var(X)} \cdot \sqrt{Var(Y)} \\ |Corr(X, Y)| &\leq 1 \end{aligned}$$

3.7.3 Projection of a random variable X on Y .

Consider the random variable

$$\begin{aligned} X^\perp &= X - \frac{\langle X, Y \rangle}{\|Y\|_2^2} Y \\ &= X - \frac{\mathbb{E}(XY)}{\mathbb{E}Y^2} Y \end{aligned}$$

This random variable is uncorrelated to Y or orthogonal to Y , in the sense that its inner product with Y is zero. We have:

$$\begin{aligned}
 \langle X^\perp, Y \rangle &= \mathbb{E}(X^\perp Y) \\
 &= \mathbb{E} \left[XY - \frac{\mathbb{E}(XY)}{\mathbb{E}Y^2} Y^2 \right] \\
 &= \mathbb{E}XY - \frac{\mathbb{E}(XY)}{\mathbb{E}Y^2} \cdot \mathbb{E}Y^2 \\
 &= 0
 \end{aligned}$$




Figure. A representation of the decomposition of the random variable X in terms of its projection on Y and the component X^\perp orthogonal to Y .

The random variable $\frac{\mathbb{E}[XY]}{\mathbb{E}[Y^2]}Y$ is the random variable of the form tY , $t \in \mathbf{R}$, that is closest to X in the L^2 -sense. We will make this more precise when we define the conditional expectation of a random variable shortly ahead. For now, we simply note that these considerations imply the decomposition

$$X = X^\perp + \frac{\mathbb{E}[XY]}{\mathbb{E}[Y^2]}Y$$

The random variable $X^\perp = X - \frac{\mathbb{E}[XY]}{\mathbb{E}[Y^2]}Y$ is the component of X *orthogonal* to Y .
The random variable:

$$\text{Proj}_Y(X) = \frac{\mathbb{E}[XY]}{\mathbb{E}[Y^2]}Y$$

is called the orthogonal projection of the random variable X onto Y . Put another way, this is the component of X in the direction of Y . This is the equivalent of the orthogonal projection of \mathbf{R}^n of a vector \mathbf{w} in the direction of \mathbf{v} , given by $\frac{\langle \mathbf{w}, \mathbf{v} \rangle}{\|\mathbf{v}\|^2} \mathbf{v}$.

Example 3.22. Going back to example (3.21), let's define the random variables $Y = 2\mathbf{1}_{\{(0,0)\}} + \mathbf{1}_{\{(1,0)\}}$ and $W = \mathbf{1}_{\{(0,0)\}}$. Then the orthogonal projection of Y onto W is:

$$\begin{aligned} \frac{\mathbb{E}(YW)}{\mathbb{E}W^2}W &= \frac{2\mathbb{P}(\{0,0\})}{\mathbb{P}(\{0,0\})}W \\ &= \frac{2 \cdot \frac{1}{4}}{\frac{1}{4}}W \\ &= 2W \end{aligned}$$

The orthogonal decomposition of Y is simply :

$$Y = 2W + (Y - 2W)$$

The notion of norm induces a notion of distance between the random variables in L^2 given by $\|X - Y\|_2 = \mathbb{E}[(X - Y)^2]^{1/2}$. In particular, we see that the orthogonal projection of Y onto X is the closest point from X amongst all multiples of Y . This is what the proof of Cauchy-Schwarz inequality does. The L^2 distance also gives rise to a notion of convergence.

3.8 Borel-Cantelli Lemmas.

Lemma 3.8. (*Borel-Cantelli Lemmas*)

(a) (*First Borel-Cantelli Lemma*) Let $\{A_n\}$ be a sequence of events such that the series $\sum_n \mathbb{P}(A_n)$ converges to a finite value L . Then, almost surely, only finitely many A_n 's will occur.

(b) (*Second Borel-Cantelli Lemma*) Let $\{A_n\}$ be a sequence of independent events such that $\sum_n \mathbb{P}(A_n)$ diverges to ∞ . Then, almost surely, infinitely many A_n 's will occur.

Fix a probability space $(\Omega, \mathcal{F}, \mathbb{P})$. Let A_1, A_2, A_3, \dots be an infinite sequence of events belonging to \mathcal{F} . We shall often be interested in finding out how many of the A_n occur.

The event “ A_n occurs infinitely often (A_n i.o.)” is the set of all ω that belong to infinitely many A_n ’s.

Imagine that an infinite number of A_n ’s occur. That is, $(\forall n)(\exists m \geq n)(\text{s.t. } A_m \text{ occurs})$. In other words:

$$\{A_n \text{ infinitely often}\} \triangleq \bigcap_{n=1}^{\infty} \underbrace{\bigcup_{m=n}^{\infty} A_m}_{B_n} \quad (3.53)$$

Here, B_n is the event that atleast one of A_n, A_{n+1}, \dots occur. For that reason, B_n is sometimes referred to as the n -th tail event. $\{A_n \text{ infinitely often}\}$ is the intersection of all the B_n ’s, so it is the event that all the B_n ’s occur. Therefore, no matter how far I go, no matter how big my n_0 is, beyond that n_0 , atleast one of $A_{n_0}, A_{n_0+1}, \dots$ occurs.

Taking the complement of both sides in (3.53), we get the expression for the event that A_n occurs finitely often.

$$\{A_n \text{ finitely often}\} \triangleq \bigcup_{n=1}^{\infty} \underbrace{\bigcap_{m=n}^{\infty} A_m^C}_{B_n^C} \quad (3.54)$$

It means there exists an n , such that each of the further A_i ’s fail to occur.

In order to prove the Borel-Cantelli lemmas, we require the following lemma.

Lemma 3.9. *If $\sum_{i=1}^{\infty} p_i = \infty$, then $\lim_{n \rightarrow \infty} \prod_{i=1}^n (1 - p_i) = 0$.*

Proof. We know that:

$$\ln(1 + x) \leq x$$

So,

$$\begin{aligned}
\ln(1 - p_i) &\leq -p_i \\
\sum_{i=1}^n \ln(1 - p_i) &\leq -\sum_{i=1}^n p_i \\
0 \leq \prod_{i=1}^n (1 - p_i) &\leq e^{-\sum_{i=1}^n p_i}
\end{aligned}$$

Passing to the limit on both sides, as $n \rightarrow \infty$, we have:

$$0 \leq \lim_{n \rightarrow \infty} \prod_{i=1}^n (1 - p_i) \leq \lim_{n \rightarrow \infty} e^{-\sum_{i=1}^n p_i} = 0$$

By the squeeze theorem, the limit $\lim_{n \rightarrow \infty} \prod_{i=1}^n (1 - p_i)$ exists and is equal to 0.

Consequently, the product series $\prod_{i=1}^n (1 - p_i)$ converges to 0. \square

Proof. (**First Borel-Cantelli Lemma**)

Our claim is that $\mathbb{P}(\bigcap_{n=1}^{\infty} B_n) = 0$.

Whenever we see something like $\bigcap_{n=1}^{\infty} B_n$, we can think of invoking continuity of probability. It turns out that, $B_n = \bigcup_{m \geq n} A_m$. So, $B_1 \supseteq B_2 \supseteq B_3 \supseteq \dots$, that is the B_n 's are nested decreasing sequence of sets. So, $\lim B_n = \bigcap_{n=1}^{\infty} B_n$. So, by continuity of probability measure:

$$\begin{aligned}
\mathbb{P}\left(\bigcap_{n=1}^{\infty} B_n\right) &= \lim_{n \rightarrow \infty} \mathbb{P}(B_n) \\
&= \lim_{n \rightarrow \infty} \mathbb{P}\left(\bigcup_{m=n}^{\infty} A_m\right) \\
&\leq \lim_{n \rightarrow \infty} [\mathbb{P}(A_n) + \mathbb{P}(A_{n+1}) + \dots] \\
&\quad \{ \text{Union bound on probability} \} \\
&= \lim_{n \rightarrow \infty} \sum_{i=n}^{\infty} \mathbb{P}(A_i)
\end{aligned}$$

We know that, $\sum_{n=1}^{\infty} \mathbb{P}(A_n)$ converges to some finite value L , and the above expression is the tail sum of a convergent series. The sequence of tail sums of a convergent series always converges to 0. Thus,

$$0 \leq \mathbb{P}\left(\bigcap_{n=1}^{\infty} B_n\right) \leq 0$$

so it follows that

$$\mathbb{P}\{A_n \text{ i.o.}\} = 0$$

(Second Borel-Cantelli Lemma)

Our claim is $\mathbb{P}\{A_n \text{ i.o.}\} = 1$. We must therefore prove that:

$$\begin{aligned} \mathbb{P}\left(\bigcap_{n=1}^{\infty} B_n\right) &= 1 \\ \iff \mathbb{P}\left(\bigcup_{n=1}^{\infty} B_n^C\right) &= 0 \end{aligned}$$

We have:

$$\mathbb{P}\left(\bigcup_{n=1}^{\infty} B_n^C\right) \leq \sum_{n=1}^{\infty} \mathbb{P}(B_n^C)$$

I want to prove that the above sum is zero. It means that each of these B_n^C events should have 0 probability. We will show that $\mathbb{P}(B_n^C) = 0$ for all $n \geq 1$.

Indeed fix n and $k \geq n$. Consider $\mathbb{P}\left(\bigcap_{i=n}^k A_i^C\right)$. That is I am taking finite intersection of A_i^C . I want to prove that B_n^C has probability zero.

If you look at A_i^C , these are independent events. So, $\mathbb{P}\left(\bigcap_{i=n}^k A_i^C\right) = \prod_{i=n}^k \mathbb{P}(A_i^C) = \prod_{i=n}^k [1 - \mathbb{P}(A_i)]$. Passing to the limit as $k \rightarrow \infty$,

$$\begin{aligned}
\lim_{k \rightarrow \infty} \mathbb{P}\left(\bigcap_{i=n}^k A_i^C\right) &= \lim_{k \rightarrow \infty} \prod_{i=n}^k [1 - \mathbb{P}(A_i)] \\
\mathbb{P}\left(\bigcap_{i=n}^{\infty} A_i^C\right) &= \lim_{k \rightarrow \infty} \prod_{i=n}^k [1 - \mathbb{P}(A_i)] \\
\mathbb{P}(B_n^C) &= \lim_{k \rightarrow \infty} \prod_{i=n}^k [1 - \mathbb{P}(A_i)]
\end{aligned}$$

Since, $\sum_{i=n}^{\infty} \mathbb{P}(A_i)$ diverges to ∞ , it follows from lemma (3.9), that $\prod_{i=n}^{\infty} [1 - \mathbb{P}(A_i)] = 0$. Hence, $\mathbb{P}(B_n^C) = 0$ for all $n \in \mathbf{N}$. So, $0 \leq \mathbb{P}\left(\bigcup_{n=1}^{\infty} B_n^C\right) \leq 0$. Therefore, $\mathbb{P}\left(\bigcup_{n=1}^{\infty} B_n^C\right) = 0$, or equivalently, $\mathbb{P}\left(\bigcap_{n=1}^{\infty} B_n\right) = 1$. The event $\{A_n \text{ i.o.}\}$ occurs almost surely. \square

3.9 Convergence of random variables.

Fix a probability space $(\Omega, \mathcal{F}, \mathbb{P})$ once for all. On this probability space, we will have a sequence (X_1, X_2, X_3, \dots) of random variables defined on it.

Definition 3.22. (*Point-wise convergence.*) A sequence of random variables $(X_n)_{n=1}^{\infty}$ on $(\Omega, \mathcal{F}, \mathbb{P})$ is said to converge point-wise to X , if and if, for all $\epsilon > 0$, and for all $\omega \in \Omega$, there exists $N(\epsilon, \omega) \in \mathbf{N}$ such that for all $n \geq N$, we have $|X_n(\omega) - X(\omega)| < \epsilon$.

It would be natural to say, that, for all $\omega \in \Omega$, $X_n(\omega) \rightarrow X(\omega)$. But, this is too demanding. So, we will weaken this convergence.

Definition 3.23. (*Almost-sure convergence.*) A sequence of random variables $(X_n)_{n=1}^{\infty}$ on $(\Omega, \mathcal{F}, \mathbb{P})$ is said to converge almost-surely to X , written $X_n \xrightarrow{a.s.} X$, if and only if, there exists a set $A \in \mathcal{F}$, such that $\mathbb{P}[A] = 1$ and for all $\omega \in A$, $X_n(\omega) \rightarrow X(\omega)$.

Theorem 3.21. (*Sufficient condition for almost-sure convergence.*) If $(\forall \epsilon > 0), \sum_{n=1}^{\infty} \mathbb{P}(|X_n - X| > \epsilon) < \infty$, then $X_n \xrightarrow{a.s.} X$.

Remark. If you notice just the object $\mathbb{P}(|X_n - X| > \epsilon)$; if this term goes to zero, then it is convergence in probability. So, if the term $\mathbb{P}(|X_n - X| > \epsilon)$ goes to zero, we have convergence in probability. The condition $\sum_{n=1}^{\infty} \mathbb{P}(|X_n - X| > \epsilon) < \infty$

is a little bit stronger, it says a little bit more. As n tends to infinity, not only do the terms $a_n = \mathbb{P}(|X_n - X| > \epsilon)$ go to zero, for every ϵ ; it goes to zero fast enough that the sum converges. For example, if this probability $\mathbb{P}(|X_n - X| > \epsilon)$ were to go to zero, as $\frac{1}{n}$, then you have convergence in probability, but $\sum \frac{1}{n}$ diverges. So, if the term a_n goes to zero fast enough to keep the summation finite, then we have convergence almost surely. For instance, if $a_n \approx \frac{1}{n^2}$, then we would have almost sure convergence.

This is just a sufficient condition. If it holds, we are guaranteed almost sure convergence, but even if it doesn't hold, sometimes we may almost sure convergence.

Proof. Let $A_n(\epsilon)$ be the event $\{|X_n - X| > \epsilon\}$. We are given that, for all $\epsilon > 0$, $\sum_{n=1}^{\infty} \mathbb{P}(A_n(\epsilon)) < \infty$. Using BCL1 (3.8), we see that that, for any $\epsilon > 0$, only finitely many $A_n(\epsilon)$ occur with probability 1. Thus, there exists an n_0 , such that for all $n \geq n_0$, $A_n^C(\epsilon) = \{|X_n - X| \leq \epsilon\}$ occurs with probability 1. So, X_n converges to X with probability 1. \square

Remark. If we plot the distance between the random variables, $X_n - X$, there may be some excursions. But, essentially, BCL1 says that, with probability 1, there must be an n_0 , beyond which the sequence X_n settles within an ϵ -band of X , and these excursions never occur. Because, only finitely many excursions occur. And this is true for every $\epsilon > 0$. So, with probability 1, $X_n \rightarrow X$. Thus, $X_n \xrightarrow{a.s.} X$.




Figure. Convergence of X_n to X .

Example 3.23. The converse of the theorem (3.21) does not hold. Consider the sequence of random variables:

$$X_n = \begin{cases} 1 & \text{with probability } \frac{1}{n} \\ 0 & \text{with probability } 1 - \frac{1}{n} \end{cases}$$

Then, for all $\epsilon > 0$, $\mathbb{P}(|X_n| < \epsilon) = 1 - \frac{1}{n}$. So, for all $\epsilon > 0$, $\lim \mathbb{P}(|X_n| < \epsilon) = 1$. Thus, the sequence (X_n) converges 0 with probability 1. So, $X_n \xrightarrow{a.s.} 0$. However, $\sum \mathbb{P}(|X_n| > \epsilon) = \sum \frac{1}{n} = \infty$.

Theorem 3.22. (Necessary and sufficient condition for almost-sure convergence.) Let $A_n(\epsilon)$ be the event that the excursion $\{|X_n - X| > \epsilon\}$ happens and define:

$$B_m(\epsilon) = \bigcup_{n \geq m} A_n(\epsilon)$$

Then,

$$X_n \xrightarrow{a.s.} X \quad \text{if and only if} \quad \lim_{m \rightarrow \infty} \mathbb{P}(B_m(\epsilon)) = 0 \quad \forall \epsilon > 0$$

Remark. Note that, $\mathbb{P}(A_n(\epsilon))$ going to 0 is convergence in probability. I am saying a little more. In words, $A_n(\epsilon)$ is the event that an excursion occurs at the n th term. In words, $B_m(\epsilon)$ is the event that atleast one of $A_m, A_{m+1}, A_{m+2}, \dots$ occurs, which means that atleast one excursion occurs m or after. What this theorem says is, if the probability of this event goes to 0, then you have almost sure convergence. In other words, if you find some m ; this m can be very large, but if you find some m beyond which no excursions ever occur, then you have almost sure convergence (and vice versa).

Proof. (\Rightarrow direction.)

We are given that $X_n \xrightarrow{a.s.} X$. Our claim is $\lim_{m \rightarrow \infty} \mathbb{P}(B_m(\epsilon)) = 0$.

Now, if $X_n \rightarrow X$ almost surely, then clearly, $(\forall \epsilon > 0)$, the event

$$\bigcup_{m \geq 1} \bigcap_{n \geq m} \{|X_n - X| < \epsilon\}$$

occurs with probability 1.

So, for all $\epsilon > 0$, the event

$$\bigcap_{m \geq 1} \bigcup_{n \geq m} \{|X_n - X| \geq \epsilon\} = \bigcap_{m \geq 1} B_m$$

occurs with probability 0. An excursion happens only finitely many times.

Now, the sequence events $B_1(\epsilon), B_2(\epsilon), B_3(\epsilon), \dots$ are nested decreasing. They are like Russian dolls. By continuity of probability measure, $\mathbb{P}(\bigcap_{m=1}^{\infty} B_m) = \lim_{m \rightarrow \infty} \mathbb{P}(B_m)$. Consequently, it follows that $\lim_{m \rightarrow \infty} \mathbb{P}(B_m) = 0$.

(\Leftarrow direction.)

We are given that, for all $\epsilon > 0$, $\lim_{m \rightarrow \infty} \mathbb{P}(B_m(\epsilon)) = 0$. We are interested to prove that $X_n \xrightarrow{a.s.} X$.

Let C be the event:

$$C = \{\omega | X_n(\omega) \rightarrow X(\omega)\}$$

Define the event:

$$A(\epsilon) = \bigcap_{m \geq 1} \bigcup_{n \geq m} \{|X_n - X| \geq \epsilon\}$$

I would like to prove that $\mathbb{P}(C) = 1$. What we will prove is that $\mathbb{P}(C^C) = 0$.

C^C is the event that, no matter how big an n you look at, there is some excursion. That is there are infinitely many excursions.

This means, there must be some $\epsilon_0 > 0$ for which $A(\epsilon)$ occurs.

$$\begin{aligned} \mathbb{P}(C^C) &= \mathbb{P}\left(\bigcup_{\epsilon > 0} A(\epsilon)\right) \\ &= \mathbb{P}\left(\bigcup_{k=0}^{\infty} A\left(\frac{1}{k}\right)\right) \\ &\leq \sum_{k=0}^{\infty} \mathbb{P}\left(A\left(\frac{1}{k}\right)\right) \end{aligned}$$

Now,

$$\begin{aligned} \lim_{m \rightarrow \infty} \mathbb{P}(B_m(\frac{1}{k})) &= \mathbb{P}\left(\bigcap_{m=1}^{\infty} B_m(\frac{1}{k})\right) \\ &\quad \{ \text{Continuity of probability measure} \} \\ &= \mathbb{P}\left(A\left(\frac{1}{k}\right)\right) \end{aligned}$$

So, $\mathbb{P}(A(1/k)) = 0$. Consequently, $\mathbb{P}(C^C) = 0$ and $\mathbb{P}(C) = 1$. Thus, $X_n \xrightarrow{a.s.} X$. \square

Remark. Therein, lies the difference between convergence in probability and convergence almost surely. Convergence in probability just says the probability of $A_n(\epsilon)$ (an excursion occurs at n) goes to zero. It just looks at one n ; it forgets about the rest of the sequence.

For convergence almost surely, you are not looking at a particular n . You fix a particular m and you're saying that the probability that beyond m an excursion occurs goes to zero.

This should convince you intuitively, that almost sure convergence implies convergence in probability.

Definition 3.24. (*Convergence in Probability.*) A sequence of random variables $(X_n)_{n=1}^\infty$ on $(\Omega, \mathcal{F}, \mathbb{P})$ is said to converge in probability to X , written $X_n \xrightarrow{P} X$, if and only if

$$\forall \epsilon > 0, \quad \lim_{n \rightarrow \infty} \mathbb{P}(|X_n - X| < \epsilon) = 0$$

Definition 3.25. (*Convergence in L^p*) A sequence of random variables $(X_n)_{n=1}^\infty$ on $(\Omega, \mathcal{F}, \mathbb{P})$ is said to converge in the p th mean to X , if and only if

$$\lim_{n \rightarrow \infty} \mathbb{E}[|X_n - X|^p] = 0$$

Definition 3.26. (*Convergence in Distribution.*) A sequence of random variables $(X_n)_{n=1}^\infty$ on $(\Omega, \mathcal{F}, \mathbb{P})$ is said to converge in distribution to X , if and only if:

$$\mathbb{P}(X_n \leq x) \rightarrow \mathbb{P}(X \leq x)$$

For example, let $\Omega = \{1, 2\}$ and $\mathbb{P}(1) = \mathbb{P}(2) = \frac{1}{2}$, $X_n(1) = \frac{-1}{n}$ and $X_n(2) = \frac{1}{n}$. We have:

- 1) $X_n \xrightarrow{a.s.} X$ because $X_n(\omega) \rightarrow 0$ for all $\omega \in \Omega$.
- 2) $X_n \xrightarrow{L^2} X$ because $\mathbb{E}(X_n^2) = \frac{1}{n^2} \rightarrow 0$.
- 3) $X_n \xrightarrow{P} X$ because $\mathbb{P}(|X_n| > \epsilon) = \mathbb{P}\left(\frac{1}{n} > \epsilon\right) = 0$.

Theorem 3.23. (*Hierarchy of Convergence*) The following implications hold:

$$\begin{array}{c}
(X_n \xrightarrow{a.s.} X) \\
\Downarrow \\
(X_n \xrightarrow{P} X) \implies (X_n \xrightarrow{D} X) \\
\Uparrow \\
(X_n \xrightarrow{L^p} X)
\end{array}$$

Proof. (i) *Claim.* $X_n \xrightarrow{L^p} X$ implies $X_n \xrightarrow{P} X$.

This is a very easy proposition to prove. By definition, we have:

$$\lim_{n \rightarrow \infty} \mathbb{E}[|X_n - X|^p] = 0$$

By Markov's inequality:

$$\begin{aligned}
0 \leq \mathbb{P}(|X_n - X| > \epsilon) &= \mathbb{P}(|X_n - X|^p > \epsilon^p) \\
&\leq \frac{1}{\epsilon^p} \mathbb{E}[|X_n - X|^p]
\end{aligned}$$

Passing to the limit on both sides, as $n \rightarrow \infty$, by the squeeze limit theorem,

$$\lim_{n \rightarrow \infty} \mathbb{P}(|X_n - X| > \epsilon) = 0$$

(ii) *Claim.* $X_n \xrightarrow{P} X$ implies that $X_n \xrightarrow{D} X$.

Fix an $\epsilon > 0$.

We have:

$$\begin{aligned}
F_{X_n}(x) &= \mathbb{P}(X_n \leq x) \\
&= \mathbb{P}(X_n \leq x, X \leq x + \epsilon) \\
&\quad + \mathbb{P}(X_n \leq x, X > x + \epsilon) \\
&\leq \mathbb{P}(X \leq x + \epsilon) + \mathbb{P}(|X_n - X| > \epsilon) \\
&\quad \because \{X_n \leq x, X \leq x + \epsilon\} \subseteq \{X \leq x + \epsilon\} \\
&= F_X(x + \epsilon) + \mathbb{P}(|X_n - X| > \epsilon)
\end{aligned}$$

Similarly,

$$\begin{aligned}
F_X(x - \epsilon) &= \mathbb{P}(X \leq x - \epsilon) \\
&= \mathbb{P}(X \leq x - \epsilon, X_n \leq x) + \mathbb{P}(X \leq x - \epsilon, X_n > x) \\
&\leq \mathbb{P}(X_n \leq x) + \mathbb{P}(|X_n - x| > \epsilon) \\
&= F_{X_n}(x) + \mathbb{P}(|X_n - X| > \epsilon)
\end{aligned}$$

Thus, we have the inequality:

$$\forall \epsilon > 0, \quad F_X(x - \epsilon) - \mathbb{P}(|X_n - X| > \epsilon) \leq F_{X_n}(x) \leq F_X(x + \epsilon) + \mathbb{P}(|X_n - X| > \epsilon)$$

We assume that F_X is continuous for all x . Pick $\epsilon = \frac{1}{n}$ and passing to the limit as $n \rightarrow \infty$, we have:

$$\begin{aligned}
\lim F_X(x - \epsilon) &\leq \lim F_{X_n}(x) \leq \lim F_X(x + \epsilon) \\
F_X(x) &\leq \lim F_{X_n}(x) \leq F_X(x)
\end{aligned}$$

By the Squeeze Theorem, the limit $F_{X_n}(x)$ exists and $\lim F_{X_n}(x) = F_X(x)$.

(iii) **Counterexample.** (Convergence in distribution does not imply convergence in probability).

Convergence in distribution simply means that only the CDFs are converging; it doesn't mean that X_n and X are getting closer in any sense.

Let X_1, X_2, X_3, \dots be such that $X_i = X$ for all $i \geq 1$ and

$$X = \begin{cases} 0 & \text{with probability } \frac{1}{2} \\ 1 & \text{with probability } \frac{1}{2} \end{cases}$$

. The entire sequence is just (X, X, X, \dots) . Let $Y = 1 - X$. By definition, $Y \sim \text{Bernoulli}(1/2)$. We have: $X_n \xrightarrow{D} Y$ in distribution, but $|X_n - Y| = 1$, we could choose $\epsilon_0 = \frac{1}{2}$ and we get:

$$\mathbb{P}(|X_n - Y| > \frac{1}{2}) = 1$$

so (X_n) does not converge to Y in probability.

(iv) **Counterexample.** (Convergence in probability does not imply convergence in the mean square sense).

Consider

$$X_n = \begin{cases} n^3 & \text{with probability } n^{-2} \\ 0 & \text{with probability } 1 - n^{-2} \end{cases}$$

Then,

$$\mathbb{P}(|X_n| > \epsilon) = \frac{1}{n^2}$$

and as $n \rightarrow \infty$, $\frac{1}{n^2} \rightarrow 0$. So, $X_n \xrightarrow{P} 0$.

But,

$$\mathbb{E}(X_n^2) = n^2$$

and as $n \rightarrow \infty$, $n^2 \rightarrow \infty$. Hence, $X_n \not\xrightarrow{L^2} 0$.

(v) **Counterexample.** (Convergence in probability does not imply convergence almost surely).

Consider

$$X_n = \begin{cases} 1 & \text{with probability } \frac{1}{n} \\ 0 & \text{with probability } 1 - \frac{1}{n} \end{cases}$$

and X_i 's are independent. The larger the value of n , the more likely that X_n takes the value 0. We have:

$$\mathbb{P}(|X_n| > \epsilon) = \frac{1}{n}$$

So,

$$\lim_{n \rightarrow \infty} \mathbb{P}(|X_n| > \epsilon) = \lim_{n \rightarrow \infty} \frac{1}{n} = 0$$

Our claim is that X_n does not converge to 0 almost surely.

Let A_n be the event that $\{X_n = 1\}$. Then, A_n 's are independent. We have:

$$\sum_{i=1}^{\infty} \mathbb{P}(A_n) = 0 + \frac{1}{2} + \frac{1}{3} + \frac{1}{4} + \dots$$

The harmonic series $\sum_{n=1}^{\infty} \frac{1}{n}$ diverges to ∞ .

By the BCL2 (Borell-Cantelli Lemma 2) (3.8), it follows that, with probability 1, infinitely many A_n 's will occur.

$$\mathbb{P}\{X_n = 1 \text{ i.o.}\} = 1$$

So, $X_n \not\overset{a.s.}{\rightarrow} 0$.

Imagine a coin-tossing experiment, where X_n represents the outcome of the n th coin-toss, and the probability of the n th coin toss falling heads is $\frac{1}{n}$. Then, no matter how far out you go in the sequence, BCL2 says that, there will some occasional head ($X_n = 1$) popping off at some-time. Which means that X_n does not converge to zero.

(vi) **Claim.** $X_n \overset{a.s.}{\rightarrow} X$ implies $X_n \overset{P}{\rightarrow} X$.

By the necessary and sufficient condition of almost sure convergence, $X_n \overset{a.s.}{\rightarrow} X$ is equivalent to saying that:

$$\lim_{m \rightarrow \infty} \mathbb{P}(B_m(\epsilon)) = 0$$

But, $B_m(\epsilon) = \bigcup_{n \geq m} A_n(\epsilon)$. Thus, $A_m(\epsilon) \subseteq B_m(\epsilon)$. So, $(\forall \epsilon > 0), 0 \leq \mathbb{P}(A_m(\epsilon)) \leq \mathbb{P}(B_m(\epsilon))$. Passing to the limit on both sides, $0 \leq \lim \mathbb{P}(A_m(\epsilon)) \leq \lim \mathbb{P}(B_m(\epsilon)) = 0$. By the squeeze theorem, $\lim \mathbb{P}(A_m(\epsilon)) = 0$. Consequently, $X_n \xrightarrow{P} X$.

(vii) **Counterexample.** (Convergence almost surely does not imply convergence in mean square)

Let

$$X_n(\omega) = \begin{cases} n & \omega \in [0, \frac{1}{n}] \\ 0 & \text{otherwise} \end{cases}$$

In this case, we do have convergence almost surely. $\mathbb{P}(|X_n| < \epsilon) = 1 - \frac{1}{n}$ and so, $\lim \mathbb{P}(|X_n| < \epsilon) = 1$. $X_n \xrightarrow{a.s.} 0$. But, $\mathbb{E}[X_n^2] = n$. Thus, $X_n \not\xrightarrow{L^2} 0$.

(viii) **Counterexample.** (Convergence in mean square does not imply convergence almost surely)

Let

$$X_n = \begin{cases} 1 & \text{with probability } 1/n \\ 0 & \text{with probability } 1 - 1/n \end{cases}$$

where the X_n 's are independent.

Now, $\mathbb{E}[X_n^2] = \frac{1}{n}$ so $\lim \mathbb{E}[X_n^2] = 0$. Thus, $X_n \xrightarrow{L^2} 0$. Define $A_n = \{|X_n| \geq \epsilon\}$. But, by BCL2, $\sum_n \mathbb{P}(A_n) = \infty$ and the events A_n are independent. Then, A_n occurs infinitely often. In other words, X_n does not converge to 0 almost surely. \square

Theorem 3.24. *If a sequence (X_n) of random variables converges in probability to X , then there exists a subsequence $(X_{n_k})_k$ which converges to X almost surely.*

Proof. Since $X_n \xrightarrow{P} X$ it follows that:

$$\forall \epsilon > 0 \quad \lim \mathbb{P}(|X_n - X| > \epsilon) = 0$$

By the definition of the limit of a sequence, there exists n_1 such that:

$$\mathbb{P}(|X_{n_1} - X| > 1) < 1$$

There exists $n_2 \geq n_1$ such that:

$$\mathbb{P}\left(|X_{n_2} - X| > \frac{1}{2}\right) < \frac{1}{2^2}$$

There exists $n_3 \geq n_2$ such that:

$$\mathbb{P}\left(|X_{n_3} - X| > \frac{1}{3}\right) < \frac{1}{3^2}$$

In general, there exists a positive integer $n_i \geq n_{i-1}$ such that:

$$\mathbb{P}\left(|X_{n_i} - X| > \frac{1}{i}\right) < \frac{1}{i^2}$$

By the sufficient condition for almost sure convergence (3.21), $\sum_{n=1}^{\infty} \mathbb{P}(A_n) < \sum_{n=1}^{\infty} \frac{1}{n^2}$ which converges to a finite value. Hence, $X_{n_i} \xrightarrow{a.s.} X$. \square

Consider a sequence of random variables (X_n) , such that $X_n \xrightarrow{L^2} X$. it turns out that the limit random variable X of the convergent sequence in L^2 is guaranteed to be in L^2 . This is because L^2 is complete. We will prove this very important result further ahead. This property is crucial for the construction of the Ito integral.

Example 3.24. (A version of the weak law of large numbers.) Consider a sequence of random variables X_1, X_2, \dots, X_n in $L^2(\Omega, \mathcal{F}, \mathbb{P})$ such that $\mathbb{E}[X_i] = 0$, $\mathbb{E}[X_i^2] = \sigma^2 < \infty$ for all $i \geq 1$ and that they are orthogonal to each other, that is, $\mathbb{E}[X_i X_j] = 0$ for all $i \neq j$. We show that the empirical mean

$$\frac{1}{n}S_n = \frac{1}{n}(X_1 + X_2 + \dots + X_n)$$

converges to zero in the L^2 sense.

Clearly,

$$\begin{aligned} \mathbb{E} \left[\frac{S_n^2}{n^2} \right] &= \lim_{n \rightarrow \infty} \frac{1}{n^2} \sum_{i=1}^n \mathbb{E}[X_i^2] \\ &= \lim_{n \rightarrow \infty} \frac{1}{n^2} \cdot n\sigma^2 \\ &= \lim_{n \rightarrow \infty} \frac{\sigma^2}{n} \\ &= 0 \end{aligned}$$

Consequently, the empirical mean $\frac{S_n}{n}$ converges to 0 in the mean square sense.

Exercise 3.1. (Ornstein-Uhlenbeck Process.) Generate 100 paths with step size = 0.01 of the following processes on $[0, 1]$:

(a) Ornstein Uhlenbeck process: $C(s, t) = \frac{e^{-2(t-s)}}{2}(1 - e^{-2s})$ for $s \leq t$. with mean 0 (so that $Y_0 = 0$).

(b) Stationary Ornstein-Uhlenbeck process: $C(s, t) = \frac{e^{-2(t-s)}}{2}$ for $s \leq t$ with mean 0 (so Y_0 is a Gaussian random variable of mean 0 and variance 1/2).

Solution.

Listing 7: Ornstein-Uhlenbeck(OU) process

```
def ornsteinUhlenbeck(numOfPaths, N):
    C = np.zeros((N, N))

    for i in range(N):
```

```

    for j in range(N):
        s = (i + 1)/N
        t = (j + 1)/N

        if s > t:
            s, t = t, s
        C[i][j] = np.exp(-2*(t-s))/2 * (1 - np.exp(-2*s))


A = np.linalg.cholesky(C)

Y = []
for i in range(numOfPaths):
    X = sampleGaussianProcess(A, N)
    X = np.concatenate(([0], X), axis=0)
    Y.append(X)


return Y

```

10 paths of Ornstein Uhlenbeck process on $[0, 1]$



10 sample paths of stationary OU process on $[0, 1]$



4 Properties of Brownian Motion.

4.1 Properties of Brownian Motion.

Let $B(t)$ be a fixed Brownian motion. We give below some simple properties that follow directly from the definition of the Brownian Motion.

Proposition 4.1. *For any $t \geq 0$, $B(t)$ is normally distributed with mean 0 and variance t . For any $s, t \geq 0$ we have $\mathbb{E}(B_s B_t) = \min\{s, t\}$.*

Proof. From condition (1), we have that $B_0 = 0$. From condition (2), $B_t - B_0 = B_t$ is normally distributed with mean 0 and variance t .

Assume that $s < t$.

We have:

$$\begin{aligned}
 \mathbb{E}(B_s B_t) &= \mathbb{E}[B_s(B_t - B_s + B_s)] && \{\text{Write } B_t = B_t - B_s + B_s\} \\
 &= \mathbb{E}[B_s(B_t - B_s)] + \mathbb{E}[B_s^2] && \{\text{Linearity of expectations}\} \\
 &= \mathbb{E}[B_s]\mathbb{E}(B_t - B_s) + s && \{B_s, (B_t - B_s) \text{ are independent}\} \\
 &= 0 \cdot 0 + s \\
 &= s
 \end{aligned}$$

This closes the proof. \square

Proposition 4.2. (*Translation Invariance*) For fixed $t_0 \geq 0$, the stochastic process $\tilde{B}(t) = B(t + t_0) - B(t_0)$ is also a Brownian motion.

Proof. Firstly, the stochastic process $\tilde{B}(t)$ is such that:

- (1) $\tilde{B}(0) = B(t_0) - B(t_0) = 0$. Hence, it satisfies condition (1).
- (2) Let $s < t$. We have: $\tilde{B}(t) - \tilde{B}(s) = B(t + t_0) - B(s + t_0)$ which a Gaussian random variable with mean 0 and variance $t - s$. Hence, for $a \leq b$,

$$\mathbb{P}\{a \leq \tilde{B}(t) \leq b\} = \frac{1}{\sqrt{2\pi(t-s)}} \int_a^b e^{-\frac{x^2}{2(t-s)}} dx$$

Hence, it satisfies condition (2).

- (3) To check condition (3) for $\tilde{B}(t)$, we may assume $t_0 > 0$. Then, for any $0 \leq t_1 \leq t_2 \leq \dots \leq t_n$, we have:

$$0 < t_0 \leq t_0 + t_1 \leq t_0 + t_2 \leq \dots \leq t_0 + t_n$$

So, $B(t_1 + t_0) - B(t_0)$, $B(t_2 + t_0) - B(t_1 + t_0)$, \dots , $B(t_k + t_0) - B(t_{k-1} + t_0)$, \dots , $B(t_n + t_0) - B(t_{n-1} + t_0)$ are independent random variables. Consequently, $\tilde{B}(t)$ satisfies condition (3).

This closes the proof. \square

The above translation invariance property says that a Brownian motion starts afresh at any moment as a new Brownian motion.

Proposition 4.3. (*Scaling Invariance*) For any real number $\lambda > 0$, the stochastic process $\tilde{B}(t) = B(\lambda t)/\sqrt{\lambda}$ is also a Brownian motion.

Proof. The scaled stochastic process $\tilde{B}(t)$ is such that:

- (1) $\tilde{B}(0) = 0$. Hence it satisfies condition (1).
- (2) Let $s < t$. Then, $\lambda s < \lambda t$. We have:

$$\tilde{B}(t) - \tilde{B}(s) = \frac{1}{\sqrt{\lambda}}(B(\lambda t) - B(\lambda s))$$

Now, $B(\lambda t) - B(\lambda s)$ is a Gaussian random variable with mean 0 and variance $\lambda(t - s)$. We know that, if X is a random variable with mean μ and variance σ^2 , $Z = \left(\frac{X - \mu}{\sigma}\right)$ has mean 0 and variance 1. Consequently, $\frac{B(\lambda t) - B(\lambda s)}{\sqrt{\lambda}}$ is a Gaussian random variable with mean 0 and variance $(t - s)$.

Hence, $\tilde{B}(t) - \tilde{B}(s)$ is normal distributed with mean 0 and variance $t - s$ and it satisfies condition (2).

(3) To check condition (3) for $\tilde{B}(t)$, we may assume $t_0 > 0$. Then, for any $0 \leq t_1 \leq t_2 \leq \dots \leq t_n$, we have:

$$0 \leq \sqrt{\lambda}t_1 \leq \sqrt{\lambda}t_2 \leq \dots \leq \sqrt{\lambda}t_n$$

Consequently, the random variables $B(\sqrt{\lambda}t_k) - B(\sqrt{\lambda}t_{k-1})$, $k = 1, 2, 3, \dots, n$ are independent. Hence it follows that $\frac{1}{\sqrt{\lambda}}[B(\sqrt{\lambda}t_k) - B(\sqrt{\lambda}t_{k-1})]$ for $k = 1, 2, \dots, n$ are also independent random variables.

This closes the proof. □

It follows from the scaling invariance property that for any $\lambda > 0$ and $0 \leq t_1 \leq t_2 \leq \dots \leq t_n$, the random vectors:

$$(B(\lambda t_1), B(\lambda t_2), \dots, B(\lambda t_n)) \quad (\sqrt{\lambda}B(t_1), \sqrt{\lambda}B(t_2), \dots, \sqrt{\lambda}B(t_n))$$

have the same distribution.

The scaling property shows that Brownian motion is *self-similar*, much like a fractal. To see this, suppose we zoom into a Brownian motion path very close to zero, say on the interval $[0, 10^{-6}]$. If the Brownian motion path were smooth and differentiable, the closer we zoom in around the origin, the flatter the function will look. In the limit, we would essentially see a straight line given by the derivative at 0. However, what we see with the Brownian motion is very different. The scaling property means that for $a = 10^{-6}$,

$$(B_{10^{-6}t}, t \in [0, 1]) \stackrel{\text{distrib.}}{=} (10^{-3}B_t, t \in [0, 1])$$

where $\stackrel{\text{distrib.}}{=}$ means equality of the distribution of the two processes. In other words, Brownian motion on $[0, 10^{-6}]$ looks like a Brownian motion on $[0, 1]$, but with its amplitude multiplied by a factor of 10^{-3} . In particular, it will remain rugged as we zoom in, unlike a smooth function.

Proposition 4.4. (*Reflection at time s*) *The process $(-B_t, t \geq 0)$ is a Brownian motion. More generally, for any $s \geq 0$, the process $(\tilde{B}(t), t \geq 0)$ defined by:*

$$\tilde{B}(t) = \begin{cases} B_t & \text{if } t \leq s \\ B_s - (B_t - B_s) & \text{if } t > s \end{cases} \quad (4.1)$$

is a Brownian motion.

Proof. (a) Consider the process $\tilde{B}(t) = (-B_t, t \geq 0)$.

(1) $\tilde{B}(0) = 0$.

(2) If X is a Gaussian random variable with mean 0 and variance $t - s$, $-X$ is also Gaussian with mean 0 and variance $t - s$. Thus, $\tilde{B}(t) - \tilde{B}(s) = -(B(t) - B(s))$ is also Gaussian with mean 0 and variance $(t - s)$. Hence condition (2) is satisfied.

(3) Assume that $0 \leq t_0 \leq t_1 \leq \dots \leq t_n$. Then, the random variables $-(B(t_k) - B(t_{k-1}))$ are independent for $k = 1, 2, 3, \dots, n$. Hence, condition (3) is satisfied.

(b) Consider the process $\tilde{B}(t)$ as defined in (4.1).

Fix an $s \geq 0$.

(1) Let $t = 0$. Then, $t \leq s$. $\tilde{B}(t) = \tilde{B}(0) = B(0) = 0$.

(2) Let $t_1 < t_2 \leq s$. Then, $\tilde{B}(t_2) - \tilde{B}(t_1) = B(t_2) - B(t_1)$. This is a Gaussian random variable with mean 0 and variance $t_2 - t_1$.

Let $t_1 < s < t_2$. Then, $\tilde{B}(t_2) - \tilde{B}(t_1) = B(s) - (B(t_2) - B(s)) - B(t_1) = (B(s) - B(t_1)) - (B(t_2) - B(s))$. Since, $B(s) - B(t_1)$ and $B(t_2) - B(s)$ are independent Gaussian random variables, any linear combination of these is Gaussian. Moreover, its mean is zero. The variance is given by:

$$\begin{aligned} \text{Var}[\tilde{B}(t_2) - \tilde{B}(t_1)] &= \text{Var}[B(s) - B(t_1)] + \text{Var}[B(t_2) - B(s)] \\ &= (s - t_1) + (t_2 - s) \\ &= t_2 - t_1 \end{aligned}$$

Let $s < t_1 < t_2$. Then,

$$\begin{aligned}\tilde{B}(t_2) - \tilde{B}(t_1) &= B_s - (B_{t_2} - B_s) - (B_s - (B_{t_1} - B_s)) \\ &= B_s - (B_{t_2} - B_s) - (B_s - (B_{t_1} - B_s)) \\ &= -(B_{t_2} - B_{t_1})\end{aligned}$$

Hence, $\tilde{B}(t_2) - \tilde{B}(t_1)$ is again a Gaussian random variable with mean 0 and variance $t_2 - t_1$. Hence, condition (3) is satisfied.

(3) Assume that $0 \leq t_1 \leq \dots \leq t_{k-1} \leq s \leq t_k \leq \dots \leq t_n$. From the above discussion, the increments $\tilde{B}(t_2) - \tilde{B}(t_1), \dots, \tilde{B}(s) - \tilde{B}(t_{k-1}), \tilde{B}(t_k) - \tilde{B}(s), \dots, \tilde{B}(t_n) - \tilde{B}(t_{n-1})$ are independent increments. The increment $\tilde{B}(t_k) - \tilde{B}(t_{k-1})$ only depends on the random variables $\tilde{B}(s) - \tilde{B}(t_{k-1})$ and $\tilde{B}(t_k) - \tilde{B}(s)$. Thus, $\tilde{B}(t_2) - \tilde{B}(t_1), \dots, \tilde{B}(t_k) - \tilde{B}(t_{k-1}), \dots, \tilde{B}(t_n) - \tilde{B}(t_{n-1})$ are independent. \square

Proposition 4.5. (*Time Reversal*). Let $(B_t, t \geq 0)$ be a Brownian motion. Show that the process $(B_1 - B_{1-t}, t \in [0, 1])$ has the distribution of a standard brownian motion on $[0, 1]$.

Proof. (1) At $t = 0$, $B(1) - B(1 - t) = B(1) - B(1) = 0$.

(2) Let $s < t$. Then, $1 - t < 1 - s$. So, the increment :

$$(B(1) - B(1 - t)) - (B(1) - B(1 - s)) = B(1 - s) - B(1 - t)$$

has a Gaussian distribution. It's mean is 0 and variance is $(1 - s) - (1 - t) = t - s$.

(3) Let $0 \leq t_1 \leq t_2 \leq \dots \leq t_n$. Then:

$$1 - t_n \leq \dots \leq 1 - t_k \leq 1 - t_{k-1} \leq \dots \leq 1 - t_2 \leq 1 - t_1$$

Consider the increments of the process for $k = 1, 2, \dots, n$:

$$(B(1) - B(1 - t_k)) - (B(1) - B(1 - t_{k-1})) = B(1 - t_{k-1}) - B(1 - t_k)$$

They are independent random variables. Hence, condition (3) is satisfied. \square

Example 4.1. (Evaluating Brownian Probabilities). Let's compute the probability that $B_1 > 0$ and $B_2 > 0$. We know from the definition that (B_1, B_2) is a Gaussian vector with mean 0 and covariance matrix:

$$C = \begin{bmatrix} 1 & 1 \\ 1 & 2 \end{bmatrix}$$

The determinant of C is 1. By performing row operations on the augmented matrix $[C|I]$ we find that:

$$C^{-1} = \begin{bmatrix} 2 & -1 \\ -1 & 1 \end{bmatrix}$$

Thus, the probability $\mathbb{P}(B_1 > 0, B_2 > 0)$ can be expressed as:

$$\mathbb{P}(B_1 > 0, B_2 > 0) = \frac{1}{\sqrt{(2\pi)^2}} \int_0^\infty \int_0^\infty \exp \left[-\frac{1}{2}(2x_1^2 - 2x_1x_2 + x_2^2) \right] dx_2 dx_1$$

This integral can be evaluated using a calculator or software and is equal to $3/8$. The probability can also be computed using the independence of increments. The increments $(B_1, B_2 - B_1)$ are IID standard Gaussians. We know their joint PDF. It remains to integrate over the correct region of \mathbf{R}^2 which in this case will be:

$$D^* = \{(z_1, z_2) : (z_1 > 0, z_1 + z_2 > 0)\}$$

We have:

$$\mathbb{P}(B_1 > 0, B_2 > 0) = \frac{1}{2\pi} \int_0^\infty \int_{z_2=-z_1}^{z_2=\infty} e^{-(z_1^2+z_2^2)/2} dz_2 dz_1$$

It turns out that this integral can be evaluated exactly. Indeed by writing $B_1 = Z_1$ and $Z_2 = B_2 - B_1$ and splitting the probability on the event $\{Z_2 \geq 0\}$ and its complement, we have that $\mathbb{P}(B_1 \geq 0, B_2 \geq 0)$ equals:

$$\begin{aligned}
\mathbb{P}(B_1 \geq 0, B_2 \geq 0) &= \mathbb{P}(Z_1 \geq 0, Z_1 + Z_2 > 0, Z_2 \geq 0) + \mathbb{P}(Z_1 \geq 0, Z_1 + Z_2 > 0, Z_2 < 0) \\
&= \mathbb{P}(Z_1 \geq 0, Z_2 \geq 0) + \mathbb{P}(Z_1 \geq 0, Z_1 > -Z_2, -Z_2 > 0) \\
&= \mathbb{P}(Z_1 \geq 0, Z_2 \geq 0) + \mathbb{P}(Z_1 \geq 0, Z_1 > Z_2, Z_2 > 0) \\
&= \frac{1}{4} + \frac{1}{8} \\
&= \frac{3}{8}
\end{aligned}$$

Note that, by symmetry, $\mathbb{P}(Z_1 \geq 0, Z_1 > Z_2, Z_2 > 0) = \mathbb{P}(Z_1 \geq 0, Z_1 \leq Z_2, Z_2 > 0) = \frac{1}{8}$.

Example 4.2. (Another look at Ornstein Uhlenbeck process.) Consider the process $(X_t, t \in \mathbf{R})$ defined by :

$$X_t = \frac{e^{-2t}}{\sqrt{2}} B(e^{4t}), \quad t \in \mathbf{R}$$

Here the process $(B_{e^{4t}}, t \geq 0)$ is called a time change of Brownian motion, since the time is now quantified by an increasing function of t namely e^{4t} . The example $(B(\lambda t), t \geq 0)$ in the scaling property is another example of time change.

It turns out that $(X_t, t \in \mathbf{R})$ is a stationary Ornstein-Uhlenbeck process. (Here the index of time is \mathbf{R} instead of $[0, \infty)$, but the definition also applies as the process is stationary. Since the original brownian motion $B(t)$ is a Gaussian process, any finite dimensional vector $(B(t_1), \dots, B(t_n))$ is Gaussian. It follows that:

$$(B(T_1), \dots, B(T_n)) = \frac{1}{\sqrt{2}} (e^{-2t_1} B(e^{4t_1}), \dots, e^{-2t_n} B(e^{4t_n}))$$

is also a Gaussian vector. Hence, $(X_t, t \in \mathbf{R})$ is a Gaussian process.

The mean of $(X_t, t \in \mathbf{R})$ is:

$$\mathbb{E}[X_t] = \frac{e^{-2t}}{\sqrt{2}} \mathbb{E}[B(e^{4t})] = 0$$

And if $s < t$,

$$\begin{aligned}\mathbb{E}[X_s X_t] &= \frac{e^{-2(s+t)}}{2} \mathbb{E}[B(e^{4s}) B(e^{4t})] \\ &= \frac{e^{-2(s+t)}}{2} e^{4s} \\ &= \frac{e^{-2(t-s)}}{2}\end{aligned}$$

Two Gaussian processes having the same mean and covariance have the same distribution. Hence, it proves the claim that (X_t) is a stationary OU process.

4.2 Properties of the paths.

First we review the definitions of the Riemann integral and the Riemann-Stieltjes integral in Calculus.

Definition 4.1. A partition P of $[a, b]$ is a *finite* set of points from $[a, b]$ that includes both a, b . The notational convention is to always list the points of a partition $P = \{a = x_0, x_1, x_2, \dots, x_n = b\}$ in increasing order. Thus:

$$a = x_0 < x_1 < \dots < x_{k-1} < x_k < \dots < x_n = b$$

For each subinterval $[x_{k-1}, x_k]$ of P , let

$$\begin{aligned}m_k &= \inf\{f(x) : x \in [x_{k-1}, x_k]\} \\ M_k &= \sup\{f(x) : x \in [x_{k-1}, x_k]\}\end{aligned}$$

The lower sum of f with respect to P is given by :

$$L(f, P) = \sum_{k=1}^n m_k (x_k - x_{k-1})$$

The upper sum of f with respect to P is given by:

$$U(f, P) = \sum_{k=1}^n M_k(x_k - x_{k-1})$$

For a particular partition P , it is clear that $U(f, P) \geq L(f, P)$ because $M_k \geq m_k$ for all $k = 0, 1, 2, \dots, n$.

Definition 4.2. A partition Q is called a *refinement* of P if Q contains all of the points of P ; that is $Q \subseteq P$.

Lemma 4.1. If $P \subseteq Q$, then $L(f, P) \leq L(f, Q)$ and $U(f, Q) \leq U(f, P)$.

Proof. Consider what happens when we refine P by adding a single point z to some subinterval $[x_{k-1}, x_k]$ of P . We have:

$$\begin{aligned} m_k(x_k - x_{k-1}) &= m_k(x_k - z) + m_k(z - x_{k-1}) \\ &\leq m'_k(x_k - z) + m''_k(z - x_{k-1}) \end{aligned}$$

where

$$\begin{aligned} m'_k &= \inf\{f(x) : x \in [z, x_k]\} \\ m''_k &= \inf\{f(x) : x \in [x_{k-1}, z]\} \end{aligned}$$

By induction we have:

$$\begin{aligned} L(f, P) &\leq L(f, Q) \\ U(f, Q) &\leq U(f, P) \end{aligned}$$

□

Lemma 4.2. If P_1 and P_2 are any two partitions of $[a, b]$, then $L(f, P_1) \leq U(f, P_2)$.

Proof. Let $Q = P_1 \cup P_2$. Then, $P_1 \subseteq Q$ and $P_2 \subseteq Q$. Thus, $L(f, P_1) \leq L(f, Q) \leq U(f, Q) \leq U(f, P_2)$. \square

Definition 4.3. Let \mathcal{P} be the collection of all possible partitions of the interval $[a, b]$. The upper integral of f is defined to be:

$$U(f) = \inf\{U(f, P) : P \in \mathcal{P}\}$$

The lower integral of f is defined by:

$$L(f) = \sup\{L(f, P) : P \in \mathcal{P}\}$$

Consider the set of all upper sums of f - $\{U(f, P) : P \in \mathcal{P}\}$. Take an arbitrary partition $P' \in \mathcal{P}$. Since $L(f, P') \leq U(f, P)$ for all $P \in \mathcal{P}$, by the Axiom of Completeness(AoC), $\inf\{U(f, P) : P \in \mathcal{P}\}$ exists. We can similarly argue for the supremum of all lower Riemann sums.

Lemma 4.3. For any bounded function f on $[a, b]$, it is always the case that $U(f) \geq L(f)$.

Proof. By the properties of the infimum of a set, $(\forall \epsilon > 0), \exists P(\epsilon)$ such that $U(f) < U(f, P(\epsilon)) < U(f) + \epsilon$. Pick $\epsilon = 1, \frac{1}{2}, \frac{1}{3}, \dots, \frac{1}{n}, \dots$. Thus, we can produce a sequence of partitions P_n such that :

$$U(f) < \dots < U(f, P_n) < U(f) + \frac{1}{n}$$

Consequently, $\lim U(f, P_n) = U(f)$. Similarly, we can produce a sequence of partitions (Q_m) such that :

$$L(f) - \frac{1}{m} < \dots < L(f, Q_m) < L(f)$$

We know that:

$$L(f, Q_m) \leq U(f, P_n)$$

Keeping m fixed and passing to the limit, as $n \rightarrow \infty$ on both sides, we have:

$$\begin{aligned} \lim_{n \rightarrow \infty} L(f, Q_m) &\leq \lim_{n \rightarrow \infty} U(f, P_n) \quad \{\text{Order Limit Theorem}\} \\ L(f, Q_m) &\leq U(f) \end{aligned}$$

Now, passing to the limit, as $m \rightarrow \infty$ on both sides, we have:

$$\begin{aligned} \lim_{m \rightarrow \infty} L(f, Q_m) &\leq \lim_{m \rightarrow \infty} U(f) \quad \{\text{Order Limit Theorem}\} \\ L(f) &\leq U(f) \end{aligned}$$

□

Definition 4.4. (Riemann Integrability). A bounded function f on the interval $[a, b]$ is said to be Riemann integrable if $U(f) = L(f)$. In this case, we define $\int_a^b f$ or $\int_a^b f(x)dx$ to be the common value:

$$\int_a^b f(x)dx = U(f) = L(f)$$

Theorem 4.1. (Integrability Criterion) A bounded function f is integrable on $[a, b]$ if and only if, for every $\epsilon > 0$, there exists a partition P_ϵ of $[a, b]$ such that:

$$U(f, P_\epsilon) - L(f, P_\epsilon) < \epsilon$$

Proof. (\Leftarrow direction.) Let $\epsilon > 0$. If such a partition P_ϵ exists, then:

$$U(f) - L(f) \leq U(f, P_\epsilon) - L(f, P_\epsilon) < \epsilon$$

Because ϵ is arbitrary, it follows that $U(f) = L(f)$ and hence f is Riemann integrable.

(\Rightarrow direction.) Let f be a bounded function on $[a, b]$ such that f is Riemann integrable.

Pick an arbitrary $\epsilon > 0$.

Then, since $U(f) = \inf\{U(f, P) : P \in \mathcal{P}\}$, there exists $P_\epsilon \in \mathcal{P}$, such that $U(f) < U(f, P_\epsilon) < U(f) + \frac{\epsilon}{2}$. Since $L(f) = \sup\{L(f, P) : P \in \mathcal{P}\}$, there exists $P_\epsilon \in \mathcal{P}$, such that $L(f) - \frac{\epsilon}{2} < L(f, P_\epsilon) < L(f)$. Consequently,

$$\begin{aligned} U(f, P_\epsilon) - L(f, P_\epsilon) &< U(f) + \frac{\epsilon}{2} - \left(L(f) - \frac{\epsilon}{2}\right) \\ &= U(f) - L(f) + \epsilon \\ &= \epsilon \end{aligned}$$

□

4.2.1 Functions considered in Stochastic Calculus.

Definition 4.5. A point c is called a discontinuity of the first kind or jump point if both limits $g(c+) = \lim_{t \uparrow c} g(t)$ and $g(c-) = \lim_{t \downarrow c} g(t)$ exist and are not equal. The jump at c is defined as $\Delta g(c) = g(c+) - g(c-)$. Any other discontinuity is said to be of the second kind.

Example 4.3. Consider the function

$$f(x) = \sin\left(\frac{1}{x}\right)$$

Let $x_n = \frac{1}{2n\pi}$. Then, $f(x_n) = (0, 0, 0, \dots)$. Next, consider $y_n = \frac{1}{\pi/2 + 2n\pi}$. Then, $f(y_n) = (1, 1, 1, \dots)$. Consequently, f is not continuous at 0. Hence, limits from the left or right don't exist. Consequently, this is a discontinuity of the second kind.

Functions in stochastic calculus are functions without discontinuities of the second kind, that is functions have both left and right hand limits at any point of the domain and have one-sided limits at the boundary. These functions are called *regular* functions. It is often agreed to identify functions if they have the same right and left limits at any point.

The class $D = D[0, T]$ of right-continuous functions on $[0, T]$ with left limits has a special name, *cadlag* functions (which is the abbreviation of right continuous with left limits in French). Sometimes these processes are called R.R.C. for regular right

continuous. Notice that this class of processes includes C , the class of continuous functions.

Let $g \in D$ be a cadlag function, then, by definition, all the discontinuities of g are jumps. An important result in analysis is that, a function can have no more than a countable number of discontinuities.

4.2.2 Variation of a function.

If g is a function of a real variable, its variation over the interval $[a, b]$ is defined as:

$$V_g([a, b]) = \sup \left\{ \sum_{i=1}^n |g(t_i) - g(t_{i-1})| \right\} \quad (4.2)$$

where the supremum is taken over all partitions $P \in \mathcal{P}$.

Clearly, by the Triangle Inequality, the sums in (4.2) increase as new points are added to the partitions. Therefore, the variation of g is:

$$V_g([a, b]) = \lim_{\|\Delta_n\| \rightarrow 0} \sum_{i=1}^n |g(t_i) - g(t_{i-1})|$$

where $\|\Delta_n\| = \max_{1 \leq i \leq n} (t_i - t_{i-1})$. If $V_g([a, b])$ is finite, then g is said to be a function of finite variation on $[a, b]$. If g is a function of $t \geq 0$, then the variation of g as a function of t is defined by:

$$V_g(t) = V_g([0, t])$$

Clearly, $V_g(t)$ is an increasing function of t .

Definition 4.6. g is a function of finite variation if $V_g(t) < \infty$ for all $t \in [0, \infty)$. g is of bounded variation if $\sup_t V_g(t) < \infty$, in other words there exists C , for all t , such that $V_g(t) < C$. Here C is independent of t .

Example 4.4. (1) If $g(t)$ is increasing then for any i , $g(t_i) \geq g(t_{i-1})$, resulting in a telescopic sum, where all terms excluding the first and the last cancel out, leaving

$$V_g(t) = g(t) - g(0)$$

(2) If $g(t)$ is decreasing, then similarly,

$$V_g(t) = g(0) - g(t)$$

Example 4.5. If $g(t)$ is differentiable with continuous derivative $g'(t)$, $g(t) = \int_0^t g'(s)ds$ then

$$V_g(t) = \int_0^t |g'(s)|ds$$

Proof. By definition,

$$V_g(t) = \lim_{\|\Delta_n \rightarrow 0\|} \sum_{i=1}^n |g(t_i) - g(t_{i-1})|$$

Since g is continuous and differentiable on $[t_{i-1}, t_i]$, there exists $z_i \in (t_{i-1}, t_i)$ such, that $g(t_i) - g(t_{i-1}) = g'(z_i)(t_i - t_{i-1})$. Therefore, we can write:

$$\begin{aligned} V_g(t) &= \lim_{\|\Delta_n \rightarrow 0\|} \sum_{i=1}^n |g'(z_i)|(t_i - t_{i-1}) \\ &= \int_0^t |g'(s)|ds \end{aligned}$$

□

Theorem 4.2. If g is continuous, g' exists and $\int_0^t |g'(s)|ds$ is finite, then g is of finite variation.

Example 4.6. The function $g(t) = t \sin(1/t)$ for $t > 0$ and $g(0) = 0$ is continuous on $[0, 1]$ and differentiable at all points except zero, but has infinite variation on any interval that includes 0. Consider the partition $\{x_n\} = \left\{ \frac{1}{\pi/2+n\pi} \right\}$. Thus,

$$\sin(x_n) = \begin{cases} 1 & \text{if } n \text{ is even} \\ -1 & \text{if } n \text{ is odd} \end{cases}$$

Thus,

$$f(x_n) = \begin{cases} x_n & n \text{ is even} \\ -x_n & n \text{ is odd} \end{cases}$$

Therefore,

$$\begin{aligned} \sum_{n=1}^m |f(x_n) - f(x_{n-1})| &= \sum_{n=1}^m (x_n + x_{n-1}) \\ &= x_0 + x_m + 2 \sum_{n=1}^{m-1} x_n \\ &\geq \sum_{n=1}^{m-1} x_n \end{aligned}$$

Now, passing to the limit as m approaches infinity, $\sum \frac{1}{\pi/2+n\pi}$ is a divergent series. Consequently, $V_g(t)$ has unbounded variation.

4.2.3 Jordan Decomposition.

Theorem 4.3. *Any function $g : [0, \infty) \rightarrow \mathbf{R}$ is of bounded variation if and only if it can be expressed as the difference of two increasing functions:*

$$g(t) = a(t) - b(t)$$

Proof. (\Rightarrow direction). If g is of finite variation, $V_g(t) < \infty$ for all t , and we can write:

$$g(t) = V_g(t) - (V_g(t) - g(t))$$

Let $a(t) = V_g(t)$ and $b(t) = V_g(t) - g(t)$. Clearly, both $a(t)$ and $b(t)$ are increasing functions.

(\Leftarrow direction). Suppose a function g can be expressed as a difference of two bounded increasing functions. Then,

$$\begin{aligned} V_g(t) &= \lim_{\|\Delta_n\| \rightarrow 0} \sum_{i=1}^n |(a(t_i) - b(t_i)) - (a(t_{i-1}) - b(t_{i-1}))| \\ &\quad \{ \text{Telescoping sum} \} \\ &= a(t) - b(t) - (a(0) - b(0)) \end{aligned}$$

Since both $a(t)$ and $b(t)$ are bounded, g has bounded variation. \square

4.2.4 Riemann-Stieltjes Integral.

Let g be a monotonically increasing function on a finite closed interval $[a, b]$. A bounded function f defined on $[a, b]$ is said to *Riemann-Stieltjes integrable* with respect to g if the following limit exists:

$$\int_a^b f(t) dg(t) = \lim_{\|\Delta_n\| \rightarrow 0} \sum_{i=1}^n f(\tau_i)(g(t_i) - g(t_{i-1})) \quad (4.3)$$

where τ_i is an evaluation point in the interval $[t_{i-1}, t_i]$. It is a well-known fact that continuous functions are Riemann integrable and Riemann-Stieltjes integrable with respect to any monotonically increasing function on $[a, b]$.

We ask the following question. For any continuous functions f and g on $[a, b]$, can we define the integral $\int_a^b f(t) dg(t)$ by Equation (4.3)?

Consider the special case $f = g$, namely, the integral:

$$\int_a^b f(t)df(t)$$

Let $\Delta_n = \{a = t_0, t_1, \dots, t_n = b\}$ be a partition of $[a, b]$. Let L_n and R_n denote the corresponding Riemann sums with the evaluation points $\tau_i = t_{i-1}$ and $\tau_i = t_i$, respectively, namely,

$$L_n = \sum_{i=1}^n f(t_{i-1})(f(t_i) - f(t_{i-1})) \quad (4.4)$$

$$R_n = \sum_{i=1}^n f(t_i)(f(t_i) - f(t_{i-1})) \quad (4.5)$$

Is it true that, $\lim L_n = \lim R_n$ as $||\Delta_n|| \rightarrow 0$? Observe that:

$$R_n - L_n = \sum_{i=1}^n (f(t_i) - f(t_{i-1}))^2 \quad (4.6)$$

$$R_n + L_n = \sum_{i=1}^n (f(t_i)^2 - f(t_{i-1})^2) = f(b)^2 - f(a)^2 \quad (4.7)$$

Therefore, R_n and L_n are given by:

$$R_n = \frac{1}{2} \left(f(b)^2 - f(a)^2 + \sum_{i=1}^n (f(t_i) - f(t_{i-1}))^2 \right) \quad (4.8)$$

$$L_n = \frac{1}{2} \left(f(b)^2 - f(a)^2 - \sum_{i=1}^n (f(t_i) - f(t_{i-1}))^2 \right) \quad (4.9)$$

The limit of the right-hand side of equation (4.6) is called the *quadratic variation* of the function f on $[a, b]$. Obviously, $\lim_{||\Delta_n|| \rightarrow 0} R_n \neq \lim_{||\Delta_n|| \rightarrow 0} L_n$ if and only the quadratic variation of the function f is non-zero.

Example 4.7. Let f be a C^1 -function that is $f'(t)$ is a continuous function. Then, by the mean value theorem:

$$\begin{aligned}
|R_n - L_n| &= \sum_{i=1}^n (f(t_i) - f(t_{i-1}))^2 \\
&= \sum_{i=1}^n (f'(t_i^*)(t_i - t_{i-1}))^2 \\
&\quad \{\text{Mean Value Theorem}\} \\
&\leq \sum_{i=1}^n \|f'\|_\infty^2 (t_i - t_{i-1})^2 \\
&\quad \{\text{Interior Extremum Theorem}\} \\
&\leq \|f'\|_\infty^2 \|\Delta_n\| \sum_{i=1}^n (t_i - t_{i-1}) \\
&= \|f'\|_\infty^2 \|\Delta_n\| (b - a)
\end{aligned}$$

where $\|f'\|_\infty = \sup_{x \in [a, b]} f'(x)$. Thus, the limit as $\|\Delta_n\| \rightarrow 0$ of the distance $|R_n - L_n|$ also approaches zero. Thus, $\lim L_n = \lim R_n$ as $\|\Delta_n\| \rightarrow 0$ and the Riemann-Stieltjes integral exists. By equation (4.7), we have:

$$\lim_{\|\Delta_n\| \rightarrow 0} L_n = \lim_{\|\Delta_n\| \rightarrow 0} R_n = \frac{1}{2}(f(b)^2 - f(a)^2) \quad (4.10)$$

On the other hand, for such a C^1 -function f , we may simply define the integral $\int_a^b f(t)df(t)$ by:

$$\int_a^b f(t)df(t) = \int_a^b f(t)f'(t)dt$$

Then, by the fundamental theorem of Calculus:

$$\int_a^b f(t)df(t) = \int_a^b f(t)f'(t)dt = \frac{1}{2}f(t)^2 \Big|_a^b = \frac{1}{2}(f(b)^2 - f(a)^2)$$

Remark. There is a very close relationship between functions with bounded variation and functions for which the classical integral makes sense. For the Ito integral, the quadratic variation plays a similar role. The quadratic variation of a smooth function $f \in C^1([0, t])$ is zero.

Example 4.8. Suppose f is a continuous function satisfying the condition

$$|f(t) - f(s)| \leq C|t - s|^{1/2}$$

where $0 < C < 1$.

In this case we have:

$$0 \leq |R_n - L_n| \leq C^2 \sum_{i=1}^n (t_i - t_{i-1}) = C^2(b - a)$$

Hence, $\lim R_n \neq \lim L_n$ as $\|\Delta_n\| \rightarrow 0$ when $a \neq b$. Consequently, the integral $\int_a^b f(t)df(t)$ cannot be defined for such a function f . Observe that the quadratic variation of the function is $b - a$ (non-zero).

We see from the above examples, that defining the integral $\int_a^b f(t)dg(t)$ even when $f = g$ is a non-trivial problem. Consider the question posed earlier - if f and g are continuous functions on $[a, b]$, can we define the integral $\int_a^b f(t)dg(t)$? There is no simple answer to this question. But then in view of example (4.8), we can ask another question:

Question. Are there continuous functions f satisfying the condition

$$|f(t) - f(s)| \leq C|t - s|^{1/2}$$

4.2.5 Brownian motion as the limit of a symmetric random walk.

Consider a random walk starting at 0 with jumps h and $-h$ equally at times $\delta, 2\delta, \dots$ where h and δ are positive numbers. More precisely, let $\{X_n\}_{n=1}^\infty$ be a sequence of independent and identically distributed random variables with :

$$\mathbb{P}\{X_j = h\} = \mathbb{P}\{X_j = -h\} = \frac{1}{2}$$

Let $Y_{\delta,h}(0) = 0$ and put:

$$Y_{\delta,h}(n\delta) = X_1 + X_2 + \dots + X_n$$

For $t > 0$, define $Y_{\delta,h}(t)$ by linearization that is, for $n\delta < t < (n+1)\delta$, define:

$$Y_{\delta,h}(t) = \frac{(n+1)\delta - t}{\delta} Y_{\delta,h}(n\delta) + \frac{t - n\delta}{\delta} Y_{\delta,h}((n+1)\delta)$$

We can think of $Y_{\delta,h}(t)$ as the position of the random walk at time t . In particular, $X_1 + X_2 + \dots + X_n$ is the position of this random walk at time $n\delta$.

Question. What is the limit of the random walk $Y_{\delta,h}$ as $\delta, h \rightarrow 0$?

Recall that the characteristic function of a random variable X is $\phi_X(\lambda) = \mathbb{E} \exp[i\lambda X]$. In order to find out the answer, let us compute the following limit of the characteristic function of $Y_{\delta,h}(t)$:

$$\lim_{\delta, h \rightarrow 0} \mathbb{E} \exp [i\lambda Y_{\delta,h}(t)]$$

where $\lambda \in \mathbf{R}$ is fixed. For heuristic derivation, let $t = n\delta$ and so $n = t/\delta$. Then we have:

$$\begin{aligned} \mathbb{E} \exp [i\lambda Y_{\delta,h}(t)] &= \prod_{j=1}^n \mathbb{E} e^{i\lambda X_j} \\ &= \prod_{j=1}^n \left(\frac{1}{2} e^{i\lambda h} + \frac{1}{2} e^{-i\lambda h} \right) \\ &= \left(\frac{1}{2} e^{i\lambda h} + \frac{1}{2} e^{-i\lambda h} \right)^n \\ &= (\cos \lambda h)^n \\ &= (\cos \lambda h)^{t/\delta} \end{aligned}$$

For fixed t and λ , when δ and h independently approach 0, the limit of $\mathbb{E} \exp [i\lambda Y_{\delta,h}(t)]$ may not exist. For example, holding h constant, letting $\delta \rightarrow 0$, since $-1 \leq \cos \theta \leq 1$, the function $(\cos \lambda h)^{t/\delta} \rightarrow 0$. Holding δ constant, letting $h \rightarrow 0$, the function $(\cos \lambda h)^{t/\delta} \rightarrow 1$. In order for the limit to exist, we impose a certain relationship between δ and h . However, depending on the relationship, we may obtain different limits.

Let $u = \cos(\lambda h)^{1/\delta}$. Then $\ln u = \frac{1}{\delta} \ln \cos(\lambda h)$. Note that:

$$\cos(\lambda h) \approx 1 - \frac{1}{2} \lambda^2 h^2$$

And $\ln(1+x) \approx x$. Hence,

$$\ln \cos(\lambda h) \approx \ln \left(1 - \frac{1}{2} \lambda^2 h^2 \right) \approx -\frac{1}{2} \lambda^2 h^2$$

Therefore for small λ and h , we have $\ln u \approx -\frac{1}{2\delta} \lambda^2 h^2$ and so:

$$u \approx \exp \left[-\frac{1}{2\delta} \lambda^2 h^2 \right]$$

In particular, if δ and h are related by $h^2 = \delta$, then

$$\lim_{\delta \rightarrow 0} \mathbb{E} \exp [i\lambda Y_{\delta,h}(t)] = e^{-\frac{1}{2} \lambda^2 t}$$

But, $e^{-\frac{1}{2} \lambda^2 t}$ is the characteristic function of a Gaussian random variable with mean 0 and variance t . Thus, we have derived the following theorem about the limit of the random walk $Y_{\delta,h}$ as $\delta, h \rightarrow 0$ in such a way that $h^2 = \delta$.

Theorem 4.4. *Let $Y_{\delta,h}(t)$ be the random walk starting at 0 with jumps h and $-h$ equally likely at times $\delta, 2\delta, 3\delta, \dots$. Assume that $h^2 = \delta$. Then, for each $t \geq 0$, the limit:*

$$\lim_{\delta \rightarrow 0} Y_{\delta,h}(t) = B(t)$$

exists in distribution. Moreover, we have:

$$\mathbb{E}e^{i\lambda B(t)} = e^{-\frac{1}{2}\lambda^2 t}$$

Theorem 4.5. (*Quadratic Variation of a Brownian motion*). Let $(B_t, t \geq 0)$ be a standard brownian motion. Then, for any sequence of partitions $(t_j, j \leq n)$ of $[0, t]$ we have:

$$\langle B \rangle_t = \sum_{j=1}^n (B_{t_{j+1}} - B_{t_j})^2 \xrightarrow{L^2} t$$

where the convergence is in the L^2 sense.

Remark. It is reasonable to have some sort of convergence as we are dealing with a sum of independent random variables. However, the conclusion would not hold if the increments were not squared. So there is something more at play here.

Proof. We have:

$$\begin{aligned} \mathbb{E} \left[\left(\sum_{j=0}^{n-1} (B(t_{j+1}) - B(t_j))^2 - t \right)^2 \right] &= \mathbb{E} \left[\left(\sum_{j=0}^{n-1} (B(t_{j+1}) - B(t_j))^2 - \sum_{j=0}^{n-1} (t_{j+1} - t_j) \right)^2 \right] \\ &= \mathbb{E} \left[\left(\sum_{j=0}^{n-1} \{ (B(t_{j+1}) - B(t_j))^2 - (t_{j+1} - t_j) \} \right)^2 \right] \end{aligned}$$

For simplicity, we define the variables $X_j = (B(t_{j+1}) - B(t_j))^2 - (t_{j+1} - t_j)$. Then, we may write:

$$\begin{aligned} \mathbb{E} \left[\left(\sum_{j=0}^{n-1} (B(t_{j+1}) - B(t_j))^2 - t \right)^2 \right] &= \mathbb{E} \left[\left(\sum_{j=0}^{n-1} X_j \right)^2 \right] \\ &= \mathbb{E} \left[\sum_{i=0}^{n-1} \sum_{j=0}^{n-1} X_i X_j \right] \\ &= \sum_{i=0}^{n-1} \sum_{j=0}^{n-1} \mathbb{E}[X_i X_j] \end{aligned}$$

Now, the random variables X_j are independent.

The expectation of X_j is $\mathbb{E}[X_j] = \mathbb{E}(B(t_{j+1}) - B(t_j))^2 - (t_{j+1} - t_j) = 0$.

Since, X_i and X_j are independent, for $i \neq j$, $\mathbb{E}[X_i X_j] = \mathbb{E}X_i \cdot \mathbb{E}X_j = 0$.

Hence, we have:

$$\mathbb{E} \left[\left(\sum_{j=0}^{n-1} (B(t_{j+1}) - B(t_j))^2 - t \right)^2 \right] = \sum_{i=0}^{n-1} \mathbb{E}[X_i^2]$$

We now develop the expectation of the square of X_i . We have:

$$\begin{aligned} \mathbb{E}[X_i^2] &= \mathbb{E} \left[\left((B(t_{i+1}) - B(t_i))^2 - (t_{i+1} - t_i) \right)^2 \right] \\ &= \mathbb{E} \left[((B(t_{i+1}) - B(t_i))^4 - 2(B(t_{i+1}) - B(t_i))^2(t_{i+1} - t_i) + (t_{i+1} - t_i)^2) \right] \end{aligned}$$

The MGF of the random variable $B(t_{i+1}) - B(t_i)$ is :

$$\begin{aligned} \phi(\lambda) &= \exp \left[\frac{\lambda^2(t_{i+1} - t_i)}{2} \right] \\ \phi'(\lambda) &= \lambda(t_{i+1} - t_i) \exp \left[\frac{\lambda^2(t_{i+1} - t_i)}{2} \right] \\ \phi''(\lambda) &= [(t_{i+1} - t_i) + \lambda^2(t_{i+1} - t_i)^2] \exp \left[\frac{\lambda^2(t_{i+1} - t_i)}{2} \right] \\ \phi^{(3)}(\lambda) &= [3\lambda(t_{i+1} - t_i)^2 + \lambda^3(t_{i+1} - t_i)^3] \exp \left[\frac{\lambda^2(t_{i+1} - t_i)}{2} \right] \\ \phi^{(4)}(\lambda) &= [3(t_{i+1} - t_i)^2 + 6\lambda^2(t_{i+1} - t_i)^3 + \lambda^4(t_{i+1} - t_i)^4] \exp \left[\frac{\lambda^2(t_{i+1} - t_i)}{2} \right] \end{aligned}$$

Thus, $\mathbb{E}[(B(t_{i+1}) - B(t_i))^4] = 3(t_{i+1} - t_i)^2$. Consequently,

$$\begin{aligned} \mathbb{E}[X_i^2] &= \mathbb{E}[(B(t_{i+1}) - B(t_i))^4] - 2(t_{i+1} - t_i)\mathbb{E}[(B(t_{i+1}) - B(t_i))^2] + (t_{i+1} - t_i)^2 \\ &= 3(t_{i+1} - t_i)^2 - 2(t_{i+1} - t_i)^2 + (t_{i+1} - t_i)^2 \\ &= 2(t_{i+1} - t_i)^2 \end{aligned}$$

Putting all this together, we finally have that:

$$\begin{aligned}
\mathbb{E} \left[\left(\sum_{j=0}^{n-1} (B(t_{j+1}) - B(t_j))^2 - t \right)^2 \right] &= 2 \sum_{i=0}^{n-1} (t_{i+1} - t_i)^2 \\
&\leq 2 \|\Delta_n\| \sum_{i=0}^{n-1} (t_{i+1} - t_i) \\
&= 2 \|\Delta_n\| \cdot t
\end{aligned} \tag{4.11}$$

As $n \rightarrow \infty$, $\|\Delta_n\| \rightarrow 0$. Hence,

$$\lim_{n \rightarrow \infty} \mathbb{E} \left[\left(\sum_{j=0}^{n-1} (B(t_{j+1}) - B(t_j))^2 - t \right)^2 \right] = 0$$

Hence, the sequence of random variables

$$\sum_{j=0}^{n-1} (B(t_{j+1}) - B(t_j))^2 \xrightarrow{L^2} t$$

□

Corollary 4.1. (*Quadratic Variation of a Brownian Motion Path*). Let $(B_s, s \geq 0)$ be a Brownian motion. For every $n \in \mathbf{N}$, consider the dyadic partition $(t_j, j \leq 2^n)$ of $[0, t]$ where $t_j = \frac{j}{2^n}t$. Then we have that:

$$\langle B \rangle_t = \sum_{j=1}^{2^n-1} (B_{t_{j+1}} - B_{t_j})^2 \xrightarrow{a.s.} t$$

Proof. We have $(t_{i+1} - t_i) = \frac{t}{2^n}$. Borrowing equation (4.11) from the proof of theorem (4.5), we have that:

$$\begin{aligned}
\mathbb{E} \left[\left(\sum_{j=0}^{2^n-1} (B(t_{j+1}) - B(t_j))^2 - t \right)^2 \right] &= 2 \sum_{i=0}^{2^n-1} \left(\frac{t}{2^n} \right)^2 \\
&= 2 \cdot (2^n) \cdot \frac{t^2}{2^{2n}} \\
&= \frac{2t^2}{2^n}
\end{aligned}$$

By Chebyshev's inequality,

$$\begin{aligned}
\mathbb{P} \left(\left| \sum_{j=0}^{2^n-1} (B(t_{j+1}) - B(t_j))^2 - t \right| > \epsilon \right) &\leq \frac{1}{\epsilon^2} \mathbb{E} \left[\left(\sum_{j=0}^{2^n-1} (B(t_{j+1}) - B(t_j))^2 - t \right)^2 \right] \\
&\leq \frac{1}{\epsilon^2} \cdot \frac{2t^2}{2^n}
\end{aligned}$$

Define $A_n := \left\{ \left| \sum_{j=0}^{2^n-1} (B(t_{j+1}) - B(t_j))^2 - t \right| > \epsilon \right\}$. Since, $\sum \frac{1}{2^n}$ is a convergent series, any multiple of it, $(2t^2/\epsilon^2) \sum \frac{1}{2^n}$ also converges. Now, $0 \leq \mathbb{P}(A_n) \leq \frac{(2t^2/\epsilon^2)}{2^n}$. By the comparison test, $\sum \mathbb{P}(A_n)$ converges to a finite value. By Theorem (3.21),

$$\sum_{j=0}^{2^n-1} (B(t_{j+1}) - B(t_j))^2 \xrightarrow{a.s.} t$$

□

We are now ready to show that every Brownian motion path has infinite variation.

If g is a C^1 function,

$$\begin{aligned}
\int_0^t |g'(t)| dt &= \int_0^t \sqrt{g'(t)^2} dt \\
&\leq \int_0^t \sqrt{1 + g'(t)^2} dt \\
&= l_g(t)
\end{aligned}$$

where $l_g(t)$ is the arclength of the function g between $[0, t]$. So, $V_g(t) \leq l_g(t)$ and further:

$$\begin{aligned} l_g(t) &= \int_0^t \sqrt{1 + g'(t)^2} dt \\ &\leq \int_0^t \left(1 + \sqrt{g'(t)^2}\right) dt \\ &\leq t + V_g(t) \end{aligned}$$

Consequently,

$$V_g(t) \leq l_g(t) \leq t + V_g(t)$$

The total variation of the function is finite if and only if it's arclength is.

Hence, intuitively, our claim is that a Brownian motion path on $[0, T]$ has infinite arclength. Since $g \in C^1([a, b]) \implies (V_g(t) < \infty)$, it follows that $(V_g(t) \rightarrow \infty) \implies g \notin C^1$.

Corollary 4.2. *(Brownian Motion paths have unbounded total variation.) Let $(B_s, s \geq 0)$ be a Brownian motion. Then, the random functions $B(s, \omega)$ on the interval $[0, t]$ have unbounded variation almost surely.*

Proof. Take the sequence of dyadic partitions of $[0, t]$: $t_j = \frac{j}{2^n}t, n \in \mathbf{N}, j \leq 2^n$. By pulling out the worst increment, we have the trivial bound for every ω :

$$\sum_{j=0}^{2^n-1} (B_{t_{j+1}}(\omega) - B_{t_j}(\omega))^2 \leq \max_{0 \leq j \leq 2^n} |B_{t_{j+1}}(\omega) - B_{t_j}(\omega)| \cdot \sum_{j=0}^{2^n-1} (B_{t_{j+1}}(\omega) - B_{t_j}(\omega)) \quad (4.12)$$

We proceed by contradiction. Let A' be the set of all ω , for which the Brownian motion paths have bounded total variation. Let A be event that the Brownian motion paths have unbounded variation.

By the definition of total variation, that would imply, $\exists M \in \mathbf{N}$:

$$(\forall \omega \in A') \quad \lim_{n \rightarrow \infty} \sum_{j=0}^{2^n-1} |(B_{t_{j+1}}(\omega) - B_{t_j}(\omega))| < M$$

Since Brownian Motion paths are continuous on the compact set $[\frac{j}{2^n}t, \frac{j+1}{2^n}t]$, they are uniformly continuous. So, as $n \rightarrow \infty$, $|t_{j+1} - t_j| \rightarrow 0$ and therefore $|B_{t_{j+1}}(\omega) - B_{t_j}(\omega)| \rightarrow 0$. And consequently, $\max_{0 \leq j \leq 2^n} |B_{t_{j+1}}(\omega) - B_{t_j}(\omega)| \rightarrow 0$.

Thus, for every $\omega \in A'$, the right hand side of the inequality (4.12), converges to 0 and therefore the left hand side converges to 0. But, this contradicts the fact that $\langle B \rangle_t \xrightarrow{a.s.} t$. So, A' is a null set, and $\mathbb{P}(A') = 0$ and $\mathbb{P}(A) = 1$. This closes the proof. \square

4.3 What exactly is $(\Omega, \mathcal{F}, \mathbb{P})$ in mathematical finance?

If we make the simplifying assumption that the process paths are continuous, we obtain the set of all continuous functions on $[0, T]$, denoted by $C[0, T]$. This is a very rich space. In a more general model, it is assumed that the process paths are right continuous with left limits (regular right-continuous RRC, cadlag) functions.

Let the sample space $\Omega = D[0, T]$ be the set of all RRC functions on $[0, T]$. An element of this set is a RRC function from $[0, T]$ into \mathbf{R} . First we must decide what kind of sets of these functions are measurable? The simplest set for which we would like to calculate the probabilities are sets of the form $\{a \leq S(t_1) \leq b\}$ for some t_1 . If $S(t)$ represents the price of a stock at time t , then the probability of such a set gives the probability that the stock price at time t_1 is between a and b . We are also interested in how the price of the stock at time t_1 affects the price at another time t_2 . Thus, we need to talk about the joint distribution of stock prices $S(t_1)$ and $S(t_2)$. This means that we need to define probability on the sets of the form $\{S(t_1) \in B_1, S(t_2) \in B_2\}$ where B_1 and B_2 are intervals on the line. More generally, we would like to have all the finite-dimensional distributions of the process $S(t)$, that is, the probabilities of the sets: $\{S(t_1) \in B_1, S(t_2) \in B_2, \dots, S(t_n) \in B_n\}$ for any choice of $0 \leq t_1 \leq \dots \leq t_n \leq T$.

The sets of the form $A = \{\omega(\cdot) \in D[0, T] : \omega(t_1) \in B_1, \dots, \omega(t_n) \in B_n\}$, where B_i 's are borel subsets of \mathbf{R} , are called cylinder sets or finite-dimensional rectangles.

The stochastic process $S(t)$ is just a (function-valued) random variable on this sample space, which takes some value $\omega(t)$ - the value of the function ω at t .

Let \mathcal{R} be the collection of all cylindrical subsets of $D[0, 1]$. Obviously \mathcal{R} is not a σ -field.

Probability is first defined by on the elements of \mathcal{R} . Let $A \subseteq \mathcal{R}$.

$$\mathbb{P}(A) = \int_{B_1} \cdots \int_{B_n} \prod_{i=1}^n \frac{1}{\sqrt{(2\pi)(t_i - t_{i-1})}} \exp \left[-\frac{(u_i - u_{i-1})^2}{2(t_i - t_{i-1})} \right] du_1 \cdots du_n$$

and then extended to the σ -field generated by taking unions, complements and intersections of cylinders. We take the smallest σ -algebra containing all the cylindrical subsets of $D[0, 1]$. Thus, $\mathcal{F} = \mathcal{B}(D[0, 1])$.

Hence, $(\Omega, \mathcal{F}, \mathbb{P}) = (D[0, 1], \mathcal{B}(D[0, 1]), \mathbb{P})$ is a probability space. It is called the *Wiener space* and \mathbb{P} here is called the *Wiener measure*.

4.4 Continuity and Regularity of paths.

As discussed in the previous section, a stochastic process is determined by its finite-dimensional distribution. In studying stochastic processes, it is often natural to think of them as function-valued random variables in t . Let $S(t)$ be defined for $0 \leq t \leq T$, then for a fixed ω , it is a function in t , called the sample path or a realization of S . Finite-dimensional distributions do not determine the continuity property of sample paths. The following example illustrates this.

Example 4.9. Let $X(t) = 0$ for all t , $0 \leq t \leq 1$ and τ be a uniformly distributed random variable on $[0, 1]$. Let $Y(t) = 0$ for $t \neq \tau$ and $Y(t) = 1$ if $t = \tau$. Then, for any fixed t , $\mathbb{P}(Y(t) \neq 0) = \mathbb{P}(\tau = t) = 0$, and hence $\mathbb{P}(Y(t) = 0) = 1$. So, that all one-dimensional distributions of $X(t)$ and $Y(t)$ are the same. Similarly, all finite-dimensional distributions of X and Y are the same. However, the sample paths of the process X , that is, the functions $X(t)_{0 \leq t \leq 1}$ are continuous in t , whereas every sample path $Y(t)_{0 \leq t \leq 1}$ has a jump at the (random) point τ . Notice that, $\mathbb{P}(X(t) = Y(t)) = 1$ for all t , $0 \leq t \leq 1$.

Definition 4.7. Two stochastic processes are called *versions* (modifications) of one another if

$$\mathbb{P}(X(t) = Y(t)) = 1 \quad \text{for all } 0 \leq t \leq T$$

Thus, the two processes in the example (4.9) are versions of one another, one has continuous sample paths, the other does not. If we agree to pick any version of the process we want, then we can pick the continuous version when it exists. In general, we choose the smoothest possible version of the process.

For two processes, X and Y , denote by $N_t = \{X(t) \neq Y(t)\}$, $0 \leq t \leq T$. In the above example, $\mathbb{P}(N_t) = \mathbb{P}(\tau = t) = 0$ for any t , $0 \leq t \leq 1$. However, $\mathbb{P}(\bigcup_{0 \leq t \leq 1} N_t) = \mathbb{P}(\tau = t \text{ for some } t \text{ in } [0, 1]) = 1$. Although, each of N_t is a \mathbb{P} -null set, the union $N = \bigcup_{0 \leq t \leq 1} N_t$ contains uncountably many null sets, and in this particular case it is a set of probability one.

If it happens that $\mathbb{P}(N) = 0$, then N is called an *evanescent set*, and the processes X and Y are called *indistinguishable*. Note that in this case, $\mathbb{P}(\{\omega : \exists t : X(t) \neq Y(t)\}) = \mathbb{P}(\bigcup_{0 \leq t \leq 1} \{X(t) \neq Y(t)\}) = 0$ and $\mathbb{P}(\bigcap_{0 \leq t \leq 1} \{X(t) = Y(t)\}) = 1$. It is clear, that if the time is discrete, then any two versions of the process are indistinguishable. It is also not hard to see, that if $X(t)$ and $Y(t)$ are versions of one another and they are both right-continuous, they are indistinguishable.

Theorem 4.6. (*Paul Levy's construction of Brownian Motion*). *Standard Brownian motion exists.*

Proof. I reproduce the standard proof as present in *Brownian Motion* by Morters and Peres. I added some remarks for greater clarity.

Let

$$\mathcal{D}_n = \left\{ \frac{k}{2^n} : k = 0, 1, 2, \dots, 2^n \right\}$$

be a finite set of dyadic points.

Let

$$\mathcal{D} = \bigcup_{n=0}^{\infty} \mathcal{D}_n$$

Let $\{Z_t : t \in \mathcal{D}\}$ be a collection of independent, standard normally distributed random variables. This is a countable set of random variables.

Let $B(0) := 0$ and $B(1) := Z_1$.

For each $n \in \mathbf{N}$, we define the random variables $B(d)$, $d \in \mathcal{D}_n$ such that, the following invariant holds:

- (1) for all $r < s < t$ in \mathcal{D}_n the random variable $B(t) - B(s)$ is normally distributed with mean zero and variance $t - s$ and is independent of $B(s) - B(r)$.
- (2) the vectors $(B(d) : d \in \mathcal{D}_n)$ and $(Z_t : t \in \mathcal{D} \setminus \mathcal{D}_n)$ are independent.

Note that we have already done this for $\mathcal{D}_0 = \{0, 1\}$. Proceeding inductively, let's assume that the above holds for some $n - 1$. We are interested to prove that the invariant also holds for n .

We define $B(d)$ for $d \in \mathcal{D}_n \setminus \mathcal{D}_{n-1}$ by:

$$B(d) = \frac{B(d - 2^{-n}) + B(d + 2^{-n})}{2} + \frac{Z_d}{2^{(n+1)/2}}$$

Note that, the points $0, \frac{1}{2^{n-1}}, \dots, \frac{k}{2^{n-1}}, \frac{k+1}{2^{n-1}}, \dots, 1$ belong to \mathcal{D}_{n-1} . The first summand is the linear interpolation of the values of B at the neighbouring points of d in \mathcal{D}_{n-1} . That is,

$$B\left(\frac{2k+1}{2^n}\right) = \frac{B\left(\frac{k}{2^{n-1}}\right) + B\left(\frac{k+1}{2^{n-1}}\right)}{2} + \frac{Z_d}{2^{(n+1)/2}}$$

Since $P(n - 1)$ holds, $B(d - 2^{-n})$ and $B(d + 2^{-n})$ have no dependence on $(Z_t : t \in \mathcal{D} \setminus \mathcal{D}_{n-1})$. Consequently, $B(d)$ has no dependence on $(Z_t : t \in \mathcal{D} \setminus \mathcal{D}_n)$ and the second property is fulfilled.

Moreover, as $\frac{1}{2}[B(d + 2^{-n}) - B(d - 2^{-n})]$ depends only on $(Z_t : t \in \mathcal{D}_{n-1})$, it is independent of $\frac{Z_d}{2^{(n+1)/2}}$. By our induction assumptions, they are both normally distributed with mean 0 and variance $\frac{1}{2^{(n+1)}}$.

So, their sum and difference random variables

$$\begin{aligned} B(d) - B(d - 2^{-n}) &= \frac{B(d + 2^{-n}) - B(d - 2^{-n})}{2} + \frac{Z_d}{2^{(n+1)/2}} \\ B(d + 2^{-n}) - B(d) &= \frac{B(d + 2^{-n}) - B(d - 2^{-n})}{2} - \frac{Z_d}{2^{(n+1)/2}} \end{aligned}$$

are also independent, with mean 0 and variance $\frac{1}{2^n}$ (the variance of independent random variables is the sum of the variances).

Indeed all increments $B(d) - B(d - 2^{-n})$ for $d \in \mathcal{D}_n \setminus \{0\}$ are independent. To see this, it suffices to show that they are pairwise independent. We have seen in the previous paragraph that the pairs $B(d) - B(d - 2^{-n})$ and $B(d + 2^{-n}) - B(d)$ with $d \in \mathcal{D}_n \setminus \mathcal{D}_{n-1}$ are independent. The other possibility is that the increments are over the intervals separated by some $d \in \mathcal{D}_{n-1}$. For concreteness, if n were 3, then the increments, $B_{7/8} - B_{6/8}$ and $B_{5/8} - B_{4/8}$ are separated by $d = \frac{3}{4} \in \mathcal{D}_2$. Choose $d \in \mathcal{D}_j$ with this property and minimal j , so, the two intervals are contained in $[d - 2^{-j}, d]$ and $[d, d + 2^{-j}]$ respectively. By induction, the increments over these two intervals of length 2^{-j} are independent and the increments over the intervals of length 2^{-n} are constructed from the independent increments $B(d) - B(d - 2^{-j})$ and $B(d + 2^{-j}) - B(d)$ using a disjoint set of variables $(Z_t : t \in \mathcal{D}_n)$. Hence, they are independent and this implies pairwise independence. This implies the first property. Consequently, the vector of increments $(B(d) - B(d - 2^{-n}))$ for all $d \in \mathcal{D}_n$ is Gaussian.

Having thus chosen the value of the process on all the dyadic points, we interpolate between them. Formally, we define:

$$F_0(t) = \begin{cases} Z_1 & \text{for } t = 1 \\ 0 & \text{for } t = 0 \\ \text{linear in between} & \end{cases}$$

and for each $n \geq 1$,

$$F_n(t) = \begin{cases} \frac{Z_t}{2^{(n+1)/2}} & \text{for } t \in \mathcal{D} \setminus \mathcal{D}_{n-1} \\ 0 & \text{for } t \in \mathcal{D}_{n-1} \\ \text{linear between consecutive points in } \mathcal{D}_n & \end{cases}$$

These functions are continuous on $[0, 1]$ and for all n and $d \in \mathcal{D}_n$, we have:

$$B(d) = \sum_{i=0}^n F_i(d) = \sum_{i=0}^{\infty} F_i(d) \quad (4.13)$$

To see this, assume that above equation holds for all $d \in \mathcal{D}_{n-1}$.

Let's consider the point $d \in \mathcal{D}_n \setminus \mathcal{D}_{n-1}$.

$$\begin{aligned} B(d) &= \frac{B(d - 2^{-n}) + B(d + 2^{-n})}{2} + \frac{Z_d}{2^{(n+1)/2}} \\ &= \sum_{i=0}^{n-1} \frac{F_i(d - 2^{-n}) + F_i(d + 2^{-n})}{2} + \frac{Z_d}{2^{(n+1)/2}} \end{aligned} \quad (4.14)$$

Now, $d - 2^{-n}$ and $d + 2^{-n}$ belong to \mathcal{D}_{n-1} and are not in $\bigcup_{i < n-1} \mathcal{D}_i$. Therefore, for $i = 0, 1, \dots, n-2$, the points $(d - 2^{-n}, F_i(d - 2^{-n}))$ and $(d + 2^{-n}, F_i(d + 2^{-n}))$ lie on some straight line and have $(d, F_i(d))$ as their midpoint. Moreover, $d - 2^{-n}$ and $d + 2^{-n}$ are vertices in \mathcal{D}_{n-1} . So, by definition of $F_{n-1}(d)$, we have $F_{n-1}(d) = [F_{n-1}(d - 2^{-n}) + F_{n-1}(d + 2^{-n})]/2$.

To summarize, the first term on the right hand side of expression (4.14) is equal to $\sum_{i=0}^{n-1} F_i(d)$. By mathematical induction, it follows that the claim (4.13) is true for all $n \in \mathbf{N}$.

It's extremely easy to find an upper bound on the probability contained in the Gaussian tails. Suppose $X \sim N(0, 1)$ and let $x > 0$. We are interested in the tail probability $\mathbb{P}(X > x)$. We have:

$$\mathbb{P}(X > x) = \int_x^\infty e^{-x^2/2} dx = \int_x^\infty \frac{x e^{-x^2/2} dx}{x}$$

Let $u = \frac{1}{x}$ and $dv = x e^{-x^2/2} dx$. We have:

$$\frac{u = \frac{1}{x}}{du = -\frac{1}{x^2} dx} \left| \begin{array}{l} dv = x e^{-x^2/2} dx \\ v = -e^{-x^2/2} \end{array} \right.$$

Thus,

$$\begin{aligned}
\mathbb{P}(X > x) &= -\frac{1}{x}e^{-x^2/2}\Big|_x^\infty - \int_x^\infty \frac{e^{-x^2/2}}{x^2}dx \\
&= \frac{e^{-x^2/2}}{x} - \int_x^\infty \frac{e^{-x^2/2}}{x^2}dx \\
&\quad \left\{ I(x) = \int_x^\infty \frac{e^{-x^2/2}}{x^2} \geq 0 \right\} \\
&\leq \frac{e^{-x^2/2}}{x}
\end{aligned}$$

Thus, for $c > 1$ and large n , we have:

$$\mathbb{P}(|Z_d| \geq c\sqrt{n}) \leq \frac{1}{c\sqrt{n}}e^{-c^2n/2} \leq \exp\left(-\frac{c^2n}{2}\right)$$

So, the series:

$$\begin{aligned}
\sum_{n=0}^{\infty} \mathbb{P}\left\{\text{There exists atleast one } d \in \mathcal{D}_n \text{ with } |Z_d| \geq c\sqrt{n}\right\} &\leq \sum_{n=0}^{\infty} \sum_{d \in \mathcal{D}_n} \mathbb{P}\left\{|Z_d| \geq c\sqrt{n}\right\} \\
&\leq \sum_{n=0}^{\infty} (2^n + 1) \exp\left(-\frac{c^2n}{2}\right)
\end{aligned}$$

Now, the series (a_n) given by, $a_n := (2^n + 1)e^{-c^2n/2}$ has the ratio between successive terms:

$$\begin{aligned}
\lim \left| \frac{a_{n+1}}{a_n} \right| &= \lim_{n \rightarrow \infty} \frac{2^{n+1} + 1}{2^n + 1} \cdot \frac{e^{(c^2n)/2}}{e^{c^2(n+1)/2}} \\
&= \lim_{n \rightarrow \infty} \frac{\frac{1}{2} + \frac{1}{2^n}}{1 + \frac{1}{2^n}} \cdot \frac{1}{e^{c^2/2}} \\
&= \frac{1}{2e^{c^2/2}}
\end{aligned}$$

If this ratio is less than unity, that is $c > \sqrt{2 \log 2}$, then by the ratio test, $\sum (2^n + 1)e^{-c^2 n/2}$ converges to a finite value. Fix such a c .

By BCL1 (Borel-Cantelli Lemma), if $A_n := \{\text{There exists atleast one } d \in \mathcal{D}_n \text{ with } |Z_d| \geq c\sqrt{n}\}$ and $\sum_{n=0}^{\infty} \mathbb{P}(A_n)$ converges to a finite value, then the event A_n occurs finitely many times with probability 1. There exists $N \in \mathbf{N}$, such that for all $n \geq N$, A_n fails to occur with probability 1. Thus, for all $n \geq N$, $\{Z_d \leq c\sqrt{n}\}$ occurs with probability 1. It follows that:

$$\sup_{t \in [0,1]} F_n(t) \leq \frac{c\sqrt{n}}{2^{(n+1)/2}}$$

Define

$$M_n = \frac{c\sqrt{n}}{2^{(n+1)/2}}$$

Since $\sum M_n$ converges, by the Weierstrass M -test, the infinite series of functions $\sum_{n=0}^{\infty} F_n(t)$ converges uniformly on $[0, 1]$. Since, each $F_n(t)$ is piecewise linear and continuous, by the Term-by-Term continuity theorem, $\sum_{n=0}^{\infty} F_n(t)$ is continuous on $[0, 1]$. \square

4.5 A point of comparison: The Poisson Process.

Like the Brownian motion, the Poisson process is defined as a process with stationary and independent increments.

Definition 4.8. A process $(N_t, t \geq 0)$ defined on $(\Omega, \mathcal{F}, \mathbb{P})$ has the distribution of the Poisson process with rate $\lambda > 0$, if and only if the following hold:

- (1) $N_0 = 0$.
- (2) For any $s < t$, the increment $N_t - N_s$ is a Poisson random variable with parameter $\lambda(t - s)$.
- (3) For any $n \in \mathbf{N}$ and any choice $0 < t_1 < t_2 < \dots < t_n < \infty$, the increments $N_{t_2} - N_{t_1}, N_{t_3} - N_{t_2}, \dots, N_{t_n} - N_{t_{n-1}}$ are independent.


Poisson paths can be sampled using this definition. By construction, it is not hard to see that the paths of Poisson processes are piecewise, constant, integer-valued and non-decreasing. In particular, the paths of Poisson processes have finite variation. Poisson paths are much simpler than the ones of Brownian motion in many ways!

Example 4.10. (Simulating the Poisson Process.) Use the definition (4.8) to generate 10 paths of the Poisson process with rate 1 on the interval $[0, 10]$ with step-size 0.01.

Listing 8: Generating 10 paths of a Poisson process

```
def generatePoissonProcess(lam,T,stepSize):
    N = int(T/stepSize)
    x = np.random.poisson(lam=lam,size=N)
    y = np.cumsum(x)
    y = np.concatenate([[0.0],y])
    return y
```

10 paths of the Poisson process on $[0, 10]$



We can construct a Poisson process as follows. Consider $(\tau_j, j \in \mathbf{N})$ IID exponential random variables with parameter $1/\lambda$. One should think of τ_j as the waiting time from the $(j - 1)$ st to the j th jump. Then, one defines :

$$N_t = \#\{k : \tau_1 + \tau_2 + \dots + \tau_k \leq t\}$$

= Number of jumps upto and including time t

Now, here is an idea! What about defining a new process with stationary and independent increments using a given distribution other than Poisson and Gaussian? Is this even possible? The answer is yes, but only if the distribution satisfies the property of being *infinitely divisible*. To see this, consider the value of the process at time 1, N_1 . Then, no matter how many subintervals we chop the interval $[0, 1]$ into, we must have the increments add up to N_1 . In other words, we must be able to write N_1 as a sum of n IID random variables for every possible n . This is certainly true for Poisson random variables and Gaussian random variables. Another example is the Cauchy distribution. In general, processes that can be constructed using independent, stationary increments are called Levy processes.

Example 4.11. Time Inversion. Let $(B_t, t \geq 0)$ be a standard brownian motion. We consider the process:

$$X_t = tB_{1/t} \quad \text{for } t > 0$$

This property relates the behavior of t large to the behavior of t small.

(a) Show that $(X_t, t > 0)$ has the distribution of Brownian motion on $t > 0$.

Proof.

Like $B(t)$, it is an easy exercise to prove that $X(t)$ is also a Gaussian process.

We have, $\mathbb{E}[X_s] = 0$.

Let $s < t$. We have:

$$\begin{aligned} \text{Cov}(X_s, X_t) &= \mathbb{E}[sB(1/s) \cdot tB(1/t)] \\ &= st\mathbb{E}[B(1/s) \cdot B(1/t)] \\ &= st \cdot \frac{1}{t} \\ &\quad \left\{ \because \frac{1}{t} < \frac{1}{s} \right\} \\ &= s \end{aligned}$$

Consequently, $X(t)$ has the distribution of a Brownian motion.

(b) Argue that $X(t)$ converges to 0 as $t \rightarrow 0$ in the sense of L^2 -convergence. It is possible to show convergence almost surely so that $(X_t, t \geq 0)$ is really a Brownian motion for $t \geq 0$.

Solution.

Let (t_n) be any arbitrary sequence of positive real numbers approaching 0 and consider the sequence of random variables $(X(t_n))_{n=1}^\infty$. We have:

$$\begin{aligned}\mathbb{E}[X(t_n)^2] &= \mathbb{E}[t_n^2 B(1/t_n)^2] \\ &= t_n^2 \mathbb{E}[B(1/t_n)^2] \\ &= t_n^2 \cdot \frac{1}{t_n} \\ &= t_n\end{aligned}$$

Hence,

$$\lim \mathbb{E}[X(t_n)^2] = \lim t_n = 0$$

Since (t_n) was an arbitrary sequence, it follows that $\lim_{t \rightarrow 0} \mathbb{E}[(X(t))^2] = 0$.

(c) Use this property of Brownian motion to show the law of large numbers for Brownian motion:

$$\lim_{t \rightarrow \infty} \frac{X(t)}{t} = 0 \quad \text{almost surely}$$

Solution.

What we need to do is to show that $X(t) \rightarrow 0$ as $t \rightarrow 0$ almost surely. That would show that $\frac{B(1/t)}{1/t} \rightarrow 0$ as $t \rightarrow 0$ almost surely, which is the same as showing $\frac{B(t)}{t} \rightarrow 0$ as $t \rightarrow \infty$, which is the law of large numbers for Brownian motion.

What we have done in part (b), is to prove the claim that $\mathbb{E}[X(t)^2] \rightarrow 0$ as $t \rightarrow 0$, which shows convergence in the L^2 sense and hence convergence in probability. This is in fact the weak law of large numbers. $\frac{B(t)}{t} \xrightarrow{\mathbf{P}} 0$ as $t \rightarrow \infty$.

For $t > 0$, continuity is clear. However, it is the proof that as $t \rightarrow 0$, $X(t) \rightarrow 0$ almost surely which we have not done.

Note that, the limit $X(t) \rightarrow 0$ as $t \rightarrow 0$ if and only if $(\forall n \geq 1), (\exists m \geq 1)$, such that $\forall r \in \mathbb{Q} \cap (0, \frac{1}{m}]$, we have $|X(r)| = |rB(\frac{1}{r})| \leq \frac{1}{n}$.

To understand the above, we just recall the $\epsilon - \delta$ definition of continuity. Note that $\frac{1}{n}$ plays the role of ϵ and $\frac{1}{m}$ works as δ .

That is,

$$\Omega^X := \left\{ \lim_{t \rightarrow 0} X(t) = 0 \right\} = \bigcap_{n \geq 1} \bigcup_{m \geq 1} \bigcap_{r \in \mathbb{Q} \cap (0, \frac{1}{m}]} \left\{ |X(r)| \leq \frac{1}{n} \right\}$$

Also, note that $X(t)$ is continuous on all $[a, 1]$ for all $a > 0$, thus, uniformly continuous on $[a, 1]$, and hence uniformly continuous on $\mathbb{Q} \cap (0, 1]$. So, there exists a continuous extension of $X(t)$ on $[0, 1]$. We already know from part (a), that $(X(t))_{t>0}$ and $(B(t))_{t>0}$ have the same finite dimensional distributions. Therefore, the RHS event has the same probability as $\Omega^B := \bigcap_{n \geq 1} \bigcup_{m \geq 1} \bigcap_{r \in \mathbb{Q} \cap (0, \frac{1}{m}]} \left\{ |B(r)| \leq \frac{1}{n} \right\}$. Since $B(t) \rightarrow 0$ as $t \rightarrow 0$ almost surely, the event Ω^B has probability 1. Thus, $\mathbb{P} \{ \lim_{t \rightarrow 0} X(t) = 0 \} = 1$.

This actually shows that $X(t)$ is a bonafide standard brownian motion, as we have established continuity as well.

5 Martingales.

5.1 Elementary conditional expectation.

In elementary probability, the conditional expectation of a variable Y given another random variable X refers to the expectation of Y given the conditional distribution $f_{Y|X}(y|x)$ of Y given X . To illustrate this, let's go through a simple example. Consider $\mathcal{B}_1, \mathcal{B}_2$ to be two independent Bernoulli-distributed random variables with $p = 1/2$. Then, construct:

$$X = \mathcal{B}_1, \quad Y = \mathcal{B}_1 + \mathcal{B}_2$$

It is easy to compute $\mathbb{E}[Y|X = 0]$ and $\mathbb{E}[Y|X = 1]$. By definition, it is given by:

$$\begin{aligned}
 \mathbb{E}[Y|X = 0] &= \sum_{j=0}^2 j\mathbb{P}(Y = j|X = 0) \\
 &= \sum_{j=0}^2 j \cdot \frac{\mathbb{P}(Y = j, X = 0)}{P(X = 0)} \\
 &= 0 + 1 \cdot \frac{(1/4)}{(1/2)} + 2 \cdot \frac{0}{(1/2)} \\
 &= \frac{1}{2}
 \end{aligned}$$

and

$$\begin{aligned}
 \mathbb{E}[Y|X = 1] &= \sum_{j=0}^2 j\mathbb{P}(Y = j|X = 1) \\
 &= \sum_{j=0}^2 j \cdot \frac{\mathbb{P}(Y = j, X = 1)}{P(X = 1)} \\
 &= 0 + 1 \cdot \frac{(1/4)}{(1/2)} + 2 \cdot \frac{(1/4)}{(1/2)} \\
 &= \frac{3}{2}
 \end{aligned}$$

With this point of view, the conditional expectation is computed given the information that the event $\{X = 0\}$ occurred or the event $\{X = 1\}$ occurred. It is possible to regroup both conditional expectations in a single object, if we think of the conditional expectation as a random variable and denote it by $\mathbb{E}[Y|X]$. Namely, we take:

$$\mathbb{E}[Y|X](\omega) = \begin{cases} \frac{1}{2} & \text{if } X(\omega) = 0 \\ \frac{3}{2} & \text{if } X(\omega) = 1 \end{cases} \quad (5.1)$$

This random variable is called the *conditional expectation* of Y given X . We make two important observations:

- (i) If the value of X is known, then the value of $\mathbb{E}[Y|X]$ is determined.
- (ii) If we have another random variable $g(X)$ constructed from X , then we have:

$$\mathbb{E}[g(X)Y] = \mathbb{E}[g(X)\mathbb{E}[Y|X]]$$

In other words, as far as X is concerned, the conditional expectation $\mathbb{E}[Y|X]$ is a proxy for Y in the expectation. We sometimes say that $\mathbb{E}[Y|X]$ is the best estimate of Y given the information of X .

The last observation is easy to verify since:

$$\begin{aligned}\mathbb{E}[g(X)Y] &= \sum_{i=0}^1 \sum_{j=0}^2 g(i) \cdot j \cdot \mathbb{P}(X = i, Y = j) \\ &= \sum_{i=0}^1 \mathbb{P}(X = i) g(i) \left\{ \sum_{j=0}^2 j \cdot \frac{\mathbb{P}(X = i, Y = j)}{\mathbb{P}(X = i)} \right\} \\ &= \mathbb{E}[g(X)\mathbb{E}[Y|X]]\end{aligned}$$

Example 5.1. (Elementary Definitions of Conditional Expectation).

(1) (X, Y) discrete. The treatment is similar to the above. If a random variable X takes values $(x_i, i \geq 1)$ and Y takes values $(y_j, j \geq 1)$, we have by definition that the conditional expectation as a random variable is:

$$\mathbb{E}[Y|X](\omega) = \sum_{j \geq 1} y_j \mathbb{P}(Y = y_j | X = x_i) \quad \text{for } \omega \text{ such that } X(\omega) = x_i$$

(2) (X, Y) continuous with joint PDF $f_{X,Y}(x, y)$: In this case, the conditional expectation is the random variable given by

$$\mathbb{E}[Y|X] = h(X)$$

where

$$h(x) = \int_{\mathbf{R}} y f_{Y|X}(y|x) dy = \int_{\mathbf{R}} y \frac{f_{X,Y}(x, y)}{f_X(x)} dy = \frac{\int_{\mathbf{R}} y f_{X,Y}(x, y) dy}{\int_{\mathbf{R}} f_{X,Y}(x, y) dy}$$

In the two examples above, the expectation of the random variable $\mathbb{E}[Y|X]$ is equal to $\mathbb{E}[Y]$. Indeed in the discrete case, we have:

$$\begin{aligned} \mathbb{E}[\mathbb{E}[Y|X]] &= \sum_{i=0}^1 P(X = x_i) \cdot \sum_{j=0}^2 y_j \mathbb{P}(Y = y_j | X = x_i) \\ &= \sum_{i=0}^1 \sum_{j=0}^2 y_j \mathbb{P}(Y = y_j, X = x_i) \\ &= \sum_{j=0}^2 y_j \mathbb{P}(Y = y_j) \\ &= \mathbb{E}[Y] \end{aligned}$$

Example 5.2. (Conditional Probability vs Conditional expectation). The conditional probability of the event A given B can be recast in terms of conditional expectation using indicator functions. If $0 < \mathbb{P}(B) < 1$, it is not hard to check that: $\mathbb{P}(A|B) = \mathbb{E}[\mathbf{1}_A | \mathbf{1}_B = 1]$ and $\mathbb{P}(A|B^C) = \mathbb{E}[\mathbf{1}_A | \mathbf{1}_B = 0]$. Indeed the random variables $\mathbf{1}_A$ and $\mathbf{1}_B$ are discrete. If we proceed as in the discrete case above, we have:

$$\begin{aligned} \mathbb{E}[\mathbf{1}_A | \mathbf{1}_B = 1] &= 1 \cdot \mathbb{P}(\mathbf{1}_A = 1 | \mathbf{1}_B = 1) \\ &= \frac{\mathbb{P}(\mathbf{1}_A = 1, \mathbf{1}_B = 1)}{\mathbb{P}(\mathbf{1}_B = 1)} \\ &= \frac{\mathbb{P}(A \cap B)}{\mathbb{P}(B)} \\ &= \mathbb{P}(A|B) \end{aligned}$$

A similar calculation gives $\mathbb{P}(A|B^C)$. In particular, the formula for total probability for A is a rewriting of the expectation of the random variable $\mathbb{E}[\mathbf{1}_A | \mathbf{1}_B]$:

$$\begin{aligned}
\mathbb{E}[\mathbb{E}[\mathbf{1}_A|\mathbf{1}_B]] &= \mathbb{E}[\mathbf{1}_A|\mathbf{1}_B = 1]\mathbb{P}(\mathbf{1}_B = 1) + \mathbb{E}[\mathbf{1}_A|\mathbf{1}_B = 0]\mathbb{P}(\mathbf{1}_B = 0) \\
&= \mathbb{P}(A|B) \cdot \mathbb{P}(B) + \mathbb{P}(A|B^C) \cdot \mathbb{P}(B^C) \\
&= \mathbb{P}(A)
\end{aligned}$$

5.2 Conditional Expectation as a projection.

Conditioning on one variable. We start by giving the definition of conditional expectation given a single variable. This relates to the two observations (A) and (B) made previously. We assume that the random variable is integrable for the expectations to be well-defined.

Definition 5.1. Let X and Y be integrable random variables on $(\Omega, \mathcal{F}, \mathbb{P})$. The conditional expectation of Y given X is the random variable denoted by $\mathbb{E}[Y|X]$ with the following two properties:

- (A) There exists a function $h : \mathbf{R} \rightarrow \mathbf{R}$ such that $\mathbb{E}[Y|X] = h(X)$.
- (B) For any bounded random variable of the form $g(X)$ for some function g ,

$$\mathbb{E}[g(X)Y] = \mathbb{E}[g(X)\mathbb{E}[Y|X]] \quad (5.2)$$

We can interpret the second property as follows. The conditional expectation $\mathbb{E}[Y|X]$ serves as a proxy for Y as far as X is concerned. Note that in equation (5.2), the expectation on the left can be seen as an average over the joint values of (X, Y) , whereas the one on the right is an average over the values of X only! Another way to see this property is to write it as:

$$\mathbb{E}[g(X)(Y - \mathbb{E}[Y|X])] = 0 \quad (5.3)$$

In other words, the *random variable* $Y - \mathbb{E}[Y|X]$ is *orthogonal to any random variable constructed from X* .

Finally, it is important to notice that if we take $g(X) = 1$, then the second property implies :

$$\mathbb{E}[Y] = \mathbb{E}[\mathbb{E}[Y|X]]$$

In other words, the expectation of the conditional expectation of Y is simply the expectation of Y .

The existence of the conditional expectation $\mathbb{E}[Y|X]$ is not obvious. We know, it exists in particular cases given in example (5.1). We will show more generally, that it exists, it is unique whenever Y is in $L^2(\Omega, \mathcal{F}, \mathbb{P})$ (In fact, it can be shown to exist whenever Y is integrable). Before doing so, let's warm up by looking at the case of Gaussian vectors.

Example 5.3. (Conditional expectation of Gaussian vectors - I). Let (X, Y) be a Gaussian vector of mean 0. Then:

$$\mathbb{E}[Y|X] = \frac{\mathbb{E}[XY]}{\mathbb{E}[X^2]}X \quad (5.4)$$

This candidate satisfies the two defining properties of conditional expectation : (A) It is clearly a function of X ; in fact it is a simple multiple of X . (B) We have that the random variable $\left(Y - \frac{\mathbb{E}[XY]}{\mathbb{E}[X^2]}X\right)$ is orthogonal and thus independent to X . This is a consequence of the proposition (3.5), since:

$$\begin{aligned} &= \mathbb{E} \left[X \left(Y - \frac{\mathbb{E}[XY]}{\mathbb{E}[X^2]}X \right) \right] = \mathbb{E}XY - \frac{\mathbb{E}[XY]}{\mathbb{E}[X^2]} \mathbb{E}X^2 \\ &= \mathbb{E}XY - \frac{\mathbb{E}[XY]}{\mathbb{E}[X^2]} \mathbb{E}X^2 \\ &= 0 \end{aligned}$$

Therefore, we have for any bounded function $g(X)$ of X :

$$\mathbb{E}[g(X)(Y - \mathbb{E}[Y|X])] = \mathbb{E}[g(X)]\mathbb{E}[Y - \mathbb{E}[Y|X]] = 0$$

Example 5.4. (Brownian conditioning-I) Let $(B_t, t \geq 0)$ be a standard Brownian motion. Consider the Gaussian vector $(B_{1/2}, B_1)$. Its covariance matrix is:

$$C = \begin{bmatrix} 1/2 & 1/2 \\ 1/2 & 1 \end{bmatrix}$$

Let's compute $\mathbb{E}[B_1|B_{1/2}]$ and $\mathbb{E}[B_{1/2}|B_1]$. This is easy using the equation (5.4). We have:

$$\begin{aligned}\mathbb{E}[B_1|B_{1/2}] &= \frac{\mathbb{E}[B_1 B_{1/2}]}{\mathbb{E}[B_{1/2}^2]} B_{1/2} \\ &= \frac{(1/2)}{(1/2)} B_{1/2} \\ &= B_{1/2}\end{aligned}$$

In other words, the best approximation of B_1 given the information of $B_{1/2}$ is $B_{1/2}$. There is no problem in computing $\mathbb{E}[B_{1/2}|B_1]$, even though we are conditioning on a future position. Indeed the same formula gives

$$\mathbb{E}[B_{1/2}|B_1] = \frac{\mathbb{E}[B_1 B_{1/2}]}{\mathbb{E}[B_1^2]} B_1 = \frac{1}{2} B_1$$

This means that the best approximation of $B_{1/2}$ given the position at time 1, is $\frac{1}{2} B_1$ which makes a whole lot of sense!

In example (5.4) for the Gaussian vector (X, Y) , the conditional expectation was equal to the *orthogonal projection* of Y onto X in L^2 . In particular, the conditional expectation was a multiple of X . Is this always the case? Unfortunately, it is not. For example, in the equation (5.1), the conditional expectation is clearly not a multiple of the random variable X . However, it is a function of X , as is always the case by definition (5.1).

The idea to construct the conditional expectation $\mathbb{E}[Y|X]$ in general is to *project* Y on the space of all random variables that can be constructed from X . To make this precise, consider the following subspace of $L^2(\Omega, \mathcal{F}, \mathbb{P})$:

Definition 5.2. Let $(\Omega, \mathcal{F}, \mathbb{P})$ be a probability space and X a random variable defined on it. The space $L^2(\Omega, \sigma(X), \mathbb{P})$ is the linear subspace of $L^2(\Omega, \mathcal{F}, \mathbb{P})$ consisting of the square-integrable random variables of the form $g(X)$ for some function $g : \mathbf{R} \rightarrow \mathbf{R}$.

This is a linear subspace of $L^2(\Omega, \mathcal{F}, \mathbb{P})$: It contains the random variable 0, and any linear combination of random variables of this kind is also a function of X and must have a finite second moment. We note the following:

Remark. $L^2(\Omega, \sigma(X), \mathbb{P})$ is a subspace of $L^2(\Omega, \mathcal{F}, \mathbb{P})$, very much how a plane or line (going through the origin) is a subspace of \mathbf{R}^3 .

In particular, as in the case of a line or a plane, we can project an element of Y of $L^2(\Omega, \mathcal{F}, \mathbb{P})$ onto $L^2(\Omega, \sigma(X), \mathbb{P})$. The resulting projection is an element of $L^2(\Omega, \sigma(X), \mathbb{P})$, a square-integrable random-variable that is a function of X . For a subspace \mathcal{S} of \mathbf{R}^3 (e.g. a line or a plane), the projection of the vector $\mathbf{v} \in \mathbf{R}^3$ onto the subspace \mathcal{S} , denoted $\text{Proj}_{\mathcal{S}}(\mathbf{v})$ is the closest point to \mathbf{v} lying in the subspace \mathcal{S} . Moreover, $\mathbf{v} - \text{Proj}_{\mathcal{S}}(\mathbf{v})$ is orthogonal to the subspace. This picture of orthogonal projection also holds in L^2 . Let Y be a random variable in $L^2(\Omega, \mathcal{F}, \mathbb{P})$ and let $L^2(\Omega, \sigma(X), \mathbb{P})$ be the subspace of those random variables that are functions of X . We write Y^* for the random variable in $L^2(\Omega, \sigma(X), \mathbb{P})$ that is *closest* to Y . In other words, we have (using the definition of the L^2 -distance square):

$$\inf_{Z \in L^2(\Omega, \sigma(X), \mathbb{P})} \mathbb{E}[(Y - Z)^2] = \mathbb{E}[(Y - Y^*)^2] \quad (5.5)$$

It turns out that Y^* is the right candidate for the conditional expectation.




Figure. An illustration of the conditional expectation $\mathbb{E}[Y|X]$ as an orthogonal projection of Y onto the subspace $L^2(\Omega, \sigma(X), \mathbb{P})$.

Theorem 5.1. (*Existence and uniqueness of the conditional expectation*) Let X be a random variable on $(\Omega, \mathcal{F}, \mathbb{P})$. Let Y be a random variable in $L^2(\Omega, \mathcal{F}, \mathbb{P})$. Then the conditional expectation $\mathbb{E}[Y|X]$ is the random variable Y^* given in the equation (5.5). Namely, it is the random variable in $L^2(\Omega, \sigma(X), \mathbb{P})$ that is closest to Y in the L^2 -distance.

In particular we have the following:

- 1) It is the orthogonal projection of Y onto $L^2(\Omega, \sigma(X), \mathbb{P})$, that is $Y - Y^*$ is orthogonal to any random variables in the subspace $L^2(\Omega, \sigma(X), \mathbb{P})$.
- 2) It is unique.

Remark. This result reinforces the meaning of the conditional expectation $\mathbb{E}[Y|X]$ as the best estimation of Y given the information of X : it is the closest random variable to Y among all the functions of X in the sense of L^2 .

Proof. We write for short $L^2(X)$ for the subspace $L^2(\Omega, \sigma(X), \mathbb{P})$. Let Y^* be as in equation (5.5). We show successively that (1) $Y - Y^*$ is orthogonal to any element of $L^2(X)$, so it is the orthogonal projection (2) Y^* has the properties of conditional expectation in definition (5.2) (3) Y^* is unique.

(1) Let $W = g(X)$ be a random variable in $L^2(X)$. We show that W is orthogonal to $Y - Y^*$; that is $\mathbb{E}[(Y - Y^*)W] = 0$. This should be intuitively clear from figure above. On the one hand, we have by developing the square:

$$\begin{aligned}\mathbb{E}[(W - (Y - Y^*))^2] &= \mathbb{E}[W^2 - 2W(Y - Y^*) + (Y - Y^*)^2] \\ &= \mathbb{E}[W^2] - 2\mathbb{E}[W(Y - Y^*)] + \mathbb{E}(Y - Y^*)^2\end{aligned}\quad (5.6)$$

On the other hand, $Y^* + W$ is an arbitrary vector in $L^2(X)$ (it is a linear combination of the elements in $L^2(X)$), we must have from equation (5.5):

$$\begin{aligned}\mathbb{E}[(W - (Y - Y^*))^2] &= \mathbb{E}[(Y - (Y^* + W))^2] \\ &\geq \inf_{Z \in L^2(X)} \mathbb{E}[(Y - Z)^2] \\ &= \mathbb{E}[(Y - Y^*)^2]\end{aligned}\quad (5.7)$$

Putting the last two equations (5.6), (5.7) together, we get that for any $W \in L^2(X)$:

$$\mathbb{E}[W^2] - 2\mathbb{E}[W(Y - Y^*)] \geq 0$$

In particular, this also holds for aW , in which case we get:

$$\begin{aligned} & a^2\mathbb{E}[W^2] - 2a\mathbb{E}[W(Y - Y^*)] \geq 0 \\ \Rightarrow & a \{a\mathbb{E}[W^2] - 2\mathbb{E}[W(Y - Y^*)]\} \geq 0 \end{aligned}$$

If $a > 0$, then:

$$a\mathbb{E}[W^2] - 2\mathbb{E}[W(Y - Y^*)] \geq 0 \quad (5.8)$$

whereas if $a < 0$, then the sign changes upon dividing throughout by a , and we have:

$$a\mathbb{E}[W^2] - 2\mathbb{E}[W(Y - Y^*)] \leq 0 \quad (5.9)$$

Rearranging (5.8) yields:

$$\mathbb{E}[W(Y - Y^*)] \leq a\mathbb{E}[W^2]/2 \quad (5.10)$$

Rearranging (5.9) yields:

$$\mathbb{E}[W(Y - Y^*)] \geq a\mathbb{E}[W^2]/2 \quad (5.11)$$

Since (5.10) holds for $a > 0$, it follows that, $\mathbb{E}[W(Y - Y^*)] \leq 0$. Since, (5.11) holds for all $a < 0$, it follows that $\mathbb{E}[W(Y - Y^*)] \geq 0$. Consequently,

$$\mathbb{E}[W(Y - Y^*)] = 0 \quad (5.12)$$

(2) It is clear that Y^* is a function of X by construction, since it is in $L^2(X)$. Moreover, for any $W \in L^2(X)$, we have from (1) that:

$$\mathbb{E}[W(Y - Y^*)] = 0$$

which is the second defining property of conditional expectations.

(3) Lastly, suppose there is another element Y' that is in $L^2(X)$ that minimizes the distance to Y . Then we would get:

$$\begin{aligned}\mathbb{E}[(Y - Y')^2] &= \mathbb{E}[(Y - Y^* + Y^* - Y')^2] \\ &= \mathbb{E}[(Y - Y^*)^2] + 2\mathbb{E}[(Y - Y^*)(Y^* - Y')] + \mathbb{E}[(Y^* - Y')^2] \\ &= \mathbb{E}[(Y - Y^*)^2] + 0 + \mathbb{E}[(Y^* - Y')^2] \\ &\quad \{ (Y^* - Y') \in L^2(X) \perp (Y - Y^*) \}\end{aligned}$$

where we used the fact, that $Y^* - Y'$ is a vector in $L^2(X)$ and the orthogonality of $Y - Y^*$ with $L^2(X)$ as in (1). But, this implies that:

$$\begin{aligned}\cancel{\mathbb{E}[(Y - Y')^2]} &= \cancel{\mathbb{E}[(Y - Y^*)^2]} + \mathbb{E}[(Y^* - Y')^2] \\ \mathbb{E}[(Y^* - Y')^2] &= 0\end{aligned}$$

So, $Y^* = Y'$ almost surely. □

Example 5.5. [Arguin-4.1] **Conditional Expectation of continuous random variables.** Let (X, Y) be two random variables with joint density $f_{X,Y}(x, y)$ on \mathbf{R}^2 . Suppose for simplicity, that $\int_{\mathbf{R}} f(x, y)dx > 0$ for every y belonging to \mathbf{R} . Show that the conditional expectation $\mathbf{E}[Y|X]$ equals $h(X)$ where h is the function:

$$h(x) = \frac{\int_{\mathbf{R}} y f_{X,Y}(x, y) dy}{\int_{\mathbf{R}} f_{X,Y}(x, y) dy} \quad (5.13)$$

In particular, verify that $\mathbf{E}[\mathbf{E}[Y|X]] = \mathbf{E}[Y]$.

Hint: To prove this, verify that the above formula satisfies both the properties of conditional expectations; then invoke uniqueness to finish it off.

Solution. (i) The density function $f_{X,Y}(x, y)$ is a map $f : \mathbf{R}^2 \rightarrow \mathbf{R}$. The integral $\int_{y=-\infty}^{y=+\infty} y f_{X,Y}(x_0, y) dy$ is the area under the curve $y f(x, y)$ at the point $x = x_0$. Let's call it $A(x_0)$. If instead, we have an arbitrary x , $\int_{y=-\infty}^{y=+\infty} y f_{X,Y}(x, y) dy$ represents the area $A(x)$ of an arbitrary slice of the surface $y f_{X,Y}$ at the point x . Hence,

it is a function of x . The denominator $\int_{\mathbf{R}} f_{X,Y}(x, y) dy = f_X(x)$, the density of X , which is a function of x . Hence, the ratio is a function of x .

(ii) Let $g(X)$ is a bounded random variable. We have:

$$\begin{aligned}
\mathbf{E}[g(X)(Y - h(X))] &= \mathbf{E}[Yg(X)] - \mathbf{E}[g(X)h(X)] \\
&= \int \int_{\mathbf{R}^2} yg(x)f_{X,Y}(x, y)dydx - \int_{\mathbf{R}} g(x)h(x)f(x)dx \\
&= \int \int_{\mathbf{R}^2} yg(x)f_{X,Y}(x, y)dydx \\
&\quad - \int_{\mathbf{R}} g(x) \cdot \frac{\int_{\mathbf{R}} yf_{X,Y}(x, y)dy}{\int_{\mathbf{R}} f_{X,Y}(x, y)dy} \cdot \int_{\mathbf{R}} f_{X,Y}(x, y)dy dx \\
&= \int \int_{\mathbf{R}^2} yg(x)f_{X,Y}(x, y)dydx \\
&\quad - \int_{\mathbf{R}} g(x) \cdot \frac{\int_{\mathbf{R}} yf_{X,Y}(x, y)dy}{\int_{\mathbf{R}} f_{X,Y}(x, y)dy} \cdot \int_{\mathbf{R}} \cancel{f_{X,Y}(x, y)dy} dx \\
&= \int \int_{\mathbf{R}^2} yg(x)f_{X,Y}(x, y)dydx - \int_{\mathbf{R}^2} yg(x)f_{X,Y}(x, y) \cdot dx \cdot dy \\
&= 0
\end{aligned}$$

Thus, $h(X)$ is a valid candidate for the conditional expectation $\mathbf{E}[Y|X]$. Moreover, by the existence and uniqueness theorem (5.1), $\mathbf{E}[Y|X]$ is unique and equals $h(X)$.

Conditioning on several random variables. We would like to generalize the conditional expectation to the case when we condition on the information of more than one random variable. Taking the L^2 point of view, we should expect that the conditional expectation is the orthogonal projection of the given random variable on the subspace generated by square integrable functions of all the variables on which we condition.

It is now useful to study sigma-fields, an object that was defined in chapter 1.

Definition 5.3. (Sigma-Field) A sigma-field or sigma-algebra \mathcal{F} of a sample space Ω is a collection of all measurable events with the following properties:

- (1) Ω is in \mathcal{F} .
- (2) Closure under complement. If $A \in \mathcal{F}$, then $A^C \in \mathcal{F}$.
- (3) Closure under countable unions. If $A_1, A_2, \dots \in \mathcal{F}$, then $\bigcup_{n=1}^{\infty} A_n \in \mathcal{F}$.

Such objects play a fundamental role in the rigorous study of probability and real analysis in general. We will focus on the intuition behind them. First let's mention some examples of sigma-fields of a given sample space Ω to get acquainted with the concept.

Example 5.6. (Examples of sigma-fields).

- (1) *The trivial sigma-field.* Note that the collection of events $\{\emptyset, \Omega\}$ is a sigma-field of Ω . We generally denote it by \mathcal{F}_0 .
- (2) *The σ -field generated by an event A .* Let A be an event that is not \emptyset and not the entire Ω . Then the smallest sigma-field containing A ought to be:

$$\mathcal{F}_1 = \{\emptyset, A, A^C, \Omega\}$$

This sigma-field is denoted by $\sigma(A)$.

- (3) *The sigma-field generated by a random variable X .*

We now define the \mathcal{F}_X as follows:

$$\mathcal{F}_X = X^{-1}(\mathcal{B}) := \{\omega : X(\omega) \in B\}, \forall B \in \mathcal{B}(\mathbf{R})$$

where \mathcal{B} is the Borel σ -algebra on \mathbf{R} . \mathcal{F}_X is sometimes denoted as $\sigma(X)$. \mathcal{F}_X is the set of all events pertaining to X . It is a sigma-algebra because:

- (i) $\Omega \in \sigma(X)$ because $\Omega = \{\omega : X(\omega) \in \mathbf{R}\}$ and $\mathbf{R} \in \mathcal{B}(\mathbf{R})$.
- (ii) Let any event $C \in \sigma(X)$. We need to show that $\Omega \setminus C \in \sigma(X)$.

Since $C \in \sigma(X)$, there exists $A \in \mathcal{B}(\mathbf{R})$, such that:

$$C = \{\omega \in \Omega : X(\omega) \in A\}$$

Now, we calculate:

$$\Omega \setminus C = \{\omega \in \Omega : X(\omega) \in \mathbf{R} \setminus A\}$$

Since $\mathcal{B}(\mathbf{R})$ is a sigma-algebra, it is closed under complementation. Hence, if $A \in \mathcal{B}(\mathbf{R})$, it implies that $\mathbf{R} \setminus A \in \mathcal{B}(\mathbf{R})$. So, $\Omega \setminus C \in \sigma(X)$.

(iii) Consider a sequence of events $C_1, C_2, \dots, C_n, \dots \in \sigma(X)$. We need to prove that $\bigcup_{n=1}^{\infty} C_n \in \sigma(X)$.

Since $C_n \in \sigma(X)$, there exists $A_n \in \mathcal{B}(\mathbf{R})$ such that:

$$C_n = \{\omega \in \Omega : X(\omega) \in A_n\}$$

Now, we calculate:

$$\bigcup_{n=1}^{\infty} C_n = \{\omega \in \Omega : X(\omega) \in \bigcup_{n=1}^{\infty} A_n\}$$

But, $\bigcup_{n=1}^{\infty} A_n \in \mathcal{B}(\mathbf{R})$. So, $\bigcup_{n=1}^{\infty} C_n \in \sigma(X)$.

Consequently, $\sigma(X)$ is indeed a σ -algebra.

Intuitively, we think of $\sigma(X)$ as containing all information about X .

(4) *The sigma-field generated by a stochastic process* $(X_s, s \leq t)$. Let $(X_s, s \geq 0)$ be a stochastic process. Consider the process restricted to $[0, t]$, $(X_s, s \leq t)$. We consider the smallest sigma-field containing all events pertaining to the random variables $X_s, s \leq t$. We denote it by $\sigma(X_s, s \leq t)$ or \mathcal{F}_t .

The sigma-fields on Ω have a natural (partial) ordering: two sigma-fields \mathcal{G} and \mathcal{F} of Ω are such that $\mathcal{G} \subseteq \mathcal{F}$ if all the events in \mathcal{G} are in \mathcal{F} . For example, the trivial σ -field $\mathcal{F}_0 = \{\emptyset, \Omega\}$ is contained in all the σ -fields of Ω . Clearly, the σ -field $\mathcal{F}_t = \sigma(X_s, s \leq t)$ is contained in $\mathcal{F}_{t'}$ if $t \leq t'$.

If all the events pertaining to a random variable X are in the σ -field \mathcal{G} (and thus we can compute $\mu(X^{-1}((a, b]))$), we will say that X is \mathcal{G} -measurable. This means that all information about X is contained in \mathcal{G} .

Definition 5.4. Let X be a random variable defined on $(\Omega, \mathcal{F}, \mathbb{P})$. Consider another $\mathcal{G} \subseteq \mathcal{F}$. Then X is said to be \mathcal{G} -measurable, if and only if:

$$\{\omega : X(\omega) \in (a, b]\} \in \mathcal{G} \text{ for all intervals } (a, b] \in \mathbf{R}$$

Example 5.7. (\mathcal{F}_0 -measurable random variables). Consider the trivial sigma-field $\mathcal{F}_0 = \{\emptyset, \Omega\}$. A random variable that is \mathcal{F}_0 -measurable must be a constant. Indeed, we have that for any interval $(a, b]$, $\{\omega : X(\omega) \in (a, b]\} = \emptyset$ or $\{\omega : X(\omega) \in (a, b]\} = \Omega$. This can only hold if X takes a single value.

Example 5.8. ($\sigma(X)$ -measurable random variables). Let X be a given random variable on $(\Omega, \mathcal{F}, \mathbb{P})$. Roughly speaking, a $\sigma(X)$ -measurable random variable is determined by the information of X only. Here is the simplest example of a $\sigma(X)$ -measurable random variable. Take the indicator function $Y = \mathbf{1}_{\{X \in B\}}$ for some event $\{X \in B\}$ pertaining to X . Then the pre-images $\{\omega : Y(\omega) \in (a, b]\}$ are either \emptyset , $\{X \in B\}$, $\{X \in B^C\}$ or Ω depending on whether 0, 1 are in $(a, b]$ or not. All of these events are in $\sigma(X)$. More generally, one can construct a $\sigma(X)$ -measurable random variable by taking linear combinations of indicator functions of events of the form $\{X \in B\}$.

It turns out that any (Borel measurable) function of X can be approximated by taking limits of such simple functions.

Concretely, this translates to the following statement:

$$\text{If } Y \text{ is } \sigma(X)\text{-measurable, then } Y = g(X) \text{ for some function } g \quad (5.14)$$

In the same way, if Z is $\sigma(X, Y)$ -measurable, then $Z = h(X, Y)$ for some h . These facts can be proved rigorously using measure theory.

We are ready to give the general definition of conditional expectation.

Example 5.9. (Coin-Tossing Space). Suppose a coin is tossed infinitely many times. Let Ω be the set of all infinite sequences of H s and T s. A generic element of Ω is denoted by $\omega_1\omega_2\ldots$, where ω_n indicates the result of the n th coin toss. Ω is an uncountable sample space. The trivial sigma-field $\mathcal{F}_0 = \{\emptyset, \Omega\}$. Assume that we don't know anything about the outcome of the experiment. Even without any information, we know that the true ω belongs to Ω and does not belong to \emptyset . It is the information learned at time 0.

Next, assume that we know the outcome of the first coin toss. Define $A_H = \{\omega : \omega_1 = H\}$ = set of all sequences beginning with H and $A_T = \{\omega : \omega_1 = T\}$ = set of all sequences beginning with T . The four sets resolved by the first coin-toss form the σ -field $\mathcal{F}_1 = \{\emptyset, A_H, A_T, \Omega\}$. We shall think of this σ -field as containing the information learned by knowing the outcome of the first coin toss. More precisely, if

instead of being told about the first coin toss, we are told for each set in \mathcal{F}_1 , whether or not the true ω belongs to that set, then we know the outcome of the first coin toss and nothing more.

If we are told the first two coin tosses, we obtain a finer resolution. In particular, the four sets:

$$\begin{aligned} A_{HH} &= \{\omega : \omega_1 = H, \omega_2 = H\} \\ A_{HT} &= \{\omega : \omega_1 = H, \omega_2 = T\} \\ A_{TH} &= \{\omega : \omega_1 = T, \omega_2 = H\} \\ A_{TT} &= \{\omega : \omega_1 = T, \omega_2 = T\} \end{aligned}$$

are resolved. Of course, the sets in \mathcal{F}_1 are resolved. Whenever a set is resolved, so is its complement, which means that $A_{HH}^C, A_{HT}^C, A_{TH}^C$ and A_{TT}^C are resolved, so is their union which means that $A_{HH} \cup A_{TH}, A_{HH} \cup A_{TT}, A_{HT} \cup A_{TH}$ and $A_{HT} \cup A_{TT}$ are resolved. The other two pair-wise unions $A_{HH} \cup A_{HT} = A_H$ and $A_{TH} \cup A_{TT} = A_T$ are already resolved. Finally, the triple unions are also resolved, because $A_{HH} \cup A_{HT} \cup A_{TH} = A_{TT}^C$ and so forth. Hence, the information pertaining to the second coin-toss is contained in:

$$\begin{aligned} \mathcal{F}_2 &= \{\emptyset, \Omega, \\ &A_H, A_T, \\ &A_{HH}, A_{HT}, A_{TH}, A_{TT}, \\ &A_{HH}^C, A_{HT}^C, A_{TH}^C, A_{TT}^C, \\ &A_{HH} \cup A_{TH}, A_{HH} \cup A_{TT}, A_{HT} \cup A_{TH}, A_{HT} \cup A_{TT}\} \end{aligned}$$

Hence, if the outcome of the first two coin tosses is known, all of the events in \mathcal{F}_2 are resolved - we exactly know, if each event has occurred or not. \mathcal{F}_2 is the information learned by observing the first two coin tosses.

Example 5.10. [Arguin-4.2] (**Exercises on sigma-fields**).

(a) Let A, B be two proper subsets of Ω such that $A \cap B \neq \emptyset$ and $A \cup B \neq \Omega$. Write down $\sigma(\{A, B\})$, the smallest sigma-field containing A and B explicitly. What if $A \cap B = \emptyset$?

(b) The Borel sigma-field is the smallest sigma-field containing intervals of the form $(a, b]$ in \mathbf{R} . Show that all singletons $\{b\}$ are in $\mathcal{B}(\mathbf{R})$ by writing $\{b\}$ as a countable intersection of intervals $(a, b]$. Conclude that all open intervals (a, b) and all closed intervals $[a, b]$ are in $\mathcal{B}(\mathbf{R})$. Is the subset \mathbf{Q} of rational numbers a Borel set?

Proof. (a) The sigma-field generated by the two events A, B is given by:

$$\begin{aligned}\sigma(\{A, B\}) = \{ & \emptyset, \Omega, \\ & A, B, A^C, B^C, \\ & A \cup B, A \cap B, \\ & A \cup B^C, A^C \cup B, A^C \cup B^C, \\ & A \cap B^C, A^C \cap B, A^C \cap B^C, \\ & (A \cup B) \cap (A \cap B)^C, \\ & (A \cup B)^C \cup (A \cap B)\}\end{aligned}$$

(b) Firstly, recall that:

$$\mathcal{B}(\mathbf{R}) = \bigcap_{\alpha \in \Lambda} \mathcal{F}_\alpha = \bigcap \sigma(\{I : I \text{ is an interval } (a, b] \subseteq \mathbf{R}\})$$

We can write:

$$\{b\} = \bigcap_{n=1}^{\infty} \left(b - \frac{1}{n}, b\right]$$

As $\mathcal{B}(\mathbf{R})$ is a sigma-field, it is closed under countable intersections. Hence, the singleton set $\{b\}$ is a Borel set.

Similarly, we can write, any open interval as the countable union:

$$(a, b) = \bigcup_{n=1}^{\infty} \left(a, b - \frac{1}{n}\right]$$

Hence, open intervals are Borel sets.

Similarly, we may write:

$$[a, b] = \bigcap_{n=1}^{\infty} \left(a - \frac{1}{n}, b + \frac{1}{n} \right)$$

Consequently, closed intervals are Borel sets. Since \mathbf{Q} is countable, it is a Borel set. Moreover, the empty set \emptyset and \mathbf{R} are Borel sets. So, $\mathbf{R} \setminus \mathbf{Q}$ is also a Borel set. \square

Example 5.11. [Arguin-4.4] Let (X, Y) be a Gaussian vector with mean 0 and covariance matrix

$$C = \begin{bmatrix} 1 & \rho \\ \rho & 1 \end{bmatrix}$$

for $\rho \in (-1, 1)$. We verify that the example (5.3) and exercise (5.5) yield the same conditional expectation.

- (a) Use equation (5.4) to show that $\mathbf{E}[Y|X] = \rho X$.
- (b) Write down the joint PDF $f(x, y)$ of (X, Y) .
- (c) Show that $\int_{\mathbf{R}} y f(x, y) dy = \rho x$ and that $\int_{\mathbf{R}} f(x, y) dy = 1$.
- (d) Deduce that $\mathbf{E}[Y|X] = \rho X$ using the equation (5.13).

Proof. (a) Since (X, Y) have mean 0 and variance 1, it follows that:

$$\begin{aligned} \mathbf{E}[(X - EX)(Y - EY)] &= \mathbf{E}(XY) \\ \sqrt{(\mathbf{E}[X^2] - (\mathbf{E}X)^2)} \cdot \sqrt{(\mathbf{E}[Y^2] - (\mathbf{E}Y)^2)} &= \sqrt{(1 - 0)(1 - 0)} \\ &= 1 \end{aligned}$$

and therefore,

$$\rho = \frac{\mathbf{E}(XY)}{1} = \frac{\mathbf{E}[XY]}{\mathbf{E}[X^2]}$$

Since (X, Y) is a Gaussian vector, using (5.4), we have:

$$\mathbf{E}[Y|X] = \frac{\mathbf{E}[XY]}{\mathbf{E}[X^2]}X = \rho X$$

(b) Consider the augmented matrix $[C|I]$. We have:

$$[C|I] = \left[\begin{array}{cc|cc} 1 & \rho & 1 & 0 \\ \rho & 1 & 0 & 1 \end{array} \right]$$

Performing $R_2 = R_2 - \rho R_1$, the above system is row-equivalent to:

$$\left[\begin{array}{cc|cc} 1 & \rho & 1 & 0 \\ 0 & 1 - \rho^2 & -\rho & 1 \end{array} \right]$$

Performing $R_2 = \frac{1}{1-\rho^2}R_2$, the above system is row-equivalent to:

$$\left[\begin{array}{cc|cc} 1 & \rho & 1 & 0 \\ 0 & 1 & \frac{-\rho}{1-\rho^2} & \frac{1}{1-\rho^2} \end{array} \right]$$

Performing $R_1 = R_1 - \rho R_2$, we have:

$$\left[\begin{array}{cc|cc} 1 & 0 & \frac{1}{1-\rho^2} & -\frac{\rho}{1-\rho^2} \\ 0 & 1 & \frac{-\rho}{1-\rho^2} & \frac{1}{1-\rho^2} \end{array} \right]$$

Thus,

$$C^{-1} = \frac{1}{1-\rho^2} \begin{bmatrix} 1 & -\rho \\ -\rho & 1 \end{bmatrix}$$

Moreover, $\det C = 1 - \rho^2$.

Therefore, the joint density of (X, Y) is given by:

$$\begin{aligned}
f(x, y) &= \frac{1}{2\pi\sqrt{1-\rho^2}} \exp \left[-\frac{1}{2(1-\rho^2)} \begin{bmatrix} x & y \end{bmatrix} \begin{bmatrix} 1 & -\rho \\ -\rho & 1 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} \right] \\
&= \frac{1}{2\pi\sqrt{1-\rho^2}} \exp \left[-\frac{1}{2(1-\rho^2)} \begin{bmatrix} x - \rho y & -\rho x + y \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} \right] \\
&= \frac{1}{2\pi\sqrt{1-\rho^2}} \exp \left[-\frac{1}{2(1-\rho^2)} (x^2 - 2\rho xy + y^2) \right]
\end{aligned}$$

(c) Claim I. $\int_{\mathbf{R}} y f(x, y) dy = \rho x$.

Completing the square, we have:

$$(x^2 - 2\rho xy + y^2) = (y - \rho x)^2 + x^2(1 - \rho^2)$$

Thus, we can write:

$$\int_{\mathbf{R}} y f(x, y) dy = \frac{1}{2\pi\sqrt{1-\rho^2}} e^{-\frac{1}{2}x^2} \int_{\mathbf{R}} y e^{-\frac{1}{2} \frac{(y-\rho x)^2}{(1-\rho^2)}} dy$$

Let's substitute

$$\begin{aligned}
z &= \frac{(y - \rho x)}{\sqrt{1 - \rho^2}} \\
dz &= \frac{dy}{\sqrt{1 - \rho^2}}
\end{aligned}$$

Therefore,

$$\begin{aligned}
\int_{\mathbf{R}} y e^{-\frac{1}{2} \frac{(y-\rho x)^2}{(1-\rho^2)}} dy &= \sqrt{1 - \rho^2} \int_{\mathbf{R}} (\rho x + \sqrt{1 - \rho^2} z) e^{-\frac{z^2}{2}} dz \\
&= \rho x \cdot \sqrt{1 - \rho^2} \int_{\mathbf{R}} e^{-\frac{z^2}{2}} dz + (1 - \rho^2) \int_{\mathbf{R}} z e^{-\frac{z^2}{2}} dz \\
&= \rho x \cdot \sqrt{1 - \rho^2} \cdot \sqrt{2\pi} + (1 - \rho^2) \cdot 0 \\
&= \rho x \cdot \sqrt{1 - \rho^2} \cdot \sqrt{2\pi}
\end{aligned}$$

Consequently,

$$\begin{aligned}
\int_{\mathbf{R}} y f(x, y) dy &= \frac{1}{2\pi\sqrt{1-\rho^2}} e^{-\frac{1}{2}x^2} \rho x \cdot \sqrt{1-\rho^2} \cdot \sqrt{2\pi} \\
&= \rho x \cdot \frac{1}{\sqrt{2\pi}} e^{-\frac{1}{2}x^2} \\
&= \rho x \cdot f_X(x) \\
\frac{\int_{\mathbf{R}} y f(x, y) dy}{f_X(x)} &= \frac{\int_{\mathbf{R}} y f(x, y) dy}{\int_{\mathbf{R}} f(x, y)} = \rho x
\end{aligned}$$

(d) For a Gaussian vector (X, Y) , the conditional expectation $\mathbf{E}[Y|X] = h(X)$. Hence, $\mathbf{E}[Y|X] = \rho X$. \square

Definition 5.5. (Conditional Expectation) Let Y be an integrable random variable on $(\Omega, \mathcal{F}, \mathbb{P})$ and let $\mathcal{G} \subseteq \mathcal{F}$ be a sigma-field of Ω . The conditional expectation of Y given \mathcal{G} is the random variable denoted by $\mathbf{E}[Y|\mathcal{G}]$ such that the following hold:

(a) $\mathbf{E}[Y|\mathcal{G}]$ is \mathcal{G} -measurable.

In other words, all events pertaining to the random variable $\mathbf{E}[Y|\mathcal{G}]$ are in \mathcal{G} .

(b) For any (bounded) random variable W , that is \mathcal{G} -measurable,

$$\mathbf{E}[WY] = \mathbf{E}[W\mathbf{E}[Y|\mathcal{G}]]$$

In other words, $\mathbf{E}[Y|\mathcal{G}]$ is a proxy for Y as far as the events in \mathcal{G} are concerned.

Note that, by taking $W = 1$ in the property (B), we recover:

$$\mathbf{E}[\mathbf{E}[Y|\mathcal{G}]] = \mathbf{E}[Y]$$

Remark. Beware of the notation! If $\mathcal{G} = \sigma(X)$, then the conditional expectation $\mathbf{E}[Y|\sigma(X)]$ is usually denoted by $\mathbf{E}[Y|X]$ for short. However, one should always keep in mind that conditioning on X is in fact projecting on the linear subspace generated by all variables constructed from X and not on the linear space generated by X alone. In the same way, the conditional expectation $\mathbf{E}[Z|\sigma(X, Y)]$ is often written $\mathbf{E}[Z|X, Y]$ for short.

As expected, if Y is in $L^2(\Omega, \mathcal{F}, \mathbb{P})$, then $\mathbf{E}[Y|\mathcal{G}]$ is given by the orthogonal projection of Y onto the subspace $L^2(\Omega, \mathcal{G}, \mathbb{P})$, the subspace of square integrable random variables that are \mathcal{G} -measurable. We write Y^* for the random variable in $L^2(\Omega, \mathcal{G}, \mathbb{P})$ that is closest to Y that is:

$$\min_{Z \in L^2(\Omega, \mathcal{G}, \mathbb{P})} \mathbf{E}[(Y - Z)^2] = \mathbf{E}[(Y - Y^*)^2] \quad (5.15)$$

Theorem 5.2. (*Existence and Uniqueness of Conditional Expectations*) Let $\mathcal{G} \subset \mathcal{F}$ be a sigma-field of Ω . Let Y be a random variable in $L^2(\Omega, \mathcal{F}, \mathbb{P})$. Then, the conditional expectation $\mathbf{E}[Y|\mathcal{G}]$ is the random variable Y^* given in the equation (5.15). Namely, it is the random variable in $L^2(\Omega, \mathcal{G}, \mathbb{P})$ that is closest to Y in the L^2 -distance. In particular we have the following:

- It is the orthogonal projection of Y onto $L^2(\Omega, \mathcal{G}, \mathbb{P})$, that is, $Y - Y^*$ is orthogonal to the random variables in $L^2(\Omega, \mathcal{G}, \mathbb{P})$.
- It is unique.

Again, the result should be interpreted as follows: The conditional expectation $\mathbf{E}[Y|\mathcal{G}]$ is the best approximation of Y given the information included in \mathcal{G} .

Remark. The conditional expectation in fact exists and is unique for any integrable random variable Y (i.e. $Y \in L^1(\Omega, \mathcal{F}, \mathbb{P})$) as the definition suggests. However, there is no orthogonal projection in L^1 , so the intuitive geometric picture is lost.




Figure. An illustration of the conditional expectation $\mathbb{E}[Y|\mathcal{G}]$ as an orthogonal projection of Y onto the subspace $L^2(\Omega, \mathcal{G}, \mathbb{P})$.

Example 5.12. (Conditional Expectation for Gaussian Vectors. II.) Consider the Gaussian vector (X_1, \dots, X_n) . Without loss of generality, suppose it has mean 0 and is non-degenerate. What is the best approximation of X_n given the information X_1, \dots, X_{n-1} ? In other words, what is:

$$\mathbb{E}[X_n | \sigma(X_1, \dots, X_{n-1})]$$

With example (5.8) in mind, let's write $\mathbb{E}[X_n | X_1 \dots X_{n-1}]$ for short. From example (5.5), we know that if (X, Y) is a Gaussian vector with mean 0, then $\mathbb{E}[Y|X]$ is a multiple of X . Thus, we expect, that $\mathbb{E}[X_n | X_1 X_2 \dots X_{n-1}]$ is a linear combination of X_1, X_2, \dots, X_{n-1} . That is, there exists a_1, \dots, a_{n-1} such that:

$$\mathbb{E}[X_n | X_1 X_2 \dots X_{n-1}] = a_1 X_1 + a_2 X_2 + \dots + a_{n-1} X_{n-1}$$

In particular, since the conditional expectation is a linear combination of the X 's, it is itself a Gaussian random variable. The best way to find the coefficient a 's is to go back to IID decomposition of Gaussian vectors.

Let $(Z_1, Z_2, \dots, Z_{n-1})$ be IID standard Gaussians constructed from the linear combination of $(X_1, X_2, \dots, X_{n-1})$. Then, we have:

$$\mathbf{E}[X_n | X_1 X_2 \dots X_{n-1}] = b_1 Z_1 + \dots + b_{n-1} Z_{n-1}$$

Now, recall, that we construct the random variables Z_1, Z_2, \dots, Z_n using Gram-Schmidt orthogonalization:

$$\begin{aligned} \tilde{Z}_1 &= X_1, & Z_1 &= \frac{\tilde{Z}_1}{\mathbf{E}(\tilde{Z}_1^2)} \\ \tilde{Z}_2 &= X_2 - \mathbf{E}(X_2 Z_1) Z_1 & Z_2 &= \frac{\tilde{Z}_2}{\mathbf{E}(\tilde{Z}_2^2)} \\ \tilde{Z}_3 &= X_3 - \sum_{i=1}^2 \mathbf{E}(X_3 Z_i) Z_i & Z_3 &= \frac{\tilde{Z}_3}{\mathbf{E}(\tilde{Z}_3^2)} \\ &\vdots & & \end{aligned}$$

The simple case for $n = 2$ random variables.

We have already seen before:

$$\begin{aligned} \mathbf{E}[X_1(X_2 - \mathbf{E}(X_2 Z_1) Z_1)] &= \frac{1}{\mathbf{E}(\tilde{Z}_1^2)} \mathbf{E}[Z_1(X_2 - \mathbf{E}(X_2 Z_1) Z_1)] \\ &= \frac{1}{\mathbf{E}(\tilde{Z}_1^2)} \{ \mathbf{E}[X_2 Z_1] - \mathbf{E}[X_2 Z_1] \mathbf{E}(Z_1^2) \} \\ &= 0 \end{aligned}$$

So, $X_2 - \mathbf{E}(X_2 Z_1) Z_1$ is orthogonal to X_1 .

Moreover, $\mathbf{E}(X_2 Z_1) Z_1$ is a function of X_1 . Thus, both the properties of conditional expectation are satisfied. Since conditional expectations are unique, we must have, $\mathbf{E}[X_2 | X_1] = \mathbf{E}(X_2 Z_1) Z_1$.

The case for $n = 3$ random variables.

We have seen that:

$$\begin{aligned}
\mathbf{E}[X_1(X_3 - \mathbf{E}(X_3 Z_1)Z_1 - \mathbf{E}(X_3 Z_2)Z_2)] &= \frac{1}{\mathbf{E}(\tilde{Z}_1^2)} \mathbf{E}[Z_1(X_3 - \mathbf{E}(X_3 Z_1)Z_1 - \mathbf{E}(X_3 Z_2)Z_2)] \\
&= \frac{1}{\mathbf{E}(\tilde{Z}_1^2)} \left\{ \mathbf{E}(X_3 Z_1) - \mathbf{E}(X_3 Z_1) \underbrace{\mathbf{E}(Z_1^2)}_{\text{equals 1}} \right\} \\
&\quad - \frac{1}{\mathbf{E}(\tilde{Z}_1^2)} \mathbf{E}(X_3 Z_2) \underbrace{\mathbf{E}(Z_2 Z_1)}_{\text{equals 0}} \\
&= 0
\end{aligned}$$

It is an easy exercise to show that it is orthogonal to X_2 .

Hence, $X_3 - \mathbf{E}(X_3 Z_1)Z_1 - \mathbf{E}(X_3 Z_2)Z_2$ is orthogonal to X_1 and X_2 . Moreover, $\mathbf{E}(X_3 Z_1)Z_1 + \mathbf{E}(X_3 Z_2)Z_2$ is a function of X_1, X_2 . Thus, we must have:

$$\mathbf{E}[X_3 | X_1 X_2] = \mathbf{E}(X_3 Z_1)Z_1 + \mathbf{E}(X_3 Z_2)Z_2$$

In general, $X_n - \sum_{i=1}^{n-1} \mathbf{E}(X_n Z_i)Z_i$ is orthogonal to X_1, X_2, \dots, X_{n-1} . Hence,

$$\mathbf{E}[X_n | X_1 X_2 \dots X_{n-1}] = \sum_{i=1}^{n-1} \mathbf{E}(X_n Z_i)Z_i$$

5.2.1 Properties of Conditional Expectation.

We now list the properties of conditional expectation that follow from the two defining properties (A), (B) in the definition. They are extremely useful, when doing explicit computations on martingales. A good way to remember them is to understand how they related to the interpretation of conditional expectation as an orthogonal projection onto a subspace or, equivalently, as the best approximation of the variable given the information available.

Proposition 5.1. *Let Y be an integrable random variable on $(\Omega, \mathcal{F}, \mathbb{P})$. Let $\mathcal{G} \subseteq \mathcal{F}$ be another sigma-field of Ω . Then, the conditional expectation $\mathbf{E}[Y|\mathcal{G}]$ has the following properties:*

(1) *If Y is \mathcal{G} -measurable, then :*

$$\mathbf{E}[Y|\mathcal{G}] = Y$$

(2) *Taking out what is known. More generally, if Y is \mathcal{G} -measurable and X is another integrable random variable (with XY also integrable), then :*

$$\mathbf{E}[XY|\mathcal{G}] = Y\mathbf{E}[X|\mathcal{G}]$$

This makes sense, since Y is determined by \mathcal{G} , so we can take out what is known; it can be treated as a constant for the conditional expectation.

(3) *Independence. If Y is independent of \mathcal{G} , that is, for any events $\{Y \in (a, b]\}$ and $A \in \mathcal{G}$:*

$$\mathbb{P}(\{Y \in I\} \cap A) = \mathbb{P}(\{Y \in I\}) \cdot \mathbb{P}(A)$$

then

$$\mathbf{E}[Y|\mathcal{G}] = \mathbf{E}[Y]$$

In other words, if you have no information on Y , your best guess for its value is simply plain expectation.

(4) *Linearity of conditional expectations. Let X be another integrable random variable on $(\Omega, \mathcal{F}, \mathbb{P})$. Then,*

$$\mathbf{E}[aX + bY|\mathcal{G}] = a\mathbf{E}[X|\mathcal{G}] + b\mathbf{E}[Y|\mathcal{G}], \quad \text{for any } a, b \in \mathbf{R}$$

The linearity justifies the cumbersome choice of notation $\mathbf{E}[Y|\mathcal{G}]$ for the random variable.

(5) *Tower Property : If $\mathcal{H} \subseteq \mathcal{G}$ is another sigma-field of Ω , then:*

$$\mathbf{E}[Y|\mathcal{H}] = \mathbf{E}[\mathbf{E}[Y|\mathcal{G}]|\mathcal{H}]$$

Think in terms of two successive projections: first on a plane, then on a line in the plane.

(6) *Pythagoras Theorem.* We have:

$$\mathbf{E}[Y^2] = \mathbf{E}[(\mathbf{E}[Y|\mathcal{G}])^2] + \mathbf{E}[(Y - \mathbf{E}[Y|\mathcal{G}])^2]$$

In particular:

$$\mathbf{E}[(\mathbf{E}[Y|\mathcal{G}])^2] \leq \mathbf{E}[Y^2]$$

In words, the L^2 norm of $\mathbf{E}[X|\mathcal{G}]$ is smaller than the one of X , which is clear if you think in terms of orthogonal projection.

(7) *Expectation of the conditional expectation.*

$$\mathbf{E}[\mathbf{E}[Y|\mathcal{G}]] = \mathbf{E}[Y]$$

Proof.

The uniqueness property of conditional expectations in theorem (5.2) might appear to be an academic curiosity. On the contrary, it is very practical, since it ensures, that if we find a candidate for the conditional expectation that has the two properties in Definition (5.1), then it must be *the* conditional expectation. To see this, let's prove property (1).

Claim 5.1. If Y is \mathcal{G} -measurable, then $\mathbf{E}[Y|\mathcal{G}] = Y$.

It suffices to show that Y has the two defining properties of conditional expectation.

- (1) We are given that, Y is \mathcal{G} -measurable. So, property (A) is satisfied.
- (2) For any bounded random variable W that is \mathcal{G} -measurable, we have:

$$\mathbf{E}[W(Y - Y)] = \mathbf{E}[0] = 0$$

So, property (B) is also a triviality.

Claim 5.2. (Taking out what is known.) If Y is \mathcal{G} -measurable and X is another integrable random variable, then:

$$\mathbf{E}[XY|\mathcal{G}] = Y\mathbf{E}[X|\mathcal{G}]$$

In a similar vein, it suffices to show that, $Y\mathbf{E}[X|\mathcal{G}]$ has the two defining properties of conditional expectation.

(1) We are given that Y is \mathcal{G} -measurable; from property (1), $\mathbf{E}[X|\mathcal{G}]$ is \mathcal{G} -measurable. It follows that, $Y\mathbf{E}[X|\mathcal{G}]$ is \mathcal{G} -measurable.

(2) From theorem (5.2), $X - \mathbf{E}[X|\mathcal{G}]$ is orthogonal to the random variables $L^2(\Omega, \mathcal{G}, \mathbb{P})$. So, if W is any bounded \mathcal{G} -measurable random variable, it follows that:

$$\begin{aligned} \mathbf{E}[WY(X - \mathbf{E}[X|\mathcal{G}])] &= 0 \\ \implies \mathbf{E}[W \cdot XY] &= \mathbf{E}[WY\mathbf{E}[X|\mathcal{G}]] \end{aligned}$$

This closes the proof.

Claim 5.3. (Independence.) If Y is independent of \mathcal{G} , that is, for all events $\{Y \in (a, b]\}$ and $A \in \mathcal{G}$,

$$\mathbb{P}\{Y \in (a, b] \cap A\} = \mathbb{P}\{Y \in (a, b]\} \cdot \mathbb{P}(A)$$

then

$$\mathbf{E}[Y|\mathcal{G}] = \mathbf{E}[Y]$$

Let us show that $\mathbf{E}[Y]$ has the two defining properties of conditional expectations.

Claim 5.4. (Linearity of conditional expectations) Let X be another integrable random variable on $(\Omega, \mathcal{F}, \mathbb{P})$. Then,

$$\mathbf{E}[aX + bY|\mathcal{G}] = a\mathbf{E}[X|\mathcal{G}] + b\mathbf{E}[Y|\mathcal{G}], \quad \text{for any } a, b \in \mathbf{R}$$

Since $\mathbf{E}[X|\mathcal{G}]$ and $\mathbf{E}[Y|\mathcal{G}]$ are \mathcal{G} -measurable, any linear combination of these two random variables is also \mathcal{G} -measurable.

Also, if W is any bounded \mathcal{G} -measurable random variable, we have:

$$\begin{aligned}\mathbf{E}[W(aX + bY - (a\mathbf{E}[X|\mathcal{G}] + b\mathbf{E}[Y|\mathcal{G}]))] &= a\mathbf{E}[W(X - \mathbf{E}[X|\mathcal{G}])] \\ &\quad + b\mathbf{E}[W(Y - \mathbf{E}[Y|\mathcal{G}])]\end{aligned}$$

By definition, $X - \mathbf{E}(X|\mathcal{G})$ is orthogonal to the subspace $L^2(\Omega, \mathcal{G}, \mathbb{P})$ and hence to all \mathcal{G} -measurable random-variables. Hence, the two expectations on the right hand side of the above expression are 0. Since, conditional expectations are unique, we have the desired result.

Claim 5.5. If $\mathcal{H} \subseteq \mathcal{G}$ is another sigma-field of Ω , then

$$\mathbf{E}[Y|\mathcal{H}] = \mathbf{E}[\mathbf{E}[Y|\mathcal{G}]|\mathcal{H}]$$

Define $U := \mathbf{E}[Y|\mathcal{G}]$. By definition, $\mathbf{E}[U|\mathcal{H}]$ is \mathcal{H} -measurable.

Let W be any bounded \mathcal{H} -measurable random variable. We have:

$$\mathbf{E}[W\{\mathbf{E}(Y|\mathcal{G}) - \mathbf{E}(\mathbf{E}(Y|\mathcal{G})|\mathcal{H})\}] = \mathbf{E}[W(U - \mathbf{E}(U|\mathcal{H}))]$$

But, by definition $U - \mathbf{E}(U|\mathcal{H})$ is always orthogonal to the subspace $L^2(\Omega, \mathcal{H}, \mathbb{P})$ and hence, $\mathbf{E}[W(U - \mathbf{E}(U|\mathcal{H}))] = 0$. Since, conditional expectations are unique, we have the desired result.

Claim 5.6. Pythagoras's theorem. We have:

$$\mathbf{E}[Y^2] = \mathbf{E}[(\mathbf{E}[Y|\mathcal{G}])^2] + \mathbf{E}[(Y - \mathbf{E}(Y|\mathcal{G}))^2]$$

In particular,

$$\mathbf{E}[(\mathbf{E}[Y|\mathcal{G}])^2] \leq \mathbf{E}[Y^2]$$

Consider the orthogonal decomposition:

$$Y = \mathbf{E}[Y|\mathcal{G}] + (Y - \mathbf{E}[Y|\mathcal{G}])$$

Squaring on both sides and taking expectations, we have:

$$\mathbf{E}[Y^2] = \mathbf{E}[(\mathbf{E}[Y|\mathcal{G}])^2] + \mathbf{E}[(Y - \mathbf{E}[Y|\mathcal{G}])^2] + 2\mathbf{E}[\mathbf{E}[Y|\mathcal{G}](Y - \mathbf{E}[Y|\mathcal{G}])]$$

By definition of conditional expectation, $(Y - \mathbf{E}[Y|\mathcal{G}])$ is orthogonal to the subspace $L^2(\Omega, \mathcal{G}, \mathbb{P})$. By the properties of conditional expectation, $\mathbf{E}[Y|\mathcal{G}]$ is \mathcal{G} -measurable, so it belongs to $L^2(\Omega, \mathcal{G}, \mathbb{P})$. Hence, the dot-product on the right-hand side is 0. Consequently, we have the desired result.

Moreover, since $(Y - \mathbf{E}[Y|\mathcal{G}])^2$ is a non-negative random variable, $\mathbf{E}[(Y - \mathbf{E}[Y|\mathcal{G}])^2] \geq 0$. It follows that: $\mathbf{E}[Y^2] \geq \mathbf{E}[(\mathbf{E}[Y|\mathcal{G}])^2]$.

Claim 5.7. Our claim is:

$$\mathbf{E}[\mathbf{E}[Y|\mathcal{G}]] = \mathbf{E}[Y]$$

We know that, if W is any bounded \mathcal{G} -measurable random variable:

$$\mathbf{E}[WY] = \mathbf{E}[W\mathbf{E}[Y|\mathcal{G}]]$$

Taking $W = 1$, we have:

$$\mathbf{E}[Y] = \mathbf{E}[\mathbf{E}[Y|\mathcal{G}]]$$

Example 5.13. (Brownian Conditioning II). We continue the example (5.4). Let's now compute the conditional expectations $\mathbf{E}[e^{aB_1}|B_{1/2}]$ and $\mathbf{E}[e^{aB_{1/2}}|B_1]$ for some parameter a . We shall need the properties of conditional expectation in proposition (5.1). For the first one we use the fact that $B_{1/2}$ is independent of $B_1 - B_{1/2}$ to get:

$$\begin{aligned}
\mathbf{E}[e^{aB_1} | B_{1/2}] &= \mathbf{E}[e^{a((B_1 - B_{1/2}) + B_{1/2})} | B_{1/2}] \\
&= \mathbf{E}[e^{a(B_1 - B_{1/2})} \cdot e^{aB_{1/2}} | B_{1/2}] \\
&\quad \{\text{Taking out what is known}\} \\
&= e^{aB_{1/2}} \mathbf{E}[e^{a(B_1 - B_{1/2})} | B_{1/2}] \\
&= e^{aB_{1/2}} \cdot \mathbf{E}[e^{a(B_1 - B_{1/2})}] \\
&\quad \{\text{Independence}\}
\end{aligned}$$

We know that, $a(B_1 - B_{1/2})$ is a gaussian random variable with mean 0 and variance $a^2/2$. We also know that, $\mathbf{E}[e^{tZ}] = e^{t^2/2}$. So, $\mathbf{E}[e^{a(B_1 - B_{1/2})}] = e^{a^2/4}$. Consequently, $\mathbf{E}[e^{aB_1} | B_{1/2}] = e^{aB_{1/2} + a^2/4}$.

The result itself has the form of the MGF of a Gaussian with mean $B_{1/2}$ and variance $1/2$. (The MGF of $X = \mu + \sigma Z$, $Z = N(0, 1)$ is $M_X(a) = \exp[\mu a + \frac{1}{2}\sigma^2 a^2]$.) In fact, this shows that the conditional distribution of B_1 given $B_{1/2}$ is Gaussian of mean $B_{1/2}$ and variance $1/2$.

For the other expectation, note that $B_{1/2} - \frac{1}{2}B_1$ is independent of B_1 . We have:

$$\begin{aligned}
\mathbf{E}\left[\left(B_{1/2} - \frac{1}{2}B_1\right) B_1\right] &= \mathbf{E}(B_{1/2}B_1) - \frac{1}{2}\mathbf{E}[B_1^2] \\
&= \frac{1}{2} - \frac{1}{2} \cdot 1 \\
&= 0
\end{aligned}$$

Therefore, we have:

$$\begin{aligned}
\mathbf{E}[e^{aB_{1/2}} | B_1] &= \mathbf{E}[e^{a(B_{1/2} - \frac{1}{2}B_1) + \frac{a}{2}B_1} | B_1] \\
&= \mathbf{E}[e^{a(B_{1/2} - \frac{1}{2}B_1)} \cdot e^{\frac{a}{2}B_1} | B_1] \\
&= e^{\frac{a}{2}B_1} \mathbf{E}[e^{a(B_{1/2} - \frac{1}{2}B_1)} | B_1] \\
&\quad \{\text{Taking out what is known}\} \\
&= e^{\frac{a}{2}B_1} \mathbf{E}[e^{a(B_{1/2} - \frac{1}{2}B_1)}] \\
&\quad \{\text{Independence}\}
\end{aligned}$$

Now, $a(B_{1/2} - \frac{1}{2}B_1)$ is a random variable with mean 0 and variance $a^2(\frac{1}{2} - \frac{1}{4}) = \frac{a^2}{4}$. Consequently, $\mathbf{E}[e^{(a/2)Z}] = e^{\frac{a^2}{8}}$. Thus, $\mathbf{E}[e^{aB_{1/2}}|B_1] = e^{\frac{a}{2}B_1 + \frac{a^2}{8}}$.

Example 5.14. (Brownian bridge is conditioned Brownian motion). We know that the Brownian bridge $M_t = B_t - tB_1$, $t \in [0, 1]$ is independent of B_1 . We use this to show that the conditional distribution of the Brownian motion given the value at the end-point B_1 is the one of a Brownian bridge shifted by the straight line going from 0 to B_1 . To see this, we compute the conditional MGF of $(B_{t_1}, B_{t_2}, \dots, B_{t_n})$ given B_1 for some arbitrary choices of t_1, t_2, \dots, t_n in $[0, 1]$. We get the following by adding and subtracting $t_j B_1$:

$$\begin{aligned} \mathbf{E}[e^{a_1 B_{t_1} + \dots + a_n B_{t_n}} | B_1] &= \mathbf{E}[e^{a_1(B_{t_1} - t_1 B_1) + \dots + a_n(B_{t_n} - t_n B_1)} \cdot e^{(a_1 t_1 B_1 + \dots + a_n t_n B_1)} | B_1] \\ &= e^{(a_1 t_1 B_1 + \dots + a_n t_n B_1)} \mathbf{E}[e^{a_1 M_{t_1} + \dots + a_n M_{t_n}} | B_1] \\ &\quad \{\text{Taking out what is known}\} \\ &= e^{(a_1 t_1 B_1 + \dots + a_n t_n B_1)} \mathbf{E}[e^{a_1 M_{t_1} + \dots + a_n M_{t_n}}] \\ &\quad \{\text{Independence}\} \end{aligned}$$

The right side is exactly the MGF of the process $M_t + tB_1$, $t \in [0, 1]$ (for a fixed value B_1), where $(M_t, t \in [0, 1])$ is a Brownian bridge. This proves the claim.

Lemma 5.1. (*Conditional Jensen's Inequality*) If c is a convex function on \mathbf{R} and X is a random variable on $(\Omega, \mathcal{F}, \mathbb{P})$, then:

$$\mathbf{E}[c(X)] \geq c(\mathbf{E}[X])$$

More generally, if $\mathcal{G} \subseteq \mathcal{F}$ is a sigma-field, then:

$$\mathbf{E}[c(X)|\mathcal{G}] \geq c(\mathbf{E}[X|\mathcal{G}]) \tag{5.16}$$

Proof. We know that, if $c(x)$ is a convex function, the tangent to the curve c at any point lies below the curve. Fix a point $(t, c(t))$. The tangent to the curve at this point, is a straight-line of the form:

$$c(t) = y = mt + c$$

where $m(t) = c'(t)$. At any arbitrary point x we have:

$$c(x) \geq y = mx + c$$

Therefore, we have:

$$c(x) - c(t) \geq m(t)(x - t)$$

$$c(X) - c(Y) \geq m(Y)(X - Y)$$

Substituting $Y = \mathbf{E}[X|\mathcal{G}]$, we get:

$$c(X) - c(\mathbf{E}[X|\mathcal{G}]) \geq m(\mathbf{E}[X|\mathcal{G}]) (X - \mathbf{E}[X|\mathcal{G}])$$

Taking expectations on both sides, we get:

$$\mathbf{E}[(c(X) - c(\mathbf{E}[X|\mathcal{G}]))|\mathcal{G}] \geq \mathbf{E}[m(\mathbf{E}[X|\mathcal{G}]) (X - \mathbf{E}[X|\mathcal{G}])|\mathcal{G}]$$

The left-hand side simplifies as:

$$\begin{aligned} \mathbf{E}[(c(X) - c(\mathbf{E}[X|\mathcal{G}]))|\mathcal{G}] &= \mathbf{E}[c(X)|\mathcal{G}] - \mathbf{E}[c(\mathbf{E}[X|\mathcal{G}])|\mathcal{G}] \\ &\quad \{\text{Linearity}\} \\ &= \mathbf{E}[c(X)|\mathcal{G}] - c(\mathbf{E}[X|\mathcal{G}]) \\ &\quad \{c(\mathbf{E}[X|\mathcal{G}]) \text{ is } \mathcal{G}\text{-measurable}\} \end{aligned}$$

On the right hand side, we have:

$$\begin{aligned} \mathbf{E}[m(\mathbf{E}[X|\mathcal{G}]) (X - \mathbf{E}[X|\mathcal{G}])|\mathcal{G}] &= \mathbf{E}[m(\mathbf{E}[X|\mathcal{G}]) \cdot X|\mathcal{G}] - \mathbf{E}[m(\mathbf{E}[X|\mathcal{G}]) \cdot \mathbf{E}[X|\mathcal{G}]|\mathcal{G}] \\ &= \mathbf{E}[X|\mathcal{G}]m(\mathbf{E}[X|\mathcal{G}]) - m(\mathbf{E}[X|\mathcal{G}]) \cdot \mathbf{E}[X|\mathcal{G}] \\ &= 0 \end{aligned}$$

Consequently, it follows that $\mathbf{E}[c(X)|\mathcal{G}] \geq c(\mathbf{E}[X|\mathcal{G}])$. □

Example 5.15. (Embeddings of L^p spaces) Square-integrable random variables are in fact integrable. In other words, there is always the inclusion $L^2(\Omega, \mathcal{F}, \mathbb{P}) \subseteq L^1(\Omega, \mathcal{F}, \mathbb{P})$. In particular, square integrable random variables always have a well-defined variance. This embedding is a simple consequence of Jensen's inequality since:

$$|\mathbf{E}[X]|^2 \leq \mathbf{E}[|X|^2]$$

By taking the square root on both sides, we get:

$$\|X\|_1 \leq \|X\|_2$$

More generally, for any $1 < p < \infty$, we can define $L^p(\Omega, \mathcal{F}, \mathbb{P})$ to be the linear space of random variables such that $\mathbf{E}[|X|^p] < \infty$. Then for $p < q$, since $x^{q/p}$ is convex, we get by Jensen's inequality :

$$\mathbf{E}[|X|^q] = \mathbf{E}[(|X|^p)^{\frac{q}{p}}] \geq (\mathbf{E}[|X|^p])^{\frac{q}{p}}$$

Taking the q -th root on both sides:

$$\mathbf{E}[|X|^p]^{1/p} \leq \mathbf{E}[|X|^q]^{1/q}$$

So, if $X \in L^q$, then it must also be in L^p . Concretely, this means that any random variable with a finite q -moment will also have a finite p -moment, for $q > p$.

5.3 Martingales.

We now have all the tools to define martingales.

Definition 5.6. (Filtration). A *filtration* $(\mathcal{F}_t : t \geq 0)$ of Ω is an increasing sequence of σ -fields of Ω . That is,

$$\mathcal{F}_s \subseteq \mathcal{F}_t, \quad \forall s \leq t$$

We will usually take $\mathcal{F}_0 = \{\emptyset, \Omega\}$. The canonical example of a filtration is the natural filtration of a given process $(M_s : s \geq 0)$. This is the filtration given by $\mathcal{F}_t = \sigma(M_s, s \leq t)$. The inclusions of the σ -fields are then clear. For a given Brownian motion $(B_t, t \geq 0)$, the filtration $\mathcal{F}_t = \sigma(B_s, s \leq t)$ is sometimes called the *Brownian filtration*. We think of the filtration as the *flow of information of the process*.

Definition 5.7. A stochastic process $(X_t : t \geq 0)$ is said to be adapted to $(\mathcal{F}_t : t \geq 0)$, if for each t , the random variable X_t is \mathcal{F}_t -measurable.

Definition 5.8. (Martingale). A process $(M_t : t \geq 0)$ is a martingale for the filtration $(\mathcal{F}_t : t \geq 0)$ if the following hold:

- (1) The process is *adapted*, that is M_t is \mathcal{F}_t -measurable for all $t \geq 0$.
- (2) $\mathbf{E}[|M_t|] < \infty$ for all $t \geq 0$. (This ensures that the conditional expectation is well defined.)
- (3) *Martingale property*:

$$\mathbf{E}[M_t | \mathcal{F}_s] = M_s \quad \forall s \leq t$$

Roughly, speaking this means that the best approximation of a process at a future time t is its value at the present.

In particular, the martingale property implies that:

$$\begin{aligned} \mathbf{E}[M_t | \mathcal{F}_0] &= M_0 \\ \mathbf{E}[\mathbf{E}[M_t | \mathcal{F}_0]] &= \mathbf{E}[M_0] \\ \mathbf{E}[M_t] &= \mathbf{E}[M_0] \end{aligned} \tag{5.17}$$

{Tower Property}

Usually, we take \mathcal{F}_0 to be the trivial sigma-field $\{\emptyset, \Omega\}$. A random variable that is \mathcal{F}_0 -measurable must be a constant, so M_0 is a constant. In this case, $\mathbf{E}[M_t] = M_0$

for all t . If properties (1) and (2) are satisfied, but the best approximation is larger, $\mathbf{E}[M_t|\mathcal{F}_s] \geq M_s$, the process is called a *submartingale*. If it is smaller on average, $\mathbf{E}[M_t|\mathcal{F}_s] \leq \mathbf{E}[M_s]$, we say it is a supermartingale.

We will be mostly interested in martingales that are continuous and square-integrable. Continuous martingales are martingales whose paths $t \mapsto M_t(\omega)$ are continuous almost surely. Square-integrable martingales are such that $\mathbf{E}[|M_t|^2] < \infty$ for all t 's. This condition is stronger than $\mathbf{E}[|M_t|] < \infty$ due to Jensen's inequality.

Remark. (Martingales in Discrete-time). Martingales can be defined the same way if the index set of the process is discrete. For example, the filtration $(\mathcal{F}_n : n \in \mathbf{N})$ is a countable set and the martingale property is then replaced by $\mathbf{E}[M_{n+1}|\mathcal{F}_n] = M_n$ as expected. The tower-property then yields the martingale property $\mathbf{E}[M_{n+k}|\mathcal{F}_n] = M_n$ for $k \geq 1$.

Remark. (Continuous Filtrations). Filtrations with continuous time can be tricky to handle rigorously. For example, one has to make sense of what it means for \mathcal{F}_s as s approaches t from the left. Is it equal to \mathcal{F}_t ? Or is there actually less information in $\lim_{s \rightarrow t^-} \mathcal{F}_s$ than in \mathcal{F}_t ? This is a bit of headache when dealing with processes with jumps, like the Poisson process. However, if the paths are continuous, the technical problems are not as heavy.

Let's look at some of the important examples of martingales constructed from Brownian Motion.

Example 5.16. (Examples of Brownian Martingales)

(i) *Standard Brownian Motion.* Let $(B_t : t \geq 0)$ be a standard Brownian motion and let $(\mathcal{F}_t : t \geq 0)$ be a *Brownian filtration*. Then $(B_t : t \geq 0)$ is a square integrable martingale for the filtration $(\mathcal{F}_t : t \geq 0)$. Property (1) is obvious, because all the sets in \mathcal{F}_t are resolved, upon observing the outcome of B_t . Similarly, $\mathbf{E}[|B_t|] = 0$. As for the martingale property, note that, by the properties of conditional expectation in proposition (5.1), we have:

$$\begin{aligned}
\mathbf{E}[B_t|\mathcal{F}_s] &= \mathbf{E}[B_t|B_s] \\
&= \mathbf{E}[B_t - B_s + B_s|B_s] \\
&= \mathbf{E}[B_t - B_s|B_s] + \mathbf{E}[B_s|B_s] \\
&\quad \{\text{Linearity}\} \\
&= \mathbf{E}[B_t - B_s] + B_s \\
&\quad \{\text{Independence}\} \\
&= B_s
\end{aligned}$$

(ii) *Geometric Brownian Motion.* Let $(B_t, t \geq 0)$ be a standard brownian motion, and $\mathcal{F}_t = \sigma(B_s, s \leq t)$. A *geometric brownian motion* is a process $(S_t, t \geq 0)$ defined by:

$$S_t = S_0 \exp(\sigma B_t + \mu t)$$

for some parameter $\sigma > 0$ and $\mu \in \mathbf{R}$. This is simply the exponential of the Brownian motion with drift. This is not a martingale for most choices of μ ! In fact, one must take

$$\mu = -\frac{1}{2}\sigma^2$$

for the process to be a martingale for the Brownian filtration. Let's verify this. Property (1) is obvious since S_t is a function of B_t for each t . So, it is \mathcal{F}_t measurable. Moreover, property (2) is clear: $\mathbf{E}[\exp(\sigma B_t + \mu t)] = \mathbf{E}[\exp(\sigma\sqrt{t}Z + \mu t)] = \exp(\mu t + \frac{1}{2}\sigma^2 t)$. So, it's a finite quantity. As for the martingale property, note that by the properties of conditional expectation, and the MGF of Gaussians, we have for $s \leq t$:

$$\begin{aligned}
\mathbf{E}[S_t|\mathcal{F}_s] &= \mathbf{E} \left[S_0 \exp \left(\sigma B_t - \frac{1}{2} \sigma^2 t \right) | \mathcal{F}_s \right] \\
&= S_0 \exp \left(-\frac{1}{2} \sigma^2 t \right) \mathbf{E}[\exp(\sigma(B_t - B_s + B_s)) | \mathcal{F}_s] \\
&= S_0 \exp \left(-\frac{1}{2} \sigma^2 t \right) \exp(\sigma B_s) \mathbf{E}[\exp(\sigma(B_t - B_s)) | \mathcal{F}_s] \\
&\quad \{\text{Taking out what is known}\} \\
&= S_0 \exp \left(\sigma B_s - \frac{1}{2} \sigma^2 t \right) \mathbf{E}[\exp(\sigma(B_t - B_s))] \\
&\quad \{\text{Independence}\} \\
&= S_0 \exp \left(\sigma B_s - \frac{1}{2} \sigma^2 t + \frac{1}{2} \sigma^2 (t - s) \right) \\
&= S_0 \exp(\sigma B_s - \frac{1}{2} \sigma^2 s) \\
&= S_s
\end{aligned}$$

We will sometimes abuse terminology and refer to the martingale case of geometric brownian motion simply as geometric Brownian Motion when the context is clear.

(iii) *The square of the Brownian motion, compensated.* It is easy to check $(B_t^2, t \geq 0)$ is a submartingale by direct computation using increments or by Jensen's inequality: $\mathbf{E}[B_t^2 | \mathcal{F}_s] > (\mathbf{E}[B_t | \mathcal{F}_s])^2 = B_s^2, s < t$. It is nevertheless possible to compensate to get a martingale:

$$M_t = B_t^2 - t$$

It is an easy exercise to verify that $(M_t : t \geq 0)$ is a martingale for the Brownian filtration $(\mathcal{F}_t : t \geq 0)$.

$$\begin{aligned}
\mathbf{E}[M_t|\mathcal{F}_s] &= \mathbf{E}[B_t^2 - t|\mathcal{F}_s] \\
&= \mathbf{E}[B_t^2|\mathcal{F}_s] - t \\
&= \mathbf{E}[(B_t - B_s + B_s)^2|\mathcal{F}_s] - t \\
&= \mathbf{E}[(B_t - B_s)^2|\mathcal{F}_s] + 2\mathbf{E}[(B_t - B_s)B_s|\mathcal{F}_s] + \mathbf{E}[B_s^2|\mathcal{F}_s] - t \\
&= \mathbf{E}[(B_t - B_s)^2] + 2B_s\mathbf{E}[(B_t - B_s)|\mathcal{F}_s] + B_s^2 - t \\
&= \mathbf{E}[(B_t - B_s)^2] + 2B_s\mathbf{E}[(B_t - B_s)] + B_s^2 - t \\
&\quad \left\{ \begin{array}{l} (B_t - B_s) \text{ is independent of } \mathcal{F}_s \\ \text{Also, } B_s \text{ is known at time } s \end{array} \right\} \\
&= (t - s) + 2B_s \cdot 0 + B_s^2 - t \\
&= B_s^2 - s \\
&= M_s
\end{aligned}$$

Example 5.17. (Other important martingales).

(1) *Symmetric random walks.* This is an example of a martingale in discrete time. Take $(X_i : i \in \mathbf{N})$ to be IID random variables with $\mathbf{E}[X_1] = 0$ and $\mathbf{E}[|X_1|] < \infty$. Take $\mathcal{F}_n = \sigma(X_i, i \leq n)$ and

$$S_n = X_1 + X_2 + \dots + X_n, \quad S_0 = 0$$

Firstly, the information learned by observing the outcomes of X_1, \dots, X_n is enough to completely determine S_n . Hence, S_n is \mathcal{F}_n -measurable.

Next,

$$\begin{aligned}
|S_n| &= \left| \sum_{i=1}^n X_i \right| \\
&\leq \sum_{i=1}^n |X_i|
\end{aligned}$$

Consequently, by the monotonicity of expectations, we have:

$$\mathbf{E}[|S_n|] \leq \sum_{i=1}^n \mathbf{E}[|X_i|] < \infty$$

The martingale property is also satisfied. We have:

$$\begin{aligned}
\mathbf{E}[S_{n+1}|\mathcal{F}_n] &= \mathbf{E}[S_n + X_{n+1}|\mathcal{F}_n] \\
&= \mathbf{E}[S_n|\mathcal{F}_n] + \mathbf{E}[X_{n+1}|\mathcal{F}_n] \\
&= S_n + \mathbf{E}[X_{n+1}] \\
&\quad \left\{ \begin{array}{l} S_n \text{ is } \mathcal{F}_n\text{-measurable} \\ X_{n+1} \text{ is independent of } \mathcal{F}_n \end{array} \right\} \\
&= S_n + 0 \\
&= S_n
\end{aligned}$$

(2) *Compensated Poisson process.* Let $(N_t : t \geq 0)$ be a Poisson process with rate λ and $\mathcal{F}_t = \sigma(N_s, s \leq t)$. Then, N_t is a submartingale for its natural filtration. Again, properties (1) and (2) are easily checked. N_t is \mathcal{F}_t measurable. Moreover, $\mathbf{E}[|N_t|] = \mathbf{E}[N_t] = \frac{1}{\lambda}t < \infty$. The submartingale property follows by the independence of increments : for $s \leq t$,

$$\begin{aligned}
\mathbf{E}[N_t|\mathcal{F}_s] &= \mathbf{E}[N_t - N_s + N_s|\mathcal{F}_s] \\
&= \mathbf{E}[N_t - N_s|\mathcal{F}_s] + \mathbf{E}[N_s|\mathcal{F}_s] \\
&= \mathbf{E}[N_t - N_s] + N_s \\
&= \lambda(t - s) + N_s \\
&= \{\cdot : \mathbf{E}[N_t] = \lambda t\}
\end{aligned}$$

More importantly, we get a martingale by slightly modifying the process. Indeed, if we subtract λt , we have that the process :

$$M_t = N_t - \lambda t$$

is a martingale. We have:

$$\begin{aligned}
\mathbf{E}[M_t|\mathcal{F}_s] &= \mathbf{E}[N_t - \lambda t|\mathcal{F}_s] \\
&= \lambda t - \lambda s + N_s - \lambda t \\
&= N_s - \lambda s \\
&= M_s
\end{aligned}$$

This is called the *compensated Poisson process*. Let us simulate 10 paths of the compensated poisson process on $[0, 10]$.

Listing 9: Generating 10 paths of a compensated Poisson process

```
import numpy as np
import matplotlib.pyplot as plt

# Generates a sample path of a compensated poisson process
# with rate : `lambda_` per unit time
# on the interval [0,T], and subintervals of size `stepSize`.

def generateCompensatedPoissonPath(lambda_,T,stepSize):
    N = int(T/stepSize)

    poissonParam = lambda_ * stepSize

    x = np.random.poisson(lam=poissonParam,size=N)
    x = np.concatenate([[0.0], x])
    N_t = np.cumsum(x)
    t = np.linspace(start=0.0,stop=10.0,num=1001)

    M_t = np.subtract(N_t,lambda_ * t)
    return M_t


t = np.linspace(0,10,1001)
plt.grid(True)

plt.xlabel(r'Time $t$')
plt.ylabel(r'Compensated poisson process $M(t)$')
plt.grid(True)
plt.title(r'$10$ paths of the compensated Poisson process on $[0,10]$')

for i in range(10):
    # Generate a poisson path with rate 1 /sec = 0.01 /millisec
    n_t = generateCompensatedPoissonPath(lambda_=1.0, T=10,
        stepSize=0.01)
    plt.plot(t, n_t)

plt.show()
plt.close()
```

10 paths of the compensated Poisson process on $[0, 10]$



We saw in the two examples, that, even though a process is not itself a martingale, we can sometimes *compensate* to obtain a martingale! Ito Calculus will greatly extend this perspective. We will have systematic rules that show when a function of Brownian motion is a martingale and if not, how to modify it to get one.

For now, we observe that a convex function of a martingale is always a submartingale by Jensen's inequality.

Corollary 5.1. *If c is a convex function on \mathbf{R} and $(M_t : t \geq 0)$ is a martingale for $(\mathcal{F}_t : t \geq 0)$, then the process $(c(M_t) : t \geq 0)$ is a submartingale for the same filtration, granted that $\mathbf{E}[|c(M_t)|] < \infty$.*

Proof. The fact that $c(M_t)$ is adapted to the filtration is clear since it is an explicit function of M_t . The integrability is by assumption. The submartingale property is checked as follows:

$$\mathbf{E}[c(M_t)|\mathcal{F}_s] \geq c(\mathbf{E}[M_t|\mathcal{F}_s]) = c(M_s)$$

□

Remark. (The Doob-Meyer Decomposition Theorem). Let $(X_n : n \in \mathbf{N})$ be a submartingale with respect to a filtration $(\mathcal{F}_n : n \in \mathbf{N})$. Define a sequence of random variables $(A_n : n \in \mathbf{N})$ by $A_0 = 0$ and

$$A_n = \sum_{i=1}^n (\mathbf{E}[X_i | \mathcal{F}_{i-1}] - X_{i-1}), \quad n \geq 1$$

Note that A_n is \mathcal{F}_{n-1} -measurable. Moreover, since $(X_n : n \in \mathbf{N})$ is a submartingale, we have $\mathbf{E}[X_i | \mathcal{F}_{i-1}] - X_{i-1} \geq 0$ almost surely. Hence, $(A_n : n \in \mathbf{N})$ is an increasing sequence almost surely. Let $M_n = X_n - A_n$.

We have:

$$\begin{aligned} \mathbf{E}[M_n | \mathcal{F}_{n-1}] &= \mathbf{E}[X_n - A_n | \mathcal{F}_{n-1}] \\ &= \mathbf{E}[X_n | \mathcal{F}_{n-1}] - \mathbf{E}[A_n | \mathcal{F}_{n-1}] \\ &= \mathbf{E}[X_n | \mathcal{F}_{n-1}] - \mathbf{E}[\mathbf{E}[X_n | \mathcal{F}_{n-1}] - X_{n-1} + A_{n-1} | \mathcal{F}_{n-1}] \\ &= \mathbf{E}[X_n | \mathcal{F}_{n-1}] - \mathbf{E}[X_n | \mathcal{F}_{n-1}] + \mathbf{E}[X_{n-1} | \mathcal{F}_{n-1}] - \mathbf{E}[A_{n-1} | \mathcal{F}_{n-1}] \\ &= \cancel{\mathbf{E}[X_n | \mathcal{F}_{n-1}]} - \cancel{\mathbf{E}[X_n | \mathcal{F}_{n-1}]} + X_{n-1} - A_{n-1} \\ &= M_{n-1} \end{aligned}$$

Thus, $(M_n : n \in \mathbf{N})$ is a martingale. Thus, we have obtained the Doob decomposition:

$$X_n = M_n + A_n \tag{5.18}$$

This decomposition of a submartingale as a sum of a martingale and an adapted increasing sequence is unique, if we require that $A_0 = 0$ and that A_n is \mathcal{F}_{n-1} -measurable.

For the continuous-time case, the situation is much more complicated. The analogue of equation (5.18) is called the *Doob-Meyer decomposition*. We briefly describe this decomposition and avoid the technical details. All stochastic processes $X(t)$ are assumed to be right-continuous with left-hand limits $X(t-)$.

Let $X(t)$, $a \leq t \leq b$ be a submartingale with respect to a right-continuous filtration $(\mathcal{F}_t : a \leq t \leq b)$. If $X(t)$ satisfies certain conditions, then it can be uniquely decomposed as:

$$X(t) = M(t) + C(t), \quad a \leq t \leq b$$

where $M(t)$, $a \leq t \leq b$ is a martingale with respect to $(\mathcal{F}_t; a \leq t \leq b)$, $C(t)$ is right-continuous and increasing almost surely with $\mathbf{E}[C(t)] < \infty$.

Example 5.18. (Square of a Poisson Process). Let $(N_t : t \geq 0)$ be a Poisson process with rate λ . We consider the compensated process $M_t = N_t - \lambda t$. By (5.1), the process $(M_t^2 : t \geq 0)$ is a submartingale for the filtration $(\mathcal{F}_t : t \geq 0)$ of the Poisson process. How should we compensated M_t^2 to get a martingale? A direct computation using the properties of conditional expectation yields:

$$\begin{aligned} \mathbf{E}[M_t^2 | \mathcal{F}_s] &= \mathbf{E}[(M_t - M_s + M_s)^2 | \mathcal{F}_s] \\ &= \mathbf{E}[(M_t - M_s)^2 + 2(M_t - M_s)M_s + M_s^2 | \mathcal{F}_s] \\ &= \mathbf{E}[(M_t - M_s)^2 | \mathcal{F}_s] + 2\mathbf{E}[(M_t - M_s)M_s | \mathcal{F}_s] + \mathbf{E}[M_s^2 | \mathcal{F}_s] \\ &= \mathbf{E}[(M_t - M_s)^2] + 2M_s \underbrace{\mathbf{E}[M_t - M_s]}_{\text{equals 0}} + M_s^2 \\ &= \mathbf{E}[(M_t - M_s)^2] + M_s^2 \end{aligned}$$

Now, if $X \sim \text{Poisson}(\lambda t)$, then $\mathbf{E}[X] = \lambda t$ and $\mathbf{E}[X^2] = \lambda t(\lambda t + 1)$.

$$\begin{aligned} \mathbf{E}[(M_t - M_s)^2] &= \mathbf{E}[\{(N_t - N_s) - \lambda(t - s)\}^2] \\ &= \mathbf{E}[(N_t - N_s)^2] - 2\lambda(t - s)\mathbf{E}[(N_t - N_s)] + \lambda^2(t - s)^2 \\ &= \lambda^2(t - s)^2 + \lambda(t - s) - 2\lambda(t - s) \cdot \lambda(t - s) + \lambda^2(t - s)^2 \\ &= \lambda(t - s) \end{aligned}$$

Thus,

$$\mathbf{E}[M_t^2 - \lambda t | \mathcal{F}_s] = M_s^2 - \lambda s$$

We conclude that the process $(M_t^2 - \lambda t : t \geq 0)$ is a martingale. The Doob-Meyer decomposition of the submartingale M_t^2 is then:

$$M_t^2 = (M_t^2 - \lambda t) + \lambda t$$

Example 5.19. Consider a Brownian motion $B(t)$. The quadratic variation of the process $(B(t) : t \geq 0)$ over the interval $[0, t]$ is given by $[B]_t = t$. On the other hand, we saw, that the square of Brownian motion compensated, $(B_t^2 - t : t \geq 0)$ is a martingale. Hence, the Doob-Meyer decomposition of $B(t)^2$ is given by:

$$B(t)^2 = (B(t)^2 - t) + t$$

5.4 Computations with Martingales.

Martingales are not only conceptually interesting, they are also formidable tools to compute probabilities and expectations of processes. For example, in this section, we will solve the *gambler's ruin* problem for Brownian motion. For convenience, we introduce the notion of *stopping time* before doing so.

Definition 5.9. A random variable $\tau : \Omega \rightarrow \mathbf{N} \cup \{+\infty\}$ is said to be a *stopping time* for the filtration $(\mathcal{F}_t : t \geq 0)$ if and only if:

$$\{\omega : \tau(\omega) \leq t\} \in \mathcal{F}_t, \quad \forall t \geq 0$$

Note that since \mathcal{F}_t is a sigma-field, if τ is a stopping time, then we must also have that $\{\omega : \tau(\omega) > t\} \in \mathcal{F}_t$.

In other words, τ is a stopping time, if we can decide if the events $\{\tau \leq t\}$ occurred or not based on the information available at time t .

The term *stopping time* comes from gambling: a gambler can decide to stop playing at a random time (depending for example on previous gains or losses), but when he or she decides to stop, his/her decision is based solely upon the knowledge of what happened before, and does not depend on future outcomes. In other words, the stopping policy/strategy can only depend on past outcomes. Otherwise, it would mean that he/she has a crystall ball.

Example 5.20. (Examples of stopping times).

(i) *First passage time.* This is the first time when a process reaches a certain value. To be precise, let $X = (X_t : t \geq 0)$ be a process and $(\mathcal{F}_t : t \geq 0)$ be its natural filtration. For $a > 0$, we define the first passage time at a to be:

$$\tau(\omega) = \inf\{s \geq 0 : X_s(\omega) \geq a\}$$

If the path ω never reaches a , we set $\tau(\omega) = \infty$. Now, for t fixed and for a given path $X(\omega)$, it is possible to know if $\{\tau(\omega) \leq t\}$ (the path has reached a before time t) or $\{\tau(\omega) > t\}$ (the path has not reached a before time t) with the information available at time t , since we are looking at the first time the process reaches a . Hence, we conclude that τ is a stopping time.

(ii) *Hitting time.* More generally, we can consider the first time (if ever) that the path of a process $(X_t : t \geq 0)$ enters or hits a subset B of \mathbf{R} :

$$\tau(\omega) = \min\{s \geq 0 : X_s(\omega) \in B\}$$

The first passage time is the particular case in which $B = [a, \infty)$.

(iii) *Minimum of two stopping times.* If τ and τ' are two stopping times for the same filtration $(\mathcal{F}_t : t \geq 0)$, then so is the minimum $\tau \wedge \tau'$ between the two, where

$$(\tau \wedge \tau')(\omega) = \min\{\tau(\omega), \tau'(\omega)\}$$

This is because for any $t \geq 0$:

$$\{\omega : (\tau \wedge \tau')(\omega) \leq t\} = \{\omega : \tau(\omega) \leq t\} \cup \{\omega : \tau'(\omega) \leq t\}$$

Since the right hand side is the union of two events in \mathcal{F}_t , it must also be in \mathcal{F}_t by the properties of a sigma-field. We conclude that $\tau \wedge \tau'$ is a stopping time. Is it also the case that the maximum $\tau \vee \tau'$ is a stopping time?

For any fixed $t \geq 0$, we have:

$$\{\omega : (\tau \vee \tau')(\omega) \leq t\} = \{\omega : \tau(\omega) \leq t\} \cap \{\omega : \tau'(\omega) \leq t\}$$

Since the right hand side is the intersection of two events in \mathcal{F}_t , it must also be in \mathcal{F}_t by the properties of a sigma-field. We conclude that $\tau \vee \tau'$ is a stopping time.

Example 5.21. (Last passage time is not a stopping time). What if we look at the last time the process reaches a , that is:

$$\rho(\omega) = \sup\{t \geq 0 : X_t(\omega) \geq a\}$$

This is a well-defined random variable, but it is not a stopping time. Based on the information available at time t , we are not able to decide whether or not $\{\rho(\omega) \leq t\}$ occurred or not, as the path can always reach a one more time after t .

It turns out that a martingale that is stopped when the stopping time is attained remains a martingale.

Proposition 5.2. (*Stopped Martingale*). If $(M_t : t \geq 0)$ is a continuous martingale for the filtration $(\mathcal{F}_t : t \geq 0)$ and τ is a stopping time for the same filtration, then the stopped process defined by

$$M_{t \wedge \tau} = \begin{cases} M_t & t \leq \tau \\ M_\tau & t > \tau \end{cases}$$

is also a continuous martingale for the same filtration.

Theorem 5.3. (*Doob's Optional sampling theorem*). If $(M_t : t \geq 0)$ is a continuous martingale for the filtration $(\mathcal{F}_t : t \geq 0)$ and τ is a stopping time such that $\tau < \infty$ and the stopped process $(M_{t \wedge \tau} : t \geq 0)$ is bounded, then:

$$\mathbf{E}[M_\tau] = M_0$$

Proof. Since $(M_{\tau \wedge t} : t \geq 0)$ is a martingale, we always have:

$$\mathbf{E}[M_{\tau \wedge t}] = M_0$$

Now, since $\tau(\omega) < \infty$, we must have that $\lim_{t \rightarrow \infty} M_{\tau \wedge t} = M_\tau$ almost surely. In particular, we have:

$$\mathbf{E}[M_\tau] = \mathbf{E} \left[\lim_{t \rightarrow \infty} M_{\tau \wedge t} \right] = \lim_{t \rightarrow \infty} \mathbf{E}[M_{\tau \wedge t}] = \lim_{t \rightarrow \infty} \mathbf{E}[M_0] = \mathbf{E}[M_0]$$

where we passed to the limit, using the dominated convergence theorem (2.6). \square

Example 5.22. (Gambler's ruin with Brownian motion). The *gambler's ruin problem* is known in different forms. Roughly speaking, it refers to the problem of computing the probability of a gambler making a series of bets reaching a certain amount before going broke. In terms of Brownian motion (and stochastic processes in general), it translates to the following questions: Let $(B_t : t \geq 0)$ be a standard brownian motion starting at $B_0 = 0$ and $a, b > 0$.

- (1) What is the probability that a Brownian path reaches a before $-b$?
- (2) What is the expected waiting time for the path to reach a or $-b$?

For the first question, it is a simple computation using stopping time and martingale properties. Define the hitting time:

$$\tau(\omega) = \inf\{t \geq 0 : B_t(\omega) \geq a \text{ or } B_t(\omega) \leq -b\}$$

Note that τ is the minimum between the first passage time at a and the one at $-b$.

We first show that $\tau < \infty$ almost surely. In other words, all Brownian paths reach a or $-b$ eventually. To see this, consider the event E_n that the n -th increment exceeds $a + b$

$$E_n := \{|B_n - B_{n-1}| > a + b\}$$

Note that, if E_n occurs, then we must have that the Brownian motion path exits the interval $[-b, a]$. Moreover, we have $\mathbb{P}(E_n) = \mathbb{P}(E_1)$ for all n . Since the events E_n are independent, we have:

$$\mathbb{P}(E_1^C \cap E_2^C \cap \dots \cap E_n^C) = (1 - p)^n$$

As $n \rightarrow \infty$ we have:

$$\lim_{n \rightarrow \infty} \mathbb{P}(E_1^C \cap E_2^C \cap \dots \cap E_n^C) = 0$$

The sequence of events (F_n) where $F_n = E_1^C \cap E_2^C \cap \dots \cap E_n^C$ is a decreasing sequence of events. By the continuity of probability measure lemma (1.8), we conclude that:

$$\lim_{n \rightarrow \infty} \mathbb{P}(F_n) = \mathbb{P}\left(\bigcap_{n=1}^{\infty} F_n\right) = 0$$

Therefore, it must be the case $\mathbb{P}(\cup_{n=1}^{\infty} E_n) = 1$. So, E_n must occur for some n , so all brownian motion paths reach a or $-b$.

Since $\tau < \infty$ with probability one, the random variable B_τ is well-defined: $B_\tau(\omega) = B_t(\omega)$ if $\tau(\omega) = t$. It can only take two values: a or $-b$. Question (1) above translates into computing $\mathbb{P}(B_\tau = a)$. On one hand, we have:

$$\mathbf{E}[B_\tau] = a\mathbb{P}(B_\tau = a) + (-b)(1 - \mathbb{P}(B_\tau = a))$$

On the other hand, by corollary (5.3), we have $\mathbf{E}[B_\tau] = \mathbf{E}[B_0] = 0$. (Note that the stopped process $(B_{t \wedge \tau} : t \geq 0)$ is bounded above by a and by $-b$ below). Putting these two observations together, we get:

$$\mathbb{P}(B_\tau = a) = \frac{b}{a + b}$$

A very simple and elegant answer!

We will revisit this problem again and again. In particular, we will answer the question above for Brownian motion with a drift at length further ahead.

Example 5.23. (Expected Waiting Time). Let τ be as in the last example. We now answer question (2) of the gambler's ruin problem:

$$\mathbf{E}[B_\tau] = ab$$

Note that the expected waiting time is consistent with the rough heuristic that Brownian motion travels a distance \sqrt{t} by time t . We now use the martingale $M_t = B_t^2 - t$. On the one hand, if we apply optional stopping in corollary (5.3), we get:

$$\mathbf{E}[M_\tau] = M_0 = 0$$

Moreover, we know the distribution of B_τ , thanks to the probability calculated in the last example. We can therefore compute $\mathbf{E}[M_\tau]$ directly:

$$\begin{aligned} 0 &= \mathbf{E}[M_\tau] \\ &= \mathbf{E}[B_\tau^2 - \tau] \\ &= \mathbf{E}[B_\tau^2] - \mathbf{E}[\tau] \\ &= a^2 \cdot \frac{b}{a+b} + b^2 \cdot \frac{a}{a+b} - \mathbf{E}[\tau] \\ \mathbf{E}[\tau] &= \frac{a^2b + b^2a}{a+b} \\ &= \frac{ab(a+b)}{(a+b)} = ab \end{aligned}$$

Why can we apply optional stopping here? The random variable τ is finite with probability 1 as before. However, the stopped martingale is not necessarily bounded as before: $B_{\tau \wedge t}$ is bounded but τ is not. However, the conclusion of optional stopping still holds. Indeed, we have:

$$\mathbf{E}[M_{t \wedge \tau}] = \mathbf{E}[B_{t \wedge \tau}^2] - \mathbf{E}[t \wedge \tau]$$

By the bounded convergence theorem, we get $\lim_{t \rightarrow \infty} \mathbf{E}[B_{t \wedge \tau}^2] = \mathbf{E}[\lim_{t \rightarrow \infty} B_{t \wedge \tau}^2] = \mathbf{E}[B_\tau^2]$. Since $\tau \wedge t$ is a non-decreasing sequence and as $t \rightarrow \infty$, $t \wedge \tau \rightarrow \tau$ almost surely, as $\tau < \infty$, by the monotone convergence theorem, $\lim_{t \rightarrow \infty} \mathbf{E}[t \wedge \tau] = \mathbf{E}[\tau]$.

Example 5.24. (First passage time of Brownian Motion.) We can use the previous two examples to get some very interesting information on the first passage time:

$$\tau_a = \inf\{t \geq 0 : B_t \geq a\}$$

Let $\tau = \tau_a \wedge \tau_{-b}$ be as in the previous examples with $\tau_{-b} = \inf\{t \geq 0 : B_t \leq -b\}$. Note that τ_{-b} is a sequence of random variables that is increasing in b . A brownian motion path must cross through -1 before it hits -2 for the first time and in general $\tau_{-n}(\omega) \leq \tau_{-(n+1)}(\omega)$. Moreover, we have $\tau_{-b} \rightarrow \infty$ almost surely as $b \rightarrow \infty$. That's because, $\mathbb{P}\{\tau < \infty\} = 1$. Moreover, the event $\{B_\tau = a\}$ is the same as $\{\tau_a < \tau_{-b}\}$. Now, the events $\{\tau_a < \tau_{-b}\}$ are increasing in b , since if a path reaches a before $-b$, it will do so as well for a more negative value of $-b$. On one hand, this means by the continuity of probability measure lemma (1.8) that:

$$\begin{aligned} \lim_{b \rightarrow \infty} \mathbb{P}\{\tau_a < \tau_{-b}\} &= \mathbb{P}\{\lim_{b \rightarrow \infty} \tau_a < \tau_{-b}\} \\ &= \mathbb{P}\{\tau_a < \infty\} \end{aligned}$$

On the other hand, we have by example (5.22)

$$\begin{aligned} \lim_{b \rightarrow \infty} \mathbb{P}\{\tau_a < \tau_{-b}\} &= \lim_{b \rightarrow \infty} \mathbb{P}\{B_\tau = a\} \\ &= \lim_{b \rightarrow \infty} \frac{b}{b+a} \\ &= 1 \end{aligned}$$

We just showed that:

$$\mathbb{P}\{\tau_a < \infty\} = 1 \quad (5.19)$$

In other words, every Brownian path will reach a , no matter how large a is!

How long will it take to reach a on average? Well, we know from example (5.23) that $\mathbf{E}[\tau_a \wedge \tau_{-b}] = ab$. On one hand this means,

$$\lim_{b \rightarrow \infty} \mathbf{E}[\tau_a \wedge \tau_{-b}] = \lim_{b \rightarrow \infty} ab = \infty$$

On the other hand, since the random variables τ_{-b} are increasing,

$$\lim_{b \rightarrow \infty} \mathbf{E}[\tau_a \wedge \tau_{-b}] = \mathbf{E}\left[\lim_{b \rightarrow \infty} \tau_a \wedge \tau_{-b}\right] = \mathbf{E}[\tau_a]$$

by the monotone convergence theorem (2.3). We just proved that:

$$\mathbf{E}[\tau_a] = \infty$$

In other words, any Brownian motion path will reach a , but the expected waiting time for this to occur is infinite, no matter, how small a is! What is happening here? No matter, how small a is, there is always paths that reach very large negative values before hitting a . These paths might be unlikely. However, the first passage time for these paths is so large that they affect the value of the expectation substantially. In other words, τ_a is a *heavy-tailed random variable*. We look at the distribution of τ_a in more detail in the next section.

Example 5.25. (When option stopping fails). Consider τ_a , the first passage time at $a > 0$. The random variable B_{τ_a} is well-defined since $\tau_a < \infty$. In fact, we have $B_{\tau_a} = a$ with probability one. Therefore, the following must hold:

$$\mathbf{E}[B_{\tau_a}] = a \neq B_0$$

Optional stopping theorem corollary (5.3) does not apply here, since the stopped process $(B_{t \wedge \tau_a} : t \geq 0)$ is not bounded. $B_{t \wedge \tau_a}$ can become infinitely negative before hitting a .

5.5 Reflection principle for Brownian motion.

Proposition 5.3. (*Bachelier's formula*). Let $(B_t : t \leq T)$ be a standard brownian motion on $[0, T]$. Then, the CDF of the random variable $\sup_{0 \leq t \leq T} B_t$ is:

$$\mathbb{P} \left(\sup_{0 \leq t \leq T} B_t \leq a \right) = \mathbb{P} (|B_T| \leq a)$$

In particular, its PDF is:

$$f_{\max}(a) = \frac{2}{\sqrt{2\pi T}} e^{-\frac{a^2}{2T}}$$

Remark. We can verify these results empirically. Note that the paths of the random variables $\max_{0 \leq s \leq t} B_s$ and $|B_t|$ are very different as t varies for a given ω . One is increasing and the other is not. The equality holds in distribution for a fixed t . As a bonus corollary, we get the distribution of the first passage time at a .

Corollary 5.2. Let $a \geq 0$ and $\tau_a = \inf\{t \geq 0 : B_t \geq a\}$. Then:

$$\mathbb{P}(\tau_a \leq T) = \mathbb{P} \left(\max_{0 \leq t \leq T} B_t \geq a \right) = \int_a^\infty \frac{2}{\sqrt{2\pi T}} e^{-\frac{x^2}{2T}} dx$$

In particular, the random variable τ_a has the PDF:

$$f_{\tau_a}(t) = \frac{a}{\sqrt{2\pi}} \frac{e^{-\frac{a^2}{2t}}}{t^{3/2}}, \quad t > 0$$

This implies that it is heavy-tailed with $\mathbf{E}[\tau_a] = \infty$.

Proof. The maximum on $[0, T]$ is larger than or equal to a if and only if $\tau_a \leq T$. Therefore, the events $\{\max_{0 \leq t \leq T} B_t \geq a\}$ and $\{\tau_a \leq T\}$ are the same. So, the CDF $\mathbb{P}(\tau_a \leq t)$ of τ_a , by proposition (5.3) $\int_a^\infty f_{\max}(x) dx = \int_a^\infty \frac{2}{\sqrt{2\pi T}} e^{-\frac{x^2}{2T}} dx$. To get the

PDF, it remains to differentiate the integral with respect to t . This is easy to do once we realize by a change of variable $u = x/\sqrt{t}$ that:

$$\begin{aligned} \int_a^\infty \frac{2}{\sqrt{2\pi t}} e^{-\frac{x^2}{2t}} dx &= \int_{a/\sqrt{t}}^\infty \frac{2}{\sqrt{2\pi t}} e^{-\frac{u^2}{2}} \sqrt{t} \cdot du \\ &= \int_{a/\sqrt{t}}^\infty \frac{2}{\sqrt{2\pi}} e^{-u^2/2} du \\ F_{\tau_a}(t) &= 2(1 - \Phi(a/\sqrt{t})) \end{aligned}$$

Differentiating with respect to t , we get:

$$\begin{aligned} f_{\tau_a}(t) &= -2\phi(a/\sqrt{t}) \cdot a \cdot \left(-\frac{1}{2t^{3/2}}\right) \\ &= \frac{a}{t^{3/2}} \phi\left(\frac{a}{\sqrt{t}}\right) \\ &= \frac{a}{t^{3/2}} \cdot \frac{1}{\sqrt{2\pi}} e^{-\frac{a^2}{2t}} \end{aligned}$$

To estimate the expectation, it suffices to realize that for $t \geq 1$, $e^{-\frac{a^2}{2t}}$ is larger than $e^{-\frac{a^2}{2}}$. Therefore, we have:

$$\mathbf{E}[\tau_a] = \int_0^\infty t \frac{a}{\sqrt{2\pi}} \frac{e^{-a^2/2t}}{t^{3/2}} dt \geq \frac{ae^{-a^2/2}}{\sqrt{2\pi}} \int_1^\infty t^{-1/2} dt$$

This is an improper integral and it diverges like \sqrt{t} and is infinite as claimed. \square

To prove proposition (5.3), we will need an important property of Brownian motion called the *reflection principle*. To motivate it, recall the reflection symmetry of Brownian motion at time s in proposition (4.4). It turns out that this reflection property also holds if s is replaced by a stopping time.

Lemma 5.2. (*Reflection principle*). *Let $(B_t : t \geq 0)$ be a standard Brownian motion and let τ be a stopping time for its filtration. Then, the process $(\tilde{B}_t : t \geq 0)$ defined by the reflection at time τ :*

$$\tilde{B}_t = \begin{cases} B_t & \text{if } t \leq \tau \\ B_\tau - (B_t - B_\tau) & \text{if } t > \tau \end{cases}$$

is also a standard brownian motion.

Remark. We defer the proof of the reflection property of Brownian motion to a further section. It is intuitive and instructive to quickly picture this in the discrete-time setting. I adopt the approach as in Shreve-I.

We repeatedly toss a fair coin (p , the probability of H on each toss, and $q = 1 - p$, the probability of T on each toss, are both equal to $\frac{1}{2}$). We denote the successive outcomes of the tosses by $\omega_1 \omega_2 \omega_3 \dots$. Let

$$X_j = \begin{cases} -1 & \text{if } \omega_j = H \\ +1 & \text{if } \omega_j = T \end{cases}$$

and define $M_0 = 0$, $M_n = \sum_{j=1}^n X_j$. The process $(M_n : n \in \mathbf{N})$ is a symmetric random walk.

Suppose we toss a coin an odd number $(2j - 1)$ of times. Some of the paths will reach level 1 in the first $2j - 1$ steps and other will not reach. In the case of 3 tosses, there are $2^3 = 8$ possible paths and 5 of these reach level 1 at some time $\tau_1 \leq 2j - 1$. From that moment on, we can create a reflected path, which steps up each time the original path steps down and steps down each time the original path steps up. If the original path ends above 1 at the final time $2j - 1$, the reflected path ends below 1 and vice versa. If the original path ends at 1, the reflected path does also. In fact, the reflection at the first hitting time has the same distribution as the original random walk.

The key here is, out of the 5 paths that reach level 1 at some time, there are as many reflected paths that exceed 1 at time $(2j - 1)$ as there are original paths that exceed 1 at time $(2j - 1)$. So, to count the total number of paths that reach level 1 by time $(2j - 1)$, we can count the paths that are at 1 at time $(2j - 1)$ and then add on *twice* the number of paths that exceed 1 at time $(2j - 1)$.

With this new tool, we can now prove proposition (5.3).

Proof. Consider $\mathbb{P}(\max_{t \leq T} B_t \geq a)$. By splitting this probability over the event of the endpoint, we have:

$$\mathbb{P}\left(\max_{t \leq T} B_t \geq a\right) = \mathbb{P}\left(\max_{t \leq T} B_t \geq a, B_T > a\right) + \mathbb{P}\left(\max_{t \leq T} B_t \geq a, B_T \leq a\right)$$

Note also, that $\mathbb{P}(B_T = a) = 0$. Hence, the first probability equals $\mathbb{P}(B_T \geq a)$. As for the second, consider the time τ_a . On the event considered, we have $\tau_a \leq T$ and using lemma (5.2) at that time, we get

$$\mathbb{P}\left(\max_{t \leq T} B_t \geq a, B_T \leq a\right) = \mathbb{P}\left(\max_{t \leq T} B_t \geq a, \tilde{B}_T \geq a\right)$$

Observe that the event $\{\max_{t \leq T} B_t \geq a\}$ is the same as $\{\max_{t \leq T} \tilde{B}_t \geq a\}$. (A rough picture might help here.) Therefore the above probability is

$$\mathbb{P}\left(\max_{t \leq T} B_t \geq a, B_T \leq a\right) = \mathbb{P}\left(\max_{t \leq T} \tilde{B}_t \geq a, \tilde{B}_T \geq a\right) = \mathbb{P}\left(\max_{t \leq T} B_t \geq a, B_T \geq a\right)$$

where the last equality follows from the reflection principle (\tilde{B}_t is also a standard brownian motion, and B_T and \tilde{B}_T have the same distribution.) But, as above, the last probability is equal to $\mathbb{P}(B_T \geq a)$. We conclude that:

$$\mathbb{P}\left(\max_{t \leq T} B_t \geq a\right) = 2\mathbb{P}(B_T \geq a) = \frac{2}{\sqrt{2\pi T}} \int_a^\infty e^{-\frac{x^2}{2T}} dx = \mathbb{P}(|B_T| \geq a)$$

This implies in particular that $\mathbb{P}(\max_{t \leq T} B_t = a) = 0$. Thus, we also have $\mathbb{P}(\max_{t \leq T} B_t \leq a) = \mathbb{P}(|B_T| \leq a)$ as claimed. \square

Example 5.26. (Simulating Martingales) Sample 10 paths of the following process with a step-size of 0.01:

- (a) $B_t^2 - t, t \in [0, 1]$
- (b) Geometric Brownian motion : $S_t = \exp(B_t - t/2), t \in [0, 1]$.

Let's write a simple `BrownianMotion` class, that we shall use to generate sample paths.

Listing 10: 10 paths of $B_t^2 - t$

```
import numpy as np
import matplotlib.pyplot as plt

import numpy as np
import matplotlib.pyplot as plt

class BrownianMotion:
    def __init__(self, stepSize, T):
        self.stepSize = stepSize
        self.T = T
        self.N = int(T/stepSize)

        # Construct the covariance matrix C for
        # the gaussian vector B(1/N), ..., B(1)
        self.C = np.zeros((self.N, self.N))

        for i in range(self.N):
            for j in range(self.N):
                s = (i + 1) / self.N
                t = (j + 1) / self.N
                self.C[i, j] = np.min([s, t])

        self.A = np.linalg.cholesky(self.C)

    def samplePath(self):
        Z = np.random.normal(0, self.T, self.N)
        B_t = np.matmul(self.A, Z)
        B_t = np.concatenate([[0.0], B_t])
        return B_t
```

Now, the process $B_t^2 - t$ can be sampled as follows:

Listing 11: 10 paths of $B_t^2 - t$

```
def generateSquareOfBMCompensated(numOfPaths, stepSize, T):
    N = int(T/stepSize)

    X = []
    brownianMotion = BrownianMotion(stepSize, T)
    for n in range(numOfPaths):

        B_t = brownianMotion.samplePath()

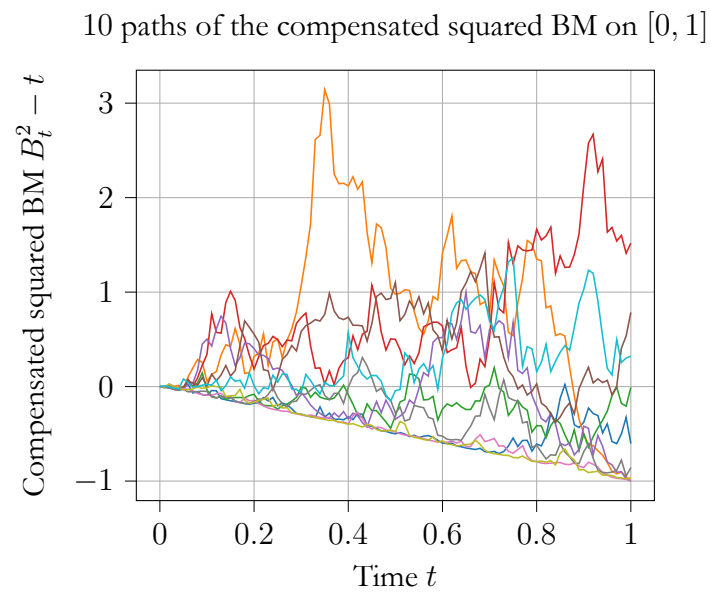
        B_t_sq = np.square(B_t)
```

```

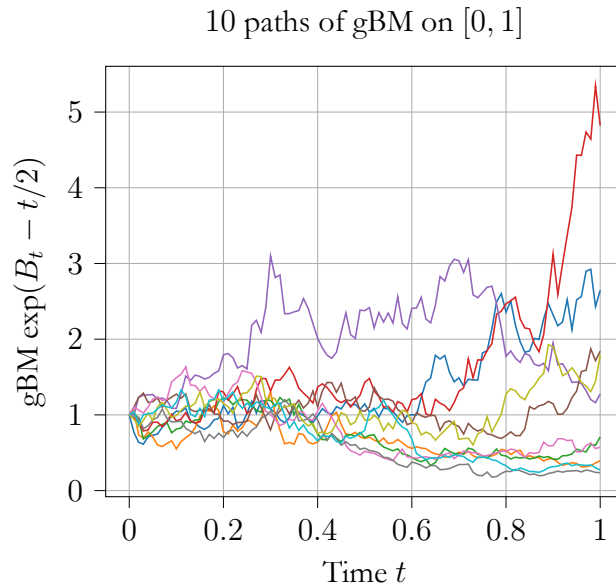
t = np.linspace(start=0.0, stop=1.0, num=N+1)
M_t = np.subtract(B_t_sq, t)
X.append(M_t)

return X

```



The gBM process can be sampled similarly, with $M_t = \text{np.exp}(\text{np.subtract}(B_t, t/2))$.



Example 5.27. (Maximum of Brownian Motion.) Consider the maximum of Brownian motion on $[0, 1]$: $\max_{s \leq 1} B_s$.

- (a) Draw the histogram of the random variable $\max_{s \leq 1} B_s$ using 10,000 sampled Brownian paths with a step size of 0.01.
- (b) Compare this to the PDF of the random variable $|B_1|$.

Solution.

I use the `itertools` python library to compute the running maximum of a brownian motion path.


Listing 12: The process $\sup_{s \leq 1} B_s$

```

brownianMotion = BrownianMotion(stepSize=0.01,T=1)
data = []

for i in range(10000):
    B_t = brownianMotion.samplePath()
    max_B_t = list(itertools.accumulate(B_t,max))
    data.append(max_B_t[100])

```



Analytically, we know that the density function of $|B_1|$ is $f_{|B_1|}(x) = f_{|Z|}(x)$